



Red Hat Enterprise Linux 8

Developing C and C++ applications in RHEL 8

Setting up a developer workstation, and developing and debugging C and C++ applications in Red Hat Enterprise Linux 8

Red Hat Enterprise Linux 8 Developing C and C++ applications in RHEL 8

Setting up a developer workstation, and developing and debugging C and C++ applications in Red Hat Enterprise Linux 8

Olga Tikhomirova
otikhomi@redhat.com

Levi Valeeva
lvaleeva@redhat.com

Vladimír Slávik
Red Hat Customer Content Services

Legal Notice

Copyright © 2020 Red Hat, Inc.

The text of and illustrations in this document are licensed by Red Hat under a Creative Commons Attribution–Share Alike 3.0 Unported license ("CC-BY-SA"). An explanation of CC-BY-SA is available at

<http://creativecommons.org/licenses/by-sa/3.0/>

. In accordance with CC-BY-SA, if you distribute this document or an adaptation of it, you must provide the URL for the original version.

Red Hat, as the licensor of this document, waives the right to enforce, and agrees not to assert, Section 4d of CC-BY-SA to the fullest extent permitted by applicable law.

Red Hat, Red Hat Enterprise Linux, the Shadowman logo, the Red Hat logo, JBoss, OpenShift, Fedora, the Infinity logo, and RHCE are trademarks of Red Hat, Inc., registered in the United States and other countries.

Linux[®] is the registered trademark of Linus Torvalds in the United States and other countries.

Java[®] is a registered trademark of Oracle and/or its affiliates.

XFS[®] is a trademark of Silicon Graphics International Corp. or its subsidiaries in the United States and/or other countries.

MySQL[®] is a registered trademark of MySQL AB in the United States, the European Union and other countries.

Node.js[®] is an official trademark of Joyent. Red Hat is not formally related to or endorsed by the official Joyent Node.js open source or commercial project.

The OpenStack[®] Word Mark and OpenStack logo are either registered trademarks/service marks or trademarks/service marks of the OpenStack Foundation, in the United States and other countries and are used with the OpenStack Foundation's permission. We are not affiliated with, endorsed or sponsored by the OpenStack Foundation, or the OpenStack community.

All other trademarks are the property of their respective owners.

Abstract

This document describes the different features and utilities that make Red Hat Enterprise Linux 8 an ideal enterprise platform for application development.

Table of Contents

PROVIDING FEEDBACK ON RED HAT DOCUMENTATION	5
CHAPTER 1. SETTING UP A DEVELOPMENT WORKSTATION	6
1.1. PREREQUISITES	6
1.2. ENABLING DEBUG AND SOURCE REPOSITORIES	6
1.3. SETTING UP TO MANAGE APPLICATION VERSIONS	6
1.4. SETTING UP TO DEVELOP APPLICATIONS USING C AND C++	7
1.5. SETTING UP TO DEBUG APPLICATIONS	8
1.6. SETTING UP TO MEASURE PERFORMANCE OF APPLICATIONS	8
PART I. CREATING C OR C++ APPLICATIONS	10
CHAPTER 2. BUILDING CODE WITH GCC	11
2.1. RELATIONSHIP BETWEEN CODE FORMS	11
2.2. COMPILING SOURCE FILES TO OBJECT CODE	12
2.3. ENABLING DEBUGGING OF C AND C++ APPLICATIONS WITH GCC	12
2.4. CODE OPTIMIZATION WITH GCC	13
2.5. OPTIONS FOR HARDENING CODE WITH GCC	14
2.6. LINKING CODE TO CREATE EXECUTABLE FILES	14
2.7. EXAMPLE: BUILDING A C PROGRAM WITH GCC	15
2.8. EXAMPLE: BUILDING A C++ PROGRAM WITH GCC	16
CHAPTER 3. USING LIBRARIES WITH GCC	17
3.1. LIBRARY NAMING CONVENTIONS	17
3.2. STATIC AND DYNAMIC LINKING	17
3.3. USING A LIBRARY WITH GCC	18
3.4. USING A STATIC LIBRARY WITH GCC	19
3.5. USING A DYNAMIC LIBRARY WITH GCC	20
3.6. USING BOTH STATIC AND DYNAMIC LIBRARIES WITH GCC	22
CHAPTER 4. CREATING LIBRARIES WITH GCC	24
4.1. LIBRARY NAMING CONVENTIONS	24
4.2. THE SONAME MECHANISM	24
4.3. CREATING DYNAMIC LIBRARIES WITH GCC	25
4.4. CREATING STATIC LIBRARIES WITH GCC AND AR	26
CHAPTER 5. MANAGING MORE CODE WITH MAKE	28
5.1. GNU MAKE AND MAKEFILE OVERVIEW	28
5.2. EXAMPLE: BUILDING A C PROGRAM USING A MAKEFILE	29
5.3. DOCUMENTATION RESOURCES FOR MAKE	30
CHAPTER 6. CHANGES IN TOOLCHAIN SINCE RHEL 7	32
6.1. CHANGES IN GCC IN RHEL 8	32
6.2. SECURITY ENHANCEMENTS IN GCC IN RHEL 8	33
6.3. COMPATIBILITY-BREAKING CHANGES IN GCC IN RHEL 8	36
C++ ABI change in <code>std::string</code> and <code>std::list</code>	36
GCC no longer builds Ada, Go, and Objective C/C++ code	36
PART II. DEBUGGING APPLICATIONS	38
CHAPTER 7. ENABLING DEBUGGING WITH DEBUGGING INFORMATION	39
7.1. DEBUGGING INFORMATION	39
7.2. ENABLING DEBUGGING OF C AND C++ APPLICATIONS WITH GCC	39
7.3. DEBUGINFO AND DEBUGSOURCE PACKAGES	40

7.4. GETTING DEBUGINFO PACKAGES FOR AN APPLICATION OR LIBRARY USING GDB	41
7.5. GETTING DEBUGINFO PACKAGES FOR AN APPLICATION OR LIBRARY MANUALLY	42
CHAPTER 8. INSPECTING APPLICATION INTERNAL STATE WITH GDB	44
8.1. GNU DEBUGGER (GDB)	44
8.2. ATTACHING GDB TO A PROCESS	44
8.3. STEPPING THROUGH PROGRAM CODE WITH GDB	45
8.4. SHOWING PROGRAM INTERNAL VALUES WITH GDB	47
8.5. USING GDB BREAKPOINTS TO STOP EXECUTION AT DEFINED CODE LOCATIONS	48
8.6. USING GDB WATCHPOINTS TO STOP EXECUTION ON DATA ACCESS AND CHANGES	49
8.7. DEBUGGING FORKING OR THREADED PROGRAMS WITH GDB	50
CHAPTER 9. RECORDING APPLICATION INTERACTIONS	52
9.1. TOOLS USEFUL FOR RECORDING APPLICATION INTERACTIONS	52
9.2. MONITORING AN APPLICATION'S SYSTEM CALLS WITH STRACE	53
9.3. MONITORING APPLICATION'S LIBRARY FUNCTION CALLS WITH LTRACE	54
9.4. MONITORING APPLICATION'S SYSTEM CALLS WITH SYSTEMTAP	56
9.5. USING GDB TO INTERCEPT APPLICATION SYSTEM CALLS	56
9.6. USING GDB TO INTERCEPT HANDLING OF SIGNALS BY APPLICATIONS	57
CHAPTER 10. DEBUGGING A CRASHED APPLICATION	59
10.1. CORE DUMPS: WHAT THEY ARE AND HOW TO USE THEM	59
10.2. RECORDING APPLICATION CRASHES WITH CORE DUMPS	59
10.3. INSPECTING APPLICATION CRASH STATES WITH CORE DUMPS	60
10.4. CREATING AND ACCESSING A CORE DUMP WITH COREDUMPCTL	62
10.5. DUMPING PROCESS MEMORY WITH GCORE	64
10.6. DUMPING PROTECTED PROCESS MEMORY WITH GDB	65
CHAPTER 11. COMPATIBILITY-BREAKING CHANGES IN GDB	66
GDBserver now starts inferiors with shell	66
gcj support removed	66
New syntax for symbol dumping maintenance commands	66
Thread numbers are no longer global	67
Memory for value contents can be limited	67
Sun version of stabs format no longer supported	68
Sysroot handling changes	68
HISTSIZE no longer controls GDB command history size	68
Completion limiting added	68
HP-UX XDB compatibility mode removed	68
Handling signals for threads	68
Breakpoint modes always-inserted off and auto merged	69
remotebaud commands no longer supported	69
PART III. ADDITIONAL TOOLSETS FOR DEVELOPMENT	70
CHAPTER 12. USING GCC TOOLSET	71
12.1. WHAT IS GCC TOOLSET	71
12.2. INSTALLING GCC TOOLSET	71
12.3. INSTALLING INDIVIDUAL PACKAGES FROM GCC TOOLSET	71
12.4. UNINSTALLING GCC TOOLSET	71
12.5. RUNNING A TOOL FROM GCC TOOLSET	72
12.6. RUNNING A SHELL SESSION WITH GCC TOOLSET	72
12.7. RELATED INFORMATION	72
CHAPTER 13. GCC TOOLSET 9	73

13.1. TOOLS AND VERSIONS PROVIDED BY GCC TOOLSET 9	73
13.2. C++ COMPATIBILITY IN GCC TOOLSET 9	73
13.3. SPECIFICS OF GCC IN GCC TOOLSET 9	74
13.4. SPECIFICS OF BINUTILS IN GCC TOOLSET 9	75
CHAPTER 14. USING THE GCC TOOLSET CONTAINER IMAGES	76
14.1. GCC TOOLSET CONTAINER IMAGES CONTENTS	76
14.2. ACCESSING AND RUNNING THE GCC TOOLSET CONTAINER IMAGES	77
14.3. EXAMPLE: USING THE GCC TOOLSET 9 TOOLCHAIN CONTAINER IMAGE	78
CHAPTER 15. COMPILER TOOLSETS	79
CHAPTER 16. THE ANNOBIN PROJECT	80
16.1. USING THE ANNOBIN PLUGIN	80
16.1.1. Enabling the annobin plugin	80
16.1.2. Passing options to the annobin plugin	81
16.2. USING THE ANNOCHECK PROGRAM	81
16.2.1. Using annoscheck to examine files	82
16.2.2. Using annoscheck to examine directories	82
16.2.3. Using annoscheck to examine RPM packages	83
16.2.4. Using annoscheck extra tools	83
16.2.4.1. Enabling the built-by tool	83
16.2.4.2. Enabling the notes tool	84
16.2.4.3. Enabling the section-size tool	84
16.2.4.4. Hardening checker basics	84
16.2.4.4.1. Hardening checker options	84
16.2.4.4.2. Disabling the hardening checker	85
16.3. REMOVING REDUNDANT ANNOBIN NOTES	85
PART IV. SUPPLEMENTARY TOPICS	87
CHAPTER 17. COMPATIBILITY-BREAKING CHANGES IN COMPILERS AND DEVELOPMENT TOOLS	88
librtkai removed	88
Sun RPC and NIS interfaces removed from glibc	88
The nose neg libraries for 32-bit Xen have been removed	88
make new operator != causes a different interpretation of certain existing makefile syntax	88
Valgrind library for MPI debugging support removed	89
Development headers and static libraries removed from valgrind-devel	89
CHAPTER 18. OPTIONS FOR RUNNING A RHEL 6 OR 7 APPLICATION ON RHEL 8	90

PROVIDING FEEDBACK ON RED HAT DOCUMENTATION

We appreciate your input on our documentation. Please let us know how we could make it better. To do so:

- For simple comments on specific passages:
 1. Make sure you are viewing the documentation in the *Multi-page HTML* format. In addition, ensure you see the **Feedback** button in the upper right corner of the document.
 2. Use your mouse cursor to highlight the part of text that you want to comment on.
 3. Click the **Add Feedback** pop-up that appears below the highlighted text.
 4. Follow the displayed instructions.
- For submitting more complex feedback, create a Bugzilla ticket:
 1. Go to the [Bugzilla](#) website.
 2. As the Component, use **Documentation**.
 3. Fill in the **Description** field with your suggestion for improvement. Include a link to the relevant part(s) of documentation.
 4. Click **Submit Bug**.

CHAPTER 1. SETTING UP A DEVELOPMENT WORKSTATION

Red Hat Enterprise Linux 8 supports development of custom applications. To allow developers to do so, the system must be set up with the required tools and utilities. This chapter lists the most common use cases for development and the items to install.

1.1. PREREQUISITES

- The system must be installed, including a graphical environment, and subscribed.

1.2. ENABLING DEBUG AND SOURCE REPOSITORIES

A standard installation of Red Hat Enterprise Linux does not enable the debug and source repositories. These repositories contain information needed to debug the system components and measure their performance.

Procedure

- Enable the source and debug information package channels:

```
# subscription-manager repos --enable rhel-8-for-$(uname -i)-baseos-debug-rpms
# subscription-manager repos --enable rhel-8-for-$(uname -i)-baseos-source-rpms
# subscription-manager repos --enable rhel-8-for-$(uname -i)-appstream-debug-rpms
# subscription-manager repos --enable rhel-8-for-$(uname -i)-appstream-source-rpms
```

The **\$(uname -i)** part is automatically replaced with a matching value for architecture of your system:

Architecture name	Value
64-bit Intel and AMD	x86_64
64-bit ARM	aarch64
IBM POWER	ppc64le
IBM Z	s390x

1.3. SETTING UP TO MANAGE APPLICATION VERSIONS

Effective version control is essential to all multi-developer projects. Red Hat Enterprise Linux is shipped with Git, a distributed version control system.

Procedure

1. Install the **git** package:

```
# yum install git
```

2. Optional: Set the full name and email address associated with your Git commits:

```
$ git config --global user.name "Full Name"
$ git config --global user.email "email@example.com"
```

Replace *Full Name* and *email@example.com* with your actual name and email address.

- Optional: To change the default text editor started by Git, set value of the **core.editor** configuration option:

```
$ git config --global core.editor command
```

Replace *command* with the command to be used to start the selected text editor.

Additional resources

- Linux manual pages for Git and tutorials:

```
$ man git
$ man gittutorial
$ man gittutorial-2
```

Note that many Git commands have their own manual pages. As an example see *git-commit(1)*.

- Git User's Manual* – HTML documentation for Git is located at **/usr/share/doc/git/user-manual.html**.
- [Pro Git](#) – The online version of the *Pro Git* book provides a detailed description of Git, its concepts, and its usage.
- [Reference](#) – Online version of the Linux manual pages for Git

1.4. SETTING UP TO DEVELOP APPLICATIONS USING C AND C++

Red Hat Enterprise Linux includes tools for creating C and C++ applications.

Prerequisites

- The debug and source repositories must be enabled.

Procedure

- Install the **Development Tools** package group including GNU Compiler Collection (GCC), GNU Debugger (GDB), and other development tools:

```
# yum group install "Development Tools"
```

- Install the LLVM-based toolchain including the **clang** compiler and **lldb** debugger:

```
# yum install llvm-toolset
```

- Optional: For Fortran dependencies, install the GNU Fortran compiler:

```
# yum install gcc-gfortran
```

1.5. SETTING UP TO DEBUG APPLICATIONS

Red Hat Enterprise Linux offers multiple debugging and instrumentation tools to analyze and troubleshoot internal application behavior.

Prerequisites

- The debug and source repositories must be enabled.

Procedure

1. Install the tools useful for debugging:

```
# yum install gdb valgrind systemtap ltrace strace
```

2. Install the **yum-utils** package in order to use the **debuginfo-install** tool:

```
# yum install yum-utils
```

3. Run a SystemTap helper script for setting up the environment.

```
# stap-prep
```

Note that **stap-prep** installs packages relevant to the currently *running* kernel, which might not be the same as the actually installed kernel(s). To ensure **stap-prep** installs the correct **kernel-debuginfo** and **kernel-headers** packages, double-check the current kernel version by using the **uname -r** command and reboot your system if necessary.

4. Make sure **SELinux** policies allow the relevant applications to run not only normally, but in the debugging situations, too. For more information, see [Using SELinux](#).

Additional resources

- [Chapter 7, Enabling Debugging with Debugging Information](#)

1.6. SETTING UP TO MEASURE PERFORMANCE OF APPLICATIONS

Red Hat Enterprise Linux includes several applications that can help a developer identify the causes of application performance loss.

Prerequisites

- The debug and source repositories must be enabled.

Procedure

1. Install the tools for performance measurement:

```
# yum install perf papi pcp-zeroconf valgrind strace sysstat systemtap
```

2. Run a SystemTap helper script for setting up the environment.

```
# stap-prep
```

-

Note that **stap-prep** installs packages relevant to the currently *running* kernel, which might not be the same as the actually installed kernel(s). To ensure **stap-prep** installs the correct **kernel-debuginfo** and **kernel-headers** packages, double-check the current kernel version by using the **uname -r** command and reboot your system if necessary.

3. Enable and start the Performance Co-Pilot (PCP) collector service:

```
# systemctl enable pmcd && systemctl start pmcd
```

PART I. CREATING C OR C++ APPLICATIONS

Red Hat offers multiple tools for creating applications using the C and C++ languages. This part of the book lists some of the most common development tasks.

CHAPTER 2. BUILDING CODE WITH GCC

This chapter describes situations where source code must be transformed into executable code.

2.1. RELATIONSHIP BETWEEN CODE FORMS

Prerequisites

- Understanding the concepts of compiling and linking


Possible code forms

The C and C++ languages have three forms of code:

- **Source code** written in the C or C++ language, present as plain text files. The files typically use extensions such as **.c**, **.cc**, **.cpp**, **.h**, **.hpp**, **.i**, **.inc**. For a complete list of supported extensions and their interpretation, see the gcc manual pages:

```
$ man gcc
```

- **Object code**, created by *compiling* the source code with a *compiler*. This is an intermediate form. The object code files use the **.o** extension.
- **Executable code**, created by *linking* object code with a *linker*. Linux application executable files do not use any file name extension. Shared object (library) executable files use the **.so** file name extension.



NOTE

Library archive files for static linking also exist. This is a variant of object code that uses the **.a** file name extension. Static linking is not recommended. See [Section 3.2, "Static and dynamic linking"](#).

Handling of code forms in GCC

Producing executable code from source code is performed in two steps, which require different applications or tools. GCC can be used as an intelligent driver for both compilers and linkers. This allows you to use a single **gcc** command for any of the required actions (compiling and linking). GCC automatically selects the actions and their sequence:

1. Source files are compiled to object files.
2. Object files and libraries are linked (including the previously compiled sources).

It is possible to run GCC so that it performs only step 1, only step 2, or both steps 1 and 2. This is determined by the types of inputs and requested type of output(s).

Because larger projects require a build system which usually runs GCC separately for each action, it is better to always consider compilation and linking as two distinct actions, even if GCC can perform both at once.

Additional resources

- [Section 2.2, "Compiling source files to object code"](#)
- [Section 2.6, "Linking code to create executable files"](#)

2.2. COMPILING SOURCE FILES TO OBJECT CODE

To create object code files from source files and not an executable file immediately, GCC must be instructed to create only object code files as its output. This action represents the basic operation of the build process for larger projects.

Prerequisites

- C or C++ source code file(s)
- [GCC installed on the system](#)

Procedure

1. Change to the directory containing the source code file(s).
2. Run **gcc** with the **-c** option:

```
$ gcc -c source.c another_source.c
```

Object files are created, with their file names reflecting the original source code files: **source.c** results in **source.o**.



NOTE

With C++ source code, replace the **gcc** command with **g++** for convenient handling of C++ Standard Library dependencies.

Additional resources

- [Section 2.5, "Options for hardening code with GCC"](#)
- [Section 2.4, "Code optimization with GCC"](#)
- [Section 2.7, "Example: Building a C program with GCC"](#)

2.3. ENABLING DEBUGGING OF C AND C++ APPLICATIONS WITH GCC

Because debugging information is large, it is not included in executable files by default. To enable debugging of your C and C++ applications with it, you must explicitly instruct the compiler to create it.

To enable creation of debugging information with **GCC** when compiling and linking code, use the **-g** option:

```
$ gcc ... -g ...
```

- Optimizations performed by the compiler and linker can result in executable code which is hard to relate to the original source code: variables may be optimized out, loops unrolled, operations merged into the surrounding ones etc. This affects debugging negatively. For improved

debugging experience, consider setting the optimization with the **-Og** option. However, changing the optimization level changes the executable code and may change the actual behaviour including removing some bugs.

- To also include macro definitions in the debug information, use the **-g3** option instead of **-g**.
- The **-fcompare-debug** GCC option tests code compiled by GCC with debug information and without debug information. The test passes if the resulting two binary files are identical. This test ensures that executable code is not affected by any debugging options, which further ensures that there are no hidden bugs in the debug code. Note that using the **-fcompare-debug** option significantly increases compilation time. See the GCC manual page for details about this option.

Additional resources

- [Chapter 7, Enabling Debugging with Debugging Information](#)
- Using the GNU Compiler Collection (GCC) – [Options for Debugging Your Program](#)
- Debugging with GDB – [Debugging Information in Separate Files](#)
- The GCC manual page:

```
$ man gcc
```

2.4. CODE OPTIMIZATION WITH GCC

A single program can be transformed into more than one sequence of machine instructions. You can achieve a more optimal result if you allocate more resources to analyzing the code during compilation.

With GCC, you can set the optimization level using the **-Olevel** option. This option accepts a set of values in place of the *level*.

Level	Description
0	Optimize for compilation speed - no code optimization (default).
1, 2, 3	Optimize to increase code execution speed (the larger the number, the greater the speed).
s	Optimize for file size.
fast	Same as a level 3 setting, plus fast disregards strict standards compliance to allow for additional optimizations.
g	Optimize for debugging experience.

For release builds, use the optimization option **-O2**.

During development, the **-Og** option is useful for debugging the program or library in some situations. Because some bugs manifest only with certain optimization levels, test the program or library with the release optimization level.

GCC offers a large number of options to enable individual optimizations. For more information, see the following Additional resources.

Additional resources

- Using GNU Compiler Collection – [3.10 Options That Control Optimization](#)
- Linux manual page for GCC:

```
$ man gcc
```

2.5. OPTIONS FOR HARDENING CODE WITH GCC

When the compiler transforms source code to object code, it can add various checks to prevent commonly exploited situations and increase security. Choosing the right set of compiler options can help produce more secure programs and libraries, without having to change the source code.

Release version options

The following list of options is the recommended minimum for developers targeting Red Hat Enterprise Linux:

```
$ gcc ... -O2 -g -Wall -Wl,-z,now,-z,relro -fstack-protector-strong -fstack-clash-protection -D_FORTIFY_SOURCE=2 ...
```

- For programs, add the **-fPIE** and **-pie** Position Independent Executable options.
- For dynamically linked libraries, the mandatory **-fPIC** (Position Independent Code) option indirectly increases security.

Development options

Use the following options to detect security flaws during development. Use these options in conjunction with the options for the release version:

```
$ gcc ... -Walloc-zero -Walloca-larger-than -Wextra -Wformat-security -Wvla-larger-than ...
```

Additional resources

- [Defensive Coding Guide](#)
- [Memory Error Detection Using GCC](#) – Red Hat Developers Blog post

2.6. LINKING CODE TO CREATE EXECUTABLE FILES

Linking is the final step when building a C or C++ application. Linking combines all object files and libraries into an executable file.

Prerequisites

- One or more object file(s)
- [GCC must be installed on the system](#)


Procedure

1. Change to the directory containing the object code file(s).
2. Run **gcc**:

```
$ gcc ... objfile.o another_object.o ... -o executable-file
```

An executable file named **executable-file** is created from the supplied object files and libraries.

To link additional libraries, add the required options after the list of object files. For more information, see [Chapter 3, Using Libraries with GCC](#).



NOTE

With C++ source code, replace the **gcc** command with **g++** for convenient handling of C++ Standard Library dependencies.

Additional resources

- [Section 2.7, “Example: Building a C program with GCC”](#)
- [Section 3.2, “Static and dynamic linking”](#)

2.7. EXAMPLE: BUILDING A C PROGRAM WITH GCC

This example shows the exact steps to build a simple sample C program.

Prerequisites

- You must understand how to use GCC.

Procedure

1. Create a directory **hello-c** and change to it:

```
$ mkdir hello-c
$ cd hello-c
```

2. Create file **hello.c** with the following contents:

```
#include <stdio.h>

int main() {
    printf("Hello, World!\n");
    return 0;
}
```

3. Compile the code with GCC:

```
$ gcc -c hello.c
```

The object file **hello.o** is created.

4. Link an executable file **helloworld** from the object file:

```
$ gcc hello.o -o helloworld
```

5. Run the resulting executable file:

```
$ ./helloworld  
Hello, World!
```

Additional resources

- [Section 5.2, “Example: Building a C program using a Makefile”](#)

2.8. EXAMPLE: BUILDING A C++ PROGRAM WITH GCC

This example shows the exact steps to build a sample minimal C++ program.

Prerequisites

- You must understand the difference between **gcc** and **g++**.

Procedure

1. Create a directory **hello-cpp** and change to it:

```
$ mkdir hello-cpp  
$ cd hello-cpp
```

2. Create file **hello.cpp** with the following contents:

```
#include <iostream>  
  
int main() {  
    std::cout << "Hello, World!\n";  
    return 0;  
}
```

3. Compile the code with **g++**:

```
$ g++ -c hello.cpp
```

The object file **hello.o** is created.

4. Link an executable file **helloworld** from the object file:

```
$ g++ hello.o -o helloworld
```

5. Run the resulting executable file:

```
$ ./helloworld  
Hello, World!
```

CHAPTER 3. USING LIBRARIES WITH GCC

This chapter describes using libraries in code.

3.1. LIBRARY NAMING CONVENTIONS

A special file name convention is used for libraries: a library known as **foo** is expected to exist as file **libfoo.so** or **libfoo.a**. This convention is automatically understood by the linking input options of GCC, but not by the output options:

- When linking against the library, the library can be specified only by its name **foo** with the **-l** option as **-lfoo**:

```
$ gcc ... -lfoo ...
```

- When creating the library, the full file name **libfoo.so** or **libfoo.a** must be specified.

Additional resources

- [Section 4.2, “The soname mechanism”](#)

3.2. STATIC AND DYNAMIC LINKING

Developers have a choice of using static or dynamic linking when building applications with fully compiled languages. This section lists the differences, particularly in the context using the C and C++ languages on Red Hat Enterprise Linux. To summarize, Red Hat discourages the use of static linking in applications for Red Hat Enterprise Linux.

Comparison of static and dynamic linking

Static linking makes libraries part of the resulting executable file. Dynamic linking keeps these libraries as separate files.

Dynamic and static linking can be compared in a number of ways:

Resource use

Static linking results in larger executable files which contain more code. This additional code coming from libraries cannot be shared across multiple programs on the system, increasing file system usage and memory usage at run time. Multiple processes running the same statically linked program will still share the code.

On the other hand, static applications need fewer run-time relocations, leading to reduced startup time, and require less private resident set size (RSS) memory. Generated code for static linking can be more efficient than for dynamic linking due to the overhead introduced by position-independent code (PIC).

Security

Dynamically linked libraries which provide ABI compatibility can be updated without changing the executable files depending on these libraries. This is especially important for libraries provided by Red Hat as part of Red Hat Enterprise Linux, where Red Hat provides security updates. Static linking against any such libraries is strongly discouraged.

Compatibility

Static linking appears to provide executable files independent of the versions of libraries provided by the operating system. However, most libraries depend on other libraries. With static linking, this

dependency becomes inflexible and as a result, both forward and backward compatibility is lost. Static linking is guaranteed to work only on the system where the executable file was built.



WARNING

Applications linking statically libraries from the GNU C library (**glibc**) still require **glibc** to be present on the system as a dynamic library. Furthermore, the dynamic library variant of **glibc** available at the application's run time must be a bitwise identical version to that present while linking the application. As a result, static linking is guaranteed to work only on the system where the executable file was built.

Support coverage

Most static libraries provided by Red Hat are in the *CodeReady Linux Builder* channel and not supported by Red Hat.

Functionality

Some libraries, notably the GNU C Library (**glibc**), offer reduced functionality when linked statically. For example, when statically linked, **glibc** does not support threads and any form of calls to the **dlopen()** function in the same program.

As a result of the listed disadvantages, static linking should be avoided at all costs, particularly for whole applications and the **glibc** and **libstdc++** libraries.

Cases for static linking

Static linking might be a reasonable choice in some cases, such as:

- Using a library which is not enabled for dynamic linking.
- Fully static linking can be required for running code in an empty **chroot** environment or container. However, static linking using the **glibc-static** package is not supported by Red Hat.

Additional resources

- [Red Hat Enterprise Linux 8: Application Compatibility GUIDE](#)
- Description of the [The CodeReady Linux Builder repository](#) in the *Package manifest*

3.3. USING A LIBRARY WITH GCC

A library is a package of code which can be reused in your program. A C or C++ library consists of two parts:

- The library code
- Header files

Compiling code that uses a library

The header files describe the interface of the library: the functions and variables available in the library. Information from the header files is needed for compiling the code.

Typically, header files of a library will be placed in a different directory than your application's code. To tell GCC where the header files are, use the **-I** option:

```
$ gcc ... -Iinclude_path ...
```

Replace *include_path* with the actual path to the header file directory.

The **-I** option can be used multiple times to add multiple directories with header files. When looking for a header file, these directories are searched in the order of their appearance in the **-I** options.

Linking code that uses a library

When linking the executable file, both the object code of your application and the binary code of the library must be available. The code for static and dynamic libraries is present in different forms:

- Static libraries are available as archive files. They contain a group of object files. The archive file has a file name extension **.a**.
- Dynamic libraries are available as shared objects. They are a form of an executable file. A shared object has a file name extension **.so**.

To tell GCC where the archives or shared object files of a library are, use the **-L** option:

```
$ gcc ... -Llibrary_path -lfoo ...
```

Replace *library_path* with the actual path to the library directory.

The **-L** option can be used multiple times to add multiple directories. When looking for a library, these directories are searched in the order of their **-L** options.

The order of options matters: GCC cannot link against a library **foo** unless it knows the directory with this library. Therefore, use the **-L** options to specify library directories before using the **-l** options for linking against libraries.

Compiling and linking code which uses a library in one step

When the situation allows the code to be compiled and linked in one **gcc** command, use the options for both situations mentioned above at once.

Additional resources

- Using the GNU Compiler Collection (GCC) – [3.15 Options for Directory Search](#)
- Using the GNU Compiler Collection (GCC) – [3.14 Options for Linking](#)

3.4. USING A STATIC LIBRARY WITH GCC

Static libraries are available as archives containing object files. After linking, they become part of the resulting executable file.

**NOTE**

Red Hat discourages use of static linking for security reasons. See [Section 3.2, “Static and dynamic linking”](#). Use static linking only when necessary, especially against libraries provided by Red Hat.

Prerequisites

- [GCC must be installed on your system.](#)
- [You must understand static and dynamic linking.](#)
- You have a set of source or object files forming a valid program, requiring some static library **foo** and no other libraries.
- The **foo** library is available as a file **libfoo.a**, and no file **libfoo.so** is provided for dynamic linking.

**NOTE**

Most libraries which are part of Red Hat Enterprise Linux are supported for dynamic linking only. The steps below only work for libraries which are *not* enabled for dynamic linking. See [Section 3.2, “Static and dynamic linking”](#).

Procedure

To link a program from source and object files, adding a statically linked library **foo**, which is to be found as a file **libfoo.a**:

1. Change to the directory containing your code.
2. Compile the program source files with headers of the **foo** library:

```
$ gcc ... -Iheader_path -c ...
```

Replace *header_path* with a path to a directory containing the header files for the **foo** library.

3. Link the program with the **foo** library:

```
$ gcc ... -Llibrary_path -lfoo ...
```

Replace *library_path* with a path to a directory containing the file **libfoo.a**.

4. To run the program later, simply:

```
$ ./program
```

CAUTION

The **-static** GCC option related to static linking forbids all dynamic linking. Instead, use the **-Wl,-Bstatic** and **-Wl,-Bdynamic** options to control linker behavior more precisely. See [Section 3.6, “Using both static and dynamic libraries with GCC”](#).

3.5. USING A DYNAMIC LIBRARY WITH GCC

Dynamic libraries are available as standalone executable files, required at both linking time and run time. They stay independent of your application's executable file.

Prerequisites

- [GCC must be installed on the system.](#)
- A set of source or object files forming a valid program, requiring some dynamic library **foo** and no other libraries.
- The **foo** library must be available as a file *libfoo.so*.

Linking a program against a dynamic library

To link a program against a dynamic library **foo**:

```
$ gcc ... -Llibrary_path -lfoo ...
```

When a program is linked against a dynamic library, the resulting program must always load the library at run time. There are two options for locating the library:

- Using a **rpath** value stored in the executable file itself
- Using the **LD_LIBRARY_PATH** variable at run time

Using a rpath Value Stored in the Executable File

The **rpath** is a special value saved as a part of an executable file when it is being linked. Later, when the program is loaded from its executable file, the runtime linker will use the **rpath** value to locate the library files.

While linking with **GCC**, to store the path *library_path* as **rpath**:

```
$ gcc ... -Llibrary_path -lfoo -Wl,-rpath=library_path ...
```

The path *library_path* must point to a directory containing the file *libfoo.so*.

CAUTION

There is no space after the comma in the **-Wl,-rpath=** option!

To run the program later:

```
$ ./program
```

Using the LD_LIBRARY_PATH environment variable

If no **rpath** is found in the program's executable file, the runtime linker will use the **LD_LIBRARY_PATH** environment variable. The value of this variable must be changed for each program. This value should represent the path where the shared library objects are located.

To run the program without **rpath** set, with libraries present in path *library_path*:

```
$ export LD_LIBRARY_PATH=library_path:$LD_LIBRARY_PATH
$ ./program
```

Leaving out the **rpath** value offers flexibility, but requires setting the **LD_LIBRARY_PATH** variable every time the program is to run.

Placing the Library into the Default Directories

The runtime linker configuration specifies a number of directories as a default location of dynamic library files. To use this default behaviour, copy your library to the appropriate directory.

A full description of the dynamic linker behavior is out of scope of this document. For more information, see the following resources:

- Linux manual pages for the dynamic linker:

```
$ man ld.so
```

- Contents of the **/etc/ld.so.conf** configuration file:

```
$ cat /etc/ld.so.conf
```

- Report of the libraries recognized by the dynamic linker without additional configuration, which includes the directories:

```
$ ldconfig -v
```

3.6. USING BOTH STATIC AND DYNAMIC LIBRARIES WITH GCC

Sometimes it is required to link some libraries statically and some dynamically. This situation brings some challenges.

Prerequisites

- [Understanding static and dynamic linking](#)

Introduction

gcc recognizes both dynamic and static libraries. When the **-lfoo** option is encountered, **gcc** will first attempt to locate a shared object (a **.so** file) containing a dynamically linked version of the **foo** library, and then look for the archive file (**.a**) containing a static version of the library. Thus, the following situations can result from this search:

- Only the shared object is found, and **gcc** links against it dynamically.
- Only the archive is found, and **gcc** links against it statically.
- Both the shared object and archive are found, and by default, **gcc** selects dynamic linking against the shared object.
- Neither shared object nor archive is found, and linking fails.

Because of these rules, the best way to select the static or dynamic version of a library for linking is having only that version found by **gcc**. This can be controlled to some extent by using or leaving out directories containing the library versions, when specifying the **-Lpath** options.

Additionally, because dynamic linking is the default, the only situation where linking must be explicitly specified is when a library with both versions present should be linked statically. There are two possible resolutions:

- Specifying the static libraries by file path instead of the **-l** option
- Using the **-Wl** option to pass options to the linker

Specifying the static libraries by file

Usually, **gcc** is instructed to link against the **foo** library with the **-lfoo** option. However, it is possible to specify the full path to file **libfoo.a** containing the library instead:

```
$ gcc ... path/to/libfoo.a ...
```

From the file extension **.a**, **gcc** will understand that this is a library to link with the program. However, specifying the full path to the library file is a less flexible method.

Using the **-Wl** option

The **gcc** option **-Wl** is a special option for passing options to the underlying linker. Syntax of this option differs from the other **gcc** options. The **-Wl** option is followed by a comma-separated list of linker options, while other **gcc** options require space-separated list of options.

The **ld** linker used by **gcc** offers the options **-Bstatic** and **-Bdynamic** to specify whether libraries following this option should be linked statically or dynamically, respectively. After passing **-Bstatic** and a library to the linker, the default dynamic linking behaviour must be restored manually for the following libraries to be linked dynamically with the **-Bdynamic** option.

To link a program, link library **first** statically (**libfirst.a**) and **second** dynamically (**libsecond.so**):

```
$ gcc ... -Wl,-Bstatic -lfirst -Wl,-Bdynamic -lsecond ...
```



NOTE

gcc can be configured to use linkers other than the default **ld**.

Additional resources

- Using the GNU Compiler Collection (GCC) – [3.14 Options for Linking](#)
- Documentation for binutils 2.27 – [2.1 Command Line Options](#)

CHAPTER 4. CREATING LIBRARIES WITH GCC

This chapter describes steps for creating libraries and explains the necessary concepts used by the Linux operating system for libraries.

4.1. LIBRARY NAMING CONVENTIONS

A special file name convention is used for libraries: a library known as **foo** is expected to exist as file **libfoo.so** or **libfoo.a**. This convention is automatically understood by the linking input options of GCC, but not by the output options:

- When linking against the library, the library can be specified only by its name **foo** with the **-l** option as **-lfoo**:

```
$ gcc ... -lfoo ...
```

- When creating the library, the full file name **libfoo.so** or **libfoo.a** must be specified.

Additional resources

- [Section 4.2, “The soname mechanism”](#)

4.2. THE SONAME MECHANISM

Dynamically loaded libraries (shared objects) use a mechanism called *soname* to manage multiple compatible versions of a library.

Prerequisites

- [You must understand dynamic linking and libraries.](#)
- You must understand the concept of ABI compatibility.
- [You must understand library naming conventions.](#)
- You must understand symbolic links.

Problem introduction

A dynamically loaded library (shared object) exists as an independent executable file. This makes it possible to update the library without updating the applications that depend on it. However, the following problems arise with this concept:

- Identification of the actual version of the library
- Need for multiple versions of the same library present
- Signalling ABI compatibility of each of the multiple versions

The soname mechanism

To resolve this, Linux uses a mechanism called *soname*.

A **foo** library version *X.Y* is ABI-compatible with other versions with the same value of *X* in a version number. Minor changes preserving compatibility increase the number *Y*. Major changes that break compatibility increase the number *X*.

The actual **foo** library version *X.Y* exists as a file **libfoo.so.x.y**. Inside the library file, a soname is recorded with value **libfoo.so.x** to signal the compatibility.

When applications are built, the linker looks for the library by searching for the file **libfoo.so**. A symbolic link with this name must exist, pointing to the actual library file. The linker then reads the soname from the library file and records it into the application executable file. Finally, the linker creates the application that declares dependency on the library using the soname, not a name or a file name.

When the runtime dynamic linker links an application before running, it reads the soname from application's executable file. This soname is **libfoo.so.x**. A symbolic link with this name must exist, pointing to the actual library file. This allows loading the library, regardless of the *Y* component of a version, because the soname does not change.



NOTE

The *Y* component of the version number is not limited to just a single number. Additionally, some libraries encode their version in their name.

Reading soname from a file

To display the soname of a library file **somelibrary**:

```
$ objdump -p somelibrary | grep SONAME
```

Replace *somelibrary* with the actual file name of the library you wish to examine.

4.3. CREATING DYNAMIC LIBRARIES WITH GCC

Dynamically linked libraries (shared objects) allow:

- resource conservation through code reuse
- increased security by making it easier to update the library code

Follow these steps to build and install a dynamic library from source.

Prerequisites

- [You must understand the soname mechanism.](#)
- [GCC must be installed on the system.](#)
- You must have source code for a library.

Procedure

1. Change to the directory with library sources.
2. Compile each source file to an object file with the Position independent code option **-fPIC**:

```
$ gcc ... -c -fPIC some_file.c ...
```

The object files have the same file names as the original source code files, but their extension is **.o**.

3. Link the shared library from the object files:

```
$ gcc -shared -o libfoo.so.x.y -Wl,-soname,libfoo.so.x some_file.o ...
```

The used major version number is X and minor version number Y.

4. Copy the **libfoo.so.x.y** file to an appropriate location, where the system's dynamic linker can find it. On Red Hat Enterprise Linux, the directory for libraries is **/usr/lib64**:

```
# cp libfoo.so.x.y /usr/lib64
```

Note that you need root permissions to manipulate files in this directory.

5. Create the symlink structure for soname mechanism:

```
# ln -s libfoo.so.x.y libfoo.so.x
# ln -s libfoo.so.x libfoo.so
```

Additional resources

- The Linux Documentation Project – Program Library HOWTO – [3. Shared Libraries](#)

4.4. CREATING STATIC LIBRARIES WITH GCC AND AR

Creating libraries for static linking is possible through conversion of object files into a special type of archive file.



NOTE

Red Hat discourages the use of static linking for security reasons. Use static linking only when necessary, especially against libraries provided by Red Hat. See [Section 3.2, "Static and dynamic linking"](#) for more details.

Prerequisites

- [GCC and binutils must be installed on the system.](#)
- [You must understand static and dynamic linking.](#)
- Source file(s) with functions to be shared as a library are available.

Procedure

1. Create intermediate object files with GCC.

```
$ gcc -c source_file.c ...
```

Append more source files if required. The resulting object files share the file name but use the **.o** file name extension.

2. Turn the object files into a static library (archive) using the **ar** tool from the **binutils** package.

```
$ ar rcs libfoo.a source_file.o ...
```

File **libfoo.a** is created.

3. Use the **nm** command to inspect the resulting archive:

```
$ nm libfoo.a
```

4. Copy the static library file to the appropriate directory.
5. When linking against the library, GCC will automatically recognize from the **.a** file name extension that the library is an archive for static linking.

```
$ gcc ... -lfoo ...
```

Additional resources

- Linux manual page for *ar(1)*:

```
$ man ar
```

CHAPTER 5. MANAGING MORE CODE WITH MAKE

The GNU **make** utility, commonly abbreviated **make**, is a tool for controlling the generation of executables from source files. **make** automatically determines which parts of a complex program have changed and need to be recompiled. **make** uses configuration files called Makefiles to control the way programs are built.

5.1. GNU MAKE AND MAKEFILE OVERVIEW

To create a usable form (usually executable files) from the source files of a particular project, perform several necessary steps. Record the actions and their sequence to be able to repeat them later.

Red Hat Enterprise Linux contains GNU **make**, a build system designed for this purpose.

Prerequisites

- Understanding the concepts of compiling and linking

GNU **make**

GNU **make** reads Makefiles which contain the instructions describing the build process. A Makefile contains multiple *rules* that describe a way to satisfy a certain condition (*target*) with a specific action (*recipe*). Rules can hierarchically depend on another rule.

Running **make** without any options makes it look for a Makefile in the current directory and attempt to reach the default target. The actual Makefile file name can be one of **Makefile**, **makefile**, and **GNUmakefile**. The default target is determined from the Makefile contents.

Makefile details

Makefiles use a relatively simple syntax for defining *variables* and *rules*, which consists of a *target* and a *recipe*. The target specifies what is the output if a rule is executed. The lines with recipes must start with the TAB character.

Typically, a Makefile contains rules for compiling source files, a rule for linking the resulting object files, and a target that serves as the entry point at the top of the hierarchy.

Consider the following **Makefile** for building a C program which consists of a single file, **hello.c**.

```
all: hello

hello: hello.o
    gcc hello.o -o hello

hello.o: hello.c
    gcc -c hello.c -o hello.o
```

This example shows that to reach the target **all**, file **hello** is required. To get **hello**, one needs **hello.o** (linked by **gcc**), which in turn is created from **hello.c** (compiled by **gcc**).

The target **all** is the default target because it is the first target that does not start with a period (.). Running **make** without any arguments is then identical to running **make all**, when the current directory contains this **Makefile**.

Typical makefile

A more typical Makefile uses variables for generalization of the steps and adds a target "clean" - remove everything but the source files.

```
CC=gcc
CFLAGS=-c -Wall
SOURCE=hello.c
OBJ=$(SOURCE:.c=.o)
EXE=hello

all: $(SOURCE) $(EXE)

$(EXE): $(OBJ)
    $(CC) $(OBJ) -o $$@

%.o: %.c
    $(CC) $(CFLAGS) $< -o $$@

clean:
    rm -rf $(OBJ) $(EXE)
```

Adding more source files to such Makefile requires only adding them to the line where the SOURCE variable is defined.

Additional resources

- GNU make: Introduction – [2 An Introduction to Makefiles](#)
- [Chapter 2, Building Code with GCC](#)

5.2. EXAMPLE: BUILDING A C PROGRAM USING A MAKEFILE

Build a sample C program using a Makefile by following the steps in this example.

Prerequisites

- You must understand the concepts of Makefiles and **make**.

Procedure

1. Create a directory **hellomake** and change to this directory:

```
$ mkdir hellomake
$ cd hellomake
```

2. Create a file **hello.c** with the following contents:

```
#include <stdio.h>

int main(int argc, char *argv[]) {
    printf("Hello, World!\n");
    return 0;
}
```

3. Create a file **Makefile** with the following contents:

```
CC=gcc
CFLAGS=-c -Wall
SOURCE=hello.c
OBJ=$(SOURCE:.c=.o)
EXE=hello

all: $(SOURCE) $(EXE)

$(EXE): $(OBJ)
    $(CC) $(OBJ) -o $@

%.o: %.c
    $(CC) $(CFLAGS) $< -o $@

clean:
    rm -rf $(OBJ) $(EXE)
```

CAUTION

The Makefile recipe lines must start with the tab character! When copying the text above from the documentation, the cut-and-paste process may paste spaces instead of tabs. If this happens, correct the issue manually.

4. Run **make**:

```
$ make
gcc -c -Wall hello.c -o hello.o
gcc hello.o -o hello
```

This creates an executable file **hello**.

5. Run the executable file **hello**:

```
$ ./hello
Hello, World!
```

6. Run the Makefile target **clean** to remove the created files:

```
$ make clean
rm -rf hello.o hello
```

Additional resources

- [Section 2.7, “Example: Building a C program with GCC”](#)
- [Section 2.8, “Example: Building a C++ program with GCC”](#)

5.3. DOCUMENTATION RESOURCES FOR **MAKE**

For more information about **make**, see the resources listed below.

Installed documentation

- Use the **man** and **info** tools to view manual pages and information pages installed on your system:

```
$ man make
$ info make
```

Online documentation

- The [GNU Make Manual](#) hosted by the Free Software Foundation
- Red Hat Developer Toolset User Guide – [Chapter 3. GNU make](#)

CHAPTER 6. CHANGES IN TOOLCHAIN SINCE RHEL 7

The following sections list changes in toolchain since the release of the described components in Red Hat Enterprise Linux 7. See also [Release notes for Red Hat Enterprise Linux 8.0](#) .

6.1. CHANGES IN GCC IN RHEL 8

In Red Hat Enterprise Linux 8, the GCC toolchain is based on the GCC 8.2 release series. Notable changes since Red Hat Enterprise Linux 7 include:

- Numerous general optimizations have been added, such as alias analysis, vectorizer improvements, identical code folding, inter-procedural analysis, store merging optimization pass, and others.
- The Address Sanitizer has been improved.
- The Leak Sanitizer for detection of memory leaks has been added.
- The Undefined Behavior Sanitizer for detection of undefined behavior has been added.
- Debug information can now be produced in the DWARF5 format. This capability is experimental.
- The source code coverage analysis tool GCOV has been extended with various improvements.
- Support for the OpenMP 4.5 specification has been added. Additionally, the offloading features of the OpenMP 4.0 specification are now supported by the C, C++, and Fortran compilers.
- New warnings and improved diagnostics have been added for static detection of certain likely programming errors.
- Source locations are now tracked as ranges rather than points, which allows much richer diagnostics. The compiler now offers “fix-it” hints, suggesting possible code modifications. A spell checker has been added to offer alternative names and ease detecting typos.

Security

GCC has been extended to provide tools to ensure additional hardening of the generated code.

For more details, see [Section 6.2, “Security enhancements in GCC in RHEL 8”](#) .

Architecture and processor support

Improvements to architecture and processor support include:

- Multiple new architecture-specific options for the Intel AVX-512 architecture, a number of its microarchitectures, and Intel Software Guard Extensions (SGX) have been added.
- Code generation can now target the 64-bit ARM architecture LSE extensions, ARMv8.2-A 16-bit Floating-Point Extensions (FPE), and ARMv8.2-A, ARMv8.3-A, and ARMv8.4-A architecture versions.
- Handling of the **-march=native** option on the ARM and 64-bit ARM architectures has been fixed.
- Support for the z13 and z14 processors of the IBM Z architecture has been added.

Languages and standards

Notable changes related to languages and standards include:

- The default standard used when compiling code in the C language has changed to C17 with GNU extensions.
- The default standard used when compiling code in the C++ language has changed to C++14 with GNU extensions.
- The C++ runtime library now supports the C++11 and C++14 standards.
- The C++ compiler now implements the C++14 standard with many new features such as variable templates, aggregates with non-static data member initializers, the extended **constexpr** specifier, sized deallocation functions, generic lambdas, variable-length arrays, digit separators, and others.
- Support for the C language standard C11 has been improved: ISO C11 atomics, generic selections, and thread-local storage are now available.
- The new **__auto_type** GNU C extension provides a subset of the functionality of C++11 **auto** keyword in the C language.
- The **_FloatN** and **_FloatNx** type names specified by the ISO/IEC TS 18661-3:2015 standard are now recognized by the C front end.
- The default standard used when compiling code in the C language has changed to C17 with GNU extensions. This has the same effect as using the **--std=gnu17** option. Previously, the default was C89 with GNU extensions.
- GCC can now experimentally compile code using the C++17 language standard and certain features from the C++20 standard.
- Passing an empty class as an argument now takes up no space on the Intel 64 and AMD64 architectures, as required by the platform ABI. Passing or returning a class with only deleted copy and move constructors now uses the same calling convention as a class with a non-trivial copy or move constructor.
- The value returned by the C++11 **alignof** operator has been corrected to match the C **_Alignof** operator and return minimum alignment. To find the preferred alignment, use the GNU extension **__alignof__**.
- The main version of the **libgfortran** library for Fortran language code has been changed to 5.
- Support for the Ada (GNAT), GCC Go, and Objective C/C++ languages has been removed. Use the Go Toolset for Go code development.

Additional resources

- See also the [Red Hat Enterprise Linux 8 Release Notes](#).
- [Using Go Toolset](#)

6.2. SECURITY ENHANCEMENTS IN GCC IN RHEL 8

This section describes in detail the changes in GCC related to security and added since the release of Red Hat Enterprise Linux 7.0.

New warnings

These warning options have been added:

Option	Displays warnings for
-Wstringop-truncation	Calls to bounded string manipulation functions such as strncat , strncpy , and stpncpy that might either truncate the copied string or leave the destination unchanged.
-Wclass-memaccess	<p>Objects of non-trivial class types manipulated in potentially unsafe ways by raw memory functions such as memcpy or realloc.</p> <p>The warning helps detect calls that bypass user-defined constructors or copy-assignment operators, corrupt virtual table pointers, data members of const-qualified types or references, or member pointers. The warning also detects calls that would bypass access controls to data members.</p>
-Wmisleading-indentation	Places where the indentation of the code gives a misleading idea of the block structure of the code to a human reader.
-Walloc-size-larger-than=size	Calls to memory allocation functions where the amount of memory to allocate exceeds <i>size</i> . Works also with functions where the allocation is specified by multiplying two parameters and with any functions decorated with attribute alloc_size .
-Walloc-zero	Calls to memory allocation functions that attempt to allocate zero amount of memory. Works also with functions where the allocation is specified by multiplying two parameters and with any functions decorated with attribute alloc_size .
-Walloca	All calls to the alloca function.
-Walloca-larger-than=size	Calls to the alloca function where the requested memory is more than <i>size</i> .
-Wvla-larger-than=size	Definitions of Variable Length Arrays (VLA) that can either exceed the specified size or whose bound is not known to be sufficiently constrained.
-Wformat-overflow=level	Both certain and likely buffer overflow in calls to the sprintf family of formatted output functions. For more details and explanation of the <i>level</i> value, see the <i>gcc(1)</i> manual page.
-Wformat-truncation=level	Both certain and likely output truncation in calls to the snprintf family of formatted output functions. For more details and explanation of the <i>level</i> value, see the <i>gcc(1)</i> manual page.
-Wstringop-overflow=type	Buffer overflow in calls to string handling functions such as memcpy and strcpy . For more details and explanation of the <i>level</i> value, see the <i>gcc(1)</i> manual page.

Warning improvements

These GCC warnings have been improved:

- The **-Warray-bounds** option has been improved to detect more instances of out-of-bounds array indices and pointer offsets. For example, negative or excessive indices into flexible array members and string literals are detected.
- The **-Wrestrict** option introduced in GCC 7 has been enhanced to detect many more instances of overlapping accesses to objects via restrict-qualified arguments to standard memory and string manipulation functions such as **memcpy** and **strcpy**.
- The **-Wnonnull** option has been enhanced to detect a broader set of cases of passing null pointers to functions that expect a non-null argument (decorated with attribute **nonnull**).

New UndefinedBehaviorSanitizer

A new run-time sanitizer for detecting undefined behavior called UndefinedBehaviorSanitizer has been added. The following options are noteworthy:

Option	Check
-fsanitize=float-divide-by-zero	Detect floating-point division by zero.
-fsanitize=float-cast-overflow	Check that the result of floating-point type to integer conversions do not overflow.
-fsanitize=bounds	Enable instrumentation of array bounds and detect out-of-bounds accesses.
-fsanitize=alignment	Enable alignment checking and detect various misaligned objects.
-fsanitize=object-size	Enable object size checking and detect various out-of-bounds accesses.
-fsanitize=vptr	Enable checking of C++ member function calls, member accesses, and some conversions between pointers to base and derived classes. Additionally, detect when referenced objects do not have correct dynamic type.
-fsanitize=bounds-strict	Enable strict checking of array bounds. This enables -fsanitize=bounds and instrumentation of flexible array member-like arrays.
-fsanitize=signed-integer-overflow	Diagnose arithmetic overflows even on arithmetic operations with generic vectors.
-fsanitize=builtin	Diagnose at run time invalid arguments to __builtin_clz or __builtin_ctz prefixed builtins. Includes checks from -fsanitize=undefined .

Option	Check
-fsanitize=pointer-overflow	Perform cheap run-time tests for pointer wrapping. Includes checks from -fsanitize=undefined .

New options for AddressSanitizer

These options have been added to AddressSanitizer:

Option	Check
-fsanitize=pointer-compare	Warn about comparison of pointers that point to a different memory object.
-fsanitize=pointer-subtract	Warn about subtraction of pointers that point to a different memory object.
-fsanitize-address-use-after-scope	Sanitize variables whose address is taken and used after a scope where the variable is defined.

Other sanitizers and instrumentation

- The option **-fstack-clash-protection** has been added to insert probes when stack space is allocated statically or dynamically to reliably detect stack overflows and thus mitigate the attack vector that relies on jumping over a stack guard page provided by the operating system.
- A new option **-fcf-protection=[full|branch|return|none]** has been added to perform code instrumentation and increase program security by checking that target addresses of control-flow transfer instructions (such as indirect function call, function return, indirect jump) are valid.

Additional resources

- For more details and explanation of the values supplied to some of the options above, see the *gcc(1)* manual page:

```
$ man gcc
```

6.3. COMPATIBILITY-BREAKING CHANGES IN GCC IN RHEL 8

C++ ABI change in `std::string` and `std::list`

The Application Binary Interface (ABI) of the `std::string` and `std::list` classes from the `libstdc++` library changed between RHEL 7 (GCC 4.8) and RHEL 8 (GCC 8) to conform to the C++11 standard. The `libstdc++` library supports both the old and new ABI, but some other C++ system libraries do not. As a consequence, applications that dynamically link against these libraries will need to be rebuilt. This affects all C++ standard modes, including C++98. It also affects applications built with Red Hat Developer Toolset compilers for RHEL 7, which kept the old ABI to maintain compatibility with the system libraries.

GCC no longer builds Ada, Go, and Objective C/C++ code

Capability for building code in the Ada (GNAT), GCC Go, and Objective C/C++ languages has been removed from the GCC compiler.

To build Go code, use the Go Toolset instead.

PART II. DEBUGGING APPLICATIONS

Debugging applications is a very wide topic. This part provides a developer with the most common techniques for debugging in multiple situations.

CHAPTER 7. ENABLING DEBUGGING WITH DEBUGGING INFORMATION

To debug applications and libraries, debugging information is required. The following sections describe how to obtain this information.

7.1. DEBUGGING INFORMATION

While debugging any executable code, two types of information allow the tools, and by extension the programmer, to comprehend the binary code:

- the source code text
- a description of how the source code text relates to the binary code

Such information is called debugging information.

Red Hat Enterprise Linux uses the ELF format for executable binaries, shared libraries, or **debuginfo** files. Within these ELF files, the DWARF format is used to hold the debug information.

To display DWARF information stored within an ELF file, run the **readelf -w file** command.

CAUTION

STABS is an older, less capable format, occasionally used with UNIX. Its use is discouraged by Red Hat. GCC and GDB provide STABS production and consumption on a best effort basis only. Some other tools such as Valgrind and **elfutils** do not work with STABS.

Additional resources

- [The DWARF Debugging Standard](#)

7.2. ENABLING DEBUGGING OF C AND C++ APPLICATIONS WITH GCC

Because debugging information is large, it is not included in executable files by default. To enable debugging of your C and C++ applications with it, you must explicitly instruct the compiler to create it.

To enable creation of debugging information with **GCC** when compiling and linking code, use the **-g** option:

```
$ gcc ... -g ...
```

- Optimizations performed by the compiler and linker can result in executable code which is hard to relate to the original source code: variables may be optimized out, loops unrolled, operations merged into the surrounding ones etc. This affects debugging negatively. For improved debugging experience, consider setting the optimization with the **-Og** option. However, changing the optimization level changes the executable code and may change the actual behaviour including removing some bugs.
- To also include macro definitions in the debug information, use the **-g3** option instead of **-g**.
- The **-fcompare-debug** GCC option tests code compiled by GCC with debug information and without debug information. The test passes if the resulting two binary files are identical. This test ensures that executable code is not affected by any debugging options, which further

ensures that there are no hidden bugs in the debug code. Note that using the **-fcompare-debug** option significantly increases compilation time. See the GCC manual page for details about this option.

Additional resources

- [Chapter 7, Enabling Debugging with Debugging Information](#)
- Using the GNU Compiler Collection (GCC) – [Options for Debugging Your Program](#)
- Debugging with GDB – [Debugging Information in Separate Files](#)
- The GCC manual page:

```
$ man gcc
```

7.3. DEBUGINFO AND DEBUGSOURCE PACKAGES

The **debuginfo** and **debugsource** packages contain debugging information and debug source code for programs and libraries. For applications and libraries installed in packages from the Red Hat Enterprise Linux repositories, you can obtain separate **debuginfo** and **debugsource** packages from an additional channel.

Debugging information package types

There are two types of packages available for debugging:

Debuginfo packages

The **debuginfo** packages provide debugging information needed to provide human-readable names for binary code features. These packages contain **.debug** files, which contain DWARF debugging information. These files are installed to the **/usr/lib/debug** directory.

Debugsource packages

The **debugsource** packages contain the source files used for compiling the binary code. With both respective **debuginfo** and **debugsource** package installed, debuggers such as GDB or LLDB can relate the execution of binary code to the source code. The source code files are installed to the **/usr/src/debug** directory.

Differences from RHEL 7

In Red Hat Enterprise Linux 7, the **debuginfo** packages contained both kinds of information. Red Hat Enterprise Linux 8 splits the source code data needed for debugging from the **debuginfo** packages into separate **debugsource** packages.

Package names

A **debuginfo** or **debugsource** package provides debugging information valid only for a binary package with the same name, version, release, and architecture:

- Binary package: ***packagename-version-release.architecture.rpm***
- Debuginfo package: ***packagename-debuginfo-version-release.architecture.rpm***
- Debugsource package: ***packagename-debugsource-version-release.architecture.rpm***

Additional resources

- [Section 7.1, “Debugging information”](#)
- [Section 1.2, “Enabling debug and source repositories”](#)

7.4. GETTING DEBUGINFO PACKAGES FOR AN APPLICATION OR LIBRARY USING GDB

Debugging information is required to debug code. For code that is installed from a package, the GNU Debugger (GDB) automatically recognizes missing debug information, resolves the package name and provides concrete advice on how to get the package.

Prerequisites

- The application or library you want to debug must be installed on the system.
- GDB and the **debuginfo-install** tool must be installed on the system.
- Channels providing **debuginfo** and **debugsource** packages must be configured and enabled on the system.

Procedure

1. Start GDB attached to the application or library you want to debug. GDB automatically recognizes missing debugging information and suggests a command to run.

```
$ gdb -q /bin/ls
Reading symbols from /bin/ls...Reading symbols from .gnu_debugdata for /usr/bin/ls...(no
debugging symbols found)...done.
(no debugging symbols found)...done.
Missing separate debuginfos, use: dnf debuginfo-install coreutils-8.30-6.el8.x86_64
(gdb)
```

2. Exit GDB: type **q** and confirm with **Enter**.

```
(gdb) q
```

3. Run the command suggested by GDB to install the required **debuginfo** packages:

```
# dnf debuginfo-install coreutils-8.30-6.el8.x86_64
```

The **dnf** package management tool provides a summary of the changes, asks for confirmation and once you confirm, downloads and installs all the necessary files.

4. In case GDB is not able to suggest the **debuginfo** package, follow the procedure described in [Section 7.5, “Getting debuginfo packages for an application or library manually”](#).

Additional resources

- [Red Hat Developer Toolset User Guide, section Installing Debugging Information](#)
- [How can I download or install debuginfo packages for RHEL systems?](#) – Red Hat Knowledgebase solution

7.5. GETTING DEBUGINFO PACKAGES FOR AN APPLICATION OR LIBRARY MANUALLY

You can determine manually which **debuginfo** packages you need to install by locating the executable file and then finding the package that installs it.



NOTE

Red Hat recommends that you [use GDB to determine the packages for installation](#). Use this manual procedure only if GDB is not able to suggest the package to install.

Prerequisites

- The application or library must be installed on the system.
- The application or library was installed from a package.
- The **debuginfo-install** tool must be available on the system.
- Channels providing the **debuginfo** packages must be configured and enabled on the system.

Procedure

1. Find the executable file of the application or library.
 - a. Use the **which** command to find the application file.

```
$ which /less
/usr/bin/less
```

- b. Use the **locate** command to find the library file.

```
$ locate /libz | grep so
/usr/lib64/libz.so.1
/usr/lib64/libz.so.1.2.11
```

If the original reasons for debugging include error messages, pick the result where the library has the same additional numbers in its file name as those mentioned in the error messages. If in doubt, try following the rest of the procedure with the result where the library file name includes no additional numbers.



NOTE

The **locate** command is provided by the **mlocate** package. To install it and enable its use:

```
# yum install mlocate
# updatedb
```

2. Search for a name and version of the package that provided the file:

```
$ rpm -qf /usr/lib64/libz.so.1.2.7
zlib-1.2.11-10.el8.x86_64
```

The output provides details for the installed package in the *name:epoch-version.release.architecture* format.



IMPORTANT

If this step does not produce any results, it is not possible to determine which package provided the binary file. There are several possible cases:

- The file is installed from a package which is not known to package management tools in their *current* configuration.
- The file is installed from a locally downloaded and manually installed package. Determining a suitable **debuginfo** package automatically is impossible in that case.
- Your package management tools are misconfigured.
- The file is not installed from any package. In such a case, no respective **debuginfo** package exists.

Because further steps depend on this one, you must resolve this situation or abort this procedure. Describing the exact troubleshooting steps is beyond the scope of this procedure.

3. Install the **debuginfo** packages using the **debuginfo-install** utility. In the command, use the package name and other details you determined during the previous step:

```
# debuginfo-install zlib-1.2.11-10.el8.x86_64
```

Additional resources

- [Red Hat Developer Toolset User Guide, section Installing Debugging Information](#)
- [How can I download or install debuginfo packages for RHEL systems?](#) – Knowledgebase article

CHAPTER 8. INSPECTING APPLICATION INTERNAL STATE WITH GDB

To find why an application does not work properly, control its execution and examine its internal state with a debugger. This section describes how to use the GNU Debugger (GDB) for this task.

8.1. GNU DEBUGGER (GDB)

Red Hat Enterprise Linux contains the GNU debugger (GDB) which lets you investigate what is happening inside a program through a command-line user interface.

For a graphical front end to GDB, install the Eclipse integrated development environment. See [Using Eclipse](#).

GDB capabilities

A single GDB session can debug the following types of programs:

- Multithreaded and forking programs
- Multiple programs at once
- Programs on remote machines or in containers with the **gdbserver** utility connected over a TCP/IP network connection

Debugging requirements

To debug any executable code, GDB requires debugging information for that particular code:

- For programs developed by you, you can create the debugging information while building the code.
- For system programs installed from packages, you must install their debuginfo packages.

8.2. ATTACHING GDB TO A PROCESS

In order to examine a process, GDB must be *attached* to the process.

Prerequisites

- [GDB must be installed on the system](#)

Starting a program with GDB

When the program is not running as a process, start it with GDB:

```
$ gdb program
```

Replace *program* with a file name or path to the program.

GDB sets up to start execution of the program. You can set up breakpoints and the **gdb** environment before beginning the execution of the process with the **run** command.

Attaching GDB to an already running process

To attach GDB to a program already running as a process:

1. Find the process ID (*pid*) with the **ps** command:

```
$ ps -C program -o pid h
pid
```

Replace *program* with a file name or path to the program.

2. Attach GDB to this process:

```
$ gdb -p pid
```

Replace *pid* with an actual process ID number from the **ps** output.

Attaching an already running GDB to an already running process

To attach an already running GDB to an already running program:

1. Use the **shell** GDB command to run the **ps** command and find the program's process ID (*pid*):

```
(gdb) shell ps -C program -o pid h
pid
```

Replace *program* with a file name or path to the program.

2. Use the **attach** command to attach GDB to the program:

```
(gdb) attach pid
```

Replace *pid* by an actual process ID number from the **ps** output.



NOTE

In some cases, GDB might not be able to find the respective executable file. Use the **file** command to specify the path:

```
(gdb) file path/to/program
```

Additional resources

- Debugging with GDB – [2.1 Invoking GDB](#)
- Debugging with GDB – [4.7 Debugging an Already-running Process](#)

8.3. STEPPING THROUGH PROGRAM CODE WITH GDB

Once the **GDB** debugger is attached to a program, you can use a number of commands to control the execution of the program.

Prerequisites

- You must have the required debugging information available:
 - The program is compiled and built with debugging information, or

- The relevant debuginfo packages are installed
- [GDB must be attached to the program to be debugged](#)

GDB commands to step through the code

r (run)

Start the execution of the program. If **run** is executed with any arguments, those arguments are passed on to the executable as if the program has been started normally. Users normally issue this command after setting breakpoints.

start

Start the execution of the program but stop at the beginning of the program's main function. If **start** is executed with any arguments, those arguments are passed on to the executable as if the program has been started normally.

c (continue)

Continue the execution of the program from the current state. The execution of the program will continue until one of the following becomes true:

- A breakpoint is reached.
- A specified condition is satisfied.
- A signal is received by the program.
- An error occurs.
- The program terminates.

n (next)

Continue the execution of the program from the current state, until the next line of code in the current source file is reached. The execution of the program will continue until one of the following becomes true:

- A breakpoint is reached.
- A specified condition is satisfied.
- A signal is received by the program.
- An error occurs.
- The program terminates.

s (step)

The **step** command also halts execution at each sequential line of code in the current source file. However, if the execution is currently stopped at a source line containing a **function call**, GDB stops the execution after entering the function call (rather than executing it).

until location

Continue the execution until the code location specified by the *location* option is reached.

fini (finish)

Resume the execution of the program and halt when execution returns from a function. The execution of the program will continue until one of the following becomes true:

- A breakpoint is reached.
- A specified condition is satisfied.
- A signal is received by the program.
- An error occurs.
- The program terminates.

q (quit)

Terminate the execution and exit GDB.

Additional resources

- [Section 8.5, “Using GDB breakpoints to stop execution at defined code locations”](#)
- Debugging with GDB – [4.2 Starting your Program](#)
- Debugging with GDB – [5.2 Continuing and Stepping](#)

8.4. SHOWING PROGRAM INTERNAL VALUES WITH GDB

Displaying the values of a program’s internal variables is important for understanding of what the program is doing. GDB offers multiple commands that you can use to inspect the internal variables. This section describes the most useful of these commands:

p (print)

Display the value of the argument given. Usually, the argument is the name of a variable of any complexity, from a simple single value to a structure. An argument can also be an expression valid in the current language, including the use of program variables and library functions, or functions defined in the program being tested.

It is possible to extend GDB with *pretty-printer* Python or Guile scripts for customized display of data structures (such as classes, structs) using the **print** command.

bt (backtrace)

Display the chain of function calls used to reach the current execution point, or the chain of functions used up until execution was terminated. This is useful for investigating serious bugs (such as segmentation faults) with elusive causes.

Adding the **full** option to the **backtrace** command displays local variables, too.

It is possible to extend GDB with *frame filter* Python scripts for customized display of data displayed using the **bt** and **info frame** commands. The term *frame* refers to the data associated with a single function call.

info

The **info** command is a generic command to provide information about various items. It takes an option specifying the item to describe.

- The **info args** command displays options of the function call that is the currently selected frame.
- The **info locals** command displays local variables in the currently selected frame.

For a list of the possible items, run the command **help info** in a GDB session:

```
(gdb) help info
```

l (list)

Show the line in the source code where the program stopped. This command is available only when the program execution is stopped. While not strictly a command to show internal state, **list** helps the user understand what changes to the internal state will happen in the next step of the program's execution.

Additional resources

- [The GDB Python API](#) – Red Hat Developers Blog entry
- Debugging with GDB – [10.9 Pretty Printing](#)

8.5. USING GDB BREAKPOINTS TO STOP EXECUTION AT DEFINED CODE LOCATIONS

Often, only small portions of code are investigated. Breakpoints are markers that tell GDB to stop the execution of a program at a certain place in the code. Breakpoints are most commonly associated with source code lines. In that case, placing a breakpoint requires specifying the source file and line number.

- To **place a breakpoint**:

- Specify the name of the source code *file* and the *line* in that file:

```
(gdb) br file:line
```

- When *file* is not present, name of the source file at the current point of execution is used:

```
(gdb) br line
```

- Alternatively, use a function name to put the breakpoint on its start:

```
(gdb) br function_name
```

- A program might encounter an error after a certain number of iterations of a task. To specify an additional **condition** to halt execution:

```
(gdb) br file:line if condition
```

Replace *condition* with a condition in the C or C++ language. The meaning of *file* and *line* is the same as above.

- To **inspect** the status of all breakpoints and watchpoints:

```
(gdb) info br
```

- To **remove** a breakpoint by using its *number* as displayed in the output of **info br**:

```
(gdb) delete number
```

- To **remove** a breakpoint at a given location:

```
(gdb) clear file:line
```

Additional resources

- Debugging with GDB – [5.1 Breakpoints, Watchpoints, and Catchpoints](#)

8.6. USING GDB WATCHPOINTS TO STOP EXECUTION ON DATA ACCESS AND CHANGES

In many cases, it is advantageous to let the program execute until certain data changes or is accessed. This section lists the most common use cases.

Prerequisites

- Understanding **GDB**

Using watchpoints in GDB

Watchpoints are markers which tell **GDB** to stop the execution of a program. Watchpoints are associated with data: placing a watchpoint requires specifying an expression that describes a variable, multiple variables, or a memory address.

- To **place** a watchpoint for data **change** (write):

```
(gdb) watch expression
```

Replace *expression* with an expression that describes what you want to watch. For variables, *expression* is equal to the name of the variable.

- To **place** a watchpoint for data **access** (read):

```
(gdb) rwatch expression
```

- To **place** a watchpoint for **any** data access (both read and write):

```
(gdb) awatch expression
```

- To **inspect** the status of all watchpoints and breakpoints:

```
(gdb) info br
```

- To **remove** a watchpoint:

```
(gdb) delete num
```

Replace the *num* option with the number reported by the **info br** command.

Additional resources

- Debugging with GDB – [5.1.2 Setting Watchpoints](#)

8.7. DEBUGGING FORKING OR THREADED PROGRAMS WITH GDB

Some programs use forking or threads to achieve parallel code execution. Debugging multiple simultaneous execution paths requires special considerations.

Prerequisites

- You must understand the concepts of process forking and threads.

Debugging forked programs with GDB

Forking is a situation when a program (**parent**) creates an independent copy of itself (**child**). Use the following settings and commands to affect what GDB does when a fork occurs:

- The **follow-fork-mode** setting controls whether GDB follows the parent or the child after the fork.

set follow-fork-mode parent

After a fork, debug the parent process. This is the default.

set follow-fork-mode child

After a fork, debug the child process.

show follow-fork-mode

Display the current setting of **follow-fork-mode**.

- The **set detach-on-fork** setting controls whether the GDB keeps control of the other (not followed) process or leaves it to run.

set detach-on-fork on

The process which is not followed (depending on the value of **follow-fork-mode**) is detached and runs independently. This is the default.

set detach-on-fork off

GDB keeps control of both processes. The process which is followed (depending on the value of **follow-fork-mode**) is debugged as usual, while the other is suspended.

show detach-on-fork

Display the current setting of **detach-on-fork**.

Debugging Threaded Programs with GDB

GDB has the ability to debug individual threads, and to manipulate and examine them independently. To make GDB stop only the thread that is examined, use the commands **set non-stop on** and **set target-async on**. You can add these commands to the **.gdbinit** file. After that functionality is turned on, GDB is ready to conduct thread debugging.

GDB uses a concept of *current thread*. By default, commands apply to the current thread only.

info threads

Display a list of threads with their **id** and **gid** numbers, indicating the current thread.

thread id

Set the thread with the specified **id** as the current thread.

thread apply ids command

Apply the command **command** to all threads listed by **ids**. The **ids** option is a space-separated list of thread ids. A special value **all** applies the command to all threads.

break *location* thread *id* if *condition*

Set a breakpoint at a certain *location* with a certain *condition* only for the thread number *id*.

watch *expression* thread *id*

Set a watchpoint defined by *expression* only for the thread number *id*.

command&

Execute command **command** and return immediately to the gdb prompt (**(gdb)**), continuing any code execution in the background.

interrupt

Halt execution in the background.

Additional resources

- Debugging with GDB – [4.10 Debugging Programs with Multiple Threads](#)
- Debugging with GDB – [4.11 Debugging Forks](#)

CHAPTER 9. RECORDING APPLICATION INTERACTIONS

The executable code of applications interacts with the code of the operating system and shared libraries. Recording an activity log of these interactions can provide enough insight into the application's behavior without debugging the actual application code. Alternatively, analyzing an application's interactions can help pinpoint the conditions in which a bug manifests.

9.1. TOOLS USEFUL FOR RECORDING APPLICATION INTERACTIONS

Red Hat Enterprise Linux offers multiple tools for analyzing an application's interactions.

strace

The **strace** tool primarily enables logging of system calls (kernel functions) used by an application.

- The **strace** tool can provide a detailed output about calls, because **strace** interprets parameters and results with knowledge of the underlying kernel code. Numbers are turned into the respective constant names, bitwise combined flags expanded to flag list, pointers to character arrays dereferenced to provide the actual string, and more. Support for more recent kernel features may be lacking.
- You can filter the traced calls to reduce the amount of captured data.
- The use of **strace** does not require any particular setup except for setting up the log filter.
- Tracing the application code with **strace** results in significant slowdown of the application's execution. As a result, **strace** is not suitable for many production deployments. As an alternative, consider using **ltrace** or SystemTap.
- The version of **strace** available in Red Hat Developer Toolset can also perform system call tampering. This capability is useful for debugging.

ltrace

The **ltrace** tool enables logging of an application's user space calls into shared objects (dynamic libraries).

- The **ltrace** tool enables tracing calls to any library.
- You can filter the traced calls to reduce the amount of captured data.
- The use of **ltrace** does not require any particular setup except for setting up the log filter.
- The **ltrace** tool is lightweight and fast, offering an alternative to **strace**: it is possible to trace the respective interfaces in libraries such as **glibc** with **ltrace** instead of tracing kernel functions with **strace**.
- Because **ltrace** does not handle a known set of calls like **strace**, it does not attempt to explain the values passed to library functions. The **ltrace** output contains only raw numbers and pointers. The interpretation of **ltrace** output requires consulting the actual interface declarations of the libraries present in the output.



NOTE

In Red Hat Enterprise Linux 8.0, a known issue prevents **ltrace** from tracing system executable files. This limitation does not apply to executable files built by users.

SystemTap

SystemTap is an instrumentation platform for probing running processes and kernel activity on the Linux system. SystemTap uses its own scripting language for programming custom event handlers.

- Compared to using **strace** and **ltrace**, scripting the logging means more work in the initial setup phase. However, the scripting capabilities extend SystemTap's usefulness beyond just producing logs.
- SystemTap works by creating and inserting a kernel module. The use of SystemTap is efficient and does not create a significant slowdown of the system or application execution on its own.
- SystemTap comes with a set of usage examples.

GDB

The GNU Debugger (GDB) is primarily meant for debugging, not logging. However, some of its features make it useful even in the scenario where an application's interaction is the primary activity of interest.

- With GDB, it is possible to conveniently combine the capture of an interaction event with immediate debugging of the subsequent execution path.
- GDB is best suited for analyzing response to infrequent or singular events, after the initial identification of problematic situation by other tools. Using GDB in any scenario with frequent events becomes inefficient or even impossible.

Additional resources

- [Red Hat Enterprise Linux SystemTap Beginners Guide](#)
- [Red Hat Developer Toolset User Guide](#)

9.2. MONITORING AN APPLICATION'S SYSTEM CALLS WITH STRACE

The **strace** tool enables monitoring the system (kernel) calls performed by an application.

Prerequisites

- You must have **strace** installed on the system.

Procedure

1. Identify the system calls to monitor.
2. Start **strace** and attach it to the program.
 - If the program you want to monitor is not running, start **strace** and specify the *program*:

```
$ strace -fvttTyy -s 256 -e trace=call program
```

 - If the program is already running, find its process id (*pid*) and attach **strace** to it:

```
$ ps -C program
(...)
$ strace -fvttTyy -s 256 -e trace=call -ppid
```

- Replace *call* with the system calls to be displayed. You can use the **-e trace=*call*** option multiple times. If left out, **strace** will display all system call types. See the *strace(1)* manual page for more information.
 - If you do not want to trace any forked processes or threads, leave out the **-f** option.
3. The **strace** tool displays the system calls made by the application and their details. In most cases, an application and its libraries make a large number of calls and **strace** output appears immediately, if no filter for system calls is set.
 4. The **strace** tool exits when the program exits. To terminate the monitoring before the traced program exits, press **Ctrl+C**.
 - If **strace** started the program, the program terminates together with **strace**.
 - If you attached **strace** to an already running program, the program terminates together with **strace**.
 5. Analyze the list of system calls done by the application.
 - Problems with resource access or availability are present in the log as calls returning errors.
 - Values passed to the system calls and patterns of call sequences provide insight into the causes of the application's behaviour.
 - If the application crashes, the important information is probably at the end of log.
 - The output contains a lot of unnecessary information. However, you can construct a more precise filter for the system calls of interest and repeat the procedure.



NOTE

It is advantageous to both see the output and save it to a file. Use the **tee** command to achieve this:

```
$ strace ... |& tee your_log_file.log
```

Additional resources

- The *strace(1)* manual page:

```
$ man strace
```

- [How do I use strace to trace system calls made by a command?](#) – Knowledgebase article
- Red Hat Developer Toolset User Guide – [Chapter strace](#)

9.3. MONITORING APPLICATION'S LIBRARY FUNCTION CALLS WITH LTRACE

The **ltrace** tool enables monitoring an application's calls to functions available in libraries (shared objects).



NOTE

In Red Hat Enterprise Linux 8.0, a known issue prevents **ltrace** from tracing system executable files. This limitation does not apply to executable files built by users.

Prerequisites

- You must have **ltrace** installed on the system.

Procedure

1. Identify the libraries and functions of interest, if possible.
2. Start **ltrace** and attach it to the program.

- If the program you want to monitor is not running, start **ltrace** and specify *program*:

```
$ ltrace -f -l library -e function program
```

- If the program is already running, find its process id (*pid*) and attach **ltrace** to it:

```
$ ps -C program
(...)
$ ltrace -f -l library -e function program -ppid
```

- Use the **-e**, **-f** and **-l** options to filter the output:
 - Supply the function names to be displayed as *function*. The **-e function** option can be used multiple times. If left out, **ltrace** displays calls to all functions.
 - Instead of specifying functions, you can specify whole libraries with the **-l library** option. This option behaves similarly to the **-e function** option.
 - If you do not want to trace any forked processes or threads, leave out the **-f** option.

See the *ltrace(1)* manual page for more information.

3. **ltrace** displays the library calls made by the application.
In most cases, an application makes a large number of calls and **ltrace** output displays immediately, if no filter is set.
4. **ltrace** exits when the program exits.
To terminate the monitoring before the traced program exits, press **ctrl+C**.
 - If **ltrace** started the program, the program terminates together with **ltrace**.
 - If you attached **ltrace** to an already running program, the program terminates together with **ltrace**.
5. Analyze the list of library calls done by the application.
 - If the application crashes, the important information is probably at the end of log.

- The output contains a lot of unnecessary information. However, you can construct a more precise filter and repeat the procedure.



NOTE

It is advantageous to both see the output and save it to a file. Use the **tee** command to achieve this:

```
$ ltrace ... |& tee your_log_file.log
```

Additional resources

- The *ltrace(1)* manual page:

```
$ man ltrace
```

- Red Hat Developer Toolset User Guide – [Chapter ltrace](#)

9.4. MONITORING APPLICATION'S SYSTEM CALLS WITH SYSTEMTAP

The SystemTap tool enables registering custom event handlers for kernel events. In comparison with the **strace** tool, it is harder to use but more efficient and enables more complicated processing logic. A SystemTap script called **strace.stp** is installed together with SystemTap and provides an approximation of **strace** functionality using SystemTap.

Prerequisites

- [SystemTap and the respective kernel packages must be installed on the system.](#)

Procedure

1. Find the process ID (*pid*) of the process you want to monitor:

```
$ ps -aux
```

2. Run SystemTap with the **strace.stp** script:

```
# stap /usr/share/systemtap/examples/process/strace.stp -x pid
```

The value of *pid* is the process id.

The script is compiled to a kernel module, which is then loaded. This introduces a slight delay between entering the command and getting the output.

3. When the process performs a system call, the call name and its parameters are printed to the terminal.
4. The script exits when the process terminates, or when you press **Ctrl+C**.

9.5. USING GDB TO INTERCEPT APPLICATION SYSTEM CALLS

GNU Debugger (GDB) lets you stop an execution in various situations that arise during program execution. To stop the execution when the program performs a system call, use a GDB *catchpoint*.

Prerequisites

- [You must understand the usage of GDB breakpoints.](#)
- [GDB must be attached to the program.](#)

Procedure

1. Set the catchpoint:

```
(gdb) catch syscall syscall-name
```

The command **catch syscall** sets a special type of breakpoint that halts execution when the program performs a system call.

The ***syscall-name*** option specifies the name of the call. You can specify multiple catchpoints for various system calls. Leaving out the ***syscall-name*** option causes GDB to stop on any system call.

2. Start execution of the program.

- If the program has not started execution, start it:

```
(gdb) r
```

- If the program execution is halted, resume it:

```
(gdb) c
```

3. GDB halts execution after the program performs any specified system call.

Additional resources

- [Section 8.4, “Showing program internal values with GDB”](#)
- [Section 8.3, “Stepping through program code with GDB”](#)
- Debugging with GDB – [Setting Watchpoints](#)

9.6. USING GDB TO INTERCEPT HANDLING OF SIGNALS BY APPLICATIONS

GNU Debugger (GDB) lets you stop the execution in various situations that arise during program execution. To stop the execution when the program receives a signal from the operating system, use a GDB *catchpoint*.

Prerequisites

- [You must understand the usage of GDB breakpoints.](#)
- [GDB must be attached to the program.](#)

Procedure

1. Set the catchpoint:

```
(gdb) catch signal signal-type
```

The command **catch signal** sets a special type of a breakpoint that halts execution when a signal is received by the program. The ***signal-type*** option specifies the type of the signal. Use the special value **'all'** to catch all signals.

2. Let the program run.

- If the program has not started execution, start it:

```
(gdb) r
```

- If the program execution is halted, resume it:

```
(gdb) c
```

3. GDB halts execution after the program receives any specified signal.

Additional resources

- [Section 8.4, “Showing program internal values with GDB”](#)
- [Section 8.3, “Stepping through program code with GDB”](#)
- Debugging With GDB – [5.1.3 Setting Catchpoints](#)

CHAPTER 10. DEBUGGING A CRASHED APPLICATION

Sometimes, it is not possible to debug an application directly. In these situations, you can collect information about the application at the moment of its termination and analyze it afterwards.

10.1. CORE DUMPS: WHAT THEY ARE AND HOW TO USE THEM

A core dump is a copy of a part of the application's memory at the moment the application stopped working, stored in the ELF format. It contains all the application's internal variables and stack, which enables inspection of the application's final state. When augmented with the respective executable file and debugging information, it is possible to analyze a core dump file with a debugger in a way similar to analyzing a running program.

The Linux operating system kernel can record core dumps automatically, if this functionality is enabled. Alternatively, you can send a signal to any running application to generate a core dump regardless of its actual state.



WARNING

Some limits might affect the ability to generate a core dump. To see the current limits:

```
$ ulimit -a
```

10.2. RECORDING APPLICATION CRASHES WITH CORE DUMPS

To record application crashes, set up core dump saving and add information about the system.

Procedure

1. To enable core dumps, ensure that the **/etc/systemd/system.conf** file contains the following lines:

```
DumpCore=yes
DefaultLimitCORE=infinity
```

You can also add comments describing if these settings were previously present, and what the previous values were. This will enable you to reverse these changes later, if needed. Comments are lines starting with the **#** character.

Changing the file requires administrator level access.

2. Apply the new configuration:

```
# systemctl daemon-reexec
```

3. Remove the limits for core dump sizes:

```
# ulimit -c unlimited
```

To reverse this change, run the command with value **0** instead of **unlimited**.

4. Install the **sos** package which provides the **sosreport** utility for collecting system information:

```
# yum install sos
```

5. When an application crashes, a core dump is generated and handled by **systemd-coredump**.
6. Create an SOS report to provide additional information about the system:

```
# sosreport
```

This creates a **.tar** archive containing information about your system, such as copies of configuration files.

7. Locate and export the core dump:

```
$ coredumpctl list executable-name
$ coredumpctl dump executable-name > /path/to/file-for-export
```

If the application crashed multiple times, output of the first command lists more captured core dumps. In that case, construct for the second command a more precise query using the other information. See the *coredumpctl(1)* manual page for details.

8. Transfer the core dump and the SOS report to the computer where the debugging will take place. Transfer the executable file, too, if it is known.



IMPORTANT

When the executable file is not known, subsequent analysis of the core file identifies it.

9. Optional: Remove the core dump and SOS report after transferring them, to free up disk space.

Additional resources

- [Introduction to systemd](#) in the document *Configuring basic system settings*
- [How to enable core file dumps when an application crashes or segmentation faults](#) – a Knowledge Base article
- [What is a sosreport and how to create one in Red Hat Enterprise Linux 4.6 and later?](#) – a Knowledge Base article

10.3. INSPECTING APPLICATION CRASH STATES WITH CORE DUMPS

Prerequisites

- You must have a core dump file and sosreport from the system where the crash occurred.
- GDB and elfutils must be installed on your system.

Procedure

1. To identify the executable file where the crash occurred, run the **eu-unstrip** command with the core dump file:

```
$ eu-unstrip -n --core=./core.9814
0x400000+0x207000 2818b2009547f780a5639c904cded443e564973e@0x400284
/usr/bin/sleep /usr/lib/debug/bin/sleep.debug [exe]
0x7fff26fff000+0x1000 1e2a683b7d877576970e4275d41a6aaec280795e@0x7fff26fff340 . -
linux-vdso.so.1
0x35e7e00000+0x3b6000
374add1ead31ccb449779bc7ee7877de3377e5ad@0x35e7e00280 /usr/lib64/libc-2.14.90.so
/usr/lib/debug/lib64/libc-2.14.90.so.debug libc.so.6
0x35e7a00000+0x224000
3ed9e61c2b7e707ce244816335776afa2ad0307d@0x35e7a001d8 /usr/lib64/ld-2.14.90.so
/usr/lib/debug/lib64/ld-2.14.90.so.debug ld-linux-x86-64.so.2
```

The output contains details for each module on a line, separated by spaces. The information is listed in this order:

1. The memory address where the module was mapped
2. The build-id of the module and where in the memory it was found
3. The module's executable file name, displayed as **-** when unknown, or as **.** when the module has not been loaded from a file
4. The source of debugging information, displayed as a file name when available, as **.** when contained in the executable file itself, or as **-** when not present at all
5. The shared library name (*soname*) or **[exe]** for the main module

In this example, the important details are the file name **/usr/bin/sleep** and the build-id **2818b2009547f780a5639c904cded443e564973e** on the line containing the text **[exe]**. With this information, you can identify the executable file required for analyzing the core dump.

2. Get the executable file that crashed.

- If possible, copy it from the system where the crash occurred. Use the file name extracted from the core file.
- You can also use an identical executable file on your system. Each executable file built on Red Hat Enterprise Linux contains a note with a unique build-id value. Determine the build-id of the relevant locally available executable files:

```
$ eu-readelf -n executable_file
```

Use this information to match the executable file on the remote system with your local copy. The build-id of the local file and build-id listed in the core dump must match.

- Finally, if the application is installed from an RPM package, you can get the executable file from the package. Use the **sosreport** output to find the exact version of the package required.
3. Get the shared libraries used by the executable file. Use the same steps as for the executable file.

4. If the application is distributed as a package, load the executable file in GDB, to display hints for missing debuginfo packages. For more details, see [Section 7.4, "Getting debuginfo packages for an application or library using GDB"](#).
5. To examine the core file in detail, load the executable file and core dump file with GDB:

```
$ gdb -e executable_file -c core_file
```

Further messages about missing files and debugging information help you identify what is missing for the debugging session. Return to the previous step if needed.

If the application's debugging information is available as a file instead of as a package, load this file in GDB with the **symbol-file** command:

```
(gdb) symbol-file program.debug
```

Replace *program.debug* with the actual file name.



NOTE

It might not be necessary to install the debugging information for all executable files contained in the core dump. Most of these executable files are libraries used by the application code. These libraries might not directly contribute to the problem you are analyzing, and you do not need to include debugging information for them.

6. Use the GDB commands to inspect the state of the application at the moment it crashed. See [Chapter 8, *Inspecting Application Internal State with GDB*](#).



NOTE

When analyzing a core file, GDB is not attached to a running process. Commands for controlling execution have no effect.

Additional resources

- Debugging with GDB – [2.1.1 Choosing Files](#)
- Debugging with GDB – [18.1 Commands to Specify Files](#)
- Debugging with GDB – [18.3 Debugging Information in Separate Files](#)

10.4. CREATING AND ACCESSING A CORE DUMP WITH COREDUMPCTL

The **coredumpctl** tool of **systemd** can significantly streamline working with core dumps on the machine where the crash happened. This procedure outlines how to capture a core dump of unresponsive process.

Prerequisites

- The system must be configured to use **systemd-coredump** for core dump handling. To verify this is true:

```
$ sysctl kernel.core_pattern
```

The configuration is correct if the output starts with the following:

```
kernel.core_pattern = /usr/lib/systemd/systemd-coredump
```

Procedure

1. Find the PID of the hung process, based on a known part of the executable file name:

```
$ pgrep -a executable-name-fragment
```

This command will output a line in the form

```
PID command-line
```

Use the *command-line* value to verify that the *PID* belongs to the intended process.

For example:

```
$ pgrep -a bc
5459 bc
```

2. Send an abort signal to the process:

```
# kill -ABRT PID
```

3. Verify that the core has been captured by **coredumpctl**:

```
$ coredumpctl list PID
```

For example:

```
$ coredumpctl list 5459
TIME                PID  UID  GID SIG COREFILE EXE
Thu 2019-11-07 15:14:46 CET  5459 1000 1000 6 present /usr/bin/bc
```

4. Further examine or use the core file as needed.

You can specify the core dump by PID and other values. See the *coredumpctl(1)* manual page for further details.

- To show details of the core file:

```
$ coredumpctl info PID
```

- To load the core file in the GDB debugger:

```
$ coredumpctl debug PID
```

Depending on availability of debugging information, GDB will suggest commands to run, such as:

```
■
```

```
Missing separate debuginfos, use: dnf debuginfo-install bc-1.07.1-5.el8.x86_64
```

For more details on this process, see [Section 7.4, “Getting debuginfo packages for an application or library using GDB”](#).

- To export the core file for further processing elsewhere:

```
$ coredumpctl dump PID > /path/to/file_for_export
```

Replace `/path/to/file_for_export` with the file where you want to put the core dump.

10.5. DUMPING PROCESS MEMORY WITH **GCORE**

The workflow of core dump debugging enables the analysis of the program’s state offline. In some cases, you can use this workflow with a program that is still running, such as when it is hard to access the environment with the process. You can use the **gcore** command to dump memory of any process while it is still running.

Prerequisites

- [You must understand what core dumps are and how they are created.](#)
- [GDB must be installed on the system.](#)

Procedure

1. Find out the process id (*pid*). Use tools such as **ps**, **pgrep**, and **top**:

```
$ ps -C some-program
```

2. Dump the memory of this process:

```
$ gcore -o filename pid
```

This creates a file **filename** and dumps the process memory in it. While the memory is being dumped, the execution of the process is halted.

3. After the core dump is finished, the process resumes normal execution.
4. Create an SOS report to provide additional information about the system:

```
# sosreport
```

This creates a tar archive containing information about your system, such as copies of configuration files.

5. Transfer the program’s executable file, core dump, and the SOS report to the computer where the debugging will take place.
6. Optional: Remove the core dump and SOS report after transferring them, to free up disk space.

Additional resources

- [How to obtain a core file without restarting an application?](#) – Knowledgebase article

10.6. DUMPING PROTECTED PROCESS MEMORY WITH GDB

You can mark the memory of processes as not to be dumped. This can save resources and ensure additional security when the process memory contains sensitive data: for example, in banking or accounting applications or on whole virtual machines. Both kernel core dumps (**kdump**) and manual core dumps (**gcore**, GDB) do not dump memory marked this way.

In some cases, you must dump the whole contents of the process memory regardless of these protections. This procedure shows how to do this using the GDB debugger.

Prerequisites

- [You must understand what core dumps are.](#)
- [GDB must be installed on the system.](#)
- [GDB must be already attached to the process with protected memory.](#)

Procedure

1. Set GDB to ignore the settings in the **/proc/PID/coredump_filter** file:

```
(gdb) set use-coredump-filter off
```

2. Set GDB to ignore the memory page flag **VM_DONTDUMP**:

```
(gdb) set dump-excluded-mappings on
```

3. Dump the memory:

```
(gdb) gcore core-file
```

Replace *core-file* with name of file where you want to dump the memory.

Additional resources

- Debugging with GDB - [How to Produce a Core File from Your Program](#)

CHAPTER 11. COMPATIBILITY-BREAKING CHANGES IN GDB

The version of GDB provided in Red Hat Enterprise Linux 8 contains a number of changes that break compatibility, especially for cases where the GDB output is read directly from the terminal. The following sections provide more details about these changes.

Parsing output of GDB is not recommended. Prefer scripts using the Python GDB API or the GDB Machine Interface (MI).

GDBserver now starts inferiors with shell

To enable expansion and variable substitution in inferior command line arguments, GDBserver now starts the inferior in a shell, same as GDB.

To disable using the shell:

- When using the **target extended-remote** GDB command, disable shell with the **set startup-with-shell off** command.
- When using the **target remote** GDB command, disable shell with the **--no-startup-with-shell** option of GDBserver.

Example 11.1. Example of shell expansion in remote GDB inferiors

This example shows how running the `/bin/echo /*` command through GDBserver differs on Red Hat Enterprise Linux versions 7 and 8:

- On RHEL 7:

```
$ gdbserver --multi :1234
$ gdb -batch -ex 'target extended-remote :1234' -ex 'set remote exec-file /bin/echo' -ex
'file /bin/echo' -ex 'run /*'
/*
```

- On RHEL 8:

```
$ gdbserver --multi :1234
$ gdb -batch -ex 'target extended-remote :1234' -ex 'set remote exec-file /bin/echo' -ex
'file /bin/echo' -ex 'run /*'
/bin /boot (...) /tmp /usr /var
```

gcj support removed

Support for debugging Java programs compiled with the GNU Compiler for Java (**gcj**) has been removed.

New syntax for symbol dumping maintenance commands

The symbol dumping maintenance commands syntax now includes options before file names. As a result, commands that worked with GDB in RHEL 7 do not work in RHEL 8.

As an example, the following command no longer stores symbols in a file, but produces an error message:

```
(gdb) maintenance print symbols /tmp/out main.c
```

The new syntax for the symbol dumping maintenance commands is:

```

maint print symbols [-pc address] [--] [filename]
maint print symbols [-objfile objfile] [-source source] [--] [filename]
maint print psymbols [-objfile objfile] [-pc address] [--] [filename]
maint print psymbols [-objfile objfile] [-source source] [--] [filename]
maint print msymbols [-objfile objfile] [--] [filename]

```

Thread numbers are no longer global

Previously, GDB used only global thread numbering. The numbering has been extended to be displayed per inferior in the form **inferior_num.thread_num**, such as **2.1**. As a consequence, thread numbers in the **\$_thread** convenience variable and in the **InferiorThread.num** Python attribute are no longer unique between inferiors.

GDB now stores a second thread ID per thread, called the global thread ID, which is the new equivalent of thread numbers in previous releases. To access the global thread number, use the **\$_gthread** convenience variable and **InferiorThread.global_num** Python attribute.

For backwards compatibility, the Machine Interface (MI) thread IDs always contains the global IDs.

Example 11.2. Example of GDB thread number changes

On Red Hat Enterprise Linux 7:

```

# debuginfo-install coreutils
$ gdb -batch -ex 'file echo' -ex start -ex 'add-inferior' -ex 'inferior 2' -ex 'file echo' -ex start -ex 'info
threads' -ex 'pring $_thread' -ex 'inferior 1' -ex 'pring $_thread'
(...)
  Id  Target Id      Frame
* 2   process 203923 "echo" main (argc=1, argv=0x7ffffffdb88) at src/echo.c:109
  1   process 203914 "echo" main (argc=1, argv=0x7ffffffdb88) at src/echo.c:109
$1 = 2
(...)
$2 = 1

```

On Red Hat Enterprise Linux 8:

```

# dnf debuginfo-install coreutils
$ gdb -batch -ex 'file echo' -ex start -ex 'add-inferior' -ex 'inferior 2' -ex 'file echo' -ex start -ex 'info
threads' -ex 'pring $_thread' -ex 'inferior 1' -ex 'pring $_thread'
(...)
  Id  Target Id      Frame
  1.1 process 4106488 "echo" main (argc=1, argv=0x7ffffffce58) at ../src/echo.c:109
* 2.1 process 4106494 "echo" main (argc=1, argv=0x7ffffffce58) at ../src/echo.c:109
$1 = 1
(...)
$2 = 1

```

Memory for value contents can be limited

Previously, GDB did not limit the amount of memory allocated for value contents. As a consequence, debugging incorrect programs could cause GDB to allocate too much memory. The **max-value-size** setting has been added to enable limiting the amount of allocated memory. The default value of this limit is 64 KiB. As a result, GDB in Red Hat Enterprise Linux 8 will not display too large values, but report that the value is too large instead.

As an example, printing a value defined as **char s[128*1024];** produces different results:

- On Red Hat Enterprise Linux 7, **\$1 = 'A' <repeats 131072 times>**
- On Red Hat Enterprise Linux 8, **value requires 131072 bytes, which is more than max-value-size**

Sun version of stabs format no longer supported

Support for the Sun version of the **stabs** debug file format has been removed. The **stabs** format produced by GCC in RHEL with the **gcc -gstabs** option is still supported by GDB.

Sysroot handling changes

The **set sysroot path** command specifies system root when searching for files needed for debugging. Directory names supplied to this command may now be prefixed with the string **target:** to make GDB read the shared libraries from the target system (both local and remote). The formerly available **remote:** prefix is now treated as **target:**. Additionally, the default system root value has changed from an empty string to **target:** for backward compatibility.

The specified system root is prepended to the file name of the main executable, when GDB starts processes remotely, or when it attaches to already running processes (both local and remote). This means that for remote processes, the default value **target:** makes GDB always try to load the debugging information from the remote system. To prevent this, run the **set sysroot** command before the **target remote** command so that local symbol files are found before the remote ones.

HISTSIZE no longer controls GDB command history size

Previously, GDB used the **HISTSIZE** environment variable to determine how long command history should be kept. GDB has been changed to use the **GDBHISTSIZE** environment variable instead. This variable is specific only to GDB. The possible values and their effects are:

- a positive number - use command history of this size,
- **-1** or an empty string - keep history of all commands,
- non-numeric values - ignored.

Completion limiting added

The maximum number of candidates considered during completion can now be limited using the **set max-completions** command. To show the current limit, run the **show max-completions** command. The default value is 200. This limit prevents GDB from generating excessively large completion lists and becoming unresponsive.

As an example, the output after the input **p <tab><tab>** is:

- on RHEL 7: **Display all 29863 possibilities? (y or n)**
- on RHEL 8: **Display all 200 possibilities? (y or n)**

HP-UX XDB compatibility mode removed

The **-xdb** option for the HP-UX XDB compatibility mode has been removed from GDB.

Handling signals for threads

Previously, GDB could deliver a signal to the current thread instead of the thread for which the signal was actually sent. This bug has been fixed, and GDB now always passes the signal to the correct thread when resuming execution.

Additionally, the **signal** command now always correctly delivers the requested signal to the current thread. If the program is stopped for a signal and the user switched threads, GDB asks for confirmation.

Breakpoint modes always-inserted off and auto merged

The **breakpoint always-inserted** setting has been changed. The **auto** value and corresponding behavior has been removed. The default value is now **off**. Additionally, the **off** value now causes GDB to not remove breakpoints from the target until all threads stop.

remotebaud commands no longer supported

The **set remotebaud** and **show remotebaud** commands are no longer supported. Use the **set serial baud** and **show serial baud** commands instead.

PART III. ADDITIONAL TOOLSETS FOR DEVELOPMENT

Besides the development related tools that are available as part of the operating system, developers can install additional toolsets on Red Hat Enterprise Linux. These toolsets can contain tools for different languages, alternative toolchains, or alternative versions of system tools.

CHAPTER 12. USING GCC TOOLSET

12.1. WHAT IS GCC TOOLSET

Red Hat Enterprise Linux 8 introduces GCC Toolset, an Application Stream containing more up-to-date versions of development and performance analysis tools. GCC Toolset is similar to [Red Hat Developer Toolset](#) for RHEL 7.

GCC Toolset is available as an Application Stream in the form of a software collection in the **AppStream** repository. GCC Toolset is fully supported under Red Hat Enterprise Linux Subscription Level Agreements, is functionally complete, and is intended for production use. Applications and libraries provided by GCC Toolset do not replace the Red Hat Enterprise Linux system versions, do not override them, and do not automatically become default or preferred choices. Using a framework called software collections, an additional set of developer tools is installed into the **/opt/** directory and is explicitly enabled by the user on demand using the **scl** utility. Unless noted otherwise for specific tools or features, GCC Toolset is available for all architectures supported by Red Hat Enterprise Linux.

12.2. INSTALLING GCC TOOLSET

Installing GCC Toolset on a system installs the main tools and all necessary dependencies. Note that some parts of the toolset are not installed by default and must be installed separately.

Procedure

- To install GCC Toolset version *N*:

```
# yum install gcc-toolset-N
```

12.3. INSTALLING INDIVIDUAL PACKAGES FROM GCC TOOLSET

To install only certain tools from GCC Toolset instead of the whole toolset, list the available packages and install the selected ones with the **yum** package management tool. This procedure is useful also for packages that are not installed by default with the toolset.

Procedure

1. List the packages available in GCC Toolset version *N*:

```
$ yum list available gcc-toolset-N-*
```

2. To install any of these packages:

```
# yum install package_name
```

Replace *package_name* with a space-separated list of packages to install. For example, to install the **gcc-toolset-9-gdb-gdbserver** and **gcc-toolset-9-gdb-doc** packages:

```
# yum install gcc-toolset-9-gdb-gdbserver gcc-toolset-9-gdb-doc
```

12.4. UNINSTALLING GCC TOOLSET

To remove GCC Toolset from your system, uninstall it using the **yum** package management tool.

Procedure

- To uninstall GCC Toolset version *N*:

```
# yum remove gcc-toolset-M*
```

12.5. RUNNING A TOOL FROM GCC TOOLSET

To run a tool from GCC Toolset, use the **scl** utility.

Procedure

- To run a tool from GCC Toolset version *N*:

```
$ scl enable gcc-toolset-N tool
```

12.6. RUNNING A SHELL SESSION WITH GCC TOOLSET

GCC Toolset allows running a shell session where the GCC Toolset tool versions are used instead of system versions of these tools, without explicitly using the **scl** command. This is useful when you need to interactively start the tools many times, such as when setting up or testing a development setup.

Procedure

- To run a shell session where tool versions from GCC Toolset version *N* override system versions of these tools:

```
$ scl enable gcc-toolset-N bash
```

12.7. RELATED INFORMATION

- [Red Hat Developer Toolset User Guide](#)

CHAPTER 13. GCC TOOLSET 9

This chapter provides information specific to GCC Toolset version 9 and the tools contained in this version.

13.1. TOOLS AND VERSIONS PROVIDED BY GCC TOOLSET 9

GCC Toolset 9 provides the following tools and versions:

Table 13.1. Tool versions in GCC Toolset 9

Name	Version	Description
GCC	9.2.1	A portable compiler suite with support for C, C++, and Fortran.
GDB	8.3	A command-line debugger for programs written in C, C++, and Fortran.
Valgrind	3.15.0	An instrumentation framework and a number of tools to profile applications in order to detect memory errors, identify memory management problems, and report any use of improper arguments in system calls.
SystemTap	4.1	A tracing and probing tool to monitor the activities of the entire system without the need to instrument, recompile, install, and reboot.
Dyninst	10.1.0	A library for instrumenting and working with user-space executables during their execution.
binutils	2.32	A collection of binary tools and other utilities to inspect and manipulate object files and binaries.
elfutils	0.176	A collection of binary tools and other utilities to inspect and manipulate ELF files.
dwz	0.12	A tool to optimize DWARF debugging information contained in ELF shared libraries and ELF executables for size.
make	4.2.1	A dependency-tracking build automation tool.
strace	5.1	A debugging tool to monitor system calls that a program uses and signals it receives.
ltrace	0.7.91	A debugging tool to display calls to dynamic libraries that a program makes. It can also monitor system calls executed by programs.
annobin	9.08	A build security checking tool.

13.2. C++ COMPATIBILITY IN GCC TOOLSET 9



IMPORTANT

The compatibility information presented here apply only to the GCC from GCC Toolset 9.

The GCC compiler in GCC Toolset can use the following C++ standards:

C++14

This is the **default** language standard setting for GCC Toolset 9, with GNU extensions, equivalent to explicitly using option **-std=gnu++14**.

Using the C++14 language version is supported when all C++ objects compiled with the respective flag have been built using GCC version 6 or later.

C++11

This language standard is available in GCC Toolset 9.

Using the C++11 language version is supported when all C++ objects compiled with the respective flag have been built using GCC version 5 or later.

C++98

This language standard is available in GCC Toolset 9. Binaries, shared libraries and objects built using this standard can be freely mixed regardless of being built with GCC from GCC Toolset, Red Hat Developer Toolset, and RHEL 5, 6, 7 and 8.

C++17, C++2a

These language standards are available in GCC Toolset 9 only as an experimental, unstable, and unsupported capability. Additionally, compatibility of objects, binary files, and libraries built using these standards cannot be guaranteed.


All of the language standards are available in both the standard compliant variant or with GNU extensions.

When mixing objects built with GCC Toolset with those built with the RHEL toolchain (particularly **.o** or **.a** files), GCC Toolset toolchain should be used for any linkage. This ensures any newer library features provided only by GCC Toolset are resolved at link time.

13.3. SPECIFICS OF GCC IN GCC TOOLSET 9

Static linking of libraries

Certain more recent library features are statically linked into applications built with GCC Toolset to support execution on multiple versions of Red Hat Enterprise Linux. This creates an additional minor security risk as standard Red Hat Enterprise Linux errata do not change this code. If the need arises for developers to rebuild their applications due to this risk, Red Hat will communicate this using a security erratum.



IMPORTANT

Because of this additional security risk, developers are strongly advised not to statically link their entire application for the same reasons.

Specify libraries after object files when linking

In GCC Toolset, libraries are linked using linker scripts which might specify some symbols through static archives. This is required to ensure compatibility with multiple versions of Red Hat Enterprise Linux.

However, the linker scripts use the names of the respective shared object files. As a consequence, the linker uses different symbol handling rules than expected, and does not recognize symbols required by object files when the option adding the library is specified before options specifying the object files:

```
$ scl enable gcc-toolset-9 'gcc -lsomelib objfile.o'
```

Using a library from GCC Toolset in this manner results in the linker error message **undefined reference to symbol**. To prevent this problem, follow the standard linking practice and specify the option adding the library after the options specifying the object files:

```
$ scl enable gcc-toolset-9 'gcc objfile.o -lsomelib'
```

Note that this recommendation also applies when using the base Red Hat Enterprise Linux version of **GCC**.

13.4. SPECIFICS OF BINUTILS IN GCC TOOLSET 9

Static linking of libraries

Certain more recent library features are statically linked into applications built with GCC Toolset to support execution on multiple versions of Red Hat Enterprise Linux. This creates an additional minor security risk as standard Red Hat Enterprise Linux errata do not change this code. If the need arises for developers to rebuild their applications due to this risk, Red Hat will communicate this using a security erratum.



IMPORTANT

Because of this additional security risk, developers are strongly advised not to statically link their entire application for the same reasons.

Specify libraries after object files when linking

In GCC Toolset, libraries are linked using linker scripts which might specify some symbols through static archives. This is required to ensure compatibility with multiple versions of Red Hat Enterprise Linux. However, the linker scripts use the names of the respective shared object files. As a consequence, the linker uses different symbol handling rules than expected, and does not recognize symbols required by object files when the option adding the library is specified before options specifying the object files:

```
$ scl enable gcc-toolset-9 'ld -lsomelib objfile.o'
```

Using a library from GCC Toolset in this manner results in the linker error message **undefined reference to symbol**. To prevent this problem, follow the standard linking practice, and specify the option adding the library after the options specifying the object files:

```
$ scl enable gcc-toolset-9 'ld objfile.o -lsomelib'
```

Note that this recommendation also applies when using the base Red Hat Enterprise Linux version of **binutils**.

CHAPTER 14. USING THE GCC TOOLSET CONTAINER IMAGES

The GCC Toolset 9 components are available in the two container images:

- GCC Toolset 9 Toolchain
- GCC Toolset 9 Perftools

The GCC Toolset container images are based on the **rhel8** base image and are available for all architectures supported by RHEL 8:

- AMD and Intel 64-bit architectures
- The 64-bit ARM architecture
- IBM Power Systems, Little Endian
- IBM Z

14.1. GCC TOOLSET CONTAINER IMAGES CONTENTS

Tools versions provided in the GCC Toolset 9 container images match [the GCC Toolset 9 components versions](#).

The GCC Toolset 9 Toolchain contents

The **rhel8/gcc-toolset-9-toolchain** image provides the GCC compiler, the GDB debugger, and other development-related tools. The container image consists of the following components:

Component	Package
gcc	gcc-toolset-9-gcc
g++	gcc-toolset-9-gcc++
gfortran	gcc-toolset-9-gfortran
gdb	gcc-toolset-9-gdb

The GCC Toolset 9 Perftools contents

The **rhel8/gcc-toolset-9-perftools** image provides a number of tools for debugging, performance monitoring, and further analysis of the applications. The container image consists of the following components:

Component	Package
Valgrind	gcc-toolset-9-valgrind
SystemTap	gcc-toolset-9-systemtap

Component	Package
Dyninst	gcc-toolset-9-dyninst
elfutils	gcc-toolset-9-elfutils

Additional resources

- To use the GCC Toolset components on RHEL 7, use Red Hat Developer Toolset that provides similar developer tools for RHEL 7 users – [Red Hat Developer Toolset user guide](#).
- Instructions on using the Red Hat Developer Toolset container images on RHEL 7 – [Red Hat Developer Toolset images](#).

14.2. ACCESSING AND RUNNING THE GCC TOOLSET CONTAINER IMAGES

The following section describes how to access and run the GCC Toolset 9 container images.

Prerequisites

- Podman is installed.

Procedure

1. Access the [Red Hat Container Registry](#) using your Customer Portal credentials:

```
$ podman login registry.redhat.io
Username: username
Password: *****
```

2. Pull a container image you require by running a relevant command as root:

```
# podman pull registry.redhat.io/rhel8/gcc-toolset-9-toolchain
```

```
# podman pull registry.redhat.io/rhel8/gcc-toolset-9-perftools
```



NOTE

On RHEL 8.1 and later versions, you can set up your system to work with containers as a non-root user. For details, see [Running containers as root or rootless](#).

3. Optional: Check that pulling was successful by running a command that lists all container images on your local system:

```
# podman images
```

4. Run a container by launching a bash shell inside a container:

```
# podman run -it image_name /bin/bash
```

The **-i** option creates an interactive session; without this option the shell opens and instantly exits.

The **-t** option opens a terminal session; without this option you cannot type anything to the shell.

Additional resources

- [Building, running, and managing Linux containers on RHEL 8](#)
- A Red Hat blog article – [Understanding root inside and outside a container](#)
- Entries in the Red Hat Container Registry – [GCC Toolset container images](#)

14.3. EXAMPLE: USING THE GCC TOOLSET 9 TOOLCHAIN CONTAINER IMAGE

This example shows how to pull and start using the GCC Toolset 9 Toolchain container image.

Prerequisites

- Podman is installed.

Procedure

1. Access the Red Hat Container Registry using your Customer Portal credentials:

```
$ podman login registry.redhat.io
Username: username
Password: *****
```

2. Pull the container image as root:

```
# podman pull registry.redhat.io/rhel8/gcc-toolset-9-toolchain
```

3. Launch the container image with an interactive shell as root:

```
# podman run -it registry.redhat.io/rhel8/gcc-toolset-9-toolchain /bin/bash
```

4. Run the GCC Toolset tools as expected. For example, to verify the **gcc** compiler version, run:

```
bash-4.4$ gcc -v
...
gcc version 9.2.1 20191120 (Red Hat 9.2.1-2) (GCC)
```

5. To list all packages provided in the container, run:

```
bash-4.4$ rpm -qa
```

CHAPTER 15. COMPILER TOOLSETS

RHEL 8.0 provides the following compiler toolsets as Application Streams:


- LLVM Toolset 9.0.1, which provides the LLVM compiler infrastructure framework, the Clang compiler for the C and C++ languages, the LLDB debugger, and related tools for code analysis. See the [Using LLVM Toolset](#) guide.
- Rust Toolset 1.41, which provides the Rust programming language compiler **rustc**, the **cargo** build tool and dependency manager, the **cargo-vendor** plugin, and required libraries. See the [Using Rust Toolset](#) guide.
- Go Toolset 1.13, which provides the Go programming language tools and libraries. Go is alternatively known as **golang**. See the [Using Go Toolset](#) guide.

CHAPTER 16. THE ANNOBIN PROJECT

The Annobin project is an implementation of the Watermark specification project. Watermark specification project intends to add markers to Executable and Linkable Format (ELF) objects to determine their properties. The Annobin project consists of the **annobin** plugin and the **annockeck** program.

The **annobin** plugin scans the GNU Compiler Collection (GCC) command line, the compilation state, and the compilation process, and generates the ELF notes. The ELF notes record how the binary was built and provide information for the **annockeck** program to perform security hardening checks.

The security hardening checker is part of the **annockeck** program and is enabled by default. It checks the binary files to determine whether the program was built with necessary security hardening options and compiled correctly. **annockeck** is able to recursively scan directories, archives, and RPM packages for ELF object files.



NOTE

The files must be in ELF format. **annockeck** does not handle any other binary file types.

The following section describes how to:

- Use the **annobin** plugin
- Use the **annockeck** program
- Remove redundant **annobin** notes

16.1. USING THE ANNOBIN PLUGIN

The following section describes how to:

- Enable the **annobin** plugin
- Pass options to the **annobin** plugin

16.1.1. Enabling the annobin plugin

The following section describes how to enable the **annobin** plugin via **gcc** and via **clang**.

Procedure

- To enable the **annobin** plugin with **gcc**, use:

```
$ gcc -fplugin=annobin
```

- If **gcc** does not find the **annobin** plugin, use:

```
$ gcc -iplugindir=/path/to/directory/containing/annobin/
```

Replace */path/to/directory/containing/annobin/* with the absolute path to the directory that contains **annobin**.

- To find the directory containing the **annobin** plugin, use:

```
$ gcc --print-file-name=plugin
```

- To enable the **annobin** plugin with **clang**, use:

```
$ clang -fplugin=/path/to/directory/containing/annobin/
```

Replace */path/to/directory/containing/annobin/* with the absolute path to the directory that contains **annobin**.

16.1.2. Passing options to the annobin plugin

The following section describes how to pass options to the **annobin** plugin via **gcc** and via **clang**.

Procedure

- To pass options to the **annobin** plugin with **gcc**, use:

```
$ gcc -fplugin=annobin -fplugin-arg-annobin-option file-name
```

Replace *option* with the **annobin** command line arguments and replace *file-name* with the name of the file.

Example

- To display additional details about what **annobin** it is doing, use:

```
$ gcc -fplugin=annobin -fplugin-arg-annobin-verbose file-name
```

Replace *file-name* with the name of the file.

- To pass options to the **annobin** plugin with **clang**, use:

```
$ clang -fplugin=/path/to/directory/containing/annobin/ -Xclang -plugin-arg-annobin -Xclang option file-name
```

Replace *option* with the **annobin** command line arguments and replace */path/to/directory/containing/annobin/* with the absolute path to the directory containing **annobin**.

Example

- To display additional details about what **annobin** it is doing, use:

```
$ clang -fplugin=/usr/lib64/clang/10/lib/annobin.so -Xclang -plugin-arg-annobin -Xclang verbose file-name
```

Replace *file-name* with the name of the file.

16.2. USING THE ANNOCHECK PROGRAM

The following section describes how to use **annocheck** to examine:

- Files

- Directories
- RPM packages
- **annockeck** extra tools

**NOTE**

annockeck recursively scans directories, archives, and RPM packages for ELF object files. The files have to be in the ELF format. **annockeck** does not handle any other binary file types.

16.2.1. Using annockeck to examine files

The following section describes how to examine ELF files using **annockeck**.

Procedure

- To examine a file, use:

```
$ annockeck file-name
```

Replace *file-name* with the name of a file.

**NOTE**

The files must be in ELF format. **annockeck** does not handle any other binary file types. **annockeck** processes static libraries that contain ELF object files.

Additional information

- For more information about **annockeck** and possible command line options, see the **annockeck** man page.

16.2.2. Using annockeck to examine directories


The following section describes how to examine ELF files in a directory using **annockeck**.

Procedure

- To scan a directory, use:

```
$ annockeck directory-name
```

Replace *directory-name* with the name of a directory. **annockeck** automatically examines the contents of the directory, its sub-directories, and any archives and RPM packages within the directory.

**NOTE**

annockeck only looks for ELF files. Other file types are ignored.

Additional information

- For more information about **annockeck** and possible command line options, see the **annockeck** man page.

16.2.3. Using annockeck to examine RPM packages

The following section describes how to examine ELF files in an RPM package using **annockeck**.

Procedure

- To scan an RPM package, use:

```
$ annockeck rpm-package-name
```

Replace *rpm-package-name* with the name of an RPM package. **annockeck** recursively scans all the ELF files inside the RPM package.



NOTE

annockeck only looks for ELF files. Other file types are ignored.

- To scan an RPM package with provided debug info RPM, use:

```
$ annockeck rpm-package-name --debug-rpm debuginfo-rpm
```

Replace *rpm-package-name* with the name of an RPM package, and *debuginfo-rpm* with the name of a debug info RPM associated with the binary RPM.

Additional information

- For more information about **annockeck** and possible command line options, see the **annockeck** man page.


16.2.4. Using annockeck extra tools

annockeck includes multiple tools for examining binary files. You can enable these tools with the command-line options.

The following section describes how to enable the:

- **built-by** tool
- **notes** tool
- **section-size** tool

You can enable multiple tools at the same time.



NOTE

The hardening checker is enabled by default.

16.2.4.1. Enabling the built-by tool

You can use the **annockeck built-by** tool to find the name of the compiler that built the binary file.

Procedure

- To enable the **built-by** tool, use:

```
$ annockeck --enable-built-by
```

Additional information

- For more information about the **built-by** tool, see the **--help** command-line option.

16.2.4.2. Enabling the notes tool

You can use the **annockeck notes** tool to display the notes stored inside a binary file created by the **annobin** plugin.

Procedure

- To enable the **notes** tool, use:

```
$ annockeck --enable-notes
```

The notes are displayed in a sequence sorted by the address range.

Additional information

- For more information about the **notes** tool, see the **--help** command-line option.

16.2.4.3. Enabling the section-size tool

You can use the **annockeck section-size** tool display the size of the named sections.

Procedure

- To enable the **section-size** tool, use:

```
$ annockeck --section-size=name
```

Replace *name* with the name of the named section. The output is restricted to specific sections. A cumulative result is produced at the end.

Additional information

- For more information about the **section-size** tool, see the **--help** command-line option.

16.2.4.4. Hardening checker basics

The hardening checker is enabled by default. You can disable the hardening checker with the **--disable-hardened** command-line option.

16.2.4.4.1. Hardening checker options

The **annockeck** program checks the following options:

- Lazy binding is disabled using the **-z now** linker option.
- The program does not have a stack in an executable region of memory.
- The relocations for the GOT table are set to read only.
- No program segment has all three of the read, write and execute permission bits set.
- There are no relocations against executable code.
- The runpath information for locating shared libraries at runtime includes only directories rooted at `/usr`.
- The program was compiled with **annobin** notes enabled.
- The program was compiled with the **-fstack-protector-strong** option enabled.
- The program was compiled with **-D_FORTIFY_SOURCE=2**.
- The program was compiled with **-D_GLIBCXX_ASSERTIONS**.
- The program was compiled with **-fexceptions** enabled.
- The program was compiled with **-fstack-clash-protection** enabled.
- The program was compiled at **-O2** or higher.
- The program does not have any relocations held in a writeable.
- Dynamic executables have a dynamic segment.
- Shared libraries were compiled with **-fPIC** or **-fPIE**.
- Dynamic executables were compiled with **-fPIE** and linked with **-pie**.
- If available, the **-fcf-protection=full** option was used.
- If available, the **-mbranch-protection** option was used.
- If available, the **-mstackrealign** option was used.

16.2.4.4.2. Disabling the hardening checker

The following section describes how to disable the hardening checker.

Procedure

- To scan the notes in a file without the hardening checker, use:

```
$ annockeck --enable-notes --disable-hardened file-name
```

Replace *file-name* with the name of a file.

16.3. REMOVING REDUNDANT ANNOBIN NOTES

Using **annobin** increases the size of binaries. To reduce the size of the binaries compiled with **annobin** you can remove redundant **annobin** notes. To remove the redundant **annobin** notes use the **objcopy** program, which is a part of the **binutils** package.

Procedure

- To remove the redundant **annobin** notes, use:

```
$ objcopy --merge-notes file-name
```

Replace *file-name* with the name of the file.

PART IV. SUPPLEMENTARY TOPICS

CHAPTER 17. COMPATIBILITY-BREAKING CHANGES IN COMPILERS AND DEVELOPMENT TOOLS

librtkaio removed

With this update, the **librtkaio** library has been removed. This library provided high-performance real-time asynchronous I/O access for some files, which was based on Linux kernel Asynchronous I/O support (KAIO).

As a result of the removal:

- Applications using the **LD_PRELOAD** method to load **librtkaio** display a warning about a missing library, load the **librt** library instead and run correctly.
- Applications using the **LD_LIBRARY_PATH** method to load **librtkaio** load the **librt** library instead and run correctly, without any warning.
- Applications using the **dlopen()** system call to access **librtkaio** directly load the **librt** library instead.

Users of **librtkaio** have the following options:

- Use the fallback mechanism described above, without any changes to their applications.
- Change code of their applications to use the **librt** library, which offers a compatible POSIX-compliant API.
- Change code of their applications to use the **libaio** library, which offers a compatible API.

Both **librt** and **libaio** can provide comparable features and performance under specific conditions.

Note that the **libaio** package has Red Hat compatibility level of 2, while **librtk** and the removed **librtkaio** level 1.

For more details, see https://fedoraproject.org/wiki/Changes/GLIBC223_librtkaio_removal

Sun RPC and NIS interfaces removed from glibc

The **glibc** library no longer provides Sun RPC and NIS interfaces for new applications. These interfaces are now available only for running legacy applications. Developers must change their applications to use the **libtirpc** library instead of Sun RPC and **libnsl2** instead of NIS. Applications can benefit from IPv6 support in the replacement libraries.

The nasegseg libraries for 32-bit Xen have been removed

Previously, the **glibc** i686 packages contained an alternative **glibc** build, which avoided the use of the thread descriptor segment register with negative offsets (**nasegseg**). This alternative build was only used in the 32-bit version of the Xen Project hypervisor without hardware virtualization support, as an optimization to reduce the cost of full paravirtualization. These alternative builds are no longer used and they have been removed.

make new operator != causes a different interpretation of certain existing makefile syntax

The **!=** shell assignment operator has been added to GNU **make** as an alternative to the **\$(shell ...)** function to increase compatibility with BSD makefiles. As a consequence, variables with name ending in exclamation mark and immediately followed by assignment such as **variable!=value** are now interpreted as the shell assignment. To restore the previous behavior, add a space after the exclamation mark, such as **variable! =value**.

For more details and differences between the operator and the function, see the GNU **make** manual.

Valgrind library for MPI debugging support removed

The **libmpiwrap.so** wrapper library for **Valgrind** provided by the **valgrind-openmpi** package has been removed. This library enabled **Valgrind** to debug programs using the Message Passing Interface (MPI). This library was specific to the Open MPI implementation version in previous versions of Red Hat Enterprise Linux.

Users of **libmpiwrap.so** are encouraged to build their own version from upstream sources specific to their MPI implementation and version. Supply these custom-built libraries to **Valgrind** using the **LD_PRELOAD** technique.

Development headers and static libraries removed from **valgrind-devel**

Previously, the **valgrind-devel** sub-package used to include development files for developing custom valgrind tools. This update removes these files because they do not have a guaranteed API, have to be linked statically, and are unsupported. The **valgrind-devel** package still does contain the development files for valgrind-aware programs and header files such as **valgrind.h**, **callgrind.h**, **drd.h**, **helgrind.h**, and **memcheck.h**, which are stable and well-supported.

CHAPTER 18. OPTIONS FOR RUNNING A RHEL 6 OR 7 APPLICATION ON RHEL 8

To run a Red Hat Enterprise Linux 6 or 7 application on Red Hat Enterprise Linux 8, a spectrum of options is available. A system administrator needs detailed guidance from the application developer. The following list outlines the options, considerations, and resources provided by Red Hat.

Run the application in a virtual machine with a matching RHEL version guest OS

Resource costs are high for this option, but the environment is a close match to the application's requirements, and this approach does not require many additional considerations. This is the currently recommended option.

Run the application in a container based on the respective RHEL version

Resource costs are lower than in the previous cases, while configuration requirements are stricter. For details on the relationship between container hosts and guest user spaces, see the [Red Hat Enterprise Linux Container Compatibility Matrix](#).

Run the application natively on RHEL 8

This option offers the lowest resource costs, but also the most strict requirements. The application developer must determine a correct configuration of the RHEL 8 system. The following resources can help the developer in this task:

- [Red Hat Enterprise Linux 8: Application Compatibility Guide](#)
- [Red Hat Enterprise Linux 7: Application Compatibility Guide](#)
- [Release notes for Red Hat Enterprise Linux 8.0](#)
- [Considerations in adopting RHEL 8](#)

Note that this list is not a complete set of resources needed to determine application compatibility. These are only starting points with lists of known incompatible changes and Red Hat policies related to compatibility.

Additionally, the [What is Kernel Application Binary Interface \(kABI\)?](#) Knowledge Centered Support article contains information relevant to kernel and compatibility.