

Red Hat Enterprise Linux 8 Performing a standard RHEL installation

Installing Red Hat Enterprise Linux 8 using the graphical user interface

Last Updated: 2020-08-05

Red Hat Enterprise Linux 8 Performing a standard RHEL installation

Installing Red Hat Enterprise Linux 8 using the graphical user interface

Legal Notice

Copyright © 2020 Red Hat, Inc.

The text of and illustrations in this document are licensed by Red Hat under a Creative Commons Attribution–Share Alike 3.0 Unported license ("CC-BY-SA"). An explanation of CC-BY-SA is available at

http://creativecommons.org/licenses/by-sa/3.0/

. In accordance with CC-BY-SA, if you distribute this document or an adaptation of it, you must provide the URL for the original version.

Red Hat, as the licensor of this document, waives the right to enforce, and agrees not to assert, Section 4d of CC-BY-SA to the fullest extent permitted by applicable law.

Red Hat, Red Hat Enterprise Linux, the Shadowman logo, the Red Hat logo, JBoss, OpenShift, Fedora, the Infinity logo, and RHCE are trademarks of Red Hat, Inc., registered in the United States and other countries.

Linux ® is the registered trademark of Linus Torvalds in the United States and other countries.

Java [®] is a registered trademark of Oracle and/or its affiliates.

XFS [®] is a trademark of Silicon Graphics International Corp. or its subsidiaries in the United States and/or other countries.

MySQL [®] is a registered trademark of MySQL AB in the United States, the European Union and other countries.

Node.js ® is an official trademark of Joyent. Red Hat is not formally related to or endorsed by the official Joyent Node.js open source or commercial project.

The OpenStack [®] Word Mark and OpenStack logo are either registered trademarks/service marks or trademarks/service marks of the OpenStack Foundation, in the United States and other countries and are used with the OpenStack Foundation's permission. We are not affiliated with, endorsed or sponsored by the OpenStack Foundation, or the OpenStack community.

All other trademarks are the property of their respective owners.

Abstract

This document is for users who want to perform a standard Red Hat Enterprise Linux 8 installation using the graphical user interface.

Table of Contents

PROVIDING FEEDBACK ON RED HAT DOCUMENTATION	. 7
CHAPTER 1. INTRODUCTION	. 8
1.1. SUPPORTED ARCHITECTURES	8
1.2. INSTALLATION TERMINOLOGY	8
PART I. INSTALLING RED HAT ENTERPRISE LINUX ON AMD64, INTEL 64, AND 64-BIT ARM	. 9
CHAPTER 2. PREPARING FOR YOUR INSTALLATION	10
2.1. RECOMMENDED STEPS	10
2.2. AVAILABLE INSTALLATION METHODS	10
2.3. SYSTEM REQUIREMENTS	11
2.4. INSTALLATION BOOT MEDIA OPTIONS	11
2.5. TYPES OF INSTALLATION ISO IMAGES	12
2.6. DOWNLOADING THE INSTALLATION ISO IMAGE	13
2.6.1. Downloading an ISO image from the Customer Portal	13
2.6.2. Downloading an ISO image using curl	14
2.7. CREATING A BOOTABLE INSTALLATION MEDIUM	15
2.7.1. Creating a bootable DVD or CD	15
2.7.2. Creating a bootable USB device on Linux	16
2.7.3. Creating a bootable USB device on Windows	17
2.7.4. Creating a bootable USB device on Mac OS X	18
2.8. PREPARING AN INSTALLATION SOURCE	20
2.8.1. Types of installation source	20
2.8.2. Specify the installation source	21
2.8.3. Ports for network-based installation	21
2.8.4. Creating an installation source on an NFS server	21
2.8.5. Creating an installation source using HTTP or HTTPS	23
2.8.6. Creating an installation source using FTP	24
CHAPTER 3. GETTING STARTED	27
3.1. BOOTING THE INSTALLATION	27
3.1.1. Boot menu	27
3.1.2. Types of boot options	28
3.1.3. Editing boot options	29
Editing the boot: prompt in BIOS	29
Editing the > prompt Editing the GRUB2 menu	29
3.1.4. Booting the installation from a USB, CD, or DVD	30 30
3.1.5. Booting the installation from a network using PXE	31
3.1.3. BOOTING THE INSTALLATION FROM THE CUSTOMER PORTAL 3.2. INSTALLING RHEL USING AN ISO IMAGE FROM THE CUSTOMER PORTAL	33
3.3. REGISTERING AND INSTALLING RHEL FROM THE CDN USING THE GUI	35
3.3.1. What is the Content Delivery Network	35
3.3.2. Registering and installing RHEL from the CDN	35
3.3.2.1. Installation source repository after system registration	38
3.3.3. Verifying your system registration from the CDN	39
3.3.4. Unregistering your system from the CDN	40
CHAPTER 4. CUSTOMIZING YOUR INSTALLATION	42
4.1. CONFIGURING LANGUAGE AND LOCATION SETTINGS	42
4.2. CONFIGURING LOCALIZATION OPTIONS	43
4.2.1. Configuring keyboard, language, and time and date settings	43

4.3. CONFIGURING SYSTEM OPTIONS	45
4.3.1. Configuring installation destination	45
4.3.1.1. Configuring boot loader	48
4.3.2. Configuring Kdump	49
4.3.3. Configuring network and host name options	50
4.3.3.1. Configuring network and host name	50
4.3.3.2. Adding a virtual network interface	51
4.3.3.3. Editing network interface configuration	52
4.3.3.4. Enabling or Disabling the Interface Connection	52
4.3.3.5. Setting up Static IPv4 or IPv6 Settings	53
4.3.3.6. Configuring Routes	54
4.3.3.7. Additional resources	54
4.3.4. Configuring Connect to Red Hat	54
4.3.4.1. Introduction to System Purpose	55
4.3.4.2. Configuring Connect to Red Hat options	56
4.3.4.3. Installation source repository after system registration	57
4.3.4.4. Verifying your system registration from the CDN	58
4.3.4.5. Unregistering your system from the CDN	58
4.3.5. Configuring Security Policy	60
4.3.5.1. About security policy	60
4.3.5.2. Configuring a security policy	60
4.3.5.3. Related information	61
4.4. CONFIGURING SOFTWARE OPTIONS	61
4.4.1. Configuring installation source	61
4.4.2. Configuring software selection	63
4.5. CONFIGURING STORAGE DEVICES	65
4.5.1. Storage device selection	65
4.5.2. Filtering storage devices	65
4.5.3. Using advanced storage options	66
4.5.3.1. Discovering and starting an iSCSI session	67
4.5.3.2. Configuring FCoE parameters	68
4.5.3.3. Configuring DASD storage devices	69
4.5.3.4. Configuring FCP devices	70
4.5.4. Installing to an NVDIMM device	71
4.5.4.1. Criteria for using an NVDIMM device as an installation target	71
4.5.4.2. Configuring an NVDIMM device using the graphical installation mode	71
4.6. CONFIGURING MANUAL PARTITIONING	71
4.6.1. Starting manual partitioning	73
	73 74
4.6.2. Adding a mount point file system 4.6.3. Configuring a mount point file system	74 75
4.6.4. Customizing a partition or volume	75 76
	78
4.6.5. Preserving the /home directory 4.6.6. Creating software RAID	70 79
-	
4.6.7. Creating an LVM logical volume	80 81
4.6.8. Configuring an LVM logical volume	01
CHAPTER 5. STARTING AND COMPLETING THE INSTALLATION	83
5.1. STARTING THE INSTALLATION	83
5.2. CONFIGURING A ROOT PASSWORD	83
5.3. CREATING A USER ACCOUNT	84
5.3.1. Editing advanced user settings	85
5.4. COMPLETING THE INSTALLATION	86

CHAPTER 6. COMPLETING POST-INSTALLATION TASKS	87
6.1. COMPLETING INITIAL SETUP	87
6.2. REGISTERING YOUR SYSTEM USING THE COMMAND LINE	89
6.3. REGISTERING YOUR SYSTEM USING THE SUBSCRIPTION MANAGER USER INTERFACE	90
6.4. REGISTRATION ASSISTANT	91
6.5. CONFIGURING SYSTEM PURPOSE USING THE SYSPURPOSE COMMAND-LINE TOOL	91
6.6. SECURING YOUR SYSTEM	93
6.7. DEPLOYING SYSTEMS THAT ARE COMPLIANT WITH A SECURITY PROFILE IMMEDIATELY AFTER AN INSTALLATION	ا 94
6.7.1. Deploying baseline-compliant RHEL systems using the graphical installation	94
6.7.2. Deploying baseline-compliant RHEL systems using Kickstart	95
6.8. NEXT STEPS	96
APPENDIX A. TROUBLESHOOTING	97
A.1. TROUBLESHOOTING AT THE START OF THE INSTALLATION PROCESS	97
A.1.1. Dracut	97
A.1.2. Using installation log files	97
A.1.2.1. Creating pre-installation log files	98
A.1.2.2. Transferring installation log files to a USB drive	98
A.1.2.3. Transferring installation log files over the network	99
A.1.3. Detecting memory faults using the Memtest86 application	100
A.1.3.1. Running Memtest86	100
A.1.4. Verifying boot media	101
A.1.5. Consoles and logging during installation	101
A.1.6. Saving screenshots	102
A.1.7. Resuming an interrupted download attempt	102
A.1.8. Cannot boot into the graphical installation	103
A.2. TROUBLESHOOTING DURING THE INSTALLATION	104
A.2.1. Disks are not detected	104
A.2.2. Reporting error messages to Red Hat Customer Support	105
A.2.3. Partitioning issues for IBM Power Systems	106
A.3. TROUBLESHOOTING AFTER INSTALLATION	106
A.3.1. Cannot boot with a RAID card	106
A.3.2. Graphical boot sequence is not responding	106
A.3.3. X server fails after log in	107
A.3.4. RAM is not recognized	108
A.3.5. System is displaying signal 11 errors	109 109
A.3.6. Unable to IPL from network storage space A.3.7. Using XDMCP	109
A.3.8. Using rescue mode	1109
A.3.8.1. Booting into rescue mode	111
A.3.8.2. Using an SOS report in rescue mode	112
A.3.8.3. Reinstalling the GRUB2 boot loader	113
A.3.8.4. Using RPM to add or remove a driver	114
A.3.9. ip= boot option returns an error	115
APPENDIX B. SYSTEM REQUIREMENTS REFERENCE	117
B.1. HARDWARE COMPATIBILITY	117
B.2. SUPPORTED INSTALLATION TARGETS	117
B.3. SYSTEM SPECIFICATIONS	117
B.4. DISK AND MEMORY REQUIREMENTS	118
B.5. RAID REQUIREMENTS	119
APPENDIX C. PARTITIONING REFERENCE	121

C.1. SUPPORTED DEVICE TYPES C.2. SUPPORTED FILE SYSTEMS C.3. SUPPORTED RAID TYPES	121 121 122
C.4. RECOMMENDED PARTITIONING SCHEME C.5. ADVICE ON PARTITIONS	123 125
APPENDIX D. BOOT OPTIONS REFERENCE	128
D.1. INSTALLATION SOURCE BOOT OPTIONS	128
D.2. NETWORK BOOT OPTIONS	132
D.3. CONSOLE BOOT OPTIONS	134
D.4. DEBUG BOOT OPTIONS	137
D.5. STORAGE BOOT OPTIONS	138
D.6. DEPRECATED BOOT OPTIONS	139
D.7. REMOVED BOOT OPTIONS	140
APPENDIX E. CHANGING A SUBSCRIPTION SERVICE	142
E.1. UNREGISTERING FROM SUBSCRIPTION MANAGEMENT SERVER	142
E.1.1. Unregistering using command line	142
E.1.2. Unregistering using Subscription Manager user interface	143
E.2. UNREGISTERING FROM SATELLITE SERVER	143
APPENDIX F. ISCSI DISKS IN INSTALLATION PROGRAM	144
PART II. INSTALLING RED HAT ENTERPRISE LINUX ON IBM POWER SYSTEM LC SERVERS	145
CHAPTER 7. INSTALLING RED HAT ENTERPRISE LINUX ON IBM POWER SYSTEM LC SERVERS	146
7.1. OVERVIEW	146
7.1.1. Creating a bootable USB device on Linux	146
7.1.2. Creating a bootable USB device on Windows	148
7.1.3. Creating a bootable USB device on Mac OS X	149
7.2. COMPLETING THE PREREQUISITES AND BOOTING YOUR FIRMWARE	150
7.3. CONFIGURING THE IP ADDRESS IBM POWER	151
7.4. POWERING ON YOUR SERVER WITH IPMI 7.5. CHOOSE YOUR INSTALLATION METHOD ON IBM LC SERVERS	152 153
7.5.1. Configuring Petitboot for installation with USB device	153
7.5.2. Access BMC Advanced System Management interface to configure virtual media	154
7.6. COMPLETING YOUR LC SERVER INSTALLATION	155
PART III. INSTALLING RED HAT ENTERPRISE LINUX ON IBM POWER SYSTEM AC SERVERS	156
CHAPTER 8. INSTALLING RED HAT ENTERPRISE LINUX ON IBM POWER SYSTEM ACCELERATED SERV	/ERS 157
8.1. OVERVIEW	157
8.2. COMPLETING THE PREREQUISITES AND BOOTING YOUR FIRMWARE	157
8.3. CONFIGURING THE FIRMWARE IP ADDRESS	158
8.4. POWERING ON YOUR SERVER WITH OPENBMC COMMANDS	159
8.5. CHOOSE YOUR INSTALLATION METHOD ON IBM ACCELERATED SERVERS	160
8.6. CONFIGURING PETITBOOT FOR NETWORK INSTALLATION	160
8.7. CONFIGURING PETITBOOT FOR INSTALLATION WITH USB DEVICE ON ACCELERATED SERVERS	161
8.8. COMPLETING YOUR ACCELERATED SERVER INSTALLATION	162
PART IV. INSTALLING RED HAT ENTERPRISE LINUX ON IBM POWER SYSTEM L SERVERS	163
CHAPTER 9. INSTALLING RED HAT ENTERPRISE LINUX ON IBM POWER SYSTEM L SERVER	164 164
9.1. OVERVIEW 9.2. COMPLETING THE PREREQUISITES AND BOOTING YOUR FIRMWARE ON L SERVER	164

9.3. CONNECTING TO ASMI WITH DHCP	165
9.4. CONNECTING TO ASMI WITH STATIC IP ADDRESS	165
9.5. ENABLING IPMI	167
9.6. POWERING ON YOUR L SERVER WITH IPMI	167
9.7. POWERING ON YOUR SYSTEM FROM A NOTEBOOK OR PC RUNNING LINUX	168
9.8. POWERING ON YOUR SYSTEM FROM A NOTEBOOK OR PC RUNNING WINDOWS	169
9.9. CONFIGURING PETITBOOT AND INSTALLING RED HAT ENTERPRISE LINUX	169
PART V. INSTALLING RED HAT ENTERPRISE LINUX ON IBM Z	171
CHAPTER 10. PREPARING FOR INSTALLATION ON IBM Z	172
10.1. OVERVIEW OF THE IBM Z INSTALLATION PROCESS	172
10.2. CUSTOMIZING BOOT PARAMETERS	172
10.3. BOOTING THE INSTALLATION	174
10.3.1. Booting the installation	174
10.4. PLANNING FOR INSTALLATION ON IBM Z	174
10.4.1. Pre-installation	174
10.5. INSTALLING UNDER Z/VM	176
10.6. USING THE Z/VM READER	177
10.7. USING A PREPARED DASD	178
10.8. USING A PREPARED FCP ATTACHED SCSI DISK	178
10.9. USING AN FCP-ATTACHED SCSI DVD DRIVE	179
10.10. USING PARAMETER AND CONFIGURATION FILES ON IBM Z	180
10.11. REQUIRED CONFIGURATION FILE PARAMETERS ON IBM Z	180
10.12. IBM Z/VM CONFIGURATION FILE	180
10.13. INSTALLATION NETWORK PARAMETERS ON IBM Z	181
10.14. PARAMETERS FOR KICKSTART INSTALLATIONS ON IBM Z	184
10.15. MISCELLANEOUS PARAMETERS ON IBM Z	184
10.16. SAMPLE PARAMETER FILE AND CMS CONFIGURATION FILE ON IBM Z	185
CHAPTER 11. INSTALLING IN AN LPAR	187
11.1. INSTALLING IN AN LPAR	187
11.2. INSTALLING IN AN LPAR USING AN FTP SERVER	187
11.3. INSTALLING IN AN LPAR USING A PREPARED DASD	188
11.4. INSTALLING IN AN LPAR USING A PREPARED FCP ATTACHED SCSI DISK	188
11.5. INSTALLING IN AN LPAR USING AN FCP ATTACHED SCSI DVD DRIVE	188
CHAPTER 12. CONFIGURING A LINUX INSTANCE ON IBM Z	190
12.1. ADDING DASDS	190
12.2. DYNAMICALLY SETTING DASDS ONLINE	190
12.3. PREPARING A NEW DASD WITH LOW-LEVEL FORMATTING	191
12.4. PERSISTENTLY SETTING DASDS ONLINE	192
12.5. DASDS THAT ARE PART OF THE ROOT FILE SYSTEM	192
12.6. FCP LUNS THAT ARE PART OF THE ROOT FILE SYSTEM	194
12.7. FCP LUNS THAT ARE NOT PART OF THE ROOT FILE SYSTEM	196
12.8. ADDING A QETH DEVICE	197
12.9. DYNAMICALLY ADDING A QETH DEVICE	197
12.10. PERSISTENTLY ADDING A QETH DEVICE	199
12.11. CONFIGURING AN IBM Z NETWORK DEVICE FOR NETWORK ROOT FILE SYSTEM	202

PROVIDING FEEDBACK ON RED HAT DOCUMENTATION

We appreciate your input on our documentation. Please let us know how we could make it better. To do so:

- For simple comments on specific passages:
 - 1. Make sure you are viewing the documentation in the *Multi-page HTML* format. In addition, ensure you see the **Feedback** button in the upper right corner of the document.
 - 2. Use your mouse cursor to highlight the part of text that you want to comment on.
 - 3. Click the Add Feedback pop-up that appears below the highlighted text.
 - 4. Follow the displayed instructions.
- For submitting more complex feedback, create a Bugzilla ticket:
 - 1. Go to the Bugzilla website.
 - 2. As the Component, use **Documentation**.
 - 3. Fill in the **Description** field with your suggestion for improvement. Include a link to the relevant part(s) of documentation.
 - 4. Click Submit Bug.

CHAPTER 1. INTRODUCTION

Red Hat Enterprise Linux 8 delivers a stable, secure, consistent foundation across hybrid cloud deployments with the tools needed to deliver workloads faster with less effort. It can be deployed as a guest on supported hypervisors and Cloud provider environments as well as deployed on physical infrastructure, so your applications can take advantage of innovations in the leading hardware architecture platforms.

1.1. SUPPORTED ARCHITECTURES

Red Hat Enterprise Linux supports the following architectures:

- AMD and Intel 64-bit architectures
- The 64-bit ARM architecture
- IBM Power Systems, Little Endian
- IBM Z

1.2. INSTALLATION TERMINOLOGY

This section describes Red Hat Enterprise Linux installation terminology. Different terminology can be used for the same concepts, depending on its upstream or downstream origin.

Anaconda: The operating system installer used in Fedora, Red Hat Enterprise Linux, and their derivatives. Anaconda is a set of Python modules and scripts with additional files like Gtk widgets (written in C), systemd units, and dracut libraries. Together, they form a tool that allows users to set parameters of the resulting (target) system. In this document, the term **installation program** refers to the installation aspect of **Anaconda**.

PART I. INSTALLING RED HAT ENTERPRISE LINUX ON AMD64, INTEL 64, AND 64-BIT ARM

If you are new to Red Hat Enterprise Linux, it is recommended that you review Chapter 2, before you begin with the installation. However, if you are familiar with Red Hat Enterprise Linux, you can proceed with the GUI-based installation. Depending on your installation requirements, see, Section 3.2, "Installing RHEL using an ISO image from the Customer Portal" or Section 3.3, "Registering and installing RHEL from the CDN using the GUI"



NOTE

This document provides details about installing RHEL using the GUI. For more information about other installation methods, see Section 2.2, "Available installation methods".

CHAPTER 2. PREPARING FOR YOUR INSTALLATION

Before you begin to install Red Hat Enterprise Linux, review the following sections to prepare your setup for the installation.

2.1. RECOMMENDED STEPS

Preparing for your RHEL installation consists of the following steps:

Steps

- 1. Review and determine the installation method.
- 2. Check system requirements.
- 3. Review the installation boot media options.
- 4. Download the required installation ISO image.
- 5. Create a bootable installation medium.
- 6. Prepare the installation source*

2.2. AVAILABLE INSTALLATION METHODS

You can install Red Hat Enterprise Linux using any of the following methods:

- GUI-based installations
- System or cloud image-based installations
- Advanced installations



NOTE

This document provides details about installing RHEL using the user interfaces (GUI).

GUI-based installations

The following GUI-based installation methods are available:

- Install RHEL using an ISO image from the Customer Portal!nstall Red Hat Enterprise Linux by downloading the Binary DVD ISO image file from the Customer Portal. Registration is performed after the GUI installation completes. This installation method is also supported by Kickstart.
- Register and install RHEL from the Content Delivery Network: Register your system, attach
 subscriptions, and install Red Hat Enterprise Linux from the Content Delivery Network (CDN).
 This installation method is supported by the Boot ISO and Binary DVD ISO image files;
 however, it is recommended that you use the Boot ISO image file as the installation source

^{*}Only required for the Boot ISO (minimal install) image if you are not using the Content Delivery Network (CDN) to download the required software packages.

defaults to CDN for the Boot ISO image file. Registration is performed before the installation packages are downloaded and installed from the CDN. This installation method is also supported by Kickstart.



IMPORTANT

You can customize the RHEL installation for your specific requirements using the GUI. You can select additional options for specific environment requirements, for example, Connect to Red Hat, software selection, partitioning, security, and many more. For more information, see Chapter 4, Customizing your installation.

System or cloud image-based installations

You can use system or cloud image-based installation methods only in virtual and cloud environments.

To perform a system or cloud image-based installation, use Red Hat Image Builder. Image Builder creates customized system images of Red Hat Enterprise Linux, including the system images for cloud deployment.

For more information about installing RHEL using Image Builder, see the Composing a customized RHEL system image document.

Advanced installations

The following advanced installation methods are available:

- Perform an automated RHEL installation using Kickstart:Install Red Hat Enterprise Linux using Kickstart. Kickstart is an automated installation that allows you to execute unattended operating system installation tasks.
- Perform a remote RHEL installation using VNC: The RHEL installation program offers two VNC installation modes: Direct and Connect. Once a connection is established, the two modes do not differ. The mode you select depends on your environment.
- Install RHEL from the network using PXE: A network installation allows you to install Red Hat Enterprise Linux to a system that has access to an installation server. At a minimum, two systems are required for a network installation.

For more information about the advanced installation methods, see the *Performing an advanced RHEL installation* document.

2.3. SYSTEM REQUIREMENTS

If this is a first-time installation of Red Hat Enterprise Linux it is recommended that you review the guidelines provided for system, hardware, security, memory, and RAID before installing. See Appendix B, System requirements reference for more information.

Additional resources

For more information about securing Red Hat Enterprise Linux, see the Security hardening document.

For more information about the system requirements for installing Red Hat Enterprise Linux in virtual and cloud environment, using Image Builder, see *Composing a customized RHEL system image* document.

2.4. INSTALLATION BOOT MEDIA OPTIONS

There are several options available to boot the Red Hat Enterprise Linux installation program.

Full installation DVD or USB flash drive

Create a full installation DVD or USB flash drive using the **Binary DVD ISO** image. The DVD or USB flash drive can be used as a boot device and as an installation source for installing software packages. Due to the size of the Binary DVD ISO image, a DVD or USB flash drive are the recommended media types.

Minimal installation DVD, CD, or USB flash drive

Create a minimal installation CD, DVD, or USB flash drive using the **Boot ISO** image, which contains only the minimum files necessary to boot the system and start the installation program.



IMPORTANT

If you are not using the Content Delivery Network (CDN) to download the required software packages, the **Boot ISO** image requires an installation source that contains the required software packages.

PXE Server

A preboot execution environment (PXE) server allows the installation program to boot over the network. After a system boot, you must complete the installation from a different installation source, such as a local hard drive or a network location.

Image Builder

Image Builder allows to create customized system and cloud images to install Red Hat Enterprise Linux in virtual and cloud environment.

Additional resources

- For instructions on how to create an installation DVD or USB flash drive, see Section 2.7, "Creating a bootable installation medium" for more information.
- For instructions on how to create a bootable DVD, CD, and USB flash drive, see Section 2.8, "Preparing an installation source" for more information.
- For more information about PXE servers, see the *Performing an advanced RHEL installation* document.
- For more information about Image Builder, see the Composing a customized RHEL system image document.

2.5. TYPES OF INSTALLATION ISO IMAGES

Two types of Red Hat Enterprise Linux 8 installation ISO images are available from the Red Hat Customer Portal.

Binary DVD ISO image file

A full installation program that contains the BaseOS and AppStream repositories and allows you to complete the installation without additional repositories.



IMPORTANT

You can use a Binary DVD for IBM Z to boot the installation program using a SCSI DVD drive, or as an installation source.

Boot ISO image file

The Boot ISO image is a minimal installation that can be used to install RHEL in two different ways:

- a. When registering and installing RHEL from the Content Delivery Network (CDN).
- b. As a minimal image that requires access to the BaseOS and AppStream repositories to install software packages. The repositories are part of the Binary DVD ISO image that is available for download from https://access.redhat.com/home. Download and unpack the Binary DVD ISO image to access the repositories.

The following table contains information about the images that are available for the supported architectures.

Table 2.1. Boot and installation images

Architecture	Installation DVD	Boot DVD
AMD64 and Intel 64	x86_64 Binary DVD ISO image file	x86_64 Boot ISO image file
ARM 64	AArch64 Binary DVD ISO image file	AArch64 Boot ISO image file
IBM POWER	ppc64le Binary DVD ISO image file	ppc64le Boot ISO image file
IBM Z	s390x Binary DVD ISO image file	s390x Boot ISO image file

Additional resources

- For instructions on how to access the Binary DVD ISO image repositories, see Section 2.8, "Preparing an installation source" for more information.
- For more details about Appstream and BaseOS repositories, refer to, *Installing, managing, and removing user-space components* document.

2.6. DOWNLOADING THE INSTALLATION ISO IMAGE

This section contains instructions about downloading a Red Hat Enterprise Linux installation image from the Red Hat Customer Portal or by using the **curl** command.

2.6.1. Downloading an ISO image from the Customer Portal

Follow this procedure to download a Red Hat Enterprise Linux 8 ISO image file from the Red Hat Customer Portal.



NOTE

- The Boot ISO image is a minimal image file that supports registering your system, attaching subscriptions, and installing RHEL from the Content Delivery Network (CDN).
- The Binary DVD ISO image file contains all repositories and software packages and does not require any additional configuration. See Section 2.8, "Preparing an installation source" for more information.

Prerequisites

- You have an active Red Hat subscription.
- You are logged in to the **Product Downloads** section of the Red Hat Customer Portal at https://access.redhat.com/downloads.

Procedure

- 1. From the **Product Downloads** page, select the **By Category** tab.
- Click the Red Hat Enterprise Linux 8 link.
 The Download Red Hat Enterprise Linuxweb page opens.
- 3. From the **Product Variant** drop-down menu, select the variant that you require.
 - a. Optional: Select the **Packages** tab to view the packages contained in the selected variant. For information on the packages available in Red Hat Enterprise Linux 8, see the Package Manifest document.
- 4. The Version drop-down menu defaults to the latest version for the selected variant.
- 5. The **Architecture** drop-down menu displays the supported architecture. The **Product Software** tab displays the image files, which include:
 - Red Hat Enterprise Linux Binary DVDimage.
 - Red Hat Enterprise Linux Boot ISOimage.

Additional images may be available, for example, preconfigured virtual machine images, but they are beyond the scope of this document.

6. Click **Download Now** beside the ISO image that you require.

2.6.2. Downloading an ISO image using curl

Use the curl command to download installation images directly from a specific URL.

Prerequisites

- Verify the curl package is installed:
 - If your distribution uses the **yum** package manager:

yum install curl

• If your distribution uses the **dnf** package manager:

dnf install curl

• If your distribution uses the apt package manager:

apt update # apt install curl

- If your Linux distribution does not use yum, dnf, or apt, or if you do not use Linux, download the most appropriate software package from the curl web site.
- You have navigated to the Product Downloads section of the Red Hat Customer Portal at https://access.redhat.com/downloads, and selected the variant, version, and architecture that you require. You have right-clicked on the required ISO image file, and selected Copy Link Location to copy the URL of the ISO image file to your clipboard.

Procedure

1. On the command line, enter a suitable directory, and run the following command to download the file:

\$ curl --output directory-path/filename.iso 'copied_link_location'

Replace *directory-path* with a path to the location where you want to save the file; replace *filename.iso* with the ISO image name as displayed in the Customer Portal; replace *copied_link_location* with the link that you have copied from the Customer Portal.

2.7. CREATING A BOOTABLE INSTALLATION MEDIUM

This section contains information about using the ISO image file that you downloaded in Section 2.6, "Downloading the installation ISO image" to create a bootable physical installation medium, such as a USB, DVD, or CD.



NOTE

By default, the **inst.stage2=** boot option is used on the installation medium and is set to a specific label, for example, **inst.stage2=hd:LABEL=RHEL8\x86_64**. If you modify the default label of the file system containing the runtime image, or if you use a customized procedure to boot the installation system, you must verify that the label is set to the correct value.

2.7.1. Creating a bootable DVD or CD

You can create a bootable installation DVD or CD using burning software and a CD/DVD burner. The exact steps to produce a DVD or CD from an ISO image file vary greatly, depending on the operating system and disc burning software installed. Consult your system's burning software documentation for the exact steps to burn a CD or DVD from an ISO image file.



WARNING

You can create a bootable DVD or CD using either the Binary DVD ISO image (full install) or the Boot ISO image (minimal install). However, the Binary DVD ISO image is larger than 4.7 GB, and as a result, it might not fit on a single or dual-layer DVD. Check the size of the Binary DVD ISO image file before you proceed. A USB key is recommended when using the Binary DVD ISO image to create bootable installation media.

2.7.2. Creating a bootable USB device on Linux

Follow this procedure to create a bootable USB device on a Linux system.



NOTE

This procedure is destructive and data on the USB flash drive is destroyed without a warning.

Prerequisites

- You have downloaded an installation ISO image as described in Section 2.6, "Downloading the installation ISO image".
- The **Binary DVD** ISO image is larger than 4.7 GB, so a USB flash drive that is large enough to hold the ISO image is required.

Procedure

- 1. Connect the USB flash drive to the system.
- 2. Open a terminal window and run the **dmesg** command:
 - \$ dmesg|tail

The **dmesg** command returns a log that details all recent events. Messages resulting from the attached USB flash drive are displayed at the bottom of the log. Record the name of the connected device.

3. Switch to user root:

\$ su -

- 4. Enter your root password when prompted.
- 5. Find the device node assigned to the drive. In this example, the drive name is **sdd**.

```
# dmesg|tail
```

[288954.686557] usb 2-1.8: New USB device strings: Mfr=0, Product=1, SerialNumber=2

[288954.686559] usb 2-1.8: Product: USB Storage

[288954.686562] usb 2-1.8: SerialNumber: 000000009225

[288954.712590] usb-storage 2-1.8:1.0: USB Mass Storage device detected

[288954.712687] scsi host6: usb-storage 2-1.8:1.0

[288954.712809] usbcore: registered new interface driver usb-storage

[288954.716682] usbcore: registered new interface driver uas

[288955.717140] scsi 6:0:0:0: Direct-Access Generic STORAGE DEVICE 9228 PQ: 0

ANSI: 0

[288955.717745] sd 6:0:0:0: Attached scsi generic sg4 type 0 [288961.876382] sd 6:0:0:0: sdd Attached SCSI removable disk

6. Run the **dd** command to write the ISO image directly to the USB device.

dd if=/image_directory/image.iso of=/dev/device

Replace /image_directory/image.iso with the full path to the ISO image file that you downloaded, and replace device with the device name that you retrieved with the **dmesg** command. In this example, the full path to the ISO image is /home/testuser/Downloads/rhel-8-x86 64-boot.iso, and the device name is sdd:

dd if=/home/testuser/Downloads/rhel-8-x86_64-boot.iso of=/dev/sdd



NOTE

Ensure that you use the correct device name, and not the name of a partition on the device. Partition names are usually device names with a numerical suffix. For example, **sdd** is a device name, and **sdd1** is the name of a partition on the device **sdd**.

7. Wait for the **dd** command to finish writing the image to the device. The data transfer is complete when the **#** prompt appears. When the prompt is displayed, log out of the root account and unplug the USB drive. The USB drive is now ready to be used as a boot device.

2.7.3. Creating a bootable USB device on Windows

Follow the steps in this procedure to create a bootable USB device on a Windows system. The procedure varies depending on the tool. Red Hat recommends using Fedora Media Writer, available for download at https://github.com/FedoraQt/MediaWriter/releases.



NOTE

- Fedora Media Writer is a community product and is not supported by Red Hat.
 You can report any issues with the tool at
 https://github.com/FedoraQt/MediaWriter/issues.
- This procedure is destructive and data on the USB flash drive is destroyed without a warning.

Prerequisites

- You have downloaded an installation ISO image as described in Section 2.6, "Downloading the installation ISO image".
- The **Binary DVD** ISO image is larger than 4.7 GB, so a USB flash drive that is large enough to hold the ISO image is required.

Procedure

 Download and install Fedora Media Writer from https://github.com/FedoraQt/MediaWriter/releases.



NOTE

To install Fedora Media Writer on Red Hat Enterprise Linux, use the pre-built Flatpak package. You can obtain the package from the official Flatpak repository Flathub.org at https://flathub.org/apps/details/org.fedoraproject.MediaWriter.

- 2. Connect the USB flash drive to the system.
- 3. Open Fedora Media Writer.
- 4. From the main window, click **Custom Image** and select the previously downloaded Red Hat Enterprise Linux ISO image.
- 5. From Write Custom Image window, select the drive that you want to use.
- 6. Click **Write to disk**. The boot media creation process starts. Do not unplug the drive until the operation completes. The operation may take several minutes, depending on the size of the ISO image, and the write speed of the USB drive.
- 7. When the operation completes, unmount the USB drive. The USB drive is now ready to be used as a boot device.

2.7.4. Creating a bootable USB device on Mac OS X

Follow the steps in this procedure to create a bootable USB device on a Mac OS X system.



NOTE

This procedure is destructive and data on the USB flash drive is destroyed without a warning.

Prerequisites

- You have downloaded an installation ISO image as described in Section 2.6, "Downloading the installation ISO image".
- The **Binary DVD** ISO image is larger than 4.7 GB, so a USB flash drive that is large enough to hold the ISO image is required.

Procedure

- 1. Connect the USB flash drive to the system.
- 2. Identify the device path with the **diskutil list** command. The device path has the format of /dev/disknumber, where number is the number of the disk. The disks are numbered starting at zero (0). Typically, Disk 0 is the OS X recovery disk, and Disk 1 is the main OS X installation. In the following example, the USB device is **disk2**:

\$ diskutil list /dev/disk0

```
TYPE NAME
                                            IDENTIFIER
                                    SIZE
0:
     GUID_partition_scheme
                                       *500.3 GB disk0
              EFI EFI
                                209.7 MB disk0s1
1:
       Apple_CoreStorage
2:
                                     400.0 GB disk0s2
           Apple_Boot Recovery HD
3:
                                        650.0 MB disk0s3
4:
       Apple_CoreStorage
                                     98.8 GB disk0s4
           Apple Boot Recovery HD
5:
                                        650.0 MB disk0s5
/dev/disk1
              TYPE NAME
                                            IDENTIFIER
#:
                                    SIZE
           Apple HFS YosemiteHD
0:
                                        *399.6 GB disk1
Logical Volume on disk0s1
8A142795-8036-48DF-9FC5-84506DFBB7B2
Unlocked Encrypted
/dev/disk2
#:
              TYPE NAME
                                    SIZE
                                            IDENTIFIER
0:
    FDisk_partition_scheme
                                      *8.1 GB
                                               disk2
         Windows_NTFS SanDisk USB
1:
                                           8.1 GB disk2s1
```

- 3. To identify your USB flash drive, compare the NAME, TYPE and SIZE columns to your flash drive. For example, the NAME should be the title of the flash drive icon in the **Finder** tool. You can also compare these values to those in the information panel of the flash drive.
- 4. Use the **diskutil unmountDisk** command to unmount the flash drive's filesystem volumes:

\$ diskutil unmountDisk /dev/disknumber Unmount of all volumes on disknumber was successful

When the command completes, the icon for the flash drive disappears from your desktop. If the icon does not disappear, you may have selected the wrong disk. Attempting to unmount the system disk accidentally returns a **failed to unmount** error.

5. Log in as root:

\$ su -

- 6. Enter your root password when prompted.
- 7. Use the **dd** command as a parameter of the sudo command to write the ISO image to the flash drive:

sudo dd if=/path/to/image.iso of=/dev/rdisknumber



NOTE

Mac OS X provides both a block (/dev/disk*) and character device (/dev/rdisk*) file for each storage device. Writing an image to the /dev/rdisknumber character device is faster than writing to the /dev/disknumber block device.

8. To write the /Users/user_name/Downloads/rhel-8-x86_64-boot.iso file to the /dev/rdisk2 device, run the following command:

sudo dd if=/Users/user_name/Downloads/rhel-8-x86_64-boot.iso of=/dev/rdisk2

9. Wait for the **dd** command to finish writing the image to the device. The data transfer is complete when the **#** prompt appears. When the prompt is displayed, log out of the root account and unplug the USB drive. The USB drive is now ready to be used as a boot device.

2.8. PREPARING AN INSTALLATION SOURCE

The Boot ISO image file does not include any repositories or software packages; it contains only the installation program and the tools required to boot the system and start the installation. This section contains information about creating an installation source for the Boot ISO image using the Binary DVD ISO image that contains the required repositories and software packages.



IMPORTANT

An installation source is required for the Boot ISO image file only if you decide not to register and install RHEL from the Content Delivery Network (CDN).

2.8.1. Types of installation source

You can use one of the following installation sources for minimal boot images:

- **DVD:** Burn the Binary DVD ISO image to a DVD. The installation program will automatically install the software packages from the DVD.
- Hard drive or USB drive: Copy the Binary DVD ISO image to the drive and configure the
 installation program to install the software packages from the drive. If you use a USB drive,
 verify that it is connected to the system before the installation begins. The installation program
 cannot detect media after the installation begins.
 - Hard drive limitation: The Binary DVD ISO image on the hard drive must be on a partition with a file system that the installation program can mount. The supported file systems are xfs, ext2, ext4, and vfat (FAT32).



WARNING

On Microsoft Windows systems, the default file system used when formatting hard drives is NTFS. The exFAT file system is also available. However, neither of these file systems can be mounted during the installation. If you are creating a hard drive or a USB drive as an installation source on Microsoft Windows, verify that you formatted the drive as FAT32. Note that the FAT32 file system cannot store files larger than 4 GiB.

In Red Hat Enterprise Linux 8, you can enable installation from a directory on a local hard drive. To do so, you need to copy the contents of the DVD ISO image to a directory on a hard drive and then specify the directory as the installation source instead of the ISO image. For example:

inst.repo=hd:<device>:<path to the directory>

Network location: Copy the Binary DVD ISO image or the installation tree (extracted contents
of the Binary DVD ISO image) to a network location and perform the installation over the
network using the following protocols:

- NFS: The Binary DVD ISO image is in a Network File System (NFS) share.
- HTTPS, HTTP or FTP: The installation tree is on a network location that is accessible over HTTP, HTTPS or FTP.

2.8.2. Specify the installation source

You can specify the installation source using any of the following methods:

- **Graphical installation:** Select the installation source in the **Installation Source** window of the graphical installation. See Section 4.4.1, "Configuring installation source" for more information.
- **Boot option:** Configure a custom boot option to specify the installation source. See Appendix D, *Boot options reference* for more information.
- **Kickstart file:** Use the install command in a Kickstart file to specify the installation source. See the *Performing an advanced RHEL installation* document for more information.

2.8.3. Ports for network-based installation

The following table lists the ports that must be open on the server providing the files for each type of network-based installation.

Table 2.2. Ports for network-based installation

Protocol used	Ports to open
HTTP	80
HTTPS	443
FTP	21
NFS	2049, 111, 20048
TFTP	69

Additional resources

• See the Securing networks document for more information.

2.8.4. Creating an installation source on an NFS server

Follow the steps in this procedure to place the installation source on an NFS server. Use this installation method to install multiple systems from a single source, without having to connect to physical media.

Prerequisites

• You have administor level access to a server with Red Hat Enterprise Linux 8, and this server is on the same network as the system to be installed.

- You have downloaded a Binary DVD ISO image. See Section 2.6, "Downloading the installation ISO image" for more information.
- You have created a bootable CD, DVD, or USB device from the image file. See Section 2.7, "Creating a bootable installation medium" for more information.
- You have verified that your firewall allows the system you are installing to access the remote installation source. See Section 2.8.3, "Ports for network-based installation" for more information.

Procedure

- 1. Install the **nfs-utils** package:
 - # yum install nfs-utils
- 2. Copy the Binary DVD ISO image to a directory on the NFS server.
- 3. Open the /etc/exports file using a text editor and add a line with the following syntax:
 - /exported_directory/ clients
- 4. Replace /exported_directory/ with the full path to the directory with the ISO image. Replace clients with the host name or IP address of the target system, the subnetwork that all target systems can use to access the ISO image, or the asterisk sign (*) if you want to allow any system with network access to the NFS server to use the ISO image. See the **exports(5)** man page for detailed information about the format of this field.

A basic configuration that makes the /**rhel8-install**/ directory available as read-only to all clients is:

- /rhel8-install *
- 5. Save the /etc/exports file and exit the text editor.
- 6. Start the nfs service:
 - # systemctl start nfs-server.service

If the service was running before you changed the /etc/exports file, run the following command for the running NFS server to reload its configuration:

systemctl reload nfs-server.service

The ISO image is now accessible over NFS and ready to be used as an installation source.



NOTE

When configuring the installation source, use **nfs:** as the protocol, the server host name or IP address, the colon sign (:), and the directory holding the ISO image. For example, if the server host name is **myserver.example.com** and you have saved the ISO image in /rhel8-install/, specify nfs:myserver.example.com:/rhel8-install/ as the installation source.

2.8.5. Creating an installation source using HTTP or HTTPS

Follow the steps in this procedure to create an installation source for a network-based installation using an installation tree, which is a directory containing extracted contents of the Binary DVD ISO image and a valid .treeinfo file. The installation source is accessed over HTTP or HTTPS.

Prerequisites

- You have administor level access to a server with Red Hat Enterprise Linux 8, and this server is on the same network as the system to be installed.
- You have downloaded a Binary DVD ISO image. See Section 2.6, "Downloading the installation ISO image" for more information.
- You have created a bootable CD, DVD, or USB device from the image file. See Section 2.7, "Creating a bootable installation medium" for more information.
- You have verified that your firewall allows the system you are installing to access the remote installation source. See Section 2.8.3, "Ports for network-based installation" for more information.

Procedure

1. Install the httpd package:

yum install httpd



WARNING

If your Apache web server configuration enables SSL security, verify that you enable only the TLSv1 protocol, and disable SSLv2 and SSLv3. This is due to the POODLE SSL vulnerability (CVE-2014-3566). See https://access.redhat.com/solutions/1232413 for details.



IMPORTANT

If you use an HTTPS server with a self-signed certificate, you must boot the installation program with the **noverifyssl** option.

- 2. Copy the Binary DVD ISO image to the HTTP(S) server.
- 3. Mount the Binary DVD ISO image, using the **mount** command, to a suitable directory:

mkdir /mnt/rhel8-install/ # mount -o loop,ro -t iso9660 /image_directory/image.iso /mnt/rhel8-install/

Replace /image_directory/image.iso with the path to the Binary DVD ISO image.

4. Copy the files from the mounted image to the HTTP(S) server root. This command creates the /var/www/html/rhel8-install/ directory with the contents of the image.

cp -r /mnt/rhel8-install/ /var/www/html/

This command creates the /var/www/html/rhel8-install/ directory with the content of the image. Note that some copying methods can skip the .treeinfo file which is required for a valid installation source. Running the cp command for whole directories as shown in this procedure will copy .treeinfo correctly.

5. Start the **httpd** service:

systemctl start httpd.service

The installation tree is now accessible and ready to be used as the installation source.



NOTE

When configuring the installation source, use http:// as the protocol, the server host name or IP address, and the directory that contains the files from the ISO image, relative to the HTTP server root. For example, if you are using HTTP, the server host name is myserver.example.com, and you have copied the files from the image to /var/www/html/rhel8-install/, specify http://myserver.example.com/rhel8-install/ as the installation source.

Additional resources

• For more information about HTTP servers, see the *Deploying different types of servers* document.

2.8.6. Creating an installation source using FTP

Follow the steps in this procedure to create an installation source for a network-based installation using an installation tree, which is a directory containing extracted contents of the Binary DVD ISO image and a valid .treeinfo file. The installation source is accessed over FTP.

Prerequisites

- You have administor level access to a server with Red Hat Enterprise Linux 8, and this server is on the same network as the system to be installed.
- You have downloaded a Binary DVD ISO image. See Section 2.6, "Downloading the installation ISO image" for more information.
- You have created a bootable CD, DVD, or USB device from the image file. See Section 2.7, "Creating a bootable installation medium" for more information.
- You have verified that your firewall allows the system you are installing to access the remote installation source. See Section 2.8.3, "Ports for network-based installation" for more information.

Procedure

1. Install the **vsftpd** package by running the following command as root:

yum install vsftpd

- 2. Open and edit the /etc/vsftpd/vsftpd.conf configuration file in a text editor.
 - a. Change the line anonymous_enable=NO to anonymous_enable=YES
 - b. Change the line write_enable=YES to write_enable=NO.
 - c. Add lines **pasv_min_port** and **pasv_max_port**. Replace *min_port* and *max_port* with the port number range used by FTP server in passive mode, e. g. **10021** and **10031**.

This step can be necessary in network environments featuring various firewall/NAT setups.

d. Optionally, add custom changes to your configuration. For available options, see the **vsftpd.conf(5)** man page. This procedure assumes that default options are used.



WARNING

If you configured SSL/TLS security in your **vsftpd.conf** file, ensure that you enable only the TLSv1 protocol, and disable SSLv2 and SSLv3. This is due to the POODLE SSL vulnerability (CVE-2014-3566). See https://access.redhat.com/solutions/1234773 for details.

- 3. Configure the server firewall.
 - a. Enable the firewall:

systemctl enable firewalld # systemctl start firewalld

b. Enable in your firewall the FTP port and port range from previous step:

```
# firewall-cmd --add-port min_port-max_port/tcp --permanent
# firewall-cmd --add-service ftp --permanent
# firewall-cmd --reload
```

Replace *min_port-max_port* with the port numbers you entered into the /etc/vsftpd/vsftpd.conf configuration file.

- 4. Copy the Binary DVD ISO image to the FTP server.
- 5. Mount the Binary DVD ISO image, using the mount command, to a suitable directory:

mkdir /mnt/rhel8-install # mount -o loop,ro -t iso9660 /image-directory/image.iso /mnt/rhel8-install

Replace /image-directory/image.iso with the path to the Binary DVD ISO image.

6. Copy the files from the mounted image to the FTP server root:

mkdir /var/ftp/rhel8-install # cp -r /mnt/rhel8-install/ /var/ftp/ This command creates the /var/ftp/rhel8-install/ directory with the content of the image. Note that some copying methods can skip the .treeinfo file which is required for a valid installation source. Running the cp command for whole directories as shown in this procedure will copy .treeinfo correctly.

7. Make sure that the correct SELinux context and access mode is set on the copied content:

```
# restorecon -r /var/ftp/rhel8-install
# find /var/ftp/rhel8-install -type f -exec chmod 444 {} \;
# find /var/ftp/rhel8-install -type d -exec chmod 755 {} \;
```

8. Start the **vsftpd** service:

systemctl start vsftpd.service

If the service was running before you changed the /etc/vsftpd/vsftpd.conf file, restart the service to load the edited file:

systemctl restart vsftpd.service

Enable the **vsftpd** service to start during the boot process:

systemctl enable vsftpd

The installation tree is now accessible and ready to be used as the installation source.



NOTE

When configuring the installation source, use **ftp:**// as the protocol, the server host name or IP address, and the directory in which you have stored the files from the ISO image, relative to the FTP server root. For example, if the server host name is **myserver.example.com** and you have copied the files from the image to /var/ftp/rhel8-install/, specify ftp://myserver.example.com/rhel8-install/ as the installation source.

CHAPTER 3. GETTING STARTED

To get started with the installation, first review the boot menu and the available boot options. Then, depending on the choice you make, proceed to boot the installation.

3.1. BOOTING THE INSTALLATION

After you have created bootable media you are ready to boot the Red Hat Enterprise Linux installation.

3.1.1. Boot menu

The Red Hat Enterprise Linux boot menu is displayed using **GRand Unified Bootloader version 2** (GRUB2) when your system has completed loading the boot media.

Figure 3.1. Red Hat Enterprise Linux boot menu



The boot menu provides several options in addition to launching the installation program. If you do not make a selection within 60 seconds, the default boot option (highlighted in white) is run. To select a different option, use the arrow keys on your keyboard to make your selection and press the **Enter** key.

You can customize boot options for a particular menu entry:

 On BIOS-based systems: Press the Tab key and add custom boot options to the command line. You can also access the boot: prompt by pressing the Esc key but no required boot options are preset. In this scenario, you must always specify the Linux option before using any other boot options. • On UEFI-based systems: Press the **e** key and add custom boot options to the command line. When ready press **Ctrl+X** to boot the modified option.

Table 3.1. Boot menu options

Boot menu option	Description
Install Red Hat Enterprise Linux 8.x	Use this option to install Red Hat Enterprise Linux using the graphical installation program. For more information, see Section 3.2, "Installing RHEL using an ISO image from the Customer Portal"
Test this media & install Red Hat Enterprise Linux 8.x	Use this option to check the integrity of the installation media. For more information, see Section A.1.4, "Verifying boot media"
Troubleshooting >	Use this option to resolve various installation issues. Press Enter to display its contents.

Table 3.2. Troubleshooting options

Troubleshooting option	Description
Troubleshooting > Install Red Hat Enterprise Linux 8.x in basic graphics mode	Use this option to install Red Hat Enterprise Linux in graphical mode even if the installation program is unable to load the correct driver for your video card. If your screen is distorted when using the Install Red Hat Enterprise Linux 8.x option, restart your system and use this option. For more information, see Section A.1.8, "Cannot boot into the graphical installation"
Troubleshooting > Rescue a Red Hat Enterprise Linux system	Use this option to repair any issues that prevent you from booting. For more information, see Section A.3.8, "Using rescue mode"
Troubleshooting > Run a memory test	Use this option to run a memory test on your system. Press Enter to display its contents. For more information, see see Section A.1.3, "Detecting memory faults using the Memtest86 application"
Troubleshooting > Boot from local drive	Use this option to boot the system from the first installed disk. If you booted this disk accidentally, use this option to boot from the hard disk immediately without starting the installation program.

3.1.2. Types of boot options

There are two types of boot options; those with an equals "=" sign, and those without an equals "=" sign. Boot options are appended to the boot command line and multiple options must be separated by a single space. Boot options that are specific to the installation program always start with **inst**.

Options with an equals "=" sign

You must specify a value for boot options that use the **=** symbol. For example, the **inst.vncpassword=** option must contain a value, in this case, a password. The correct syntax for this example is **inst.vncpassword=password**.

Options without an equals "=" sign

This boot option does not accept any values or parameters. For example, the **rd.live.check** option forces the installation program to verify the installation media before starting the installation. If this boot option is present, the verification is performed; if the boot option is not present, the verification is skipped.

3.1.3. Editing boot options

This section contains information about the different ways that you can edit boot options from the boot menu. The boot menu opens after you boot the installation media.

Editing the boot: prompt in BIOS

When using the **boot:** prompt, the first option must always specify the installation program image file that you want to load. In most cases, you can specify the image using the keyword. You can specify additional options according to your requirements.

Prerequisites

- You have created bootable installation media (USB, CD or DVD).
- You have booted the installation from the media, and the installation boot menu is open.

Procedure

- 1. With the boot menu open, press the **Esc** key on your keyboard.
- 2. The **boot:** prompt is now accessible.
- 3. Press the **Tab** key on your keyboard to display the help commands.
- 4. Press the **Enter** key on your keyboard to start the installation with your options. To return from the **boot:** prompt to the boot menu, restart the system and boot from the installation media again.



NOTE

The **boot:** prompt also accepts **dracut** kernel options. A list of options is available in the **dracut.cmdline(7)** man page.

Editing the > prompt

You can use the > prompt to edit predefined boot options. For example, select **Test this media and install Red Hat Enterprise Linux 8.1** from the boot menu to display a full set of options.



NOTE

This procedure is for BIOS-based AMD64 and Intel 64 systems.

Prerequisites

- You have created bootable installation media (USB, CD or DVD).
- You have booted the installation from the media, and the installation boot menu is open.

Procedure

- 1. From the boot menu, select an option and press the **Tab** key on your keyboard. The > prompt is accessible and displays the available options.
- 2. Append the options that you require to the > prompt.
- 3. Press the **Enter** key on your keyboard to start the installation.
- 4. Press the **Esc** key on your keyboard to cancel editing and return to the boot menu.

Editing the GRUB2 menu

The GRUB2 menu is available on UEFI-based AMD64, Intel 64, and 64-bit ARM systems.

Prerequisites

- You have created bootable installation media (USB, CD or DVD).
- You have booted the installation from the media, and the installation boot menu is open.

Procedure

- 1. From the boot menu window, select the required option and press the **e** key on your keyboard.
- 2. Move the cursor to the kernel command line. On UEFI systems, the kernel command line starts with **linuxefi**.
- 3. Move the cursor to the end of the linuxefi kernel command line.
- 4. Edit the parameters as required. For example, to configure one or more network interfaces, add the **ip=** parameter at the end of the **linuxefi** kernel command line, followed by the required value.
- 5. When you finish editing, press **Ctrl+X** on your keyboard to start the installation using the specified options.

3.1.4. Booting the installation from a USB, CD, or DVD

Follow the steps in this procedure to boot the Red Hat Enterprise Linux installation using a USB, CD, or DVD. The following steps are generic. Consult your hardware manufacturer's documentation for specific instructions.

Prerequisite

You have created bootable installation media (USB, CD or DVD). See Section 2.7, "Creating a bootable installation medium" for more information.

Procedure

- 1. Power off the system to which you are installing Red Hat Enterprise Linux.
- 2. Disconnect any drives from the system.

- 3. Power on the system.
- 4. Insert the bootable installation media (USB, DVD, or CD).
- 5. Power off the system but do not remove the boot media.
- 6. Power on the system.



NOTE

You might need to press a specific key or combination of keys to boot from the media or configure the Basic Input/Output System (BIOS) of your system to boot from the media. For more information, see the documentation that came with your system.

- 7. The **Red Hat Enterprise Linux boot** window opens and displays information about a variety of available boot options.
- 8. Use the arrow keys on your keyboard to select the boot option that you require, and press **Enter** to select the boot option. The **Welcome to Red Hat Enterprise Linux**window opens and you can install Red Hat Enterprise Linux using the graphical user interface.



NOTE

The installation program automatically begins if no action is performed in the boot window within 60 seconds.

- 9. Optional: For UEFI-based systems, press **E** to edit the available boot options. For BIOS-based systems, press the **Tab** key on your keyboard to edit the available boot options. The boot window enters edit mode and you can change the predefined command line to add or remove boot options.
 - a. Press **Enter** to confirm your choice.

Additional Resources

- For more information about installing Red Hat Enterprise Linux using the Graphical User Interface, see Chapter 4, *Customizing your installation*.
- For more information about the list of available boot options you can use on the boot command line, see Appendix D, Boot options reference.

3.1.5. Booting the installation from a network using PXE

When installing Red Hat Enterprise Linux on a large number of systems simultaneously, the best approach is to boot from a PXE server and install from a source in a shared network location. Follow the steps in this procedure to boot the Red Hat Enterprise Linux installation from a network using PXE.



NOTE

To boot the installation process from a network using PXE, you must use a physical network connection, for example, Ethernet. You cannot boot the installation process with a wireless connection.

Prerequisites

- You have configured a TFTP server, and there is a network interface in your system that supports PXE. See **Additional resources** for more information.
- You have configured your system to boot from the network interface. This option is in the BIOS, and can be labeled **Network Boot** or **Boot Services**.
- You have verified that the BIOS is configured to boot from the specified network interface.
 Some BIOS systems specify the network interface as a possible boot device, but do not support the PXE standard. See your hardware's documentation for more information. When you have properly enabled PXE booting, the system can boot the Red Hat Enterprise Linux installation program without any other media.

Procedure

- 1. Verify that the network cable is attached. The link indicator light on the network socket should be lit, even if the computer is not switched on.
- Switch on the system.
 Depending on your hardware, some network setup and diagnostic information can be displayed before your system connects to a PXE server. When connected, a menu is displayed according to the PXE server configuration.
- 3. Press the number key that corresponds to the option that you require.



NOTE

In some instances, boot options are not displayed. If this occurs, press the **Enter** key on your keyboard or wait until the boot window opens.

The **Red Hat Enterprise Linux boot** window opens and displays information about a variety of available boot options.

4. Use the arrow keys on your keyboard to select the boot option that you require, and press **Enter** to select the boot option. The **Welcome to Red Hat Enterprise Linux**window opens and you can install Red Hat Enterprise Linux using the graphical user interface.



NOTE

The installation program automatically begins if no action is performed in the boot window within 60 seconds.

- 5. Optional: For UEFI-based systems, press **E** to edit the available boot options. For BIOS-based systems, press the **Tab** key on your keyboard to edit the available boot options. The boot window enters edit mode and you can change the predefined command line to add or remove boot options.
 - a. Press **Enter** to confirm your choice.

Additional Resources

• For information about installing Red Hat Enterprise Linux using the Graphical User Interface, see Chapter 4, Customizing your installation.

- For information about how to prepare to install Red Hat Enterprise Linux from the network using PXE, see the *Performing an advanced RHEL installation* document.
- For more information about the list of available boot options you can use on the boot command line, see Appendix D, Boot options reference.

3.2. INSTALLING RHEL USING AN ISO IMAGE FROM THE CUSTOMER PORTAL

Use this procedure to perform a graphical installation of RHEL using a Binary DVD ISO image that you downloaded from the Customer Portal.



NOTE

This section provides quick installation instructions. For more detailed information about customizing the installation settings and parameters, and completing the installation, see Chapter 4, Customizing your installation and Chapter 5, Starting and completing the installation.

Prerequisites

- You have downloaded the Binary DVD ISO image file from the Customer Portal. See Section 2.6, "Downloading the installation ISO image" for more information.
- You have created bootable installation media. See Section 2.7, "Creating a bootable installation medium" for more information.
- You have booted the installation program and the boot menu is displayed. See Section 3.1, "Booting the installation" for more information.



WARNING

When performing a GUI installation using the Binary DVD ISO image file, a race condition in the installer can sometimes prevent the installation from proceeding until you register the system using the Connect to Red Hat feature. For more information, see BZ#1823578 in the Known Issues section of the *RHEL 8.2 Release Notes* document.

Procedure

- 1. From the boot menu, select Install Red Hat Enterprise Linux 8.x
- 2. Press the **Enter** key on your keyboard.
- 3. From the **Welcome to Red Hat Enterprise Linux 8.x**window, select your language and location.
- 4. Click **Continue** to proceed to the **Installation Summary** window.



The **Installation Summary** window is the central hub to configure the Red Hat Enterprise Linux graphical user interface. The default settings assigned by the installation program are displayed under each category.

- 5. From the Installation Summary window, accept the default Localization and Software options.
- 6. Select **System > Installation Destination**.
 - a. From the Local Standard Disks pane, select the target disk.
 - b. Click **Done** to accept the selection and the default setting of automatic partitioning, and return to the **Installation Summary** window.

7. Select Network & Host Name

- a. Toggle the **Ethernet** switch to **ON** to enable network configuration.
 - i. Optional: Select a network device and click **Configure** to update the network interface configuration.
- b. Click **Done** to accept the changes and return to the **Installation Summary** window.
- 8. Optional: Select Connect to Red Hat
 - a. Configure the **Connect to Red Hat** options to attach RHEL subscriptions, install RHEL from the CDN, configure System Purpose, and enable Red Hat Insights.
 - b. Click **Done** to accept the changes and return to the **Installation Summary** window.
- 9. Optional: Select KDUMP.
 - a. Change the default KDUMP settings for your requirements.
 - b. Click **Done** to accept the changes and return to the **Installation Summary** window.
- 10. Optional: Select Security Policy.
 - a. Select the profile that you require, and click **Select profile**.
 - b. Click **Done** to accept the changes and return to the **Installation Summary** window.
- 11. Click **Begin Installation** to start the installation.
- 12. From the Configuration window, configure a root password and create a user account.
- 13. When the installation process is complete, click **Reboot** to restart the system.
- 14. From the Initial Setup window, accept the licensing agreement and register your system.

Additional resources

• To learn more about how to prepare for your installation, see Chapter 2, *Preparing for your installation*.

- To learn more about how to customize your network, connect to Red Hat, system purpose, installation destination, KDUMP, and security policy, see Chapter 4, Customizing your installation.
- To learn more about how to register your system, see Chapter 6, Completing post-installation tasks.

3.3. REGISTERING AND INSTALLING RHEL FROM THE CDN USING THE GUI

This section contains information about how to register your system, attach RHEL subscriptions, and install RHEL from the Red Hat Content Delivery Network (CDN) using the GUI.

3.3.1. What is the Content Delivery Network

The Red Hat Content Delivery Network (CDN), available from cdn.redhat.com, is a geographically distributed series of static web servers that contain content and errata that is consumed by systems. The content can be consumed directly, such as using a system registered to Red Hat Subscription Management. The CDN is protected by x.509 certificate authentication to ensure that only valid users have access. When a system is registered to Red Hat Subscription Management, the attached subscriptions govern which subset of the CDN the system can access.

Registering and installing RHEL from the CDN provides the following benefits:

- The CDN installation method supports the Boot ISO and the Binary DVD ISO image files. However, the use of the smaller Boot ISO image file is recommended as it consumes less space than the larger Binary DVD ISO image file.
- The CDN uses the latest packages resulting in a fully up-to-date system right after installation. There is no requirement to install package updates immediately after installation as is often the case when using the Binary DVD ISO image file.
- Integrated support for connecting to Red Hat Insights and enabling System Purpose.

Registering and installing RHEL from the CDN is supported by the GUI and Kickstart. For information about how to register and install RHEL from the CDN using the GUI, see Section 3.3.2, "Registering and installing RHEL from the CDN". For information about how to register and install RHEL using Kickstart, see the *Performing an advanced RHEL installation* document.

3.3.2. Registering and installing RHEL from the CDN

Use this procedure to register your system, attach RHEL subscriptions, and install RHEL from the Red Hat Content Delivery Network (CDN) using the GUI.



IMPORTANT

The CDN feature is supported by the **Boot ISO** and **Binary DVD ISO** image files. However, it is recommended that you use the **Boot ISO** image file as the installation source defaults to CDN for the Boot ISO image file.

Prerequisites

Your system is connected to a network that can access the CDN.

- You have downloaded the **Boot ISO** image file from the Customer Portal. See Section 2.6, "Downloading the installation ISO image" for more information.
- You have created bootable installation media. See Section 2.7, "Creating a bootable installation medium" for more information.
- You have booted the installation program and the boot menu is displayed. See Section 3.1, "Booting the installation" for more information.



IMPORTANT

The installation source repository used after system registration is dependent on how the system was booted. For more information, see Section 3.3.2.1, "Installation source repository after system registration".

Procedure

- 1. From the boot menu, select Install Red Hat Enterprise Linux 8.x
- 2. Press the **Enter** key on your keyboard.
- 3. From the **Welcome to Red Hat Enterprise Linux 8.x**window, select your language and location.
- 4. Click **Continue** to proceed to the **Installation Summary** window.



NOTE

The **Installation Summary** window is the central hub to configure a RHEL installation when using the GUI. The default settings assigned by the installation program are displayed under each category. The **Installation Source** defaults to *Red Hat CDN* and **Software Selection** displays a *Red Hat CDN requires registration* message.

- 5. From the **Installation Summary** window, the default **Localization** settings are displayed.
 - a. Optional: Change the default settings for your requirements.
- 6. Select System > Network & Host Name
 - a. Toggle the **Ethernet** switch to **ON** to enable network configuration.
 - i. Optional: Select a network device and click **Configure** to update the network interface configuration.
 - b. Click **Done** to accept the changes and return to the **Installation Summary** window.
- 7. Select System > Connect to Red Hat



IMPORTANT

You can register to the CDN using either your Red Hat account or your activation key details.

a. Click Account.

- i. Enter your Red Hat Customer Portal username and password details.
- b. Optional: Click Activation Key.
 - i. Enter your organization ID and activation key. You can enter more than one activation key, separated by a comma, as long as the activation keys are registered to your subscription.
- c. Select the **Set System Purpose** check box.
 - i. Select the required Role, SLA, and Usage from the corresponding drop-down lists.
- d. The **Connect to Red Hat Insights**check box is enabled by default. Clear the check box if you do not want to connect to Red Hat Insights.



Red Hat Insights is a Software-as-a-Service (SaaS) offering that provides continuous, in-depth analysis of registered Red Hat-based systems to proactively identify threats to security, performance and stability across physical, virtual and cloud environments, and container deployments.

- e. Optional: Expand Options.
 - i. Select the Use HTTP proxy check box if your network environment allows external Internet access only or access to content servers through an HTTP proxy. Clear the Use HTTP proxy check box if an HTTP proxy is not used.
 - ii. If you are running Satellite Server or performing internal testing, select the **Custom server URL** and **Custom base URL**check boxes and enter the required details.



IMPORTANT

- The Custom server URL field does not require the HTTP protocol, for example nameofhost.com. However, the Custom base URL field requires the HTTP protocol.
- To change the **Custom base URL** after registration, you must unregister, provide the new details, and then re-register.
- f. Click **Register** to register the system. When the system is successfully registered and subscriptions are attached, the **Connect to Red Hat** window displays the attached subscription details.



NOTE

Depending on the amount of subscriptions, the registration and attachment process might take up to a minute to complete.

- g. Click **Done** to return to the **Installation Summary** window.
 - i. A Registered message is displayed under Connect to Red Hat
- 8. Select System > Software Selection

- a. Choose the required options from Base Environment and Additional software for selected environment. The Server with GUI base environment is the default base environment and it launches the Initial Setup application after the installation completes and you restart the system.
- b. Click **Done** to accept the selection and return to the **Installation Summary** window.
- 9. Select **System > Installation Destination**
 - a. From the Local Standard Disks pane, select the target disk.
 - b. Click **Done** to accept the selection and the default setting of automatic partitioning and return to the **Installation Summary** window.
- 10. Optional: Select System > KDUMP.
 - a. Change the default KDUMP settings for your requirements.
 - b. Click **Done** to accept the changes and return to the **Installation Summary** window.
- 11. Optional: Select System > Security Policy.
 - a. Select the profile that you require and click **Select profile**.
 - b. Click **Done** to accept the changes and return to the **Installation Summary** window.
- 12. Click **Begin Installation** to start the installation. The required packages are downloaded from the CDN.
- 13. From the **Configuration** window, configure a root password and create a user account.
- 14. When the installation process is complete, click **Reboot** to restart the system.
- 15. From the **Initial Setup** window, accept the licensing agreement.

Additional resources

- To learn more about how to prepare for your installation, see Chapter 2, *Preparing for your installation*.
- To learn more about how to customize your network, connect to Red Hat, system purpose, installation destination, KDUMP, and security policy, see Chapter 4, Customizing your installation.
- To learn more about completing Initial Setup, see Section 6.1, "Completing initial setup".
- For information about Red Hat Insights, see the Red Hat Insights product documentation.
- For information about Activation Keys, see the *Understanding Activation Keys* chapter of the *Using Red Hat Subscription Management* document.
- For information about how to set up an HTTP proxy for Subscription Manager, see the *Using an HTTP proxy* chapter of the *Using and Configuring Red Hat Subscription Manager* document.

3.3.2.1. Installation source repository after system registration

The installation source repository used after system registration is dependent on how the system was booted.

System booted from the Boot ISO or the Binary DVD ISO image file

If you booted the RHEL installation using either the **Boot ISO** or the **Binary DVD ISO** image file with the default boot parameters, the installation program automatically switches the installation source repository to the CDN after registration.

System booted with theinst.repo=<URL> boot parameter

If you booted the RHEL installation with the **inst.repo=<URL>** boot parameter, the installation program does not automatically switch the installation source repository to the CDN after registration. If you want to use the CDN to install RHEL, you must manually switch the installation source repository to the CDN by selecting the **Red Hat CDN** option in the **Installation Source** window of the graphical installation. If you do not manually switch to the CDN, the installation program installs the packages from the repository specified on the kernel command line.



IMPORTANT

- You can switch the installation source repository to the CDN using the rhsm
 Kickstart command only if you do not specify an installation source using
 inst.repo= on the kernel command line or the url command in the Kickstart file.
 You must use inst.stage2=<URL> on the kernel command line to fetch the
 installation image, but not specify the installation source.
- An installation source URL specified using a boot option or included in a Kickstart file takes precedence over the CDN, even if the Kickstart file contains the **rhsm** command with valid credentials. The system is registered, but it is installed from the URL installation source. This ensures that earlier installation processes operate as normal.

3.3.3. Verifying your system registration from the CDN

Use this procedure to verify that your system is registered to the CDN using the GUI.



WARNING

You can only verify your registration from the CDN if you have **not** clicked the **Begin Installation** button from the **Installation Summary** window. Once the **Begin Installation** button is clicked, you cannot return to the Installation Summary window to verify your registration.

Prerequisite

• You have completed the registration process as documented in the Section 3.3.2, "Registering and installing RHEL from the CDN" and *Registered* is displayed under **Connect to Red Hat**on the **Installation Summary** window.

Procedure

1. From the Installation Summary window, select Connect to Red Hat

2. The window opens and displays a registration summary:

Method

The registered account name or activation keys are displayed.

System Purpose

If set, the role, SLA, and usage details are displayed.

Insights

If enabled, the Insights details are displayed.

Number of subscriptions

The number of subscriptions attached are displayed.

3. Verify that the registration summary matches the details that were entered.

3.3.4. Unregistering your system from the CDN

Use this procedure to unregister your system from the CDN using the GUI.



WARNING

- You can unregister from the CDN if you have not clicked the Begin
 Installation button from the Installation Summary window. Once the
 Begin Installation button is clicked, you cannot return to the Installation
 Summary window to unregister your registration.
- When unregistering, the installation program switches to the first available repository, in the following order:
 - a. The URL used in the inst.repo=<URL> boot parameter on the kernel command line.
 - b. An automatically detected repository on the installation media (USB or DVD).

Prerequisite

• You have completed the registration process as documented in the Section 3.3.2, "Registering and installing RHEL from the CDN" and *Registered* is displayed under **Connect to Red Hat** on the **Installation Summary** window.

Procedure

- 1. From the Installation Summary window, select Connect to Red Hat
- 2. The **Connect to Red Hat**window opens and displays a registration summary:

Method

The registered account name or activation keys used are displayed.

System Purpose

If set, the role, SLA, and usage details are displayed.

Insights

If enabled, the Insights details are displayed.

Number of subscriptions

The number of subscriptions attached are displayed.

- 3. Click **Unregister** to remove the registration from the CDN. The original registration details are displayed with a **Not registered** message displayed in the lower-middle part of the window.
- 4. Click **Done** to return to the **Installation Summary** window.
- 5. **Connect to Red Hat** displays a *Not registered* message, and **Software Selection** displays a *Red Hat CDN requires registration* message.



NOTE

After unregistering, it is possible to register your system again. Click **Connect to Red Hat** The previously entered details are populated. Edit the original details, or update the fields based on the account, purpose, and connection. Click **Register** to complete.

CHAPTER 4. CUSTOMIZING YOUR INSTALLATION

When installing Red Hat Enterprise Linux, you can customize location, software, and system settings and parameters, using the **Installation Summary** window.

The Installation Summary window contains three categories:

- LOCALIZATION: You can configure Keyboard, Language Support, and Time and Date.
- **SOFTWARE**: You can configure Installation Source and Software Selection.
- **SYSTEM**: You can configure Installation Destination, Connect to Red Hat, KDUMP, Network and Host Name, and Security Policy.

A category has a different status depending on where it is in the installation program.

Table 4.1. Category status

Category status	Status	Description
Warning symbol type 1	Yellow triangle with an exclamation mark and red text	Requires attention before installation. For example, Network & Host Name requires attention before you can register and download from the Content Delivery Network (CDN).
Warning symbol type 2	Grayed out and with a warning symbol (yellow triangle with an exclamation mark)	The installation program is configuring a category and you must wait for it to finish before accessing the window.



NOTE

A warning message is displayed at the bottom of the **Installation Summary** window and the **Begin Installation** button is disabled until you configure all of the required categories.

This section contains information about customizing your Red Hat Enterprise Linux installation using the Graphical User Interface (GUI). The GUI is the preferred method of installing Red Hat Enterprise Linux when you boot the system from a CD, DVD, or USB flash drive, or from a network using PXE.



NOTE

There may be some variance between the online help and the content that is published on the Customer Portal. For the latest updates, see the installation content on the Customer Portal.

4.1. CONFIGURING LANGUAGE AND LOCATION SETTINGS

The installation program uses the language that you selected during installation.

Prerequisites

- 1. You created installation media. See Section 2.7, "Creating a bootable installation medium" for more information.
- 2. You specified an installation source if you are using the Boot ISO image file. See Section 2.8, "Preparing an installation source" for more information.
- 3. You booted the installation. See Section 3.1, "Booting the installation" for more information.

Procedure

1. From the left-hand pane of the **Welcome to Red Hat Enterprise Linux**window, select a language. Alternatively, type your preferred language into the **Search** field.



NOTE

A language is pre-selected by default. If network access is configured, that is, if you booted from a network server instead of local media, the pre-selected language is determined by the automatic location detection feature of the **GeoIP** module. If you used the **inst.lang=** option on the boot command line or in your PXE server configuration, then the language that you define with the boot option is selected.

- 2. From the right-hand pane of the **Welcome to Red Hat Enterprise Linux**window, select a location specific to your region.
- 3. Click **Continue** to proceed to the Chapter 4, Customizing your installation window.



IMPORTANT

If you are installing a pre-release version of Red Hat Enterprise Linux, a warning message is displayed about the pre-release status of the installation media. Click **I want to proceed** to continue with the installation, or **I want to exit** to quit the installation and reboot the system.

Additional resources

For information about how to change language and location settings during the installation program, see Section 4.2, "Configuring localization options"

4.2. CONFIGURING LOCALIZATION OPTIONS

This section contains information about configuring your keyboard, language support, and time and date settings.



IMPORTANT

If you use a layout that cannot accept Latin characters, such as **Russian**, add the **English** (**United States**) layout and configure a keyboard combination to switch between the two layouts. If you select a layout that does not have Latin characters, you might be unable to enter a valid **root** password and user credentials later in the installation process. This might prevent you from completing the installation.

4.2.1. Configuring keyboard, language, and time and date settings



Keyboard, Language, and Time and Date Settings are configured by default as part of Section 4.1, "Configuring language and location settings". To change any of the settings, complete the following steps, otherwise proceed to Section 4.4, "Configuring software options".

Procedure: Configuring keyboard settings

- 1. From the **Installation Summary** window, click **Keyboard**. The default layout depends on the option selected in Section 4.1, "Configuring language and location settings".
 - a. Click + to open the Add a Keyboard Layout window and change to a different layout.
 - b. Select a layout by browsing the list or use the **Search** field.
 - c. Select the required layout and click **Add**. The new layout appears under the default layout.
 - d. Click **Options** to optionally configure a keyboard switch that you can use to cycle between available layouts. The **Layout Switching Options**window opens.
 - e. To configure key combinations for switching, select one or more key combinations and click **OK** to confirm your selection.



NOTE

When you select a layout, click the **Keyboard** button to open a new dialog box that displays a visual representation of the selected layout.

f. Click **Done** to apply the settings and return to Chapter 4, Customizing your installation.

Procedure: Configuring language settings

- 1. From the **Installation Summary** window, click **Language Support** The **Language Support** window opens. The left pane lists the available language groups. If at least one language from a group is configured, a check mark is displayed and the supported language is highlighted.
 - a. From the left pane, click a group to select additional languages, and from the right pane, select regional options. Repeat this process for languages that you require.
 - b. Click **Done** to apply the changes and return to Chapter 4, Customizing your installation.

Procedure: Configuring time and date settings

1. From the **Installation Summary** window, click **Time & Date**. The **Time & Date** window opens.



NOTE

The **Time & Date** settings are configured by default based on the settings you selected in Section 4.1, "Configuring language and location settings".

The list of cities and regions come from the Time Zone Database (**tzdata**) public domain that is maintained by the Internet Assigned Numbers Authority (IANA). Red Hat can not add cities or regions to this database. You can find more information at the IANA official website.

a. From the **Region** drop-down menu, select a region.



NOTE

Select **Etc** as your region to configure a time zone relative to Greenwich Mean Time (GMT) without setting your location to a specific region.

- b. From the **City** drop-down menu, select the city, or the city closest to your location in the same time zone.
- c. Toggle the **Network Time** switch to enable or disable network time synchronization using the Network Time Protocol (NTP).



NOTE

Enabling the Network Time switch keeps your system time correct as long as the system can access the internet. By default, one NTP pool is configured; you can add a new option, or disable or remove the default options by clicking the **gear wheel** button next to the **Network Time** switch.

d. Click **Done** to apply the changes and return to Chapter 4, Customizing your installation.



NOTE

If you disable network time synchronization, the controls at the bottom of the window become active, allowing you to set the time and date manually.

4.3. CONFIGURING SYSTEM OPTIONS

This section contains information about configuring Installation Destination, KDUMP, Network and Host Name, Security Policy, and System Purpose.

4.3.1. Configuring installation destination

Use the **Installation Destination** window to configure the storage options, for example, the disks that you want to use as the installation target for your Red Hat Enterprise Linux installation. You must select at least one disk.



WARNING

Back up your data if you plan to use a disk that already contains data. For example, if you want to shrink an existing Microsoft Windows partition and install Red Hat Enterprise Linux as a second system, or if you are upgrading a previous release of Red Hat Enterprise Linux. Manipulating partitions always carries a risk. For example, if the process is interrupted or fails for any reason data on the disk can be lost.



IMPORTANT

Special cases

- Some BIOS types do not support booting from a RAID card. In these instances, the /boot partition must be created on a partition outside of the RAID array, such as on a separate hard drive. It is necessary to use an internal hard drive for partition creation with problematic RAID cards. A /boot partition is also necessary for software RAID setups. If you choose to partition your system automatically, you should manually edit your /boot partition.
- To configure the Red Hat Enterprise Linux boot loader to chain load from a
 different boot loader, you must specify the boot drive manually by clicking the
 Full disk summary and bootloaderlink from the Installation Destination
 window.
- When you install Red Hat Enterprise Linux on a system with both multipath and non-multipath storage devices, the automatic partitioning layout in the installation program creates volume groups that contain a mix of multipath and non-multipath devices. This defeats the purpose of multipath storage. It is recommended that you select either multipath or non-multipath devices on the Installation Destination window. Alternatively, proceed to manual partitioning.

Prerequisite

The Installation Summary window is open.

Procedure

- From the Installation Summary window, click Installation Destination. The Installation
 Destination window opens.
 - a. From the Local Standard Disks section, select the storage device that you require; a white check mark indicates your selection. Disks without a white check mark are not used during the installation process; they are ignored if you choose automatic partitioning, and they are not available in manual partitioning.



NOTE

All locally available storage devices (SATA, IDE and SCSI hard drives, USB flash and external disks) are displayed under **Local Standard Disks**. Any storage devices connected after the installation program has started are not detected. If you use a removable drive to install Red Hat Enterprise Linux, your system is unusable if you remove the device.

b. Optional: Click the Refresh link in the lower right-hand side of the window if you want to configure additional local storage devices to connect new hard drives. The Rescan Disks dialog box opens.



NOTE

All storage changes that you make during the installation are lost when you click **Rescan Disks**.

i. Click **Rescan Disks** and wait until the scanning process completes.

- ii. Click **OK** to return to the **Installation Destination** window. All detected disks including any new ones are displayed under the **Local Standard Disks** section.
- Optional: To add a specialized storage device, click Add a disk....
 The Storage Device Selection window opens and lists all storage devices that the installation program has access to. For information about how to add a specialized disk, see Section 4.5.3, "Using advanced storage options".
- 3. Optional: Under **Storage Configuration**, select the **Automatic** radio button.



IMPORTANT

Automatic partitioning is the **recommended** method of partitioning your storage. You can also configure custom partitioning, for more details see Section 4.6, "Configuring manual partitioning"

- 4. Optional: To reclaim space from an existing partitioning layout, select the I would like to make additional space available check box. For example, if a disk you want to use already contains a different operating system and you want to make this system's partitions smaller to allow more room for Red Hat Enterprise Linux.
- 5. Optional: Select **Encrypt my data** to encrypt all partitions except the ones needed to boot the system (such as /boot) using *Linux Unified Key Setup* (LUKS). Encrypting your hard drive is recommended.
 - a. Click **Done**. The **Disk Encryption Passphrase** dialog box opens.
 - i. Type your passphrase in the Passphrase and Confirm fields.
 - ii. Click Save Passphrase to complete disk encryption.



WARNING

If you lose the LUKS passphrase, any encrypted partitions and their data is completely inaccessible. There is no way to recover a lost passphrase. However, if you perform a Kickstart installation, you can save encryption passphrases and create backup encryption passphrases during the installation. See the *Performing an advanced RHEL installation* document for information.

6. Optional: Click the **Full disk summary and bootloader** link in the lower left-hand side of the window to select which storage device contains the boot loader. For more information, see Section 4.3.1.1, "Configuring boot loader".



NOTE

In most cases it is sufficient to leave the boot loader in the default location. Some configurations, for example, systems that require chain loading from another boot loader require the boot drive to be specified manually.

7. Click Done.

a. If you selected automatic partitioning and I would like to make additional space available, or if there is not enough free space on your selected hard drives to install Red Hat Enterprise Linux, the Reclaim Disk Space dialog box opens when you click Done, and lists all configured disk devices and all partitions on those devices. The dialog box displays information about how much space the system needs for a minimal installation and how much space you have reclaimed.



WARNING

If you **delete** a partition, all data on that partition is lost. If you want to preserve your data, use the **Shrink** option, not the **Delete** option.

- b. Review the displayed list of available storage devices. The **Reclaimable Space** column shows how much space can be reclaimed from each entry.
- c. To reclaim space, select a disk or partition, and click either the **Delete** button to delete that partition, or all partitions on a selected disk, or click **Shrink** to use free space on a partition while preserving the existing data.



NOTE

Alternatively, you can click **Delete all**, this deletes all existing partitions on all disks and makes this space available to Red Hat Enterprise Linux. Existing data on all disks is lost.

d. Click **Reclaim space** to apply the changes and return to Chapter 4, *Customizing your installation*.



IMPORTANT

No disk changes are made until you click **Begin Installation** on the **Installation Summary** window. The **Reclaim Space** dialog only marks partitions for resizing or deletion; no action is performed.

4.3.1.1. Configuring boot loader

Red Hat Enterprise Linux uses GRand Unified Bootloader version 2 (**GRUB2**) as the boot loader for AMD64 and Intel 64, IBM Power Systems, and ARM. For IBM Z, the **zipl** boot loader is used.

The boot loader is the first program that runs when the system starts and is responsible for loading and transferring control to an operating system. **GRUB2** can boot any compatible operating system (including Microsoft Windows) and can also use chain loading to transfer control to other boot loaders for unsupported operating systems.



WARNING

Installing GRUB2 may overwrite your existing boot loader.

If an operating system is already installed, the Red Hat Enterprise Linux installation program attempts to automatically detect and configure the boot loader to start the other operating system. If the boot loader is not detected, you can manually configure any additional operating systems after you finish the installation.

If you are installing a Red Hat Enterprise Linux system with more than one disk, you might want to manually specify the disk where you want to install the boot loader.

Procedure

- 1. From the **Installation Destination** window, click the **Full disk summary and bootloader** link. The **Selected Disks** dialog box opens.
 - The boot loader is installed on the device of your choice, or on a UEFI system; the **EFI system** partition is created on the target device during guided partitioning.
- 2. To change the boot device, select a device from the list and click **Set as Boot Device**. You can set only one device as the boot device.
- 3. To disable a new boot loader installation, select the device currently marked for boot and click **Do not install boot loader**. This ensures **GRUB2** is not installed on any device.



WARNING

If you choose not to install a boot loader, you cannot boot the system directly and you must use another boot method, such as a standalone commercial boot loader application. Use this option only if you have another way to boot your system.

The boot loader may also require a special partition to be created, depending on if your system uses BIOS or UEFI firmware, or if the boot drive has a *GUID Partition Table* (GPT) or a **Master Boot Record** (MBR, also known as **msdos**) label. If you use automatic partitioning, the installation program creates the partition.

4.3.2. Configuring Kdump

Kdump is a kernel crash-dumping mechanism. In the event of a system crash, **Kdump** captures the contents of the system memory at the moment of failure. This captured memory can be analyzed to find the cause of the crash. If **Kdump** is enabled, it must have a small portion of the system's memory (RAM) reserved to itself. This reserved memory is not accessible to the main kernel.

Procedure

- 1. From the Installation Summary window, click Kdump. The Kdump window opens.
- 2. Select the **Enable kdump** check box.
- 3. Select either the **Automatic** or **Manual** memory reservation setting.
 - a. If you select Manual, enter the amount of memory (in megabytes) that you want to reserve in the Memory to be reserved field using the + and - buttons. The Usable System Memory readout below the reservation input field shows how much memory is accessible to your main system after reserving the amount of RAM that you select.
- 4. Click **Done** to apply the settings and return to Chapter 4, Customizing your installation.



The amount of memory that you reserve is determined by your system architecture (AMD64 and Intel 64 have different requirements than IBM Power) as well as the total amount of system memory. In most cases, automatic reservation is satisfactory.



IMPORTANT

Additional settings, such as the location where kernel crash dumps will be saved, can only be configured after the installation using either the **system-config-kdump** graphical interface, or manually in the /etc/kdump.conf configuration file.

4.3.3. Configuring network and host name options

Use the **Network and Host name** window to configure network interfaces. Options that you select here are available both during the installation for tasks such as downloading packages from a remote location, and on the installed system.

4.3.3.1. Configuring network and host name

Follow the steps in this procedure to configure your network and host name.

Procedure

- 1. From the Installation Summary window, click Network and Host Name*.
- 2. From the list in the left-hand pane, select an interface. The details are displayed in the right-hand pane.
- 3. Toggle the **ON/OFF** switch to enable or disable the selected interface.



NOTE

The installation program automatically detects locally accessible interfaces, and you cannot add or remove them manually.

- 4. Click + to add a virtual network interface, which can be either: Team, Bond, Bridge, or VLAN.
- 5. Click to remove a virtual interface.

- 6. Click **Configure** to change settings such as IP addresses, DNS servers, or routing configuration for an existing interface (both virtual and physical).
- 7. Type a host name for your system in the **Host Name** field.



- There are several types of network device naming standards used to identify network devices with persistent names, for example, em1 and wl3sp0. For information about these standards, see the Configuring and managing networking document.
- The host name can be either a fully-qualified domain name (FQDN) in the format hostname.domainname, or a short host name with no domain name. Many networks have a Dynamic Host Configuration Protocol (DHCP) service that automatically supplies connected systems with a domain name. To allow the DHCP service to assign the domain name to this machine, specify only the short host name. The value localhost.localdomain means that no specific static host name for the target system is configured, and the actual host name of the installed system is configured during the processing of the network configuration, for example, by NetworkManager using DHCP or DNS.
- 8. Click **Apply** to apply the host name to the environment.

4.3.3.2. Adding a virtual network interface

Follow the steps in this procedure to add a virtual network interface.

Procedure

- 1. From the **Network & Host name** window, click the + button to add a virtual network interface. The **Add a device** dialog opens.
- 2. Select one of the four available types of virtual interfaces:
 - **Bond**: NIC (*Network Interface Controller*) Bonding, a method to bind multiple physical network interfaces together into a single bonded channel.
 - Bridge: Represents NIC Bridging, a method to connect multiple separate networks into one aggregate network.
 - **Team**: NIC Teaming, a new implementation to aggregate links, designed to provide a small kernel driver to implement the fast handling of packet flows, and various applications to do everything else in user space.
 - Vlan (*Virtual LAN*): A method to create multiple distinct broadcast domains which are mutually isolated.
- 3. Select the interface type and click **Add**. An editing interface dialog box opens, allowing you to edit any available settings for your chosen interface type. For more information see Section 4.3.3.3, "Editing network interface configuration".
- 4. Click **Save** to confirm the virtual interface settings and return to the **Network & Host name** window.



If you need to change the settings of a virtual interface, select the interface and click **Configure**.

4.3.3.3. Editing network interface configuration

This section contains information about the most important settings for a typical wired connection used during installation. Configuration of other types of networks is broadly similar, although the specific configuration parameters might be different.



NOTE

On IBM Z, you cannot add a new connection as the network subchannels need to be grouped and set online beforehand, and this is currently done only in the booting phase.

Procedure

 To configure a network connection manually, select the interface from the Network and Host name window and click Configure.

An editing dialog specific to the selected interface opens.



NOTE

The options present depend on the connection type - the available options are slightly different depending on whether the connection type is a physical interface (wired or wireless network interface controller) or a virtual interface (Bond, Bridge, Team, or Vlan) that was previously configured in Section 4.3.3.2, "Adding a virtual network interface".

The following sections contain information about the three most common and useful options in the editing dialog:

4.3.3.4. Enabling or Disabling the Interface Connection

Follow the steps in this procedure to enable or disable an interface connection.

Procedure

- 1. Click the **General** tab.
- 2. Select the **Connect automatically with priority** check box to enable connection by default. Keep the default priority setting at **0**.



IMPORTANT

- When enabled on a wired connection, the system automatically connects
 during startup or reboot. On a wireless connection, the interface attempts to
 connect to any known wireless networks in range. For further information
 about NetworkManager, including the nm-connection-editor tool, see the
 Configuring and managing networking document.
- You can enable or disable all users on the system from connecting to this
 network using the All users may connect to this network option. If you
 disable this option, only root will be able to connect to this network.
- It is not possible to only allow a specific user other than **root** to use this
 interface, as no other users are created at this point during the installation. If
 you need a connection for a different user, you must configure it after the
 installation.
- 3. Click **Save** to apply the changes and return to the **Network and Host name** window.

4.3.3.5. Setting up Static IPv4 or IPv6 Settings

By default, both IPv4 and IPv6 are set to automatic configuration depending on current network settings. This means that addresses such as the local IP address, DNS address, and other settings are detected automatically when the interface connects to a network. In many cases, this is sufficient, but you can also provide static configuration in the **IPv4 Settings** and **IPv6 Settings** tabs. Complete the following steps to configure IPv4 or IPv6 settings:

Procedure

 To set static network configuration, navigate to one of the IPv Settings tabs and from the Method drop-down menu, select a method other than Automatic, for example, Manual. The Addresses pane is enabled.



NOTE

In the **IPv6 Settings** tab, you can also set the method to **Ignore** to disable IPv6 on this interface.

- 2. Click **Add** and enter your address settings.
- 3. Type the IP addresses in the **Additional DNS servers** field; it accepts one or more IP addresses of DNS servers, for example, **10.0.0.1,10.0.0.8**.
- 4. Select the Require IPvX addressing for this connection to complete check box.



NOTE

Select this option in the **IPv4 Settings** or **IPv6 Settings** tabs to allow this connection only if IPv4 or IPv6 was successful. If this option remains disabled for both IPv4 and IPv6, the interface is able to connect if configuration succeeds on either IP protocol.

5. Click **Save** to apply the changes and return to the **Network & Host name** window.

4.3.3.6. Configuring Routes

Complete the following steps to configure routes.

Procedure

- 1. In the IPv4 Settings and IPv6 Settings tabs, click **Routes** to configure routing settings for a specific IP protocol on an interface. An editing routes dialog specific to the interface opens.
- 2. Click Add to add a route.
- 3. Select the **Ignore automatically obtained routes** check box to configure at least one static route and to disable all routes not specifically configured.
- 4. Select the **Use this connection only for resources on its network**check box to prevent the connection from becoming the default route.



NOTE

This option can be selected even if you did not configure any static routes. This route is used only to access certain resources, such as intranet pages that require a local or VPN connection. Another (default) route is used for publicly available resources. Unlike the additional routes configured, this setting is transferred to the installed system. This option is useful only when you configure more than one interface

- 5. Click **OK** to save your settings and return to the editing routes dialog that is specific to the interface.
- 6. Click **Save** to apply the settings and return to the **Network and Host Name** window.

4.3.3.7. Additional resources

• To learn more about network configuration after installation, see the *Configuring and managing networking* document.

4.3.4. Configuring Connect to Red Hat

The Red Hat Content Delivery Network (CDN), available from cdn.redhat.com, is a geographically distributed series of static web servers that contain content and errata that is consumed by systems. The content can be consumed directly, such as using a system registered to Red Hat Subscription Management. The CDN is protected by x.509 certificate authentication to ensure that only valid users have access. When a system is registered to Red Hat Subscription Management, the attached subscriptions govern which subset of the CDN the system can access.

Registering and installing RHEL from the CDN provides the following benefits:

- The CDN installation method supports the Boot ISO and the Binary DVD ISO image files. However, the use of the smaller Boot ISO image file is recommended as it consumes less space than the larger Binary DVD ISO image file.
- The CDN uses the latest packages resulting in a fully up-to-date system right after installation. There is no requirement to install package updates immediately after installation as is often the case when using the Binary DVD ISO image file.

• Integrated support for connecting to Red Hat Insights and enabling System Purpose.

4.3.4.1. Introduction to System Purpose

System Purpose is an optional but recommended feature of the Red Hat Enterprise Linux installation. You use System Purpose to record the intended use of a Red Hat Enterprise Linux 8 system, and ensure that the entitlement server auto-attaches the most appropriate subscription to your system.

Benefits include:

- In-depth system-level information for system administrators and business operations.
- Reduced overhead when determining why a system was procured and its intended purpose.
- Improved customer experience of Subscription Manager auto-attach as well as automated discovery and reconciliation of system usage.

You can enter System Purpose data in one of the following ways:

- During image creation
- During a GUI installation when using **Connect to Red Hat** to register your system and attach your Red Hat subscription
- During a Kickstart installation when using Kickstart automation scripts
- After installation using the **syspurpose** command-line (CLI) tool

To record the intended purpose of your system, you can configure the following components of System Purpose. The selected values are used by the entitlement server upon registration to attach the most suitable subscription for your system.

Role

- Red Hat Enterprise Linux Server
- Red Hat Enterprise Linux Workstation
- Red Hat Enterprise Linux Compute Node

Service Level Agreement

- Premium
- Standard
- Self-Support

Usage

- Production
- Development/Test
- Disaster Recovery

Additional resources

- For more information about Image Builder, see the Composing a customized RHEL system image document.
- For more information about Kickstart, see the *Performing an advanced RHEL installation* document
- For more information about Subscription Manager, see the *Using and Configuring Red Hat Subscription Manager* document.

4.3.4.2. Configuring Connect to Red Hat options

Use the following procedure to configure the Connect to Red Hat options in the GUI.



NOTE

You can register to the CDN using either your Red Hat account or your activation key details.

Procedure

- 1. Click Account.
 - a. Enter your Red Hat Customer Portal username and password details.
- 2. Optional: Click **Activation Key**.
 - a. Enter your organization ID and activation key. You can enter more than one activation key, separated by a comma, as long as the activation keys are registered to your subscription.
- 3. Select the **Set System Purpose** check box. System Purpose enables the entitlement server to determine and automatically attach the most appropriate subscription to satisfy the intended use of your RHEL 8 system.
 - a. Select the required Role, SLA, and Usage from the corresponding drop-down lists.
- 4. The **Connect to Red Hat Insights**check box is enabled by default. Clear the check box if you do not want to connect to Red Hat Insights.



NOTE

Red Hat Insights is a Software-as-a-Service (SaaS) offering that provides continuous, in-depth analysis of registered Red Hat-based systems to proactively identify threats to security, performance and stability across physical, virtual and cloud environments, and container deployments.

- 5. Optional: Expand Options.
 - a. Select the **Use HTTP proxy** check box if your network environment only allows external Internet access or access to content servers through an HTTP proxy. Clear the **Use HTTP proxy** check box if an HTTP proxy is not used.
 - b. If you are running Satellite Server or performing internal testing, select the **Custom server URL** and **Custom base URL** check boxes and enter the required details.



IMPORTANT

- The Custom server URL field does not require the HTTP protocol, for example nameofhost.com. However, the Custom base URL field requires the HTTP protocol.
- To change the **Custom base URL** after registration, you must unregister, provide the new details, and then re-register.
- 6. Click **Register** to register the system. When the system is successfully registered and subscriptions are attached, the **Connect to Red Hat**window displays the attached subscription details.



NOTE

Depending on the amount of subscriptions, the registration and attachment process might take up to a minute to complete.

- 7. Click **Done** to return to the **Installation Summary** window.
 - a. A Registered message is displayed under Connect to Red Hat

4.3.4.3. Installation source repository after system registration

The installation source repository used after system registration is dependent on how the system was booted.

System booted from the Boot ISO or the Binary DVD ISO image file

If you booted the RHEL installation using either the **Boot ISO** or the **Binary DVD ISO** image file with the default boot parameters, the installation program automatically switches the installation source repository to the CDN after registration.

System booted with theinst.repo=<URL> boot parameter

If you booted the RHEL installation with the <code>inst.repo=<URL></code> boot parameter, the installation program does not automatically switch the installation source repository to the CDN after registration. If you want to use the CDN to install RHEL, you must manually switch the installation source repository to the CDN by selecting the <code>Red Hat CDN</code> option in the <code>Installation Source</code> window of the graphical installation. If you do not manually switch to the CDN, the installation program installs the packages from the repository specified on the kernel command line.



IMPORTANT

- You can switch the installation source repository to the CDN using the rhsm
 Kickstart command only if you do not specify an installation source using
 inst.repo= on the kernel command line or the url command in the Kickstart file.
 You must use inst.stage2=<URL> on the kernel command line to fetch the
 installation image, but not specify the installation source.
- An installation source URL specified using a boot option or included in a Kickstart file takes precedence over the CDN, even if the Kickstart file contains the **rhsm** command with valid credentials. The system is registered, but it is installed from the URL installation source. This ensures that earlier installation processes operate as normal.

4.3.4.4. Verifying your system registration from the CDN

Use this procedure to verify that your system is registered to the CDN using the GUI.



WARNING

You can only verify your registration from the CDN if you have **not** clicked the **Begin Installation** button from the **Installation Summary** window. Once the **Begin Installation** button is clicked, you cannot return to the Installation Summary window to verify your registration.

Prerequisite

 You have completed the registration process as documented in the Section 3.3.2, "Registering and installing RHEL from the CDN" and Registered is displayed under Connect to Red Hat on the Installation Summary window.

Procedure

- 1. From the Installation Summary window, select Connect to Red Hat
- 2. The window opens and displays a registration summary:

Method

The registered account name or activation keys are displayed.

System Purpose

If set, the role, SLA, and usage details are displayed.

Insights

If enabled, the Insights details are displayed.

Number of subscriptions

The number of subscriptions attached are displayed.

3. Verify that the registration summary matches the details that were entered.

4.3.4.5. Unregistering your system from the CDN

Use this procedure to unregister your system from the CDN using the GUI.



WARNING

- You can unregister from the CDN if you have not clicked the Begin Installation button from the Installation Summary window. Once the Begin Installation button is clicked, you cannot return to the Installation Summary window to unregister your registration.
- When unregistering, the installation program switches to the first available repository, in the following order:
 - a. The URL used in the inst.repo=<URL> boot parameter on the kernel command line.
 - b. An automatically detected repository on the installation media (USB or DVD).

Prerequisite

 You have completed the registration process as documented in the Section 3.3.2, "Registering and installing RHEL from the CDN" and Registered is displayed under Connect to Red Haton the Installation Summary window.

Procedure

- 1. From the Installation Summary window, select Connect to Red Hat
- 2. The **Connect to Red Hat** window opens and displays a registration summary:

Method

The registered account name or activation keys used are displayed.

System Purpose

If set, the role, SLA, and usage details are displayed.

Insights

If enabled, the Insights details are displayed.

Number of subscriptions

The number of subscriptions attached are displayed.

- 3. Click **Unregister** to remove the registration from the CDN. The original registration details are displayed with a **Not registered** message displayed in the lower-middle part of the window.
- 4. Click **Done** to return to the **Installation Summary** window.
- 5. **Connect to Red Hat** displays a *Not registered* message, and **Software Selection** displays a *Red Hat CDN requires registration* message.



After unregistering, it is possible to register your system again. Click **Connect to Red Hat** The previously entered details are populated. Edit the original details, or update the fields based on the account, purpose, and connection. Click **Register** to complete.

Related information

- For information about Red Hat Insights, see the Red Hat Insights product documentation.
- For information about Activation Keys, see the *Understanding Activation Keys* chapter of the *Using Red Hat Subscription Management* document.
- For information about how to set up an HTTP proxy for Subscription Manager, see the *Using an HTTP proxy* chapter of the *Using and Configuring Red Hat Subscription Manager* document.

4.3.5. Configuring Security Policy

This section contains information about the Red Hat Enterprise Linux 8 security policy and how to configure it for use on your system.

4.3.5.1. About security policy

The Red Hat Enterprise Linux security policy adheres to restrictions and recommendations (compliance policies) defined by the Security Content Automation Protocol (SCAP) standard. The packages are automatically installed. However, by default, no policies are enforced and therefore no checks are performed during or after installation unless specifically configured.

Applying a security policy is not a mandatory feature of the installation program. If you apply a security policy to the system, it is installed using restrictions and recommendations defined in the profile that you selected. The **openscap-scanner** package is added to your package selection, providing a preinstalled tool for compliance and vulnerability scanning. After the installation finishes, the system is automatically scanned to verify compliance. The results of this scan are saved to the /root/openscap_data directory on the installed system. You can also load additional profiles from an HTTP, HTTPS, or FTP server.

4.3.5.2. Configuring a security policy

Complete the following steps to configure a security policy.

Prerequisite

The **Installation Summary** window is open.

Procedure

- From the Installation Summary window, click Security Policy. The Security Policy window opens.
- 2. To enable security policies on the system, toggle the Apply security policy switch to ON.
- 3. Select one of the profiles listed in the top pane.
- 4. Click Select profile.

Profile changes that you must apply before installation appear in the bottom pane.



The default profiles do not require changes before installation. However, loading a custom profile can require pre-installation tasks.

- 5. Click **Change content** to use a custom profile. A separate window opens allowing you to enter a URL for valid security content.
 - a. Click **Fetch** to retrieve the URL.
 - b. Click **Use SCAP Security Guide** to return to the **Security Policy** window.



NOTE

You can load custom profiles from an HTTP, HTTPS, or FTP server. Use the full address of the content including the protocol, such as http://. A network connection must be active before you can load a custom profile. The installation program detects the content type automatically.

6. Click **Done** to apply the settings and return to the **Installation Summary** window.

4.3.5.3. Related information

- **scap-security-guide(8)** The manual page for the **scap-security-guide** project contains information about SCAP security profiles, including examples on how to utilize the provided benchmarks using the OpenSCAP utility.
- Red Hat Enterprise Linux security compliance information is available in the Security hardening document.

4.4. CONFIGURING SOFTWARE OPTIONS

This section contains information about configuring your installation source and software selection settings, and activating a repository.

4.4.1. Configuring installation source

Complete the steps in this procedure to configure an installation source from either auto-detected installation media, Red Hat CDN, or the network.



NOTE

When the **Installation Summary** window first opens, the installation program attempts to configure an installation source based on the type of media that was used to boot the system. The full Red Hat Enterprise Linux Server DVD configures the source as local media.

Prerequisites

- You have downloaded the Binary DVD ISO image as detailed in Section 2.6, "Downloading the installation ISO image".
- You have created bootable installation media as detailed in Section 2.7, "Creating a bootable installation medium".

• The **Installation Summary** window is open.

Procedure

- 1. From the **Installation Summary** window, click **Installation Source**. The **Installation Source** window opens.
 - a. Review the **Auto-detected installation media** section to verify the details. This option is selected by default if you started the installation program from media containing an installation source, for example, a DVD.
 - b. Click **Verify** to check the media integrity.
 - c. Review the **Additional repositories** section and note that the **AppStream** checkbox is selected by default.



IMPORTANT

- No additional configuration is necessary as the BaseOS and AppStream repositories are installed as part of the full installation image.
- Do not disable the AppStream repository check box if you want a full Red Hat Enterprise Linux 8 installation.
- 2. Optional: Select the **Red Hat CDN** option to register your system, attach RHEL subscriptions, and install RHEL from the Red Hat Content Delivery Network (CDN). For more information, see the *Registering and installing RHEL from the CDN* section.
- 3. Optional: Select the **On the network** option to download and install packages from a network location instead of local media.



NOTE

- If you do not want to download and install additional repositories from a network location, proceed to Section 4.4.2, "Configuring software selection".
- This option is available only when a network connection is active. See Section 4.3.3, "Configuring network and host name options" for information about how to configure network connections in the GUI.
- a. Select the **On the network** drop-down menu to specify the protocol for downloading packages. This setting depends on the server that you want to use.



WARNING

The AppStream repository check box is disabled if you select **On the network** and then decide to revert to **Auto-detected installation**. You must select the AppStream check box to enable the AppStream repository.

b. Type the server address (without the protocol) into the address field. If you choose NFS, a second input field opens where you can specify custom **NFS mount options**. This field accepts options listed in the **nfs(5)** man page.



IMPORTANT

When selecting an NFS installation source, you must specify the address with a colon (:) character separating the host name from the path. For example:



server.example.com:/path/to/directory



NOTE

The following steps are optional and are only required if you use a proxy for network access.

- c. Click **Proxy setup...** to configure a proxy for an HTTP or HTTPS source.
- d. Select the **Enable HTTP proxy** check box and type the URL into the **Proxy Host** field.
- e. Select the **Use Authentication** check box if the proxy server requires authentication.
- f. Type in your user name and password.
- g. Click **OK** to finish the configuration and exit the **Proxy Setup...** dialog box.



NOTE

If your HTTP or HTTPS URL refers to a repository mirror menu, select the required option from the **URL type** drop-down list. All environments and additional software packages are available for selection when you finish configuring the sources.

- 4. Click + to add a repository.
- 5. Click to delete a repository.
- 6. Click the **arrow** icon to revert the current entries to the setting when you opened the **Installation Source** window.
- 7. To activate or deactivate a repository, click the check box in the **Enabled** column for each entry in the list.



NOTE

You can name and configure your additional repository in the same way as the primary repository on the network.

8. Click **Done** to apply the settings and return to the **Installation Summary** window.

4.4.2. Configuring software selection

Use the **Software Selection** window to select the software packages that you require. The packages are organized by Base Environment and Additional Software.

- Base Environment contains predefined packages. You can select only one base environment, and availability is dependent on the installation ISO image that is used as the installation source.
- Additional Software for Selected Environment contains additional software packages for the base environment. You can select multiple software packages.

Use a predefined environment and additional software to customize your system. However, in a standard installation, you cannot select individual packages to install. To view the packages contained in a specific environment, see the *repository*/repodata/*-comps-repository.architecture.xml file on your installation source media (DVD, CD, USB). The XML file contains details of the packages installed as part of a base environment. Available environments are marked by the <environment> tag, and additional software packages are marked by the <group> tag.

If you are unsure about which packages to install, Red Hat recommends that you select the **Minimal Install** base environment. Minimal install installs a basic version of Red Hat Enterprise Linux with only a minimal amount of additional software. After the system finishes installing and you log in for the first time, you can use the **Yum package manager** to install additional software. For more information about Yum package manager, see the *Configuring basic system settings* document.



NOTE

- The **yum group list** command lists all package groups from yum repositories. See the *Configuring basic system settings* document for more information.
- If you need to control which packages are installed, you can use a Kickstart file and define the packages in the **%packages** section. See the *Performing an advanced RHEL installation* document for information about installing Red Hat Enterprise Linux using Kickstart.

Prerequisites

- You have configured the installation source.
- The installation program downloaded package metadata.
- The **Installation Summary** window is open.

Procedure

- 1. From the **Installation Summary** window, click **Software Selection**. The **Software Selection** window opens.
- 2. From the **Base Environment** pane, select a base environment. You can select only one base environment.



NOTE

The **Server with GUI** base environment is the default base environment and it launches the **Initial Setup** application after the installation completes and you restart the system.

3. From the Additional Software for Selected Environmentpane, select one or more options.

4. Click **Done** to apply the settings and return to Chapter 4, Customizing your installation.

4.5. CONFIGURING STORAGE DEVICES

You can install Red Hat Enterprise Linux on a large variety of storage devices. You can configure basic, locally accessible, storage devices in the **Installation Destination** window. Basic storage devices directly connected to the local system, such as hard disk drives and solid-state drives, are displayed in the **Local Standard Disks** section of the window. On IBM Z, this section contains activated Direct Access Storage Devices (DASDs).



WARNING

A known issue prevents DASDs configured as HyperPAV aliases from being automatically attached to the system after the installation is complete. These storage devices are available during the installation, but are not immediately accessible after you finish installing and reboot. To attach HyperPAV alias devices, add them manually to the /etc/dasd.conf configuration file of the system.

4.5.1. Storage device selection

The storage device selection window lists all storage devices that the installation program can access. Depending on your system and available hardware, some tabs might not be displayed. The devices are grouped under the following tabs:

Multipath Devices

Storage devices accessible through more than one path, such as through multiple SCSI controllers or Fiber Channel ports on the same system.



IMPORTANT

The installation program only detects multipath storage devices with serial numbers that are 16 or 32 characters long.

Other SAN Devices

Devices available on a Storage Area Network (SAN).

Firmware RAID

Storage devices attached to a firmware RAID controller.

NVDIMM Devices

Under specific circumstances, Red Hat Enterprise Linux 8 can boot and run from (NVDIMM) devices in sector mode on the Intel 64 and AMD64 architectures.

System z Devices

Storage devices, or Logical Units (LUNs), attached through the zSeries Linux FCP (Fiber Channel Protocol) driver.

4.5.2. Filtering storage devices

In the storage device selection window you can filter storage devices either by their World Wide Identifier (WWID) or by the port, target, or logical unit number (LUN).

Prerequisite

The Installation Summary window is open.

Procedure

- From the Installation Summary window, click Installation Destination. The Installation
 Destination window opens, listing all available drives.
- 2. Under the **Specialized & Network Disks**section, click **Add a disk...**. The storage devices selection window opens.
- 3. Click the **Search by** tab to search by port, target, LUN, or WWID.

 Searching by WWID or LUN requires additional values in the corresponding input text fields.
- 4. Select the option that you require from the **Search** drop-down menu.
- 5. Click **Find** to start the search. Each device is presented on a separate row with a corresponding check box.
- 6. Select the check box to enable the device that you require during the installation process. Later in the installation process you can choose to install Red Hat Enterprise Linux on any of the selected devices, and you can choose to mount any of the other selected devices as part of the installed system automatically.



NOTE

- Selected devices are not automatically erased by the installation process and selecting a device does not put the data stored on the device at risk.
- You can add devices to the system after installation by modifying the /etc/fstab file.
- 7. Click **Done** to return to the **Installation Destination** window.



IMPORTANT

Any storage devices that you do not select are hidden from the installation program entirely. To chain load the boot loader from a different boot loader, select all the devices present.

4.5.3. Using advanced storage options

To use an advanced storage device, you can configure an iSCSI (SCSI over TCP/IP) target or FCoE (Fibre Channel over Ethernet) SAN (Storage Area Network).

To use iSCSI storage devices for the installation, the installation program must be able to discover them as iSCSI targets and be able to create an iSCSI session to access them. Each of these steps might require a user name and password for Challenge Handshake Authentication Protocol (CHAP) authentication. Additionally, you can configure an iSCSI target to authenticate the iSCSI initiator on the system to which the target is attached (reverse CHAP), both for discovery and for the session. Used

together, CHAP and reverse CHAP are called mutual CHAP or two-way CHAP. Mutual CHAP provides the greatest level of security for iSCSI connections, particularly if the user name and password are different for CHAP authentication and reverse CHAP authentication.



NOTE

Repeat the iSCSI discovery and iSCSI login steps to add all required iSCSI storage. You cannot change the name of the iSCSI initiator after you attempt discovery for the first time. To change the iSCSI initiator name, you must restart the installation.

4.5.3.1. Discovering and starting an iSCSI session

Complete the following steps to discover and start an iSCSI session.

Prerequisites

• The **Installation Summary** window is open.

Procedure

- 1. From the **Installation Summary** window, click **Installation Destination**. The **Installation Destination** window opens, listing all available drives.
- 2. Under the **Specialized & Network Disks**section, click **Add a disk...**. The storage devices selection window opens.
- 3. Click **Add iSCSI target...**. The **Add iSCSI Storage Target** window opens.



IMPORTANT

You cannot place the /boot partition on iSCSI targets that you have manually added using this method - an iSCSI target containing a /boot partition must be configured for use with iBFT. However, in instances where the installed system is expected to boot from iSCSI with iBFT configuration provided by a method other than firmware iBFT, for example using iPXE, you can remove the /boot partition restriction using the inst.nonibftiscsiboot installer boot option.

- 4. Enter the IP address of the iSCSI target in the **Target IP Address** field.
- 5. Type a name in the **iSCSI Initiator Name** field for the iSCSI initiator in iSCSI qualified name (IQN) format. A valid IQN entry contains the following information:
 - The string **ign.** (note the period).
 - A date code that specifies the year and month in which your organization's Internet domain
 or subdomain name was registered, represented as four digits for the year, a dash, and two
 digits for the month, followed by a period. For example, represent September 2010 as 201009.
 - Your organization's Internet domain or subdomain name, presented in reverse order with the top-level domain first. For example, represent the subdomain **storage.example.com** as **com.example.storage**.
 - A colon followed by a string that uniquely identifies this particular iSCSI initiator within your domain or subdomain. For example, :diskarrays-sn-a8675309.

A complete IQN is as follows: **iqn.2010-09.storage.example.com:diskarrays-sn-a8675309**. The installation program prepopulates the **iSCSI Initiator Name** field with a name in this format to help you with the structure. For more information about IQNs, see 3.2.6. iSCSI Names in RFC 3720 - Internet Small Computer Systems Interface (iSCSI) available from tools.ietf.org and 1. iSCSI Names and Addresses in RFC 3721 - Internet Small Computer Systems Interface (iSCSI) Naming and Discovery available from tools.ietf.org.

- 6. Select the **Discovery Authentication Type** drop-down menu to specify the type of authentication to use for iSCSI discovery. The following options are available:
 - No credentials
 - CHAP pair
 - CHAP pair and a reverse pair
- 7. a. If you selected **CHAP pair** as the authentication type, enter the user name and password for the iSCSI target in the **CHAP Username** and **CHAP Password** fields.
 - b. If you selected **CHAP pair and a reverse pair** as the authentication type, enter the user name and password for the iSCSI target in the **CHAP Username** and **CHAP Password** field, and the user name and password for the iSCSI initiator in the **Reverse CHAP Username** and **Reverse CHAP Password** fields.
- 8. Optionally, select the **Bind targets to network interfaces** check box.
- 9. Click Start Discovery.

The installation program attempts to discover an iSCSI target based on the information provided. If discovery succeeds, the **Add iSCSI Storage Target** window displays a list of all iSCSI nodes discovered on the target.

10. Select the check boxes for the node that you want to use for installation.



NOTE

The **Node login authentication type** menu contains the same options as the **Discovery Authentication Type** menu. However, if you need credentials for discovery authentication, use the same credentials to log in to a discovered node.

- 11. Click the additional **Use the credentials from discovery** drop-down menu. When you provide the proper credentials, the **Log In** button becomes available.
- 12. Click **Log In** to initiate an iSCSI session.

4.5.3.2. Configuring FCoE parameters

Complete the following steps to configure FCoE parameters.

Prerequisite

The **Installation Summary** window is open.

Procedure

1. From the **Installation Summary** window, click **Installation Destination**. The **Installation Destination** window opens, listing all available drives.

- 2. Under the **Specialized & Network Disks**section, click **Add a disk...**. The storage devices selection window opens.
- 3. Click **Add FCoE SAN...**. A dialog box opens for you to configure network interfaces for discovering FCoE storage devices.
- 4. Select a network interface that is connected to an FCoE switch in the **NIC** drop-down menu.
- 5. Click **Add FCoE disk(s)** to scan the network for SAN devices.
- 6. Select the required check boxes:
 - Use DCB:Data Center Bridging (DCB) is a set of enhancements to the Ethernet protocols designed to increase the efficiency of Ethernet connections in storage networks and clusters. Select the check box to enable or disable the installation program's awareness of DCB. Enable this option only for network interfaces that require a host-based DCBX client. For configurations on interfaces that use a hardware DCBX client, disable the check box.
 - Use auto vlan: Auto VLAN is enabled by default and indicates whether VLAN discovery should be performed. If this check box is enabled, then the FIP (FCoE Initiation Protocol) VLAN discovery protocol runs on the Ethernet interface when the link configuration has been validated. If they are not already configured, network interfaces for any discovered FCoE VLANs are automatically created and FCoE instances are created on the VLAN interfaces.
- 7. Discovered FCoE devices are displayed under the **Other SAN Devices** tab in the **Installation Destination** window.

4.5.3.3. Configuring DASD storage devices

Complete the following steps to configure DASD storage devices.

Prerequisite

The **Installation Summary** window is open.

- 1. From the Installation Summary window, click Installation Destination. The Installation Destination window opens, listing all available drives.
- 2. Under the **Specialized & Network Disks**section, click **Add a disk...**. The storage devices selection window opens.
- 3. Click **Add DASD**. The **Add DASD Storage Target** dialog box opens and prompts you to specify a device number, such as **0.0.0204**, and attach additional DASDs that were not detected when the installation started.
- 4. Type the device number of the DASD that you want to attach in the **Device number** field.
- 5. Click Start Discovery.



- If a DASD with the specified device number is found and if it is not already attached, the dialog box closes and the newly-discovered drives appear in the list of drives. You can then select the check boxes for the required devices and click Done. The new DASDs are available for selection, marked as DASD device 0.0.xxxx in the Local Standard Disks section of the Installation Destination window.
- If you entered an invalid device number, or if the DASD with the specified device number is already attached to the system, an error message appears in the dialog box, explaining the error and prompting you to try again with a different device number.

4.5.3.4. Configuring FCP devices

FCP devices enable IBM Z to use SCSI devices rather than, or in addition to, Direct Access Storage Device (DASD) devices. FCP devices provide a switched fabric topology that enables IBM Z systems to use SCSI LUNs as disk devices in addition to traditional DASD devices.

Prerequisites

- The **Installation Summary** window is open.
- For an FCP-only installation, remove the **DASD=** option from the CMS configuration file or the **rd.dasd=** option from the parameter file to indicate that no DASD is present.

Procedure

- 1. From the **Installation Summary** window, click **Installation Destination**. The **Installation Destination** window opens, listing all available drives.
- 2. Under the **Specialized & Network Disks**section, click **Add a disk...**. The storage devices selection window opens.
- 3. Click Add ZFCP LUN. The Add zFCP Storage Target dialog box opens allowing you to add a FCP (Fibre Channel Protocol) storage device.
 IBM Z requires that you enter any FCP device manually so that the installation program can activate FCP LUNs. You can enter FCP devices either in the graphical installation, or as a unique parameter entry in the parameter or CMS configuration file. The values that you enter must be unique to each site that you configure.
- 4. Type the 4 digit hexadecimal device number in the **Device number** field.
- 5. Type the 16 digit hexadecimal World Wide Port Number (WWPN) in the **WWPN** field.
- 6. Type the 16 digit hexadecimal FCP LUN identifier in the **LUN** field.
- 7. Click **Start Discovery** to connect to the FCP device.

The newly-added devices are displayed in the **System z Devices** tab of the **Installation Destination** window.



- Interactive creation of an FCP device is only possible in graphical mode. It is not possible to configure an FCP device interactively in text mode installation.
- Use only lower-case letters in hex values. If you enter an incorrect value and click Start Discovery, the installation program displays a warning. You can edit the configuration information and retry the discovery attempt.
- For more information about these values, consult the hardware documentation and check with your system administrator.

4.5.4. Installing to an NVDIMM device

Non-Volatile Dual In-line Memory Module (NVDIMM) devices combine the performance of RAM with disk-like data persistence when no power is supplied. Under specific circumstances, Red Hat Enterprise Linux 8 can boot and run from NVDIMM devices.

4.5.4.1. Criteria for using an NVDIMM device as an installation target

You can install Red Hat Enterprise Linux 8 to Non-Volatile Dual In-line Memory Module (NVDIMM) devices in sector mode on the Intel 64 and AMD64 architectures, supported by the **nd_pmem** driver.

Conditions for using an NVDIMM device as storage

To use an NVDIMM device as storage, the following conditions must be satisfied:

- The architecture of the system is Intel 64 or AMD64.
- The NVDIMM device is configured to sector mode. The installation program can reconfigure NVDIMM devices to this mode.
- The NVDIMM device must be supported by the **nd_pmem** driver.

Conditions for booting from an NVDIMM Device

Booting from an NVDIMM device is possible under the following conditions:

- All conditions for using the NVDIMM device as storage are satisfied.
- The system uses UEFI.
- The NVDIMM device must be supported by firmware available on the system, or by an UEFI driver. The UEFI driver may be loaded from an option ROM of the device itself.
- The NVDIMM device must be made available under a namespace.

Utilize the high performance of NVDIMM devices during booting, place the /boot and /boot/efi directories on the device. The Execute-in-place (XIP) feature of NVDIMM devices is not supported during booting and the kernel is loaded into conventional memory.

4.5.4.2. Configuring an NVDIMM device using the graphical installation mode

A Non-Volatile Dual In-line Memory Module (NVDIMM) device must be properly configured for use by Red Hat Enterprise Linux 8 using the graphical installation.



WARNING

Reconfiguration of a NVDIMM device process destroys any data stored on the device.

Prerequisites

- A NVDIMM device is present on the system and satisfies all the other conditions for usage as an installation target.
- The installation has booted and the **Installation Summary** window is open.

Procedure

- From the Installation Summary window, click Installation Destination. The Installation
 Destination window opens, listing all available drives.
- 2. Under the **Specialized & Network Disks**section, click **Add a disk...**. The storage devices selection window opens.
- 3. Click the NVDIMM Devices tab.
- 4. To reconfigure a device, select it from the list. If a device is not listed, it is not in sector mode.
- 5. Click **Reconfigure NVDIMM...**. A reconfiguration dialog opens.
- 6. Enter the sector size that you require and click **Start Reconfiguration**. The supported sector sizes are 512 and 4096 bytes.
- 7. When reconfiguration completes click **OK**.
- 8. Select the device check box.
- Click **Done** to return to the **Installation Destination** window.
 The NVDIMM device that you reconfigured is displayed in the **Specialized & Network Disks** section.
- 10. Click **Done** to return to the **Installation Summary** window.

The NVDIMM device is now available for you to select as an installation target. Additionally, if the device meets the requirements for booting, you can set the device as a boot device.

4.6. CONFIGURING MANUAL PARTITIONING

You can use manual partitioning to configure your disk partitions and mount points and define the file system that Red Hat Enterprise Linux is installed on.



Before installation, you should consider whether you want to use partitioned or unpartitioned disk devices. For more information, see the Knowledgebase article at https://access.redhat.com/solutions/163853.

An installation of Red Hat Enterprise Linux requires a minimum of one partition but Red Hat recommends using at least the following partitions or volumes: **PReP**, /, /home, /boot, and swap. You can also create additional partitions and volumes as you require.



NOTE

An installation of Red Hat Enterprise Linux on IBM Power Systems servers requires a **PReP** boot partition.



WARNING

To prevent data loss it is recommended that you back up your data before proceeding. If you are upgrading or creating a dual-boot system, you should back up any data you want to keep on your storage devices.

4.6.1. Starting manual partitioning

Prerequisites

- The Installation Summary screen is currently displayed.
- All disks are available to the installation program.

- 1. Select disks for installation:
 - a. Click **Installation Destination** to open the **Installation Destination** window.
 - b. Select the disks that you require for installation by clicking the corresponding icon. A selected disk has a check-mark displayed on it.
 - c. Under **Storage Configuration**, select the **Custom** radio-button.
 - d. Optional: To enable storage encryption with LUKS, select the **Encrypt my data** check box.
 - e. Click Done.
- 2. If you selected to encrypt the storage, a dialog box for entering a disk encryption passphrase opens. Type in the LUKS passphrase:
 - a. Enter the passphrase in the two text fields. To switch keyboard layout, use the keyboard icon.



WARNING

In the dialog box for entering the passphrase, you cannot change the keyboard layout. Select the English keyboard layout to enter the passphrase in the installation program.

- b. Click **Save Passphrase**. The **Manual Partitioning** window opens.
- 3. Deleted Mount points are listed in the left-hand pane. The mount points are organized by detected operating system installations. As a result, some file systems may be displayed multiple times if a partition is shared among several installations.
 - a. Select the mount points in the left pane; the options that can be customized are displayed in the right pane.



NOTE

- If your system contains existing file systems, ensure that enough space is available for the installation. To remove any partitions, select them in the list and click the button.
 - The dialog has a check box that you can use to remove all other partitions used by the system to which the deleted partition belongs.
- If there are no existing partitions and you want to create the
 recommended set of partitions as a starting point, select your preferred
 partitioning scheme from the left pane (default for Red Hat
 Enterprise Linux is LVM) and click the Click here to create them
 automatically link.

A /boot partition, a / (root) volume, and a **swap** volume proportionate to the size of the available storage are created and listed in the left pane. These are the recommended file systems for a typical installation, but you can add additional file systems and mount points.

b. Click **Done** to confirm any changes and return to the **Installation Summary** window.

Continue with adding mount points, configuring the individual mount points, and configuring the underlying partitions or volumes.

4.6.2. Adding a mount point file system

Complete the following steps to add multiple mount point file systems.

Prerequisites

- Plan for your partitions:
 - To avoid problems with space allocation, first create small partitions with known fixed sizes, such as /boot, and then create the remaining partitions, letting the installation program allocate the remaining capacity to them.

• If you want to install the system on multiple disks, or if your disks differ in size and a particular partition must be created on the first disk detected by BIOS, then create these partitions first.

Procedure

- 1. Click + to create a new mount point file system. The Add a New Mount Point dialog opens.
- 2. Select one of the preset paths from the **Mount Point** drop-down menu or type your own; for example, select / for the root partition or /**boot** for the boot partition.
- 3. Enter the size of the file system in to the **Desired Capacity** field; for example, **2GiB**.



WARNING

If you do not specify a value in the Desired Capacity field, or if you specify a size bigger than available space, then all remaining free space is used.

4. Click **Add mount point** to create the partition and return to the **Manual Partitioning** window.

4.6.3. Configuring a mount point file system

This procedure describes how to set the partitioning scheme for each mount point that was created manually. The available options are **Standard Partition**, **LVM**, and **LVM Thin Provisioning**.



NOTE

- Btfrs support has been removed in Red Hat Enterprise Linux 8.
- The /boot partition is always located on a standard partition, regardless of the value selected.

Procedure

- 1. To change the devices that a single non-LVM mount point should be located on, select the required mount point from the left-hand pane.
- 2. Under the Device(s) heading, click Modify.... The Configure Mount Point dialog opens.
- 3. Select one or more devices and click **Select** to confirm your selection and return to the **Manual Partitioning** window.
- 4. Click **Update Settings** to apply the changes.



NOTE

Click the **Rescan** button (circular arrow button) to refresh all local disks and partitions; this is only required after performing advanced partition configuration outside the installation program. Clicking the **Rescan Disks** button resets all configuration changes made in the installation program.

5. In the lower left-hand side of the **Manual Partitioning** window, click the **storage device selected** link to open the **Selected Disks** dialog and review disk information.

4.6.4. Customizing a partition or volume

You can customize a partition or volume if you want to set specific settings.



IMPORTANT


If /usr or /var is partitioned separately from the rest of the root volume, the boot process becomes much more complex as these directories contain critical components. In some situations, such as when these directories are placed on an iSCSI drive or an FCoE location, the system is unable to boot, or hangs with a Device is busy error when powering off or rebooting.

This limitation only applies to /usr or /var, not to directories below them. For example, a separate partition for /var/www works successfully.

Procedure

1. From the left pane, select the mount point.

Figure 4.1. Customizing Partitions



- 2. From the right-hand pane, you can customize the following options:
 - a. Enter the file system mount point into the **Mount Point** field. For example, if a file system is the root file system, enter /; enter /**boot** for the /**boot** file system, and so on. For a swap file system, do not set the mount point as setting the file system type to **swap** is sufficient.

- b. Enter the size of the file system in the **Desired Capacity** field. You can use common size units such as KiB or GiB. The default is MiB if you do not set any other unit.
- c. Select the device type that you require from the drop-down **Device Type** menu: **Standard Partition**, **LVM**, or **LVM Thin Provisioning**.



WARNING

The installation program does not support overprovisioned LVM thin pools.



NOTE

RAID is available only if two or more disks are selected for partitioning. If you choose **RAID**, you can also set the **RAID Level**. Similarly, if you select **LVM**, you can specify the **Volume Group**.

- d. Select the **Encrypt** check box to encrypt the partition or volume. You must set a password later in the installation program. The **LUKS Version** drop-down menu is displayed.
- e. Select the LUKS version that you require from the drop-down menu.
- f. Select the appropriate file system type for this partition or volume from the **File system** drop-down menu.



NOTE

Support for **VFAT** file system is not available for Linux system partitions. For example, /, /var, /usr, and so on.

- g. Select the **Reformat** check box to format an existing partition, or clear the **Reformat** check box to retain your data. The newly-created partitions and volumes must be reformatted, and the check box cannot be cleared.
- h. Type a label for the partition in the **Label** field. Use labels to easily recognize and address individual partitions.
- i. Type a name in the **Name** field.



NOTE

Note that standard partitions are named automatically when they are created and you cannot edit the names of standard partitions. For example, you cannot edit the /boot name sda1.

3. Click **Update Settings** to apply your changes and if required, select another partition to customize. Changes are not applied until you click **Begin Installation** from the **Installation Summary** window.



Click **Reset All** to discard your partition changes.

- 4. Click **Done** when you have created and customized all file systems and mount points. If you choose to encrypt a file system, you are prompted to create a passphrase.
 A **Summary of Changes** dialog box opens, displaying a summary of all storage actions for the installation program.
- 5. Click **Accept Changes** to apply the changes and return to the **Installation Summary** window.

4.6.5. Preserving the /home directory

In a RHEL 8 graphical installation, you can preserve the /**home** directory that was used on your RHEL 7 system.



WARNING

Preserving /home is only possible if the /home directory is located on a separate /home partition on your RHEL 7 system.

Preserving the /home directory that includes various configuration settings, makes it possible that the GNOME Shell environment on the new RHEL 8 system is set in the same way as it was on your RHEL 7 system. Note that this applies only for users on RHEL 8 with the same user name and ID as on the previous RHEL 7 system.

Complete this procedure to preserve the /home directory from your RHEL 7 system.

Prerequisites

- RHEL 7 system is installed on your computer.
- The /home directory is located on a separate /home partition on your RHEL 7 system.
- The RHEL 8 **Installation Summary** window is currently displayed.

- 1. Click Installation Destination to open the Installation Destination window.
- 2. Under **Storage Configuration**, select the **Custom** radio button. Click **Done**.
- 3. Click **Done**, the **Manual Partitioning** window opens.
- 4. Choose the /home partition, fill in /home under Mount Point: and clear the Reformat check box.

MANUAL PARTITIONING RED HAT ENTERPRISE LINUX 8.1 INSTALLATION rhel-home ▼ New Red Hat Enterprise Linux 8.1 Installation 1ount Point Device(s): /home 1000 MiB > 0x1af4 (vda) /home Modify. 4.88 GiB 1000 MiB rhel-root /boot 1024 MiB vda1 Device Type: Volume Group swap 820 MiB LVM ▼ Encrypt (4 MiB free) ▼ rhel rhel-swap File System: Modify... Red Hat Enterprise Linux Workstation 7.7 for x86_64 xfs DATA /home 1000 MiB rhel-home Label: Name home Update Settings Note: The settings you make on this screen will not be applied until you click on the main menu's 'Begin G Installation' button. 340.97 MiB 8 GiB 1 storage device selected Reset All

Figure 4.2. Ensuring that /home is not formatted

- 5. Optional: You can also customize various aspects of the /home partition required for your RHEL 8 system as described in Section 4.6.4, "Customizing a partition or volume". However, to preserve /home from your RHEL 7 system, it is necessary to clear the Reformat check box.
- 6. After you customized all partitions according to your requirements, click **Done**. The **Summary of changes** dialog box opens.
- 7. Verify that the **Summary of changes** dialog box does not show any change for **/home**. This means that the **/home** partition is preserved.
- 8. Click **Accept Changes** to apply the changes, and return to the **Installation Summary** window.

4.6.6. Creating software RAID

Follow the steps in this procedure to create a Redundant Arrays of Independent Disks (RAID) device. RAID devices are constructed from multiple storage devices that are arranged to provide increased performance and, in some configurations, greater fault tolerance.

A RAID device is created in one step and disks are added or removed as necessary. You can configure one RAID partition for each physical disk in your system, so the number of disks available to the installation program determines the levels of RAID device available. For example, if your system has two hard drives, you cannot create a RAID 10 device, as it requires a minimum of three separate disks.



NOTE

On IBM Z, the storage subsystem uses RAID transparently. You do not have to configure software RAID manually.

Prerequisites

- You have selected two or more disks for installation before RAID configuration options are visible. At least two disks are required to create a RAID device.
- You have created a mount point. By configuring a mount point, you configure the RAID device.
- You have selected the **Custom** radio button on the **Installation Destination** window.

Procedure

- 1. From the left pane of the Manual Partitioning window, select the required partition.
- 2. Under the **Device(s)** section, click **Modify**. The **Configure Mount Point** dialog box opens.
- 3. Select the disks that you want to include in the RAID device and click **Select**.
- 4. Click the **Device Type** drop-down menu and select **RAID**.
- 5. Click the **File System** drop-down menu and select your preferred file system type.
- 6. Click the RAID Level drop-down menu and select your preferred level of RAID.
- 7. Click **Update Settings** to save your changes.
- 8. Click **Done** to apply the settings and return to the **Installation Summary** window.

A message is displayed at the bottom of the window if the specified RAID level requires more disks.

4.6.7. Creating an LVM logical volume

Logical Volume Management (LVM) presents a simple logical view of underlying physical storage space, such as hard drives or LUNs. Partitions on physical storage are represented as physical volumes that you can group together into volume groups. You can divide each volume group into multiple logical volumes, each of which is analogous to a standard disk partition. Therefore, LVM logical volumes function as partitions that can span multiple physical disks.



NOTE

LVM configuration is available only in the graphical installation program.



IMPORTANT

During text-mode installation, LVM configuration is not available. To create an LVM configuration, press **Ctrl+Alt+F2** to use a different virtual console, and run the **lvm** command. To return to the text-mode installation, press **Ctrl+Alt+F1**.

- 1. From the left-hand pane of the Manual Partitioning window, select the mount point.
- 2. Click the **Device Type** drop-down menu and select **LVM**. The **Volume Group** drop-down menu is displayed with the newly-created volume group name.



You cannot specify the size of the volume group's physical extents in the configuration dialog. The size is always set to the default value of 4 MiB. If you want to create a volume group with different physical extents, you must create it manually by switching to an interactive shell and using the **vgcreate** command, or use a Kickstart file with the **volgroup --pesize=**size command. See the Performing an advanced RHEL installation document for more information about Kickstart.

Additional resources

• For more information about LVM, see the Configuring and managing logical volumes document.

4.6.8. Configuring an LVM logical volume

Follow the steps in this procedure to configure a newly-created LVM logical volume.



WARNING

Placing the /boot partition on an LVM volume is not supported.

Procedure

- 1. From the left-hand pane of the Manual Partitioning window, select the mount point.
- 2. Click the **Device Type** drop-down menu and select **LVM**. The **Volume Group** drop-down menu is displayed with the newly-created volume group name.
- 3. Click **Modify** to configure the newly-created volume group. The **Configure Volume Group** dialog box opens.



NOTE

You cannot specify the size of the volume group's physical extents in the configuration dialog. The size is always set to the default value of 4 MiB. If you want to create a volume group with different physical extents, you must create it manually by switching to an interactive shell and using the **vgcreate** command, or use a Kickstart file with the **volgroup --pesize=**size command. See the Performing an advanced RHEL installation document for more information about Kickstart.

- 4. From the **RAID Level** drop-down menu, select the RAID level that you require. The available RAID levels are the same as with actual RAID devices.
- 5. Select the **Encrypt** check box to mark the volume group for encryption.
- 6. From the **Size policy** drop-down menu, select the size policy for the volume group. The available policy options are:

- Automatic: The size of the volume group is set automatically so that it is large enough to contain the configured logical volumes. This is optimal if you do not need free space within the volume group.
- As large as possible The volume group is created with maximum size, regardless of the size
 of the configured logical volumes it contains. This is optimal if you plan to keep most of your
 data on LVM and later need to increase the size of some existing logical volumes, or if you
 need to create additional logical volumes within this group.
- **Fixed**: You can set an exact size of the volume group. Any configured logical volumes must then fit within this fixed size. This is useful if you know exactly how large you need the volume group to be.
- 7. Click **Save** to apply the settings and return to the **Manual Partitioning** window.
- 8. Click **Update Settings** to save your changes.
- 9. Click **Done** to return to the **Installation Summary** window.

CHAPTER 5. STARTING AND COMPLETING THE INSTALLATION

After you have completed the tasks in the **Installation Summary** window, you can start the installation process.



NOTE

When the installation process has started, it is not possible to return to the Installation Summary window to change any settings. To change any settings, wait for the installation process to finish, reboot your system, log in, and then change your settings on the installed system.

5.1. STARTING THE INSTALLATION

Follow the steps in this procedure to start the installation from the **Installation Summary** window.

Procedure

• From the **Installation Summary** window, click **Begin Installation**. The **Configuration** window opens and the installation process starts.

Two user setting options, **Root Password** (mandatory) and **User Creation** (optional) are available. After you start the installation program, configure the user settings.

5.2. CONFIGURING A ROOT PASSWORD

You must configure a **root** password to finish the installation process and to log in to the administrator (also known as superuser or root) account that is used for system administration tasks. These tasks include installing and updating software packages and changing system-wide configuration such as network and firewall settings, storage options, and adding or modifying users, groups and file permissions.



IMPORTANT

- Use one or both of the following ways to gain root privileges to the installed system:
 - Use a root account.
 - Create a user account with administrative privileges (member of the wheel group). The **root** account is always created during the installation. Switch to the administrator account only when you need to perform a task that requires administrator access.



WARNING

The **root** account has complete control over the system. If unauthorized personnel gain access to the account, they can access or delete users' personal files.

Procedure

- 1. From the Configuration window, click Root Password. The Root Password window opens.
- Type your password in the Root Password field.
 The requirements and recommendations for creating a strong root password are:
 - Must be at least eight characters long
 - May contain numbers, letters (upper and lower case) and symbols
 - Is case-sensitive
- 3. Type the same password in the **Confirm** field.
- 4. Click **Done** to confirm your root password and return to Chapter 5, *Starting and completing the installation*.



NOTE

If you proceeded with a weak password, you must click **Done** twice.

5.3. CREATING A USER ACCOUNT

It is recommended that you create a user account to finish the installation. If you do not create a user account, you must log in to the system as **root** directly, which is **not** recommended.

Procedure

- 1. From the Configuration window, click User Creation. The Create User window opens.
- 2. Type the user account name in to the Full name field, for example: John Smith.
- 3. Type the username in to the **User name** field, for example: jsmith.



NOTE

The **User name** is used to log in from a command line; if you install a graphical environment, then your graphical login manager uses the **Full name**.

4. Select the **Make this user administrator** check box if the user requires administrative rights (the installation program adds the user to the **wheel** group).



IMPORTANT

An administrator user can use the **sudo** command to perform tasks that are only available to **root** using the user password, instead of the **root** password. This may be more convenient, but it can also cause a security risk.

5. Select the **Require a password to use this account**check box.



WARNING

If you give administrator privileges to a user, verify that the account is password protected. Never give a user administrator privileges without assigning a password to the account.

- 6. Type a password into the **Password** field.
- 7. Type the same password into the **Confirm password** field.
- 8. Click **Save Changes** to apply the changes and return to the **Configuration** window.
- 9. When the installation process is complete, click **Reboot** to reboot and log in to your Red Hat Enterprise Linux 8 system.

5.3.1. Editing advanced user settings

Follow the steps in this procedure to edit the default settings for the user account in the **Advanced User Configuration** dialog box.

Procedure

- 1. Edit the details in the **Home directory** field, if required. The field is populated by default with /home/username.
- 2. In the User and Groups IDs section you can:
 - a. Select the **Specify a user ID manually**check box and use + or to enter the required value.



NOTE

The default value is 1000. User IDs (UIDs) 0-999 are reserved by the system so they cannot be assigned to a user.

b. Select the **Specify a group ID manually** check box and use + or - to enter the required value.



NOTE

The default group name is the same as the user name, and the default Group ID (GID) is 1000. GIDs 0-999 are reserved by the system so they can not be assigned to a user group.

3. Specify additional groups as a comma-separated list in the **Group Membership** field. Groups that do not already exist are created; you can specify custom GIDs for additional groups in parentheses. If you do not specify a custom GID for a new group, the new group receives a GID automatically.



The user account created always has one default group membership (the user's default group with an ID set in the **Specify a group ID manually** field).

4. Click **Save Changes** to apply the updates and return to the **Configuration** window.

5.4. COMPLETING THE INSTALLATION

Remove any installation media if it is not ejected automatically upon reboot.

Red Hat Enterprise Linux 8 starts after your system's normal power-up sequence is complete. If your system was installed on a workstation with the X Window System, applications to configure your system are launched. These applications guide you through initial configuration and you can set your system time and date, register your system with Red Hat, and more. If the X Window System is not installed, a **login:** prompt is displayed.

To learn how to complete initial setup, register, and secure your system, see Chapter 6, Completing post-installation tasks.

CHAPTER 6. COMPLETING POST-INSTALLATION TASKS

This section describes how to complete the following post-installation tasks:

- Completing initial setup
- Registering your system



NOTE

Depending on your requirements, there are several methods to register your system. Most of these methods are completed as part of post-installation tasks. However, the Red Hat Content Delivery Network (CDN) registers your system and attaches RHEL subscriptions **before** the installation process starts. See Section 3.3.2, "Registering and installing RHEL from the CDN" for more information.

Securing your system

6.1. COMPLETING INITIAL SETUP

This section contains information about how to complete initial setup on a Red Hat Enterprise Linux 8 system.



IMPORTANT

- If you selected the **Server with GUI** base environment during installation, the **Initial Setup** window opens the first time you reboot your system after the installation process is complete.
- If you registered and installed RHEL from the CDN, the Subscription Manager option displays a note that all installed products are covered by valid entitlements.

The information displayed in the **Initial Setup** window might vary depending on what was configured during installation. At a minimum, the **Licensing** and **Subscription Manager** options are displayed.

Prerequisites

- You have completed the graphical installation according to the recommended workflow described on Section 3.2, "Installing RHEL using an ISO image from the Customer Portal".
- You have an active, non-evaluation Red Hat Enterprise Linux subscription.

- From the Initial Setup window, select Licensing Information.
 The License Agreement window opens and displays the licensing terms for Red Hat Enterprise Linux.
- 2. Review the license agreement and select the I accept the license agreement checkbox.



You must accept the license agreement. Exiting **Initial Setup** without completing this step causes a system restart. When the restart process is complete, you are prompted to accept the license agreement again.

3. Click **Done** to apply the settings and return to the **Initial Setup** window.



NOTE

If you did not configure network settings, you cannot register your system immediately. In this case, click **Finish Configuration**. Red Hat Enterprise Linux 8 starts and you can login, activate access to the network, and register your system. See Section 6.3, "Registering your system using the Subscription Manager User Interface" for more information. If you configured network settings, as described in Section 4.3.3, "Configuring network and host name options", you can register your system immediately, as shown in the following steps:

4. From the Initial Setup window, select Subscription Manager.



IMPORTANT

If you registered and installed RHEL from the CDN, the Subscription Manager option displays a note that all installed products are covered by valid entitlements.

- 5. The **Subscription Manager** graphical interface opens and displays the option you are going to register, which is: **subscription.rhsm.redhat.com**.
- 6. Click Next.
- 7. Enter your Login and Password details and click Register.
- 8. Confirm the Subscription details and click **Attach**. You must receive the following confirmation message: **Registration with Red Hat Subscription Management is Done!**
- 9. Click **Done**. The **Initial Setup** window opens.
- 10. Click **Finish Configuration**. The login window opens.
- 11. Configure your system. See the *Configuring basic system settings* document for more information.

Additional resources

Depending on your requirements, there are five methods to register your system:

- Using the Red Hat Content Delivery Network (CDN) to register your system, attach RHEL subscriptions, and install Red Hat Enterprise Linux. See Section 3.3.2, "Registering and installing RHEL from the CDN" for more information.
- During installation using Initial Setup.
- After installation using the command line. See Section 6.2, "Registering your system using the command line" for more information.

- After installation using the Subscription Manager user interface. See Section 6.3, "Registering your system using the Subscription Manager User Interface" for more information.
- After installation using Registration Assistant. Registration Assistant is designed to help you
 choose the most suitable registration option for your Red Hat Enterprise Linux environment.
 See https://access.redhat.com/labs/registrationassistant/ for more information.

6.2. REGISTERING YOUR SYSTEM USING THE COMMAND LINE

This section contains information about how to register your Red Hat Enterprise Linux 8 system using the command line.



NOTE

When auto-attaching a system, the subscription service checks if the system is physical or virtual, as well as how many sockets are on the system. A physical system usually consumes two entitlements, a virtual system usually consumes one. One entitlement is consumed per two sockets on a system.

Prerequisites

- You have an active, non-evaluation Red Hat Enterprise Linux subscription.
- Your Red Hat subscription status is verified.
- You have not previously received a Red Hat Enterprise Linux 8 subscription.
- You have activated your subscription before attempting to download entitlements from the Customer Portal. You need an entitlement for each instance that you plan to use. Red Hat Customer Service is available if you need help activating your subscription.
- You have successfully installed Red Hat Enterprise Linux 8 and logged into the system.

- 1. Open a terminal window and register a subscription using your Red Hat Customer Portal username and password:
 - # subscription-manager register --username [username] --password [password]
- 2. When the subscription is successfully registered, an output similar to the following is displayed:
 - # The system has been registered with ID: 123456abcdef # The registered system name is: localhost.localdomain
- 3. Set the role for the system, for example:
 - # subscription-manager role --set="Red Hat Enterprise Linux Server"



Available roles depend on the subscriptions that have been purchased by the organization and the architecture of the RHEL 8 system. You can set one of the following roles: **Red Hat Enterprise Linux Server**, **Red Hat Enterprise Linux Workstation**, or **Red Hat Enterprise Linux Compute Node**.

4. Set the service level for the system, for example:

subscription-manager service-level --set="Premium"

5. Set the usage for the system, for example:

subscription-manager usage --set="Production"

6. Attach the system to an entitlement that matches the host system architecture:

subscription-manager attach

7. When the subscription is successfully attached, an output similar to the following is displayed:

Installed Product Current Status:

Product Name: Red Hat Enterprise Linux for x86_64

Status: Subscribed



NOTE

You can also register Red Hat Enterprise Linux 8 by logging in to the system as a **root** user and using the Subscription Manager graphical user interface.

6.3. REGISTERING YOUR SYSTEM USING THE SUBSCRIPTION MANAGER USER INTERFACE

This section contains information about how to register your Red Hat Enterprise Linux 8 system using the Subscription Manager User Interface to receive updates and access package repositories.

Prerequisites

- You have completed the graphical installation as per the recommended workflow described on Section 3.2, "Installing RHEL using an ISO image from the Customer Portal" .
- You have an active, non-evaluation Red Hat Enterprise Linux subscription.
- Your Red Hat subscription status is verified.

- 1. Log in to your system.
- 2. From the top left-hand side of the window, click **Activities**.
- 3. From the menu options, click the **Show Applications** icon.

- 4. Click the **Red Hat Subscription Manager** icon, or enter **Red Hat Subscription Manager** in the search.
- 5. Enter your administrator password in the **Authentication Required** dialog box.



Authentication is required to perform privileged tasks on the system.

- 6. The **Subscriptions** window opens, displaying the current status of Subscriptions, System Purpose, and installed products. Unregistered products display a red X.
- 7. Click the **Register** button.
- 8. The **Register System** dialog box opens. Enter your **Customer Portal** credentials and click the **Register** button.

The **Register** button in the **Subscriptions** window changes to **Unregister** and installed products display a green X. You can troubleshoot an unsuccessful registration using the **subscription-manager status** command.

Additional resources

- For more information about using and configuring Subscription Manager, see the *Using and Configuring Red Hat Subscription Manager* document.
- For information about preparing subscriptions in Subscription Management, configuring virtwho, and registering virtual machines so that they inherit a subscription from their hypervisor, see the Configuring Virtual Machine Subscriptions in Red Hat Subscription Management document.

6.4. REGISTRATION ASSISTANT

Registration Assistant is designed to help you choose the most suitable registration option for your Red Hat Enterprise Linux environment. See https://access.redhat.com/labs/registrationassistant/ for more information.

6.5. CONFIGURING SYSTEM PURPOSE USING THE SYSPURPOSE COMMAND-LINE TOOL

System Purpose is an optional but recommended feature of the Red Hat Enterprise Linux installation. You use System Purpose to record the intended use of a Red Hat Enterprise Linux 8 system, and ensure that the entitlement server auto-attaches the most appropriate subscription to your system. The **syspurpose** command-line tool is part of the **python3_syspurpose.rpm** package. If System Purpose was not configured during the installation process, you can use the **syspurpose** command-line tool after installation to set the required attributes.

Prerequisites

- You installed and registered your Red Hat Enterprise Linux 8 system, but System Purpose is not configured.
- You are logged in as a **root** user.

The python3_syspurpose.rpm package is available on your system.



NOTE

If your system is registered but has subscriptions that do not satisfy the required purpose, you can run the **subscription-manager remove --all** command to remove attached subscriptions. You can then use the **syspurpose** command-line tool to set the required purpose attributes, and run **subscription-manager attach --auto** to entitle the system with the updated attributes.

Procedure

Complete the steps in this procedure to configure System Purpose after installation using the **syspurpose** command-line tool. The selected values are used by the entitlement server to attach the most suitable subscription to your system.

1. From a terminal window, run the following command to set the intended role of the system:

syspurpose set-role "VALUE"

Replace **VALUE** with the role that you want to assign:

- Red Hat Enterprise Linux Server
- Red Hat Enterprise Linux Workstation
- Red Hat Enterprise Linux Compute Node

For example:

syspurpose set-role "Red Hat Enterprise Linux Server"

a. Optional: Run the following command to unset the role:

syspurpose unset-role

2. Run the following command to set the intended Service Level Agreement (SLA) of the system:

syspurpose set-sla "VALUE"

Replace **VALUE** with the SLA that you want to assign:

- Premium
- Standard
- Self-Support

For example:

syspurpose set-sla "Standard"

a. Optional: Run the following command to unset the SLA:

syspurpose unset-sla

3. Run the following command to set the intended usage of the system:

syspurpose set-usage "VALUE"

Replace **VALUE** with the usage that you want to assign:

- Production
- Disaster Recovery
- Development/Test

For example:

syspurpose set-usage "Production"

- a. Optional: Run the following command to unset the usage:
 - # syspurpose unset-usage
- 4. Run the following command to show the current system purpose properties:
 - # syspurpose show
 - a. Optional: Run the following command to access the **syspurpose** man page:
 - # man syspurpose

6.6. SECURING YOUR SYSTEM

Complete the following security-related steps immediately after you install Red Hat Enterprise Linux.

Prerequisites

• You have completed the graphical installation according to the recommended workflow described in Section 3.2, "Installing RHEL using an ISO image from the Customer Portal" .

Procedure

- 1. To update your system, run the following command as root:
 - # yum update
- 2. Even though the firewall service, **firewalld**, is automatically enabled with the installation of Red Hat Enterprise Linux, there are scenarios where it might be explicitly disabled, for example in a Kickstart configuration. In that scenario, it is recommended that you re-enable the firewall. To start **firewalld**, run the following commands as root:

systemctl start firewalld # systemctl enable firewalld 3. To enhance security, disable services that you do not need. For example, if your system has no printers installed, disable the cups service using the following command:

systemctl mask cups

To review active services, run the following command:

\$ systemctl list-units | grep service

6.7. DEPLOYING SYSTEMS THAT ARE COMPLIANT WITH A SECURITY PROFILE IMMEDIATELY AFTER AN INSTALLATION

You can use the OpenSCAP suite to deploy RHEL systems that are compliant with a security profile, such as OSPP or PCI-DSS, immediately after the installation process. Using this deployment method, you can apply specific rules that cannot be applied later using remediation scripts, for example, a rule for password strength and partitioning.

6.7.1. Deploying baseline-compliant RHEL systems using the graphical installation

Use this procedure to deploy a RHEL system that is aligned with a specific baseline. This example uses Protection Profile for General Purpose Operating System (OSPP).

Prerequisites

- You have booted into the **graphical** installation program. Note that the **OSCAP Anaconda Add-on** does not support text-only installation.
- You have accessed the **Installation Summary** window.

Procedure

- 1. From the **Installation Summary** window, click **Software Selection**. The **Software Selection** window opens.
- 2. From the **Base Environment** pane, select the **Server** environment. You can select only one base environment.



WARNING

Do not use the **Server with GUI** base environment if you want to deploy a compliant system. Security profiles provided as part of the **SCAP Security Guide** may not be compatible with the extended package set of **Server with GUI**. For more information, see, for example, BZ#1648162, BZ#1787156, or BZ#1816199.

- 3. Click **Done** to apply the setting and return to the **Installation Summary** window.
- 4. Click **Security Policy**. The **Security Policy** window opens.

- 5. To enable security policies on the system, toggle the **Apply security policy** switch to **ON**.
- 6. Select **Protection Profile for General Purpose Operating Systems** from the profile pane.
- 7. Click **Select Profile** to confirm the selection.
- 8. Confirm the changes in the **Changes that were done or need to be done** pane that is displayed at the bottom of the window. Complete any remaining manual changes.
- 9. Because OSPP has strict partitioning requirements that must be met, create separate partitions for /boot, /home, /var, /var/log, /var/tmp, and /var/log/audit.
- 10. Complete the graphical installation process.



The graphical installation program automatically creates a corresponding Kickstart file after a successful installation. You can use the /**root/anaconda-ks.cfg** file to automatically install OSPP-compliant systems.

Verification steps

1. To check the current status of the system after installation is complete, reboot the system and start a new scan:

oscap xccdf eval --profile ospp --report eval_postinstall_report.html /usr/share/xml/scap/ssg/content/ssg-rhel8-ds.xml

Additional resources

For more details on partitioning, see Configuring manual partitioning.

6.7.2. Deploying baseline-compliant RHEL systems using Kickstart

Use this procedure to deploy RHEL systems that are aligned with a specific baseline. This example uses Protection Profile for General Purpose Operating System (OSPP).

Prerequisites

• The **scap-security-guide** package is installed on your RHEL 8 system.

- 1. Open the /usr/share/scap-security-guide/kickstarts/ssg-rhel8-ospp-ks.cfg Kickstart file in an editor of your choice.
- 2. Update the partitioning scheme to fit your configuration requirements. For OSPP compliance, the separate partitions for /boot, /home, /var, /var/log, /var/tmp, and /var/log/audit must be preserved, and you can only change the size of the partitions.



WARNING

Because the **OSCAP Anaconda Addon** plugin does not support text-only installation, do not use the **text** option in your Kickstart file. For more information, see RHBZ#1674001.

3. Start a Kickstart installation as described in Performing an automated installation using Kickstart.



IMPORTANT

Passwords in the hash form cannot be checked for OSPP requirements.

Verification steps

1. To check the current status of the system after installation is complete, reboot the system and start a new scan:

oscap xccdf eval --profile ospp --report eval_postinstall_report.html /usr/share/xml/scap/ssg/content/ssg-rhel8-ds.xml

Additional resources

• For more details, see the OSCAP Anaconda Addon project page.

6.8. NEXT STEPS

When you have completed the required post-installation steps, you can configure basic system settings. For information about completing tasks such as installing software with yum, using systemd for service management, managing users, groups, and file permissions, using chrony to configure NTP, and working with Python 3, see the *Configuring basic system settings* document.

APPENDIX A. TROUBLESHOOTING

The following sections cover various troubleshooting information that might be helpful when diagnosing issues during different stages of the installation process.

A.1. TROUBLESHOOTING AT THE START OF THE INSTALLATION PROCESS

The troubleshooting information in the following sections might be helpful when diagnosing issues at the start of the installation process. The following sections are for all supported architectures. However, if an issue is for a particular architecture, it is specified at the start of the section.

A.1.1. Dracut

Dracut is a tool that manages the **initramfs** image during the Linux operating system boot process. The **dracut** emergency shell is an interactive mode that can be initiated while the **initramfs** image is loaded. You can run basic troubleshooting commands from the **dracut** emergency shell. For more information, see the **Troubleshooting** section of the **dracut** man page.

A.1.2. Using installation log files

For debugging purposes, the installation program logs installation actions in files that are located in the /tmp directory. These log files are listed in the following table.

Table A.1. Log files generated during the installation

Log file	Contents
/tmp/anaconda.log	General messages.
/tmp/program.log	All external programs run during the installation.
/tmp/storage.log	Extensive storage module information.
/tmp/packaging.log	yum and rpm package installation messages.
/tmp/dbus.log	Information about the dbus session that is used for installation program modules.
/tmp/ifcfg.log	Information about networking scripts.
/tmp/sensitive-info.log	Configuration information that is not part of other logs and not copied to the installed system.
/tmp/syslog	Hardware-related system messages.

If the installation fails, the messages are consolidated into /tmp/anaconda-tb-identifier, where identifier is a random string. After a successful installation, these files are copied to the installed system under the directory /var/log/anaconda/. However, if the installation is unsuccessful, or if the inst.nosave=all or inst.nosave=logs options are used when booting the installation system, these logs only exist in the

installation program's RAM disk. This means that the logs are not saved permanently and are lost when the system is powered down. To store them permanently, copy the files to another system on the network or copy them to a mounted storage device such as a USB flash drive.

A.1.2.1. Creating pre-installation log files

Use this procedure to set the **inst.debug** option to create log files before the installation process starts. These log files contain, for example, the current storage configuration.

Prerequisites

• The Red Hat Enterprise Linux boot menu is displayed.

Procedure

- 1. Select the Install Red Hat Enterprise Linux option from the boot menu.
- 2. Press the **Tab** key on BIOS-based systems or the **e** key on UEFI-based systems to edit the selected boot options.
- 3. Append inst.debug to the options. For example:
 - vmlinuz ... inst.debug
- 4. Press the **Enter** key on your keyboard. The system stores the pre-installation log files in the /tmp/pre-anaconda-logs/ directory before the installation program starts.
- 5. To access the log files, switch to the console.
- 6. Change to the /tmp/pre-anaconda-logs/ directory:
 - # cd /tmp/pre-anaconda-logs/

Additional Resources

- For more information about boot options, see Appendix D, Boot options reference.
- For more information about how to switch consoles, see Section A.1.5, "Consoles and logging during installation"

A.1.2.2. Transferring installation log files to a USB drive

Use this procedure to transfer installation log files to a USB drive.

Prerequisites

- Back up any data on the USB drive before using this procedure.
- You are logged into a root account and you have access to the installation program's temporary file system.

Procedure

1. Press Ctrl + Alt + F2 to access a shell prompt on the system you are installing.

2. Connect a USB flash drive to the system and run the **dmesg** command:

dmesg

A log detailing all recent events is displayed. At the end of this log, a set of messages is displayed. For example:

[170.171135] sd 5:0:0:0: [sdb] Attached SCSI removable disk

- 3. Note the name of the connected device. In the above example, it is **sdb**.
- 4. Navigate to the /mnt directory and create a new directory that serves as the mount target for the USB drive. This example uses the name usb:

mkdir usb

- 5. Mount the USB flash drive onto the newly created directory. In most cases, you do not want to mount the whole drive, but a partition on it. Do not use the name **sdb**, use the name of the partition you want to write the log files to. In this example, the name **sdb1** is used:
 - # mount /dev/sdb1 /mnt/usb
- 6. Verify that you mounted the correct device and partition by accessing it and listing its contents:
 - # cd /mnt/usb
 - # ls
- 7. Copy the log files to the mounted device.
 - # cp /tmp/*log /mnt/usb
- 8. Unmount the USB flash drive. If you receive an error message that the target is busy, change your working directory to outside the mount (for example, /).
 - # umount /mnt/usb

A.1.2.3. Transferring installation log files over the network

Use this procedure to transfer installation log files over the network.

Prerequisites

• You are logged into a root account and you have access to the installation program's temporary file system.

- 1. Press Ctrl + Alt + F2 to access a shell prompt on the system you are installing.
- 2. Switch to the /tmp directory where the log files are located:

cd /tmp

3. Copy the log files onto another system on the network using the **scp** command:

scp *log user@address:path

a. Replace user with a valid user name on the target system, address with the target system's address or host name, and path with the path to the directory where you want to save the log files. For example, if you want to log in as john on a system with an IP address of 192.168.0.122 and place the log files into the /home/john/logs/ directory on that system, the command is as follows:

scp *log john@192.168.0.122:/home/john/logs/

When connecting to the target system for the first time, the SSH client asks you to confirm that the fingerprint of the remote system is correct and that you want to continue:

The authenticity of host '192.168.0.122 (192.168.0.122)' can't be established. ECDSA key fingerprint is a4:60:76:eb:b2:d0:aa:23:af:3d:59:5c:de:bb:c4:42. Are you sure you want to continue connecting (yes/no)?

b. Type **yes** and press **Enter** to continue. Provide a valid password when prompted. The files are transferred to the specified directory on the target system.

A.1.3. Detecting memory faults using the Memtest86 application

Faults in memory (RAM) modules can cause your system to fail unpredictably. In certain situations, memory faults might only cause errors with particular combinations of software. For this reason, you should test your system's memory before you install Red Hat Enterprise Linux.



NOTE

Red Hat Enterprise Linux includes the **Memtest86+** memory testing application for BIOS systems only. Support for UEFI systems is currently unavailable.

A.1.3.1. Running Memtest86

Use this procedure to run the **Memtest86** application to test your system's memory for faults before you install Red Hat Enterprise Linux.

Prerequisites

• You have accessed the Red Hat Enterprise Linux boot menu.

Procedure

From the Red Hat Enterprise Linux boot menu, select Troubleshooting > Run a memory test
The Memtest86 application window is displayed and testing begins immediately. By default,
Memtest86 performs ten tests in every pass. After the first pass is complete, a message is
displayed in the lower part of the window informing you of the current status. Another pass
starts automatically.

If **Memtest86+** detects an error, the error is displayed in the central pane of the window and is highlighted in red. The message includes detailed information such as which test detected a

problem, the memory location that is failing, and others. In most cases, a single successful pass of all 10 tests is sufficient to verify that your RAM is in good condition. In rare circumstances, however, errors that went undetected during the first pass might appear on subsequent passes. To perform a thorough test on important systems, run the tests overnight or for a few days to complete multiple passes.



NOTE

The amount of time it takes to complete a single full pass of **Memtest86+** varies depending on your system's configuration, notably the RAM size and speed. For example, on a system with 2 GiB of DDR2 memory at 667 MHz, a single pass takes 20 minutes to complete.

- 2. Optional: Follow the on-screen instructions to access the **Configuration** window and specify a different configuration.
- 3. To halt the tests and reboot your computer, press the **Esc** key at any time.

Additional resources

 For more information about using Memtest86, see the official website at http://www.memtest.org/.

A.1.4. Verifying boot media

Verifying ISO images helps to avoid problems that are sometimes encountered during installation. These sources include DVD and ISO images stored on a hard drive or NFS server. Use this procedure to test the integrity of an ISO-based installation source before using it to install Red Hat Enterprise Linux.

Prerequisites

• You have accessed the Red Hat Enterprise Linux boot menu.

Procedure

- 1. From the boot menu, select **Test this media & install Red Hat Enterprise Linux 8.1**to test the boot media.
- 2. The boot process tests the media and highlights any issues.
- 3. Optional: You can start the verification process by appending **rd.live.check** to the boot command line.

A.1.5. Consoles and logging during installation

The Red Hat Enterprise Linux installer uses the **tmux** terminal multiplexer to display and control several windows in addition to the main interface. Each of these windows serve a different purpose; they display several different logs, which can be used to troubleshoot issues during the installation process. One of the windows provides an interactive shell prompt with **root** privileges, unless this prompt was specifically disabled using a boot option or a Kickstart command.



In general, there is no reason to leave the default graphical installation environment unless you need to diagnose an installation problem.

The terminal multiplexer is running in virtual console 1. To switch from the actual installation environment to **tmux**, press **Ctrl+Alt+F1**. To go back to the main installation interface which runs in virtual console 6, press **Ctrl+Alt+F6**.



NOTE

If you choose text mode installation, you will start in virtual console 1 (tmux), and switching to console 6 will open a shell prompt instead of a graphical interface.

The console running **tmux** has five available windows; their contents are described in the following table, along with keyboard shortcuts. Note that the keyboard shortcuts are two-part: first press **Ctrl+b**, then release both keys, and press the number key for the window you want to use.

You can also use **Ctrl+b n**, **Alt+ Tab**, and **Ctrl+b p** to switch to the next or previous **tmux** window, respectively.

Table A.2. Available tmux windows

Shortcut	Contents
Ctrl+b 1	Main installation program window. Contains text- based prompts (during text mode installation or if you use VNC direct mode), and also some debugging information.
Ctrl+b 2	Interactive shell prompt with root privileges.
Ctrl+b 3	Installation log; displays messages stored in /tmp/anaconda.log.
Ctrl+b 4	Storage log; displays messages related to storage devices and configuration, stored in /tmp/storage.log.
Ctrl+b 5	Program log; displays messages from utilities executed during the installation process, stored in /tmp/program.log.

A.1.6. Saving screenshots

You can press **Shift+Print Screen** at any time during the graphical installation to capture the current screen. The screenshots are saved to /tmp/anaconda-screenshots.

A.1.7. Resuming an interrupted download attempt

You can resume an interrupted download using the **curl** command.

Prerequisite

- You have navigated to the Product Downloads section of the Red Hat Customer Portal at https://access.redhat.com/downloads, and selected the required variant, version, and architecture.
- You have right-clicked on the required ISO file, and selected **Copy Link Location** to copy the URL of the ISO image file to your clipboard.

Procedure

- 1. Download the ISO image from the new link. Add the **--continue-at -** option to automatically resume the download:
 - \$ curl --output directory-path/filename.iso 'new_copied_link_location' --continue-at -
- 2. Use a checksum utility such as **sha256sum** to verify the integrity of the image file after the download finishes:

```
$ sha256sum rhel-8.1-x86_64-dvd.iso
`85a...46c rhel-8.1-x86_64-dvd.iso`
```

Compare the output with reference checksums provided on the Red Hat Enterprise Linux **Product Download** web page.

Example A.1. Resuming an interrupted download attempt

The following is an example of a **curl** command for a partially downloaded ISO image:

```
$ curl --output _rhel-8.1-x86_64-dvd.iso 'https://access.cdn.redhat.com//content/origin/files/sha256/85/85a...46c/rhel-8.1-x86_64-dvd.iso? _auth=141...963' --continue-at -
```

A.1.8. Cannot boot into the graphical installation

Some video cards have trouble booting into the Red Hat Enterprise Linux graphical installation program. If the installation program does not run using its default settings, it attempts to run in a lower resolution mode. If that fails, the installation program attempts to run in text mode. There are several possible solutions to resolve display issues, most of which involve specifying custom boot options. For more information, see Section D.3, "Console boot options".

Table A.3. Solutions

Solution	Description
Use the basic graphics mode	You can attempt to perform the installation using the basic graphics driver. To do this, either select Troubleshooting > Install Red Hat Enterprise Linux 8.1 in basic graphics mode from the boot menu, or edit the installation program's boot options and append inst.xdriver=vesa at the end of the command line.

Solution	Description
Specify the display resolution manually	If the installation program fails to detect your screen resolution, you can override the automatic detection and specify it manually. To do this, append the inst.resolution=x option at the boot menu, where x is your display's resolution, for example, 1024x768.
Use an alternate video driver	You can attempt to specify a custom video driver, overriding the installation program's automatic detection. To specify a driver, use the inst.xdriver=x option, where x is the device driver you want to use (for example, nouveau)*.
Perform the installation using VNC	If the above options fail, you can use a separate system to access the graphical installation over the network, using the Virtual Network Computing (VNC) protocol. For details on installing using VNC, see the Performing a remote RHEL installation using VNC section of the Performing an advanced RHEL installation document.

^{*}If specifying a custom video driver solves your problem, you should report it as a bug at https://bugzilla.redhat.com under the **anaconda** component. The installation program should be able to detect your hardware automatically and use the appropriate driver without intervention.

A.2. TROUBLESHOOTING DURING THE INSTALLATION

The troubleshooting information in the following sections might be helpful when diagnosing issues during the installation process. The following sections are for all supported architectures. However, if an issue is for a particular architecture, it is specified at the start of the section.

A.2.1. Disks are not detected

If the installation program cannot find a writable storage device to install to, it returns the following error message in the **Installation Destination** window: **No disks detected. Please shut down the computer, connect at least one disk, and restart to complete installation.**

Check the following items:

- Your system has at least one storage device attached.
- If your system uses a hardware RAID controller; verify that the controller is properly configured and working as expected. See your controller's documentation for instructions.
- If you are installing into one or more iSCSI devices and there is no local storage present on the system, verify that all required LUNs are presented to the appropriate Host Bus Adapter (HBA).

If the error message is still displayed after rebooting the system and starting the installation process, the installation program failed to detect the storage. In many cases the error message is a result of attempting to install on an iSCSI device that is not recognized by the installation program.

In this scenario, you must perform a driver update before starting the installation. Check your hardware

vendor's website to determine if a driver update is available. For more general information on driver updates, see the *Updating drivers during installation* section of the *Performing an advanced RHEL installation* document.

You can also consult the Red Hat Hardware Compatibility List, available at https://access.redhat.com/ecosystem/search/#/category/Server.

A.2.2. Reporting error messages to Red Hat Customer Support

If the graphical installation encounters an error, it displays the **unknown error** dialog box. You can send information about the error to Red Hat Customer Support. To send a report, you must enter your Customer Portal credentials. If you do not have a Customer Portal account, you can register at https://www.redhat.com/wapps/ugc/register.html. Automated error reporting requires a network connection.

Prerequisite

The graphical installation program encountered an error and displayed the unknown error dialog box.

Procedure

- 1. From the **unknown error** dialog box, click **Report Bug** to report the problem, or **Quit** to exit the installation.
 - a. Optionally, click **More Info...** to display a detailed output that might help determine the cause of the error. If you are familiar with debugging, click **Debug**. This displays the virtual terminal **tty1**, where you can request additional information. To return to the graphical interface from **tty1**, use the **continue** command.
- 2. Click Report a bug to Red Hat Customer Support.
- 3. The **Red Hat Customer Support Reporting Configuration**dialog box is displayed. From the **Basic** tab, enter your Customer Portal user name and password. If your network settings require you to use an HTTP or HTTPS proxy, you can configure it by selecting the **Advanced** tab and entering the address of the proxy server.
- 4. Complete all fields and click **OK**.
- 5. A text box is displayed. Explain each step that was taken before the **unknown error** dialog box was displayed.
- 6. Select an option from the **How reproducible is this problem** drop-down menu and provide additional information in the text box.
- 7. Click **Forward**.
- 8. Verify that all the information you provided is in the **Comment** tab. The other tabs include information such as your system's host name and other details about your installation environment. You can remove any of the information that you do not want to send to Red Hat, but be aware that providing less detail might affect the investigation of the issue.
- 9. Click **Forward** when you have finished reviewing all tabs.
- 10. A dialog box displays all the files that will be sent to Red Hat. Clear the check boxes beside the files that you do not want to send to Red Hat. To add a file, click **Attach a file**.
- 11. Select the check box I have reviewed the data and agree with submitting it.

- 12. Click **Forward** to send the report and attachments to Red Hat.
- 13. Click **Show log** to view the details of the reporting process or click **Close** to return to the **unknown error** dialog box.
- 14. Click Quit to exit the installation.

A.2.3. Partitioning issues for IBM Power Systems



NOTE

This issue is for IBM Power Systems.

If you manually created partitions, but cannot move forward in the installation process, you might not have created all the partitions that are necessary for the installation to proceed. At a minimum, you must have the following partitions:

- / (root) partition
- **PReP** boot partition
- /boot partition (only if the root partition is an LVM logical volume)

See Section C.4, "Recommended partitioning scheme" for more information.

A.3. TROUBLESHOOTING AFTER INSTALLATION

The troubleshooting information in the following sections might be helpful when diagnosing issues after the installation process. The following sections are for all supported architectures. However, if an issue is for a particular architecture, it is specified at the start of the section.

A.3.1. Cannot boot with a RAID card

If you cannot boot your system after the installation, you might need to reinstall and repartition your system's storage. Some BIOS types do not support booting from RAID cards. After you finish the installation and reboot the system for the first time, a text-based screen displays the boot loader prompt (for example, **grub>**) and a flashing cursor might be displayed. If this is the case, you must repartition your system and move your /**boot** partition and the boot loader outside of the RAID array. The /**boot** partition and the boot loader must be on the same drive. Once these changes have been made, you should be able to finish your installation and boot the system properly.

A.3.2. Graphical boot sequence is not responding

When rebooting your system for the first time after installation, the system might be unresponsive during the graphical boot sequence. If this occurs, a reset is required. In this scenario, the boot loader menu is displayed successfully, but selecting any entry and attempting to boot the system results in a halt. This usually indicates that there is a problem with the graphical boot sequence. To resolve the issue, you must disable the graphical boot by temporarily altering the setting at boot time before changing it permanently.

Procedure: Disabling the graphical boot temporarily

1. Start your system and wait until the boot loader menu is displayed. If you set your boot timeout period to **0**, press the **Esc** key to access it.

- 2. From the boot loader menu, use your cursor keys to highlight the entry you want to boot. Press the **Tab** key on BIOS-based systems or the **e** key on UEFI-based systems to edit the selected entry options.
- 3. In the list of options, find the kernel line that is, the line beginning with the keyword **linux**. On this line, locate and delete **rhgb**.
- 4. Press **F10** or **Ctrl+X** to boot your system with the edited options.

If the system started successfully, you can log in normally. However, if you do not disable graphical boot permanently, you must perform this procedure every time the system boots.

Procedure: Disabling the graphical boot permanently

- 1. Log in to the root account on your system.
- 2. Use the grubby tool to find the default GRUB2 kernel:

```
# grubby --default-kernel
/boot/vmlinuz-4.18.0-94.el8.x86_64
```

- 3. Use the grubby tool to remove the **rhgb** boot option from the default kernel in your GRUB2 configuration. For example:
 - # grubby --remove-args="rhgb" --update-kernel /boot/vmlinuz-4.18.0-94.el8.x86_64
- 4. Reboot the system. The graphical boot sequence is no longer used. If you want to enable the graphical boot sequence, follow the same procedure, replacing the **--remove-args="rhgb"** parameter with the **--args="rhgb"** parameter. This restores the **rhgb** boot option to the default kernel in your GRUB2 configuration.

A.3.3. X server fails after log in

An X server is a program in the X Window System that runs on local machines, that is, the computers used directly by users. X server handles all access to the graphics cards, display screens and input devices, typically a keyboard and mouse on those computers. The X Window System, often referred to as X, is a complete, cross-platform and free client-server system for managing GUIs on single computers and on networks of computers. The client-server model is an architecture that divides the work between two separate but linked applications, referred to as clients and servers.*

If X server crashes after login, one or more of the file systems might be full. To troubleshoot the issue, execute the following command:

\$ df -h

The output verifies which partition is full - in most cases, the problem is on the /home partition. The following is a sample output of the df command:

 Filesystem
 Size
 Used Avail Use% Mounted on devtmpfs

 396M
 0
 396M
 0% /dev

 tmpfs
 411M
 0
 411M
 0% /dev/shm

 tmpfs
 411M
 6.7M
 405M
 2% /run

 tmpfs
 411M
 0
 411M
 0% /sys/fs/cgroup

 /dev/mapper/rhel-root
 17G
 4.1G
 13G
 25% /

/dev/sda1	1014M 173M 842M 17% /boot
tmpfs	83M 20K 83M 1% /run/user/42
tmpfs	83M 84K 83M 1% /run/user/1000
/dev/dm-4	90G 90G 0 100% /home

In the example, you can see that the /home partition is full, which causes the failure. Remove any unwanted files. After you free up some disk space, start X using the **startx** command. For additional information about **df** and an explanation of the options available, such as the **-h** option used in this example, see the **df(1)** man page.

*Source: http://www.linfo.org/x_server.html

A.3.4. RAM is not recognized

In some scenarios, the kernel does not recognize all memory (RAM), which causes the system to use less memory than is installed. You can find out how much RAM is being utilized using the free **-m** command. If the total amount of memory does not match your expectations, it is likely that at least one of your memory modules is faulty. On BIOS-based systems, you can use the **Memtest86+** utility to test your system's memory.

Some hardware configurations have part of the system's RAM reserved, and as a result, it is unavailable to the system. Some laptop computers with integrated graphics cards reserve a portion of memory for the GPU. For example, a laptop with 4 GiB of RAM and an integrated Intel graphics card shows roughly 3.7 GiB of available memory. Additionally, the kdump crash kernel dumping mechanism, which is enabled by default on most Red Hat Enterprise Linux systems, reserves some memory for the secondary kernel used in case of a primary kernel failure. This reserved memory is not displayed as available when using the free command.

Procedure: Manually configuring the memory

Use this procedure to manually set the amount of memory using the **mem=** kernel option.

- 1. Start your system and wait until the boot loader menu is displayed. If you set your boot timeout period to **0**, press the **Esc** key to access it.
- 2. From the boot loader menu, use your cursor keys to highlight the entry you want to boot, and press the **Tab** key on BIOS-based systems or the **e** key on UEFI-based systems to edit the selected entry options.
- 3. In the list of options, find the kernel line that is, the line beginning with the keyword **linux**. Append the following option to the end of this line:
 - mem=xxM
- 4. Replace xx with the amount of RAM you have in MiB.
- 5. Press **F10** or **Ctrl+X** to boot your system with the edited options.
- 6. Wait for the system to boot and then log in.
- 7. Open a command line and execute the free -m command again. If the total amount of RAM displayed by the command matches your expectations, append the following to the line beginning with GRUB_CMDLINE_LINUX in the /etc/default/grub file to make the change permanent:

grub2-mkconfig --output=/boot/grub2/grub.cfg

A.3.5. System is displaying signal 11 errors

A signal 11 error, commonly known as a segmentation fault means that a program accessed a memory location that it was not assigned. A signal 11 error can occur due to a bug in one of the software programs that are installed, or faulty hardware. If you receive a signal 11 error during the installation process, verify that you are using the most recent installation images and prompt the installation program to verify them to ensure they are not corrupt. For more information, see Section A.1.4, "Verifying boot media".

Faulty installation media (such as an improperly burned or scratched optical disk) are a common cause of signal 11 errors. Verifying the integrity of the installation media is recommended before every installation. For information about obtaining the most recent installation media, see Section 2.6, "Downloading the installation ISO image".

To perform a media check before the installation starts, append the **rd.live.check** boot option at the boot menu. If you performed a media check without any errors and you still have issues with segmentation faults, it usually indicates that your system encountered a hardware error. In this scenario, the problem is most likely in the system's memory (RAM). This can be a problem even if you previously used a different operating system on the same computer without any errors.



NOTE

For AMD and Intel 64-bit and 64-bit ARM architectures: On BIOS-based systems, you can use the **Memtest86+** memory testing module included on the installation media to perform a thorough test of your system's memory. For more information, see Section A.1.3, "Detecting memory faults using the Memtest86 application".

Other possible causes are beyond this document's scope. Consult your hardware manufacturer's documentation and also see the Red Hat Hardware Compatibility List, available online at https://access.redhat.com/ecosystem/search/#/category/Server.

A.3.6. Unable to IPL from network storage space



NOTE

This issue is for IBM Power Systems.

If you experience difficulties when trying to IPL from Network Storage Space (*NWSSTG), it is most likely due to a missing PReP partition. In this scenario, you must reinstall the system and create this partition during the partitioning phase or in the Kickstart file.

A.3.7. Using XDMCP

There are scenarios where you have installed the X Window System and want to log in to your Red Hat Enterprise Linux system using a graphical login manager. Use this procedure to enable the X Display Manager Control Protocol (XDMCP) and remotely log in to a desktop environment from any X-compatible client, such as a network-connected workstation or X11 terminal.



NOTE

XDMCP is not supported by the Wayland protocol. For more information, see the *Using the desktop environment in RHEL 8* document.



NOTE

This issue is for IBM Z.

Procedure

- 1. Open the /etc/gdm/custom.conf configuration file in a plain text editor such as vi or nano.
- 2. In the **custom.conf** file, locate the section starting with **[xdmcp]**. In this section, add the following line:
 - Enable=true
- 3. Save the file and exit the text editor.
- 4. Restart the X Window System. To do this, either reboot the system, or restart the GNOME Display Manager using the following command as root:
 - # systemctl restart gdm.service
- 5. Wait for the login prompt and log in using your user name and password. The X Window System is now configured for XDMCP. You can connect to it from another workstation (client) by starting a remote X session using the X command on the client workstation. For example:
 - \$ X :1 -query address
- 6. Replace **address** with the host name of the remote X11 server. The command connects to the remote X11 server using XDMCP and displays the remote graphical login screen on display :1 of the X11 server system (usually accessible by pressing **Ctrl-Alt-F8**). You can also access remote desktop sessions using a nested X11 server, which opens the remote desktop as a window in your current X11 session. You can use Xnest to open a remote desktop nested in a local X11 session. For example, run Xnest using the following command, replacing address with the host name of the remote X11 server:
 - \$ Xnest :1 -query address

For more information about XDMCP, see the X Window System documentation at http://www.x.org/releases/X11R7.6/doc/libXdmcp/xdmcp.html.

A.3.8. Using rescue mode

The installation program's rescue mode is a minimal Linux environment that can be booted from the Red Hat Enterprise Linux DVD or other boot media. It contains command-line utilities for repairing a wide variety of issues. Rescue mode can be accessed from the **Troubleshooting** menu of the boot menu. In this mode, you can mount file systems as read-only, blacklist or add a driver provided on a driver disc, install or upgrade system packages, or manage partitions.



NOTE

The installation program's rescue mode is different from rescue mode (an equivalent to single-user mode) and emergency mode, which are provided as parts of the **systemd** system and service manager.

To boot into rescue mode, you must be able to boot the system using one of the Red Hat Enterprise Linux boot media, such as a minimal boot disc or USB drive, or a full installation DVD.



IMPORTANT

Advanced storage, such as iSCSI or zFCP devices, must be configured either using **dracut** boot options such as **rd.zfcp=** or **root=iscsi:** options, or in the CMS configuration file on IBM Z. It is not possible to configure these storage devices interactively after booting into rescue mode. For information about **dracut** boot options, see the **dracut.cmdline(7)** man page.

A.3.8.1. Booting into rescue mode

Use this procedure to boot into rescue mode.

Procedure

- 1. Boot the system from either minimal boot media, or a full installation DVD or USB drive, and wait for the boot menu to be displayed.
- 2. From the boot menu, either select **Troubleshooting** > **Rescue** a **Red Hat Enterprise Linux** system option, or append the **inst.rescue** option to the boot command line. To enter the boot command line, press the **Tab** key on BIOS-based systems or the **e** key on UEFI-based systems.
- 3. Optional: If your system requires a third-party driver provided on a driver disc to boot, append the **inst.dd=driver name** to the boot command line:
 - inst.rescue inst.dd=driver_name
- 4. Optional: If a driver that is part of the Red Hat Enterprise Linux distribution prevents the system from booting, append the **modprobe.blacklist=** option to the boot command line:
 - inst.rescue modprobe.blacklist=driver_name
- 5. Press **Enter** (BIOS-based systems) or **Ctrl+X** (UEFI-based systems) to boot the modified option. Wait until the following message is displayed:

The rescue environment will now attempt to find your Linux installation and mount it under the directory: /mnt/sysimage/. You can then make any changes required to your system. Choose 1 to proceed with this step. You can choose to mount your file systems read-only instead of read-write by choosing 2. If for some reason this process does not work choose 3 to skip directly to a shell.

- 1) Continue
- 2) Read-only mount
- 3) Skip to shell
- 4) Quit (Reboot)

If you select 1, the installation program attempts to mount your file system under the directory /mnt/sysimage/. You are notified if it fails to mount a partition. If you select 2, it attempts to mount your file system under the directory /mnt/sysimage/, but in read-only mode. If you select 3, your file system is not mounted.

6. Select 1 to continue. Once your system is in rescue mode, a prompt appears on VC (virtual console) 1 and VC 2. Use the **Ctrl+Alt+F1** key combination to access VC 1 and **Ctrl+Alt+F2** to access VC 2:

sh-4.2#

7. Even if your file system is mounted, the default root partition while in rescue mode is a temporary root partition, not the root partition of the file system used during normal user mode (**multi-user.target** or **graphical.target**). If you selected to mount your file system and it mounted successfully, you can change the root partition of the rescue mode environment to the root partition of your file system by executing the following command:

sh-4.2# chroot /mnt/sysimage

This is useful if you need to run commands, such as **rpm**, that require your root partition to be mounted as /. To exit the chroot environment, type **exit** to return to the prompt.

8. If you selected **3**, you can still try to mount a partition or LVM2 logical volume manually inside rescue mode by creating a directory, such as /directory/, and typing the following command:

sh-4.2# mount -t xfs /dev/mapper/VolGroup00-LogVol02 /directory

In the above command, /directory/ is the directory that you created and /dev/mapper/VolGroup00-LogVol02 is the LVM2 logical volume you want to mount. If the partition is a different type than XFS, replace the xfs string with the correct type (such as ext4).

9. If you do not know the names of all physical partitions, use the following command to list them:

sh-4.2# fdisk -l

If you do not know the names of all LVM2 physical volumes, volume groups, or logical volumes, use the **pvdisplay**, **vgdisplay** or **lvdisplay** commands.

A.3.8.2. Using an SOS report in rescue mode

The **sosreport** command-line utility collects configuration and diagnostic information, such as the running kernel version, loaded modules, and system and service configuration files from the system. The utility output is stored in a tar archive in the /var/tmp/ directory. The **sosreport** utility is useful for analyzing system errors and troubleshooting. Use this procedure to capture an **sosreport** output in rescue mode.

Prerequisites

- You have booted into rescue mode.
- You have mounted the installed system / **(root)** partition in read-write mode.
- You have contacted Red Hat Support about your case and received a case number.

Procedure

1. Change the root directory to the /mnt/sysimage/ directory:

sh-4.2# chroot /mnt/sysimage/

2. Execute **sosreport** to generate an archive with system configuration and diagnostic information:

sh-4.2# sosreport



IMPORTANT

sosreport prompts you to enter your name and the case number you received from Red Hat Support. Use only letters and numbers because adding any of the following characters or spaces could render the report unusable:

3. Optional: If you want to transfer the generated archive to a new location using the network, it is necessary to have a network interface configured. In this scenario, use the dynamic IP addressing as no other steps required. However, when using static addressing, enter the following command to assign an IP address (for example 10.13.153.64/23) to a network interface, for example dev eth0:

bash-4.2# ip addr add 10.13.153.64/23 dev eth0

4. Exit the chroot environment:

sh-4.2# exit

5. Store the generated archive in a new location, from where it can be easily accessible:

sh-4.2# cp /mnt/sysimage/var/tmp/sosreport new_location

6. For transferring the archive through the network, use the **scp** utility:

 $sh-4.2\# scp /mnt/sysimage/var/tmp/sosreport \ username@hostname:sosreport$

Additional resources

- For general information about sosreport, see What is an sosreport and how to create one in Red Hat Enterprise Linux?.
- For information about using sosreport in rescue mode, see How to generate sosreport from the rescue environment.
- For information about generating an sosreport to a different location than /tmp, see How do I make sosreport write to an alternative location?.
- For information about collecting an sosreport manually, see Sosreport fails. What data should I provide in its place?.

A.3.8.3. Reinstalling the GRUB2 boot loader

In some scenarios, the GRUB2 boot loader is mistakenly deleted, corrupted, or replaced by other operating systems. Use this procedure to reinstall GRUB2 on the master boot record.

Prerequisites

- You have booted into rescue mode.
- You have mounted the installed system / (root) partition in read-write mode.

Procedure

- 1. Change the root partition:
 - sh-4.2# chroot /mnt/sysimage/
- 2. Reinstall the GRUB2 boot loader, where install device is the boot device, typically, /dev/sda:
 - sh-4.2# /sbin/grub2-install install_device
- 3. Reboot the system.

A.3.8.4. Using RPM to add or remove a driver

Missing or malfunctioning drivers cause problems when booting the system. Rescue mode provides an environment in which you can add or remove a driver even when the system fails to boot. Wherever possible, it is recommended that you use the RPM package manager to remove malfunctioning drivers or to add updated or missing drivers. Use the following procedures to add or remove a driver.



IMPORTANT

When you install a driver from a driver disc, the driver disc updates all **initramfs** images on the system to use this driver. If a problem with a driver prevents a system from booting, you cannot rely on booting the system from another **initramfs** image.

Procedure: Adding a driver using RPM

Use this procedure to add a driver.

Prerequisites

- You have booted into rescue mode.
- You have mounted the installed system in read-write mode.
 - Make the RPM package that contains the driver available. For example, mount a CD or USB flash drive and copy the RPM package to a location of your choice under /mnt/sysimage/, for example: /mnt/sysimage/root/drivers/.
 - 2. Change the root directory to /mnt/sysimage/:
 - sh-4.2# chroot /mnt/sysimage/
 - 3. Use the **rpm -ivh** command to install the driver package. For example, run the following command to install the **xorg-x11-drv-wacom** driver package from /**root/drivers**/:
 - sh-4.2# rpm -ivh /root/drivers/xorg-x11-drv-wacom-0.23.0-6.el7.x86_64.rpm



NOTE

The /**root/drivers**/ directory in this chroot environment is the /**mnt/sysimage/root/drivers**/ directory in the original rescue environment.

4. Exit the chroot environment:

sh-4.2# exit

Procedure: Removing a driver using RPM

Use this procedure to remove a driver.

Prerequisites

- You have booted into rescue mode.
- You have mounted the installed system in read-write mode.
 - 1. Change the root directory to the /mnt/sysimage/ directory:
 - sh-4.2# chroot /mnt/sysimage/
 - 2. Use the **rpm -e** command to remove the driver package. For example, to remove the **xorg- x11-drv-wacom** driver package, run:
 - sh-4.2# rpm -e xorg-x11-drv-wacom
 - 3. Exit the chroot environment:
 - sh-4.2# exit

If you cannot remove a malfunctioning driver for some reason, you can instead blacklist the driver so that it does not load at boot time.

4. When you have finished adding and removing drivers, reboot the system.

A.3.9. ip= boot option returns an error

Using the **ip=** boot option format **ip=[ip address]** for example, **ip=192.168.1.1** returns the error message **Fatal for argument 'ip=[insert ip here]'\n sorry, unknown value [ip address] refusing to continue**.

In previous releases of Red Hat Enterprise Linux, the boot option format was:

--ip=192.168.1.15 --netmask=255.255.255.0 --gateway=192.168.1.254 --nameserver=192.168.1.250 --hostname=myhost1

However, in Red Hat Enterprise Linux 8, the boot option format is:

ip=192.168.1.15::192.168.1.254:255.255.255.0:myhost1::none: nameserver=192.168.1.250

To resolve the issue, use the format: ip=ip::gateway:netmask:hostname:interface:none where:

- **ip** specifies the client ip address. You can specify IPv6 addresses in square brackets, for example, **[2001:DB8::1]**.
- gateway is the default gateway. IPv6 addresses are also accepted.
- **netmask** is the netmask to be used. This can be either a full netmask, for example, 255.255.255.0, or a prefix, for example, **64**.
- **hostname** is the host name of the client system. This parameter is optional.

For more information, see Section D.2, "Network boot options".

APPENDIX B. SYSTEM REQUIREMENTS REFERENCE

This section provides information and guidelines for hardware, installation target, system, memory, and RAID when installing Red Hat Enterprise Linux.

B.1. HARDWARE COMPATIBILITY

Red Hat works closely with hardware vendors on supported hardware.

- To verify that your hardware is supported, see the Red Hat Hardware Compatibility List, available at https://access.redhat.com/ecosystem/search/#/category/Server.
- To view supported memory sizes or CPU counts, see https://access.redhat.com/articles/rhellimits for information.

B.2. SUPPORTED INSTALLATION TARGETS

An installation target is a storage device that stores Red Hat Enterprise Linux and boots the system. Red Hat Enterprise Linux supports the following installation targets for AMD64, Intel 64, and 64-bit ARM systems:

- Storage connected by a standard internal interface, such as SCSI, SATA, or SAS
- BIOS/firmware RAID devices
- NVDIMM devices in sector mode on the Intel64 and AMD64 architectures, supported by the nd_pmem driver.
- Fibre Channel Host Bus Adapters and multipath devices. Some can require vendor-provided drivers.
- Xen block devices on Intel processors in Xen virtual machines.
- VirtlO block devices on Intel processors in KVM virtual machines.

Red Hat does not support installation to USB drives or SD memory cards. For information about support for third-party virtualization technologies, see the Red Hat Hardware Compatibility List.

B.3. SYSTEM SPECIFICATIONS

The Red Hat Enterprise Linux installation program automatically detects and installs your system's hardware, so you should not have to supply any specific system information. However, for certain Red Hat Enterprise Linux installation scenarios, it is recommended that you record system specifications for future reference. These scenarios include:

Installing RHEL with a customized partition layout

Record: The model numbers, sizes, types, and interfaces of the hard drives attached to the system. For example, Seagate ST3320613AS 320 GB on SATA0, Western Digital WD7500AAKS 750 GB on SATA1.

Installing RHEL as an additional operating system on an existing system

Record: Partitions used on the system. This information can include file system types, device node names, file system labels, and sizes, and allows you to identify specific partitions during the partitioning process. If one of the operating systems is a Unix operating system, Red Hat Enterprise Linux may

report the device names differently. Additional information can be found by executing the equivalent of the **mount** command and the **blkid** command, and in the **/etc/fstab** file.

If multiple operating systems are installed, the Red Hat Enterprise Linux installation program attempts to automatically detect them, and to configure boot loader to boot them. You can manually configure additional operating systems if they are not detected automatically. See *Configuring boot loader* in Section 4.4, "Configuring software options" for more information.

Installing RHEL from an image on a local hard drive

Record: The hard drive and directory that holds the image.

Installing RHEL from a network location

If the network has to be configured manually, that is, DHCP is not used.

Record:

- IP address
- Netmask
- Gateway IP address
- Server IP addresses, if required

Contact your network administrator if you need assistance with networking requirements.

Installing RHEL on an iSCSI target

Record: The location of the iSCSI target. Depending on your network, you may need a CHAP user name and password, and a reverse CHAP user name and password.

Installing RHEL if the system is part of a domain

Verify that the domain name is supplied by the DHCP server. If it is not, enter the domain name during installation.

B.4. DISK AND MEMORY REQUIREMENTS

If several operating systems are installed, it is important that you verify that the allocated disk space is separate from the disk space required by Red Hat Enterprise Linux.



NOTE

- For AMD64, Intel 64, and 64-bit ARM, at least two partitions (/ and **swap**) must be dedicated to Red Hat Enterprise Linux.
- For IBM Power Systems servers, at least three partitions (/, swap, and a PReP boot partition) must be dedicated to Red Hat Enterprise Linux.

You must have a minimum of 10 GiB of available disk space. See Appendix C, Partitioning reference for more information.

To install Red Hat Enterprise Linux, you must have a minimum of 10 GiB of space in either unpartitioned disk space or in partitions that can be deleted. See Appendix C, *Partitioning reference* for more information.

Table B.1. Minimum RAM requirements

Installation type	Recommended minimum RAM
Local media installation (USB, DVD)	768 MiB
NFS network installation	768 MiB
HTTP, HTTPS or FTP network installation	1.5 GiB



NOTE

It is possible to complete the installation with less memory than the recommended minimum requirements. The exact requirements depend on your environment and installation path. It is recommended that you test various configurations to determine the minimum required RAM for your environment. Installing Red Hat Enterprise Linux using a Kickstart file has the same recommended minimum RAM requirements as a standard installation. However, additional RAM may be required if your Kickstart file includes commands that require additional memory, or write data to the RAM disk. See the *Performing an advanced RHEL installation* document for more information.

B.5. RAID REQUIREMENTS

It is important to understand how storage technologies are configured and how support for them may have changed between major versions of Red Hat Enterprise Linux.

Hardware RAID

Any RAID functions provided by the mainboard of your computer, or attached controller cards, need to be configured before you begin the installation process. Each active RAID array appears as one drive within Red Hat Enterprise Linux.

Software RAID

On systems with more than one hard drive, you can use the Red Hat Enterprise Linux installation program to operate several of the drives as a Linux software RAID array. With a software RAID array, RAID functions are controlled by the operating system rather than the dedicated hardware.



NOTE

When a pre-existing RAID array's member devices are all unpartitioned disks/drives, the installation program treats the array as a disk and there is no method to remove the array.

USB Disks

You can connect and configure external USB storage after installation. Most devices are recognized by the kernel, but some devices may not be recognized. If it is not a requirement to configure these disks during installation, disconnect them to avoid potential problems.

NVDIMM devices

To use a Non-Volatile Dual In-line Memory Module (NVDIMM) device as storage, the following conditions must be satisfied:

- Version of Red Hat Enterprise Linux is 7.6 or later.
- The architecture of the system is Intel 64 or AMD64.
- The device is configured to sector mode. Anaconda can reconfigure NVDIMM devices to this mode.
- The device must be supported by the nd_pmem driver.

Booting from an NVDIMM device is possible under the following additional conditions:

- The system uses UEFI.
- The device must be supported by firmware available on the system, or by a UEFI driver. The UEFI driver may be loaded from an option ROM of the device itself.
- The device must be made available under a namespace.

To take advantage of the high performance of NVDIMM devices during booting, place the /boot and /boot/efi directories on the device.



NOTE

The Execute-in-place (XIP) feature of NVDIMM devices is not supported during booting and the kernel is loaded into conventional memory.

Considerations for Intel BIOS RAID Sets

Red Hat Enterprise Linux uses **mdraid** for installing on Intel BIOS RAID sets. These sets are automatically detected during the boot process and their device node paths can change across several booting processes. For this reason, local modifications to the /etc/fstab, /etc/crypttab or other configuration files that refer to the devices by their device node paths may not work in Red Hat Enterprise Linux. It is recommended that you replace device node paths (such as /dev/sda) with file system labels or device UUIDs. You can find the file system labels and device UUIDs using the blkid command.

APPENDIX C. PARTITIONING REFERENCE

C.1. SUPPORTED DEVICE TYPES

Standard partition

A standard partition can contain a file system or swap space. Standard partitions are most commonly used for /boot and the BIOS Boot and EFI System partitions. LVM logical volumes are recommended for most other uses.

LVM

Choosing **LVM** (or Logical Volume Management) as the device type creates an LVM logical volume. If no LVM volume group currently exists, one is automatically created to contain the new volume; if an LVM volume group already exists, the volume is assigned. LVM can improve performance when using physical disks, and it allows for advanced setups such as using multiple physical disks for one mount point, and setting up software RAID for increased performance, reliability, or both.

LVM thin provisioning

Using thin provisioning, you can manage a storage pool of free space, known as a thin pool, which can be allocated to an arbitrary number of devices when needed by applications. You can dynamically expand the pool when needed for cost-effective allocation of storage space.



WARNING

The installation program does not support overprovisioned LVM thin pools.

C.2. SUPPORTED FILE SYSTEMS

This section describes the file systems available in Red Hat Enterprise Linux.

xfs

XFS is a highly scalable, high-performance file system that supports file systems up to 16 exabytes (approximately 16 million terabytes), files up to 8 exabytes (approximately 8 million terabytes), and directory structures containing tens of millions of entries. **XFS** also supports metadata journaling, which facilitates quicker crash recovery. The maximum supported size of a single XFS file system is 500 TB. **XFS** is the default and recommended file system on Red Hat Enterprise Linux.

ext4

The **ext4** file system is based on the **ext3** file system and features a number of improvements. These include support for larger file systems and larger files, faster and more efficient allocation of disk space, no limit on the number of subdirectories within a directory, faster file system checking, and more robust journaling. The maximum supported size of a single **ext4** file system is 50 TB.

ext3

The **ext3** file system is based on the **ext2** file system and has one main advantage – journaling. Using a journaling file system reduces the time spent recovering a file system after it terminates unexpectedly, as there is no need to check the file system for metadata consistency by running the fsck utility every time.

ext2

An **ext2** file system supports standard Unix file types, including regular files, directories, or symbolic links. It provides the ability to assign long file names, up to 255 characters.

swap

Swap partitions are used to support virtual memory. In other words, data is written to a swap partition when there is not enough RAM to store the data your system is processing.

vfat

The **VFAT** file system is a Linux file system that is compatible with Microsoft Windows long file names on the FAT file system.



NOTE

Support for **VFAT** file system is not available for Linux system partitions. For example, /, /var, /usr and so on.

BIOS Boot

A very small partition required for booting from a device with a GUID partition table (GPT) on BIOS systems and UEFI systems in BIOS compatibility mode.

EFI System Partition

A small partition required for booting a device with a GUID partition table (GPT) on a UEFI system.

PReP

This small boot partition is located on the first partition of the hard drive. The **PReP** boot partition contains the GRUB2 boot loader, which allows other IBM Power Systems servers to boot Red Hat Enterprise Linux.

C.3. SUPPORTED RAID TYPES

RAID stands for Redundant Array of Independent Disks, a technology which allows you to combine multiple physical disks into logical units. Some setups are designed to enhance performance at the cost of reliability, while others improve reliability at the cost of requiring more disks for the same amount of available space.

This section describes supported software RAID types which you can use with LVM and LVM Thin Provisioning to set up storage on the installed system.

None

No RAID array is set up.

RAID 0

Performance: Distributes data across multiple disks. RAID 0 offers increased performance over standard partitions and can be used to pool the storage of multiple disks into one large virtual device. Note that RAID 0 offers no redundancy and that the failure of one device in the array destroys data in the entire array. RAID 0 requires at least two disks.

RAID 1

Redundancy: Mirrors all data from one partition onto one or more other disks. Additional devices in the array provide increasing levels of redundancy. RAID 1 requires at least two disks.

RAID 4

Error checking: Distributes data across multiple disks and uses one disk in the array to store parity information which safeguards the array in case any disk in the array fails. As all parity information is stored on one disk, access to this disk creates a "bottleneck" in the array's performance. RAID 4 requires at least three disks.

RAID 5

Distributed error checking: Distributes data and parity information across multiple disks. RAID 5 offers the performance advantages of distributing data across multiple disks, but does not share the performance bottleneck of RAID 4 as the parity information is also distributed through the array. RAID 5 requires at least three disks.

RAID 6

Redundant error checking: RAID 6 is similar to RAID 5, but instead of storing only one set of parity data, it stores two sets. RAID 6 requires at least four disks.

RAID₁₀

Performance and redundancy: RAID 10 is nested or hybrid RAID. It is constructed by distributing data over mirrored sets of disks. For example, a RAID 10 array constructed from four RAID partitions consists of two mirrored pairs of striped partitions. RAID 10 requires at least four disks.

C.4. RECOMMENDED PARTITIONING SCHEME

Red Hat recommends that you create separate file systems at the following mount points:

- /boot
- /(root)
- /home
- swap
- /boot/efi
- PReP

/boot partition - recommended size at least 1 GiB

The partition mounted on /boot contains the operating system kernel, which allows your system to boot Red Hat Enterprise Linux 8, along with files used during the bootstrap process. Due to the limitations of most firmwares, creating a small partition to hold these is recommended. In most scenarios, a 1 GiB boot partition is adequate. Unlike other mount points, using an LVM volume for /boot is not possible - /boot must be located on a separate disk partition.



WARNING

Normally, the **/boot** partition is created automatically by the installation program. However, if the **/** (root) partition is larger than 2 TiB and (U)EFI is used for booting, you need to create a separate **/boot** partition that is smaller than 2 TiB to boot the machine successfully.



NOTE

If you have a RAID card, be aware that some BIOS types do not support booting from the RAID card. In such a case, the **/boot** partition must be created on a partition outside of the RAID array, such as on a separate hard drive.

root - recommended size of 10 GiB

This is where "/", or the root directory, is located. The root directory is the top-level of the directory structure. By default, all files are written to this file system unless a different file system is mounted in the path being written to, for example, /boot or /home.

While a 5 GiB root file system allows you to install a minimal installation, it is recommended to allocate at least 10 GiB so that you can install as many package groups as you want.



IMPORTANT

Do not confuse the / directory with the /**root** directory. The /**root** directory is the home directory of the root user. The /**root** directory is sometimes referred to as *slash root* to distinguish it from the root directory.

/home - recommended size at least 1 GiB

To store user data separately from system data, create a dedicated file system for the /home directory. Base the file system size on the amount of data that is stored locally, number of users, and so on. You can upgrade or reinstall Red Hat Enterprise Linux 8 without erasing user data files. If you select automatic partitioning, it is recommended to have at least 55 GiB of disk space available for the installation, to ensure that the /home file system is created.

swap partition - recommended size at least 1 GB

Swap file systems support virtual memory; data is written to a swap file system when there is not enough RAM to store the data your system is processing. Swap size is a function of system memory workload, not total system memory and therefore is not equal to the total system memory size. It is important to analyze what applications a system will be running and the load those applications will serve in order to determine the system memory workload. Application providers and developers can provide quidance.

When the system runs out of swap space, the kernel terminates processes as the system RAM memory is exhausted. Configuring too much swap space results in storage devices being allocated but idle and is a poor use of resources. Too much swap space can also hide memory leaks. The maximum size for a swap partition and other additional information can be found in the **mkswap(8)** manual page.

The following table provides the recommended size of a swap partition depending on the amount of RAM in your system and if you want sufficient memory for your system to hibernate. If you let the installation program partition your system automatically, the swap partition size is established using these guidelines. Automatic partitioning setup assumes hibernation is not in use. The maximum size of the swap partition is limited to 10 percent of the total size of the hard drive, and the installation program cannot create swap partitions more than 128 GB in size. To set up enough swap space to allow for hibernation, or if you want to set the swap partition size to more than 10 percent of the system's storage space, or more than 128 GB, you must edit the partitioning layout manually.

/boot/efi partition - recommended size of 200 MiB

UEFI-based AMD64, Intel 64, and 64-bit ARM require a 200 MiB EFI system partition. The recommended minimum size is 200 MiB, the default size is 600 MiB, and the maximum size is 600 MiB. BIOS systems do not require an EFI system partition.

Table C.1. Recommended System Swap Space

Amount of RAM in the system	Recommended swap space	Recommended swap space if allowing for hibernation
Less than 2 GB	2 times the amount of RAM	3 times the amount of RAM
2 GB - 8 GB	Equal to the amount of RAM	2 times the amount of RAM
8 GB - 64 GB	4 GB to 0.5 times the amount of RAM	1.5 times the amount of RAM
More than 64 GB	Workload dependent (at least 4GB)	Hibernation not recommended

At the border between each range, for example, a system with 2 GB, 8 GB, or 64 GB of system RAM, discretion can be exercised with regard to chosen swap space and hibernation support. If your system resources allow for it, increasing the swap space can lead to better performance.

Distributing swap space over multiple storage devices - particularly on systems with fast drives, controllers and interfaces - also improves swap space performance.

Many systems have more partitions and volumes than the minimum required. Choose partitions based on your particular system needs.



NOTE

- Only assign storage capacity to those partitions you require immediately. You can allocate free space at any time, to meet needs as they occur.
- If you are unsure about how to configure partitions, accept the automatic default partition layout provided by the installation program.

PReP boot partition - recommended size of 4 to 8 MiB

When installing Red Hat Enterprise Linux on IBM Power System servers, the first partition of the hard drive should include a **PReP** boot partition. This contains the GRUB2 boot loader, which allows other IBM Power Systems servers to boot Red Hat Enterprise Linux.

C.5. ADVICE ON PARTITIONS

There is no best way to partition every system; the optimal setup depends on how you plan to use the system being installed. However, the following tips may help you find the optimal layout for your needs:

- Create partitions that have specific requirements first, for example, if a particular partition must be on a specific disk.
- Consider encrypting any partitions and volumes which might contain sensitive data. Encryption prevents unauthorized people from accessing the data on the partitions, even if they have

access to the physical storage device. In most cases, you should at least encrypt the /home partition, which contains user data.

- In some cases, creating separate mount points for directories other than /, /boot and /home may be useful; for example, on a server running a MySQL database, having a separate mount point for /var/lib/mysql will allow you to preserve the database during a reinstallation without having to restore it from backup afterward. However, having unnecessary separate mount points will make storage administration more difficult.
- Some special restrictions apply to certain directories with regards on which partitioning layouts can they be placed. Notably, the /boot directory must always be on a physical partition (not on an LVM volume).
- If you are new to Linux, consider reviewing the *Linux Filesystem Hierarchy Standard* at http://refspecs.linuxfoundation.org/FHS_2.3/fhs-2.3.html for information about various system directories and their contents.
- Each kernel installed on your system requires approximately 56 MB on the /**boot** partition:
 - o 32 MB initramfs
 - 14 MB kdump initramfs
 - 3.5 MB system map
 - 6.6 MB vmlinuz



NOTE

For rescue mode, **initramfs** and **vmlinuz** require 80 MB.

The default partition size of 1 GB for /**boot** should suffice for most common use cases. However, it is recommended that you increase the size of this partition if you are planning on retaining multiple kernel releases or errata kernels.

- The /var directory holds content for a number of applications, including the Apache web server, and is used by the DNF package manager to temporarily store downloaded package updates.
 Make sure that the partition or volume containing /var has at least 3 GB.
- The contents of the /var directory usually change very often. This may cause problems with older solid state drives (SSDs), as they can handle a lower number of read/write cycles before becoming unusable. If your system root is on an SSD, consider creating a separate mount point for /var on a classic (platter) HDD.
- The /usr directory holds the majority of software on a typical Red Hat Enterprise Linux installation. The partition or volume containing this directory should therefore be at least 5 GB for minimal installations, and at least 10 GB for installations with a graphical environment.
- If /usr or /var is partitioned separately from the rest of the root volume, the boot process becomes much more complex because these directories contain boot-critical components. In some situations, such as when these directories are placed on an iSCSI drive or an FCoE location, the system may either be unable to boot, or it may hang with a Device is busy error when powering off or rebooting.
 - This limitation only applies to /usr or /var, not to directories below them. For example, a separate partition for /var/www will work without issues.

- Consider leaving a portion of the space in an LVM volume group unallocated. This unallocated
 space gives you flexibility if your space requirements change but you do not wish to remove data
 from other volumes. You can also select the LVM Thin Provisioning device type for the
 partition to have the unused space handled automatically by the volume.
- The size of an XFS file system can not be reduced if you need to make a partition or volume with this file system smaller, you must back up your data, destroy the file system, and create a new, smaller one in its place. Therefore, if you expect needing to manipulate your partitioning layout later, you should use the ext4 file system instead.
- Use Logical Volume Management (LVM) if you anticipate expanding your storage by adding
 more hard drives or expanding virtual machine hard drives after the installation. With LVM, you
 can create physical volumes on the new drives, and then assign them to any volume group and
 logical volume as you see fit for example, you can easily expand your system's /home (or any
 other directory residing on a logical volume).
- Creating a BIOS Boot partition or an EFI System Partition may be necessary, depending on your system's firmware, boot drive size, and boot drive disk label. See Section C.4, "Recommended partitioning scheme" for information about these partitions. Note that graphical installation will not let you create a BIOS Boot or EFI System Partition if your system does not require one in that case, they will be hidden from the menu.
- If you need to make any changes to your storage configuration after the installation, Red Hat Enterprise Linux repositories offer several different tools which can help you do this. If you prefer a command line tool, try system-storage-manager.

APPENDIX D. BOOT OPTIONS REFERENCE

This section contains information about some of the boot options that you can use to modify the default behavior of the installation program. For Kickstart and advanced boot options, see the *Performing an advanced RHEL installation* document.

D.1. INSTALLATION SOURCE BOOT OPTIONS

This section contains information about the various installation source boot options.

inst.repo=

The **inst.repo**= boot option specifies the installation source, that is, the location providing the package repositories and a valid **.treeinfo** file that describes them. For example: **inst.repo**=**cdrom**. The target of the **inst.repo**= option must be one of the following installation media:

- an installable tree, which is a directory structure containing the installation program images, packages, and repository data as well as a valid .treeinfo file
- a DVD (a physical disk present in the system DVD drive)
- an ISO image of the full Red Hat Enterprise Linux installation DVD, placed on a hard drive or a network location accessible to the system.

You can use the **inst.repo=** boot option to configure different installation methods using different formats. The following table contains details of the **inst.repo=** boot option syntax:

Table D.1. inst.repo= installation source boot options

Source type	Boot option format	Source format	
CD/DVD drive	inst.repo=cdrom[:device]	Installation DVD as a physical disk. ^[a]	
Installable tree	inst.repo=hd:device:/path	Image file of the installation DVD, or an installation tree, which is a complete copy of the directories and files on the installation DVD.	
NFS Server	inst.repo=nfs:[options:]server:/path	Image file of the installation DVD, or an installation tree, which is a complete copy of the directories and files on the installation DVD. [b]	
HTTP Server	inst.repo=http://host/path	Installation tree, which is a complete copy of the	
HTTPS Server	inst.repo=https://host/path	directories and files on the installation DVD.	
FTP Server	inst.repo=ftp://username:password@ho st/path		

Source type	Boot option format	Source format
НМС	inst.repo=hmc	

[a] If device is left out, installation program automatically searches for a drive containing the installation DVD.

[b] The NFS Server option uses NFS protocol version 3 by default. To use a different version, add **nfsvers=** X to *options*, replacing X with the version number that you want to use.

You can set disk device names with the following formats:

- Kernel device name, for example /dev/sda1 or sdb2
- File system label, for example LABEL=Flash or LABEL=RHEL8
- File system UUID, for example **UUID=8176c7bf-04ff-403a-a832-9557f94e61db**Non-alphanumeric characters must be represented as **xNN**, where *NN* is the hexadecimal representation of the character. For example, **x20** is a white space (" ").

inst.addrepo=

Use the **inst.addrepo=** boot option to add an additional repository that can be used as another installation source along with the main repository (**inst.repo=**). You can use the **inst.addrepo=** boot option multip le times during one boot. The following table contains details of the **inst.addrepo=** boot option syntax.



NOTE

The **REPO_NAME** is the name of the repository and is required in the installation process. These repositories are only used during the installation process; they are not installed on the installed system.

Table D.2. inst.addrepo installation source boot options

Installation source	Boot option format	Additional information
Installable tree at a URL	inst.addrepo=REPO_NAME, [http,https,ftp]:// <host>/<pat h></pat </host>	Looks for the installable tree at a given URL.
Installable tree at an NFS path	inst.addrepo=REPO_NAME, nfs:// <server>:/<path></path></server>	Looks for the installable tree at a given NFS path. A colon is required after the host. The installation program passes every thing after nfs: // directly to the mount command instead of parsing URLs according to RFC 2224.

Installation source	Boot option format	Additional information
Installable tree in the installation environment	inst.addrepo=REPO_NAME, file:// <path></path>	Looks for the installable tree at the given location in the installation environment. To use this option, the repositor y must be mounted before the installation program attempts to load the available software groups. The benefit of this option is that you can have multiple repositories on one bootable ISO, and you can install both the main repository and additional repositories from the ISO. The path to the additional repositories is /run/install/source/REPO_ISO_PATH. Additional, you can mount the repository directory in the %pre secti on in the Kickstart file. The path must be absolute and start with /, for example inst.addrepo=REPO_NAME, file:/// <path></path>
Hard Drive	inst.addrepo=REPO_NAME, hd: <device>:<path></path></device>	Mounts the given <device> partition and installs from the ISO that is specified by the <path>. If the <path> is not specified, the installation p rogram looks for a valid installation ISO on the <device>. This installation method requires an ISO with a valid installable tree.</device></path></path></device>

inst.noverifyssl

The **inst.noverifyssl** boot option prevents the installation program from verifying the SSL certificate for all HTTPS connections with the exception of the additional Kickstart repositories, where **--noverifyssl** can be set per repository.

inst.stage2=

Use the **inst.stage2=** boot option to specify the location of the installation program runtime image. This option expects a path to a directory containing a valid **.treeinfo** file. The location of the runtime image is read from the **.treeinfo** file. If the **.treeinfo** file is not available, the installation program attempts to load the image from **images/install.img**.

When the **inst.stage2** option is not specified, the installation program attempts to use the location specified with **inst.repo** option.

You should specify this option only for PXE boot. The installation DVD and Boot ISO already contain a correct **inst.stage2** option to boot the installation program from themselves.



NOTE

By default, the **inst.stage2=** boot option is used on the installation media and is set to a specific label, for example, **inst.stage2=hd:LABEL=RHEL-8-0-0-BaseOS-x86_64**. If you modify the default label of the file system containing the runtime image, or if you use a customized procedure to boot the installation system, you must verify that the **inst.stage2=** boot option is set to the correct value.

inst.stage2.all

The **inst.stage2.all** boot option is used to specify several HTTP, HTTPS, or FTP sources. You can use the **inst.stage2=** boot option multiple times with the **inst.stage2.all** option to fetch the image from the sources sequentially until one succeeds. For example:

```
inst.stage2.all
inst.stage2=http://hostname1/path_to_install_tree/
inst.stage2=http://hostname2/path_to_install_tree/
inst.stage2=http://hostname3/path_to_install_tree/
```

inst.dd=

The **inst.dd=** boot option is used to perform a driver update during the installation. See the *Performing an advanced RHEL installation* document for information on how to update drivers during installation.

inst.repo=hmc

When booting from a Binary DVD, the installation program prompts you to enter additional kernel parameters. To set the DVD as an installation source, append **inst.repo=hmc** to the kernel parameters. The installation program then enables **SE** and **HMC** file access, fetches the images for stage2 from the DVD, and provides access to the packages on the DVD for software selection. This option eliminates the requirement of an external network setup and expands the installation options.

inst.proxy=

The **inst.proxy=** boot option is used when performing an installation from a HTTP, HTTPS, FTP source. For example:



inst.nosave=

Use the **inst.nosave=** boot option to control which installation logs and related files are not saved to the installed system, for example **input_ks**, **output_ks**, **all_ks**, **logs** and **all**. Multiple values can be combined as a comma-separated list, for example: **input_ks,logs**.



NOTE

The **inst.nosave** boot option is used for excluding files from the installed system that can't be removed by a Kickstart %post script, such as logs and input/output Kickstart results.

Table D.3. inst.nosave boot options

Option	Description
input_ks	Disables the ability to save the input Kickstart results.
output_ks	Disables the ability to save the output Kickstart results generated by the installation program.
all_ks	Disables the ability to save the input and output Kickstart results.
logs	Disables the ability to save all installation logs.
all	Disables the ability to save all Kickstart results, and all logs.

inst.multilib

Use the **inst.multilib** boot option to set DNF's **multilib_policy** to **all**, instead of **best**.

inst.memcheck

The **inst.memcheck** boot option performs a check to verify that the system has enough RAM to complete the installation. If there isn't enough RAM, the installation process is stopped. The system check is approximate and memory usage during installation depends on the package selection, user interface, for example graphical or text, and other parameters.

inst.nomemcheck

The **inst.nomemcheck** boot option does not perform a check to verify if the system has enough RAM to complete the installation. Any attempt to perform the installation with less than the recommended minimum amount of memory is unsupported, and might result in the installation process failing.

D.2. NETWORK BOOT OPTIONS

This section contains information about commonly used network boot options.



NOTE

Initial network initialization is handled by **dracut**. For a complete list, see the **dracut.cmdline(7)** man page.

ip=

Use the **ip=** boot option to configure one or more network interfaces. To configure multiple interfaces, you can use the **ip** option multiple times, once for each interface; to do so, you must use the **rd.neednet=1** option, and you must specify a primary boot interface using the **bootdev** option. Alternatively, you can use the **ip** option once, and then use Kickstart to set up further interfaces. This option accepts several different formats. The following tables contain information about the most common options.



NOTE

In the following tables:

- The **ip** parameter specifies the client IP address and requires square brackets, for example [2001:db8::99].
- The **gateway** parameter is the default gateway. IPv6 addresses are also accepted.
- The **netmask** parameter is the netmask to be used. This can be either a full netmask (for example, 255.255.255.0) or a prefix (for example, 64).
- The **hostname** parameter is the host name of the client system. This parameter is optional.

Table D.4. Network interface configuration boot option formats

Configuration method	Boot option format
Automatic configuration of any interface	ip=method
Automatic configuration of a specific interface	ip=interface:method
Static configuration	ip= <i>ip::gateway:netmask:hostname:interfac e</i> :none
Automatic configuration of a specific interface with an override	ip=ip::gateway:netmask:hostname:interfac e:method:mtu



NOTE

The method **automatic configuration of a specific interface with an override** brings up the interface using the specified method of automatic configuration, such as **dhcp**, but overrides the automatically-obtained IP address, gateway, netmask, host name or other specified parameters. All parameters are optional, so specify only the parameters that you want to override.

The **method** parameter can be any of the following:

Table D.5. Automatic interface configuration methods

Automatic configuration method	Value
DHCP	dhcp
IPv6 DHCP	dhcp6
IPv6 automatic configuration	auto6
iSCSI Boot Firmware Table (iBFT)	ibft



NOTE

- If you use a boot option that requires network access, such as
 inst.ks=http://host:/path, without specifying the ip option, the installation
 program uses ip=dhcp.
- To connect to an iSCSI target automatically, you must activate a network device for accessing the target. The recommended way to activate a network is to use the ip=ibft boot option.

nameserver=

The **nameserver=** option specifies the address of the name server. You can use this option multiple times.



NOTE

The **ip=** parameter requires square brackets. However, an IPv6 address does not work with square brackets. An example of the correct syntax to use for an IPv6 address is **nameserver=2001:db8::1**.

bootdev=

The **bootdev=** option specifies the boot interface. This option is mandatory if you use more than one **ip** option.

ifname=

The **ifname=** options assigns an interface name to a network device with a given MAC address. You can use this option multiple times. The syntax is **ifname=interface:MAC**. For example:



ifname=eth0:01:23:45:67:89:ab



The **ifname=** option is the only supported way to set custom network interface names during installation.

inst.dhcpclass=

The **inst.dhcpclass=** option specifies the DHCP vendor class identifier. The **dhcpd** service sees this value as **vendor-class-identifier**. The default value is **anaconda-\$(uname -srm)**.

inst.waitfornet=

Using the **inst.waitfornet=SECONDS** boot option causes the installation system to wait for network connectivity before installation. The value given in the **SECONDS** argument specifies the maximum amount of time to wait for network connectivity before timing out and continuing the installation process even if network connectivity is not present.

Additional resources

• For more information about networking, see the *Configuring and managing networking* document.

D.3. CONSOLE BOOT OPTIONS

This section contains information about configuring boot options for your console, monitor display, and keyboard.

console=

Use the **console=** option to specify a device that you want to use as the primary console. For example, to use a console on the first serial port, use **console=ttyS0**. Use this option in conjunction with the **inst.text** option. You can use the **console=** option multiple times. If you do, the boot message is displayed on all specified consoles, but only the last one is used by the installation program. For example, if you specify **console=ttyS0 console=ttyS1**, the installation program uses **ttyS1**.

inst.lang=

Use the **inst.lang=** option to set the language that you want to use during the installation. The **locale -a** | **grep** _ or **localectl list-locales** | **grep** _ commands return a list of locales.

inst.singlelang

Use the **inst.singlelang** option to install in single language mode, which results in no available interactive options for the installation language and language support configuration. If a language is specified using the **inst.lang** boot option or the **lang** Kickstart command, then it is used. If no language is specified, the installation program defaults to **en_US.UTF-8**.

inst.geoloc=

Use the **inst.geoloc=** option to configure geolocation usage in the installation program. Geolocation is used to preset the language and time zone, and uses the following syntax: **inst.geoloc=value**. The **value** can be any of the following parameters:

Table D.6. Values for the inst.geoloc boot option

Value	Boot option format
Disable geolocation	inst.geoloc=0
Use the Fedora GeoIP API	inst.geoloc=provider_fedora_geoip
Use the Hostip.info GeoIP API	inst.geoloc=provider_hostip

If you do not specify the **inst.geoloc=** option, the installation program uses **provider_fedora_geoip**.

inst.keymap=

Use the **inst.keymap=** option to specify the keyboard layout that you want to use for the installation.

inst.cmdline

Use the **inst.cmdline** option to force the installation program to run in command-line mode. This mode does not allow any interaction, and you must specify all options in a Kickstart file or on the command line.

inst.graphical

Use the **inst.graphical** option to force the installation program to run in graphical mode. This mode is the default.

inst.text

Use the **inst.text** option to force the installation program to run in text mode instead of graphical mode.

inst.noninteractive

Use the **inst.noninteractive** boot option to run the installation program in a non-interactive mode. User interaction is not permitted in the non-interactive mode, and **inst.noninteractive** can be used with a graphical or text installation. When the **inst.noninteractive** option is used in text mode it behaves the same as the **inst.cmdline** option.

inst.resolution=

Use the **inst.resolution**= option to specify the screen resolution in graphical mode. The format is **NxM**, where N is the screen width and M is the screen height (in pixels). The lowest supported resolution is 1024x768.

inst.vnc

Use the **inst.vnc** option to run the graphical installation using VNC. You must use a VNC client application to interact with the installation program. When VNC sharing is enabled, multiple clients can connect. A system installed using VNC starts in text mode.

inst.vncpassword=

Use the **inst.vncpassword=** option to set a password on the VNC server that is used by the installation program.

inst.vncconnect=

Use the **inst.vncconnect=** option to connect to a listening VNC client at the given host location. For example **inst.vncconnect=<host>[:<port>]** The default port is 5900. This option can be used with **vncviewer-listen**.

inst.xdriver=

Use the **inst.xdriver=** option to specify the name of the X driver that you want to use both during installation and on the installed system.

inst.usefbx

Use the **inst.usefbx** option to prompt the installation program to use the frame buffer X driver instead of a hardware-specific driver. This option is equivalent to **inst.xdriver=fbdev**.

modprobe.blacklist=

Use the **modprobe.blacklist=** option to blacklist or completely disable one or more drivers. Drivers (mods) that you disable using this option cannot load when the installation starts, and after the installation finishes, the installed system retains these settings. You can find a list of the blacklisted drivers in the /etc/modprobe.d/ directory. Use a comma-separated list to disable multiple drivers. For example:

modprobe.blacklist=ahci,firewire_ohci

inst.xtimeout=

Use the **inst.xtimeout**= option to specify the timeout in seconds for starting X server.

inst.sshd

Use the **inst.sshd** option to start the **sshd** service during installation, so that you can connect to the system during the installation using SSH, and monitor the installation progress. For more information about SSH, see the **ssh(1)** man page. By default, the **sshd** option is automatically started only on the IBM Z architecture. On other architectures, **sshd** is not started unless you use the **inst.sshd** option.



NOTE

During installation, the root account has no password by default. You can set a root password during installation with the **sshpw** Kickstart command.

inst.kdump_addon=

Use the **inst.kdump_addon=** option to enable or disable the Kdump configuration screen (add-on) in the installation program. This screen is enabled by default; use **inst.kdump_addon=off** to disable it. Disabling the add-on disables the Kdump screens in both the graphical and text-based interface as well as the **%addon com_redhat_kdump** Kickstart command.

D.4. DEBUG BOOT OPTIONS

This section contains information about the options that you can use when debugging issues.

inst.rescue

Use the **inst.rescue** option to run the rescue environment. The option is useful for trying to diagnose and fix systems.

inst.updates=

Use the **inst.updates=** option to specify the location of the **updates.img** file that you want to apply during installation. There are a number of sources for the updates.

Table D.7. inst.updates= source updates

Source	Description	Example
Updates from a network	The easiest way to use inst.updates= is to specify the network location of updates.img. This does not require any modification to the installation tree. To use this method, edit the kernel command line to include inst.updates.	inst.updates=http://some.w ebsite.com/path/to/updates. img.
Updates from a disk image	You can save an updates.img on a floppy drive or a USB key. This can be done only with an ext2 filesystem type of updates.img . To save the contents of the image on your floppy drive, insert the floppy disc and run the command.	dd if=updates.img of=/dev/fd0 bs=72k count=20. To use a USB key or flash media, replace /dev/fd0 with the device name of your USB key.
Updates from an installation tree	If you are using a CD, hard drive, HTTP, or FTP install, you can save the updates.img in the installation tree so that all installations can detect the .img file. Save the file in the images / directory. The file name must be updates.img .	For NFS installs, there are two options: You can either save the image in the images / directory, or in the RHupdates / directory in the installation tree.

inst.loglevel=

Use the inst.loglevel= option to specify the minimum level of messages logged on a terminal. This

concerns only terminal logging; log files always contain messages of all levels. Possible values for this option from the lowest to highest level are: **debug**, **info**, **warning**, **error** and **critical**. The default value is **info**, which means that by default, the logging terminal displays messages ranging from **info** to **critical**.

inst.syslog=

When installation starts, the **inst.syslog=** option sends log messages to the **syslog** process on the specified host. The remote **syslog** process must be configured to accept incoming connections.

inst.virtiolog=

Use the **inst.virtiolog=** option to specify the virtio port (a character device at /dev/virtio-ports/name) that you want to use for forwarding logs. The default value is **org.fedoraproject.anaconda.log.0**; if this port is present, it is used.

inst.zram=

The **inst.zram=** option controls the usage of zRAM swap during installation. The option creates a compressed block device inside the system RAM and uses it for swap space instead of the hard drive. This allows the installation program to run with less available memory than is possible without compression, and it might also make the installation faster. By default, swap on zRAM is enabled on systems with 2 GiB or less RAM, and disabled on systems with more than 2 GiB of memory. You can use this option to change this behavior; on a system with more than 2 GiB RAM, use **inst.zram=1** to enable the feature, and on systems with 2 GiB or less memory, use **inst.zram=0** to disable the feature.

rd.live.ram

If the **rd.live.ram** option is specified, the **stage 2** image is copied into RAM. Using this option when the **stage 2** image is on an NFS server increases the minimum required memory by the size of the image by roughly 500 MiB.

inst.nokill

The **inst.nokill** option is a debugging option that prevents the installation program from rebooting when a fatal error occurs, or at the end of the installation process. Use the **inst.nokill** option to capture installation logs which would be lost upon reboot.

inst.noshell

Use inst.noshell option if you do not want a shell on terminal session 2 (tty2) during installation.

inst.notmux

Use **inst.notmux** option if you do not want to use tmux during installation. The output is generated without terminal control characters and is meant for non-interactive uses.

inst.remotelog=

You can use the **inst.remotelog=** option to send all of the logs to a remote **host:port** using a TCP connection. The connection is retired if there is no listener and the installation proceeds as normal.

D.5. STORAGE BOOT OPTIONS

inst.nodmraid

Use the **inst.nodmraid** option to disable **dmraid** support.



WARNING

Use this option with caution. If you have a disk that is incorrectly identified as part of a firmware RAID array, it might have some stale RAID metadata on it that must be removed using the appropriate tool, for example, **dmraid** or **wipefs**.

inst.nompath

Use the **inst.nompath** option to disable support for multipath devices. This option can be used for systems on which a false-positive is encountered which incorrectly identifies a normal block device as a multipath device. There is no other reason to use this option.



WARNING

Use this option with caution. You should not use this option with multipath hardware. Using this option to attempt to install to a single path of a multipath is not supported.

inst.gpt

The **inst.gpt** boot option forces the installation program to install partition information to a GUID Partition Table (GPT) instead of a Master Boot Record (MBR). This option is not valid on UEFI-based systems, unless they are in BIOS compatibility mode. Normally, BIOS-based systems and UEFI-based systems in BIOS compatibility mode attempt to use the MBR schema for storing partitioning information, unless the disk is 232 sectors in size or larger. Disk sectors are typically 512 bytes in size, meaning that this is usually equivalent to 2 TiB. Using the **inst.gpt** boot option changes this behavior, allowing a GPT to be written to smaller disks.

D.6. DEPRECATED BOOT OPTIONS

This section contains information about deprecated boot options. These options are still accepted by the installation program but they are deprecated and are scheduled to be removed in a future release of Red Hat Enterprise Linux.

method

The **method** option is an alias for **inst.repo**.

dns

Use **nameserver** instead of **dns**. Note that nameserver does not accept comma-separated lists; use multiple nameserver options instead.

netmask, gateway, hostname

The **netmask**, **gateway**, and **hostname** options are provided as part of the **ip** option.

ip=bootif

A PXE-supplied **BOOTIF** option is used automatically, so there is no requirement to use **ip=bootif**.

ksdevice

Table D.8. Values for the ksdevice boot option

Value	Information
Not present	N/A
ksdevice=link	Ignored as this option is the same as the default behavior
ksdevice=bootif	Ignored as this option is the default if BOOTIF= is present
ksdevice=ibft	Replaced with ip=ibft . See ip for details
ksdevice= <mac></mac>	Replaced with BOOTIF=\${MAC/:/-}
ksdevice= <dev></dev>	Replaced with bootdev

D.7. REMOVED BOOT OPTIONS

This section contains the boot options that have been removed from Red Hat Enterprise Linux.



NOTE

dracut provides advanced boot options. For more information about **dracut**, see the **dracut.cmdline(7)** man page.

askmethod, asknetwork

initramfs is completely non-interactive, so the **askmethod** and **asknetwork** options have been removed. Instead, use **inst.repo** or specify the appropriate network options.

blacklist, nofirewire

The **modprobe** option handles blacklisting kernel modules; use **modprobe.blacklist=<mod1>**, **<mod2>**. You can blacklist the firewire module by using **modprobe.blacklist=firewire_ohci**.

inst.headless=

The **headless=** option specified that the system that is being installed to does not have any display hardware, and that the installation program is not required to look for any display hardware.

inst.decorated

The **inst.decorated** option was used to specify the graphical installation in a decorated window. By default, the window is not decorated, so it doesn't have a title bar, resize controls, and so on. This option was no longer required.

repo=nfsiso

Use the **inst.repo=nfs:** option.

serial

Use the console=ttyS0 option.

updates

Use the **inst.updates** option.

essid, wepkey, wpakey

Dracut does not support wireless networking.

ethtool

This option was no longer required.

gdb

This option was removed as there are many options available for debugging dracut-based **initramfs**.

inst.mediacheck

Use the dracut option rd.live.check option.

ks=floppy

Use the **inst.ks=hd:<device>** option.

display

For a remote display of the UI, use the **inst.vnc** option.

utf8

This option was no longer required as the default TERM setting behaves as expected.

noipv6

ipv6 is built into the kernel and cannot be removed by the installation program. You can disable ipv6 using **ipv6.disable=1**. This setting is used by the installed system.

upgradeany

This option was no longer required as the installation program no longer handles upgrades.

APPENDIX E. CHANGING A SUBSCRIPTION SERVICE

To manage the subscriptions, you can register a RHEL system with either Red Hat Subscription Management Server or Red Hat Satellite Server. If required, you can change the subscription service at a later point. To change the subscription service under which you are registered, unregister the system from the current service and then register it with a new service.

This section contains information about how to unregister your RHEL system from the Red Hat Subscription Management Server and Red Hat Satellite Server.

Prerequisites

You have registered your system with any one of the following:

- Red Hat Subscription Management Server
- Red Hat Satellite Server



NOTE

To receive the system updates, register your system with either of the management server.

E.1. UNREGISTERING FROM SUBSCRIPTION MANAGEMENT SERVER

This section contains information about how to unregister a RHEL system from Red Hat Subscription Management Server, using a command line and the Subscription Manager user interface.

E.1.1. Unregistering using command line

Use the **unregister** command to unregister a RHEL system from Red Hat Subscription Management Server.

Procedure

- 1. Run the unregister command as a root user, without any additional parameters.
 - # subscription-manager unregister
- 2. When prompted, provide a root password.

The system is unregistered from the Subscription Management Server, and the status 'The system is currently not registered' is displayed with the **Register** button enabled.



NOTE

To continue uninterrupted services, re-register the system with either of the management services. If you do not register the system with a management service, you may fail to receive the system updates. For more information about registering a system, see Registering your system using the command line

For more information about Red Hat Subscription Management server, see the *Using and Configuring Red Hat Subscription Manager* document.

E.1.2. Unregistering using Subscription Manager user interface

This section contains information about how to unregister a RHEL system from Red Hat Subscription Management Server, using Subscription Manager user interface.

Procedure

- 1. Log in to your system.
- 2. From the top left-hand side of the window, click Activities.
- 3. From the menu options, click the **Show Applications** icon.
- 4. Click the **Red Hat Subscription Manager** icon, or enter **Red Hat Subscription Manager** in the search.
- Enter your administrator password in the Authentication Required dialog box. The Subscriptions window appears and displays the current status of Subscriptions, System Purpose, and installed products. Unregistered products display a red X.



NOTE

Authentication is required to perform privileged tasks on the system.

6. Click the **Unregister** button.

The system is unregistered from the Subscription Management Server, and the status 'The system is currently not registered' is displayed with the **Register** button enabled.



NOTE

To continue uninterrupted services, re-register the system with either of the management services. If you do not register the system with a management service, you may fail to receive the system updates. For more information about registering a system, see Registering your system using the Subscription Manager User Interface

For more information about Red Hat Subscription Management server, see the *Using and Configuring Red Hat Subscription Manager* document.

E.2. UNREGISTERING FROM SATELLITE SERVER

To unregister a Red Hat Enterprise Linux system from Satellite Server, remove the system from Satellite Server.

For more information, see Removing a Host from Red Hat Satellite in the *Managing Hosts* guide from Satellite Server documentation.

APPENDIX F. ISCSI DISKS IN INSTALLATION PROGRAM

The Red Hat Enterprise Linux installer can discover and log in to iSCSI disks in two ways:

When the installer starts, it checks if the BIOS or add-on boot ROMs of the system support
iSCSI Boot Firmware Table (iBFT), a BIOS extension for systems that can boot from iSCSI. If
the BIOS supports iBFT, the installer reads the iSCSI target information for the configured boot
disk from the BIOS and logs in to this target, making it available as an installation target.



IMPORTANT

To connect automatically to an iSCSI target, activate a network device for accessing the target. To do so, use **ip=ibft** boot option. For more information, see Network boot options.

• You can discover and add iSCSI targets manually in the installer's graphical user interface. For more information, see Configuring storage devices.



IMPORTANT

You cannot place the **/boot** partition on iSCSI targets that you have manually added using this method - an iSCSI target containing a **/boot** partition must be configured for use with iBFT. However, in instances where the installed system is expected to boot from iSCSI with iBFT configuration provided by a method other than firmware iBFT, for example using iPXE, you can remove the **/boot** partition restriction using the **inst.nonibftiscsiboot** installer boot option.

While the installer uses **iscsiadm** to find and log into iSCSI targets, **iscsiadm** automatically stores any information about these targets in the **iscsiadm** iSCSI database. The installer then copies this database to the installed system and marks any iSCSI targets that are not used for root partition, so that the system automatically logs in to them when it starts. If the root partition is placed on an iSCSI target, initrd logs into this target and the installer does not include this target in start up scripts to avoid multiple attempts to log into the same target.

PART II. INSTALLING RED HAT ENTERPRISE LINUX ON IBM POWER SYSTEM LC SERVERS

This section describes how to install Red Hat Enterprise Linux on the IBM Power Systems LC server.

CHAPTER 7. INSTALLING RED HAT ENTERPRISE LINUX ON IBM POWER SYSTEM LC SERVERS

This guide helps you install Red Hat Enterprise Linux on a Linux on Power Systems LC server. Use these instructions for the following IBM Power System servers:

- 8335-GCA (IBM Power System S822LC)
- 8335-GTA (IBM Power System S822LC)
- 8335-GTB (IBM Power System S822LC)
- 8001-12C (IBM Power System S821LC)
- 8001-22C (IBM Power System S822LC for Big Data)
- 9006-12P (IBM Power System LC921)
- 9006-22P (IBM Power System LC922)

7.1. OVERVIEW

Use this information to install Red Hat Enterprise Linux 8 on a non-virtualized or bare metal IBM Power System LC server. This procedure follows these general steps:

- Create a bootable USB device
- Connect to the BMC firmware to set up network connection
- Connect to the BMC firmware with IPMI
- Choose your installation method:
 - Install Red Hat Enterprise Linux from USB device
 - Install Red Hat Enterprise Linux with virtual media Download your ISO file from the Red Hat Enterprise Linux website.

Additional resources

• For a list of virtualization options, see Supported Linux distributions and virtualization options for POWER8 and POWER9 Linux on Power systems.

7.1.1. Creating a bootable USB device on Linux

Follow this procedure to create a bootable USB device on a Linux system.



NOTE

This procedure is destructive and data on the USB flash drive is destroyed without a warning.

Prerequisites

- You have downloaded an installation ISO image as described in Section 2.6, "Downloading the installation ISO image".
- The **Binary DVD** ISO image is larger than 4.7 GB, so a USB flash drive that is large enough to hold the ISO image is required.

Procedure

- 1. Connect the USB flash drive to the system.
- 2. Open a terminal window and run the **dmesg** command:

\$ dmesg|tail

The **dmesg** command returns a log that details all recent events. Messages resulting from the attached USB flash drive are displayed at the bottom of the log. Record the name of the connected device.

3. Switch to user root:

\$ su -

- 4. Enter your root password when prompted.
- 5. Find the device node assigned to the drive. In this example, the drive name is **sdd**.

```
# dmesg|tail
[288954.686557] usb 2-1.8: New USB device strings: Mfr=0, Product=1, SerialNumber=2
[288954.686559] usb 2-1.8: Product: USB Storage
[288954.686562] usb 2-1.8: SerialNumber: 0000000009225
[288954.712590] usb-storage 2-1.8:1.0: USB Mass Storage device detected
[288954.712687] scsi host6: usb-storage 2-1.8:1.0
[288954.712809] usbcore: registered new interface driver usb-storage
[288954.716682] usbcore: registered new interface driver uas
[288955.717140] scsi 6:0:0:0: Direct-Access Generic STORAGE DEVICE 9228 PQ: 0
ANSI: 0
[288955.717745] sd 6:0:0:0: Attached scsi generic sg4 type 0
[288961.876382] sd 6:0:0:0: sdd Attached SCSI removable disk
```

6. Run the **dd** command to write the ISO image directly to the USB device.

dd if=/image_directory/image.iso of=/dev/device

Replace /image_directory/image.iso with the full path to the ISO image file that you downloaded, and replace device with the device name that you retrieved with the **dmesg** command. In this example, the full path to the ISO image is /home/testuser/Downloads/rhel-8-x86 64-boot.iso, and the device name is sdd:

dd if=/home/testuser/Downloads/rhel-8-x86_64-boot.iso of=/dev/sdd



NOTE

Ensure that you use the correct device name, and not the name of a partition on the device. Partition names are usually device names with a numerical suffix. For example, **sdd** is a device name, and **sdd1** is the name of a partition on the device **sdd**.

7. Wait for the **dd** command to finish writing the image to the device. The data transfer is complete when the **#** prompt appears. When the prompt is displayed, log out of the root account and unplug the USB drive. The USB drive is now ready to be used as a boot device.

7.1.2. Creating a bootable USB device on Windows

Follow the steps in this procedure to create a bootable USB device on a Windows system. The procedure varies depending on the tool. Red Hat recommends using Fedora Media Writer, available for download at https://github.com/FedoraQt/MediaWriter/releases.



NOTE

- Fedora Media Writer is a community product and is not supported by Red Hat. You can report any issues with the tool at https://github.com/FedoraQt/MediaWriter/issues.
- This procedure is destructive and data on the USB flash drive is destroyed without a warning.

Prerequisites

- You have downloaded an installation ISO image as described in Section 2.6, "Downloading the installation ISO image".
- The **Binary DVD** ISO image is larger than 4.7 GB, so a USB flash drive that is large enough to hold the ISO image is required.

Procedure

 Download and install Fedora Media Writer from https://github.com/FedoraQt/MediaWriter/releases.



NOTE

To install Fedora Media Writer on Red Hat Enterprise Linux, use the pre-built Flatpak package. You can obtain the package from the official Flatpak repository Flathub.org at https://flathub.org/apps/details/org.fedoraproject.MediaWriter.

- 2. Connect the USB flash drive to the system.
- 3. Open Fedora Media Writer.
- 4. From the main window, click **Custom Image** and select the previously downloaded Red Hat Enterprise Linux ISO image.
- 5. From **Write Custom Image** window, select the drive that you want to use.

- 6. Click **Write to disk**. The boot media creation process starts. Do not unplug the drive until the operation completes. The operation may take several minutes, depending on the size of the ISO image, and the write speed of the USB drive.
- 7. When the operation completes, unmount the USB drive. The USB drive is now ready to be used as a boot device.

7.1.3. Creating a bootable USB device on Mac OS X

Follow the steps in this procedure to create a bootable USB device on a Mac OS X system.



NOTE

This procedure is destructive and data on the USB flash drive is destroyed without a warning.

Prerequisites

- You have downloaded an installation ISO image as described in Section 2.6, "Downloading the installation ISO image".
- The **Binary DVD** ISO image is larger than 4.7 GB, so a USB flash drive that is large enough to hold the ISO image is required.

Procedure

- 1. Connect the USB flash drive to the system.
- 2. Identify the device path with the **diskutil list** command. The device path has the format of /dev/disknumber, where number is the number of the disk. The disks are numbered starting at zero (0). Typically, Disk 0 is the OS X recovery disk, and Disk 1 is the main OS X installation. In the following example, the USB device is **disk2**:

```
$ diskutil list
/dev/disk0
#:
              TYPE NAME
                                    SIZE
                                             IDENTIFIER
0:
     GUID partition scheme
                                       *500.3 GB disk0
              EFI EFI
                                 209.7 MB disk0s1
1:
       Apple_CoreStorage
2:
                                      400.0 GB disk0s2
3:
           Apple_Boot Recovery HD
                                        650.0 MB disk0s3
       Apple CoreStorage
4:
                                      98.8 GB disk0s4
           Apple Boot Recovery HD
5:
                                        650.0 MB disk0s5
/dev/disk1
              TYPE NAME
                                    SIZE
                                             IDENTIFIER
#:
           Apple HFS YosemiteHD
                                        *399.6 GB disk1
Logical Volume on disk0s1
8A142795-8036-48DF-9FC5-84506DFBB7B2
Unlocked Encrypted
/dev/disk2
              TYPE NAME
                                             IDENTIFIER
#:
                                    SIZE
0:
    FDisk partition scheme
                                       *8.1 GB
                                               disk2
                                           8.1 GB
          Windows_NTFS SanDisk USB
1:
                                                    disk2s1
```

- 3. To identify your USB flash drive, compare the NAME, TYPE and SIZE columns to your flash drive. For example, the NAME should be the title of the flash drive icon in the **Finder** tool. You can also compare these values to those in the information panel of the flash drive.
- 4. Use the diskutil unmountDisk command to unmount the flash drive's filesystem volumes:

\$ diskutil unmountDisk /dev/disknumber
Unmount of all volumes on disknumber was successful

When the command completes, the icon for the flash drive disappears from your desktop. If the icon does not disappear, you may have selected the wrong disk. Attempting to unmount the system disk accidentally returns a **failed to unmount** error.

- 5. Log in as root:
 - \$ su -
- 6. Enter your root password when prompted.
- 7. Use the **dd** command as a parameter of the sudo command to write the ISO image to the flash drive:

sudo dd if=/path/to/image.iso of=/dev/rdisknumber



NOTE

Mac OS X provides both a block (/dev/disk*) and character device (/dev/rdisk*) file for each storage device. Writing an image to the /dev/rdisknumber character device is faster than writing to the /dev/disknumber block device.

- 8. To write the /Users/user_name/Downloads/rhel-8-x86_64-boot.iso file to the /dev/rdisk2 device, run the following command:
 - # sudo dd if=/Users/user_name/Downloads/rhel-8-x86_64-boot.iso of=/dev/rdisk2
- 9. Wait for the **dd** command to finish writing the image to the device. The data transfer is complete when the # prompt appears. When the prompt is displayed, log out of the root account and unplug the USB drive. The USB drive is now ready to be used as a boot device.

7.2. COMPLETING THE PREREQUISITES AND BOOTING YOUR FIRMWARE

Before you power on the system, ensure that you have the following items:

- Ethernet cable
- VGA monitor. The VGA resolution must be set to 1024x768-60Hz.
- USB Keyboard
- Power cords and outlet for your system.
 - PC or notebook that has IPMItool level 1.8.15 or greater. (Verifying this piece of info)

• Bootable USB device

Complete these steps:

- 1. If your system belongs in a rack, install your system into that rack. For instructions, see IBM Power Systems information at https://www.ibm.com/support/knowledgecenter/.
- 2. Connect an Ethernet cable to the embedded Ethernet port next to the serial port on the back of your system. Connect the other end to your network.
- 3. Connect your VGA monitor to the VGA port on back of system.
- 4. Connect your USB keyboard to an available USB port.
- 5. Connect the power cords to the system and plug them into the outlets.

At this point, your firmware is booting. Wait for the green LED on the power button to start flashing, indicating that it is ready to use. If your system does not have a green LED indicator light, then wait 1 to 2 minutes.

7.3. CONFIGURING THE IP ADDRESS IBM POWER

To set up or enable your network connection to the baseboard management controller (BMC) firmware, use the Petitboot bootloader interface. Follow these steps:

- 1. Power on your server using the power button on the front of your system. Your system will power on to the Petitboot bootloader menu. This process takes about 1 2 minutes to complete. Do not walk away from your system! When Petitboot loads, your monitor will become active and you will need to push any key in order to interrupt the boot process.
- 2. At the Petitboot bootloader main menu, select Exit to Shell.
- 3. Run **ipmitool lan print 1**. If this command returns an IP address, verify that is correct and continue. To set a static IP address, follow these steps:
 - a. Set the mode to static by running this command: ipmitool lan set 1 ipsrc static
 - b. Set your IP address by running this command: **ipmitool lan set 1 ipaddr** *ip_address* where *ip_address* is the static IP address that you are assigning to this system.
 - c. Set your netmask by running this command: **ipmitool lan set 1 netmask netmask_address** where netmask_address is the netmask for the system.
 - d. Set your gateway server by running this command: **ipmitool lan set 1 defgw ipaddr gateway_server** where gateway_server is the gateway for this system.
 - e. Confirm the IP address by running the command **ipmitool lan print 1** again. This network interface is not active until after you perform the following steps:
- 4. To reset your firmware, run the following command: **ipmitool mc reset cold**.

 This command must complete before continuing the process; however, it does not return any information. To verify that this command has completed, ping your system BMC address (the same IP address used in your IPMItool command). When the ping returns successfully, continue to the next step.
 - a. If your ping does not return successfully within a reasonable amount of time (2 3 minutes), try these additional steps:

- i. Power your system off with this command: **ipmitool power off**.
- ii. Unplug the power cords from the back of the system. Wait 30 seconds and then apply power to boot BMC.

7.4. POWERING ON YOUR SERVER WITH IPMI

Intelligent Platform Management Interface (IPMI) is the default console to use when connecting to the OPAL firmware.

Use the default values for IPMI:

Default user: ADMIN

• Default password: admin



NOTE

After your system powers on, the Petitboot interface loads. If you do not interrupt the boot process by pressing any key within 10 seconds, Petitboot automatically boots the first option. To power on your server from a PC or notebook that is running Linux, follow these steps:

- 1. Open a terminal program on your PC or notebook.
- 2. To power on your server, run the following command:

ipmitool -I lanplus -H server_ip_address -U ipmi_user -P ipmi_password chassis power on

where server_ip_ipaddress is the IP address of the Power system and ipmi_password is the password set up for IPMI.



NOTE

If your system is already powered on, continue to activate your IPMI console.

3. Activate your IPMI console by running this command

ipmitool -I lanplus -H server_ip_address -U ipmi_user -P ipmi_password sol activate



NOTE

Use your keyboard up arrow to display the previous **ipmitool** command. You can edit previous commands to avoid typing the entire command again. If you need to power off or reboot your system, deactivate the console by running this command:

ipmitool -I lanplus -H server_ip_address -U user-name -P ipmi_password sol deactivate

To reboot the system, run this command:

ipmitool -I lanplus -H server_ip_address -U user-name -P ipmi_password chassis power reset

7.5. CHOOSE YOUR INSTALLATION METHOD ON IBM LC SERVERS

You can either install Red Hat Enterprise Linux from a USB device or through virtual media.

7.5.1. Configuring Petitboot for installation with USB device

After the system powers on, the Petitboot bootloader scans local boot devices and network interfaces to find boot options that are available to the system. For information about creating a bootable USB device, see Section 7.1.1, "Creating a bootable USB device on Linux".

Use one of the following USB devices:

- USB attached DVD player with a single USB cable to stay under 1.0 Amps
- 8 GB 2.0 USB flash drive

Procedure

Follow these steps to configure Petitboot:

1. Insert your bootable USB device into the front USB port. Petitboot displays the following option:

[USB: sdb1 / 2015-10-30-11-05-03-00]

Rescue a Red Hat Enterprise Linux system (64-bit kernel)

Test this media & install Red Hat Enterprise Linux 8.1 (64-bit kernel)

* Install Red Hat Enterprise Linux 8.1 (64-bit kernel)



NOTE

Select Rescan devices if the USB device does not appear. If your device is not detected, you may have to try a different type.

- 2. Record the UUID of the USB device. For example, the UUID of the USB device in the above example is **2015-10-30-11-05-03-00**.
- 3. Select Install Red Hat Enterprise Linux 8.1 (64-bit kernel) and press **e** (Edit) to open the Petitboot Option Editor window.
- 4. Move the cursor to the Boot arguments section and add the following information:

- 5. Select **OK** to save your options and return to the Main menu.
- 6. Verify that Install Red Hat Enterprise Linux 8.x (64-bit kernel) is selected and then press **Enter** to begin your installation.

7.5.2. Access BMC Advanced System Management interface to configure virtual media

Baseboard Management Controller (BMC) Advanced Systems Management is a remote management controller used to access system information, status, and other process for your server. You can use the BMC Advanced Systems Management to set up your installation and provide the CD image as virtual media to the Power System. However, the actual installation requires a serial-over-LAN (SOL) connection through IPMI.

To access the BMC Advanced Systems Management, open a web browser to **http://ip_address** where *ip_address* is the IP address for the BMC. Log in using these default values:

Default user name: ADMIN

• Default password: admin

In order to fully use the BMC Advanced Systems Management, you need to add the IP address of the BMC firmware to the Exceptions list in the Java Control Panel of your laptop or PC. On a Windows system, this is usually located by selecting Control Panel > Control Panel for Java.

On a Linux system, this is usually located by selecting the Control Center and then selecting the Java web browser plugin.

After accessing the Control Panel for Java, select Security tab. Then add the IP address of the BMC firmware to the Exceptions list, by clicking Edit Site List and then clicking Add. Enter the IP address and click OK.

To create a virtual CD/DVD, follow these steps:

- 1. Log into the BMC Advanced Systems Management interface from a PC or notebook using the default user name and password.
- 2. Select Remote Control > Console Redirection.

- 3. **Select** Java Console. As the console opens, you might need to direct your browser to open the **jviewer.jnlp** file by selecting to Open with Java Web Start and click OK. Accept the warning and click Run.
- 4. In the Console Redirection window, select Media > Virtual Media wizard from the menu.
- 5. In the Virtual Media wizard, select CD/DVD Media:1.
- Select CD Image and the path to the Linux distribution ISO file. For example, /tmp/RHEL-7.2-20151030.0-Server-ppc64el-dvd1.iso. Click Connect CD/DVD. If the connection is successful, the message Device redirected in Read Only Mode is displayed.
- 7. Verify that CD/DVD is shown as an option in Petitboot as **sr0**:

CD/DVD: sr0 Install Repair



NOTE

Select Rescan devices if CD/DVD does not appear.

8. **Select** Install. Ater selecting Install, your remote console may become inactive. Open or reactivate your IPMI console to complete the installation.



NOTE

Be patient! It can sometimes take a couple minutes for the installation to begin.

7.6. COMPLETING YOUR LC SERVER INSTALLATION

After you select to boot the Red Hat Enterprise Linux 8 (RHEL) installer, the installation program walks you through the steps.

1. Complete the installation program for Red Hat Enterprise Linux to set up disk options, your user name and password, time zones, and so on. The last step is to restart your system.



NOTE

While your system is restarting, remove the USB device.

2. After the system restarts, Petitboot displays the option to boot Red Hat Enterprise Linux 8. Select this option and press **Enter**.

PART III. INSTALLING RED HAT ENTERPRISE LINUX ON IBM POWER SYSTEM AC SERVERS

This section describes how to install Red Hat Enterprise Linux on the IBM Power Systems accelerated server.

CHAPTER 8. INSTALLING RED HAT ENTERPRISE LINUX ON IBM POWER SYSTEM ACCELERATED SERVERS

This guide helps you install Red Hat Enterprise Linux on an IBM Power Systems accelerated server (AC). Use these instructions for the following IBM Power System servers:

- 8335-GTG (IBM Power System AC922)
- 8335-GTH (IBM Power System AC922)
- 8335-GTX (IBM Power System AC922)

8.1. OVERVIEW

Use this information to install Red Hat Enterprise Linux on a non-virtualized, or bare metal IBM Power System accelerated server. This procedure follows these general steps:

- Connect to the BMC firmware to set up network connection
- Choose your installation method:
 - Install Red Hat Enterprise Linux from USB device
 - Install Red Hat Enterprise Linux from network
- Install Red Hat Enterprise Linux

Additional resources

• For a list of supported Red Hat Enterprise Linux versions, see Supported Linux distributions for POWER8 and POWER9 Linux on Power systems.

8.2. COMPLETING THE PREREQUISITES AND BOOTING YOUR FIRMWARE

Before you power on the system, ensure that you have the following items:

- Ethernet cable
- VGA monitor. The VGA resolution must be set to 1024x768-60Hz.
- USB Keyboard
- Power cords and outlet for your system

These instructions require that you have a network server set up with Red Hat Enterprise Linux 7.x. Download Red Hat Enterprise Linux 7.x LE ALT at https://access.redhat.com/products/red-hat-enterprise-linux/#addl-arch.

- 1. Select Downloads for Red Hat Enterprise Linux for Power, little endian
- 2. Log into your Red Hat account (if you have not already done so). Select **Red Hat Enterprise**Linux for Power 9 from the Product Variant list.

3. Click **Download Now** beside the **Red Hat Enterprise Linux Alternate Architectures 7.x Binary DVD** ISO file.

Complete these steps:

- If your system belongs in a rack, install your system into that rack. For instructions, see IBM
 Power Systems information at
 https://www.ibm.com/support/knowledgecenter/POWER9/p9hdx/POWER9welcome.htm.
- Connect an Ethernet cable to the embedded Ethernet port next to the serial port on the back of your system. Connect the other end to your network.
- Connect your VGA monitor to the VGA port on back of system.
- Connect your USB keyboard to an available USB port.
- Connect the power cords to the system and plug them into the outlets.

At this point, your firmware is booting. Wait for the green LED on the power button to start flashing, indicating that it is ready to use. If your system does not have a green LED indicator light, then wait 1 to 2 minutes.

8.3. CONFIGURING THE FIRMWARE IP ADDRESS

To set up or enable your network connection to the BMC firmware, use the Petitboot bootloader interface. Follow these steps:

- 1. Power on your server using the power button on the front of your system. Your system will power on to the Petitboot bootloader menu. This process usually takes about 1 2 minutes to complete, but may take 5 10 minutes on the first boot or after a firmware update. Do not walk away from your system! When Petitboot loads, your monitor will become active and you will need to push any key in order to interrupt the boot process.
- 2. At the Petitboot bootloader main menu, select Exit to Shell.
- 3. Run **ipmitool lan print 1**. If this command returns an IP address, verify that is correct and continue to step 4. If no IP addresses are returned, follow these steps:
 - a. Set the mode to static by running this command:
 - ipmitool lan set 1 ipsrc static
 - b. Set your IP address by running this command:
 - ipmitool lan set 1 ipaddr _ip_address_

Where *ip_address* is the static IP address that you are assigning to this system.

- c. Set your netmask by running this command:
 - ipmitool lan set 1 netmask _netmask_address_

Where netmask address is the netmask for the system.

d. Set your gateway server by running this command:

ipmitool lan set 1 defgw ipaddr _gateway_server_

Where gateway_server is the gateway for this system.

e. Confirm the IP address by running the command ipmitool lan print 1 again.



NOTE

This interface is not active until after you perform the following steps.

f. To reset your firmware, run the following command:

ipmitool raw 0x06 0x40.

This command must complete before continuing the process; however, it does not return any information. To verify that this command has completed, ping your system BMC address (the same IP address used in your IPMItool command). When the ping returns successfully, continue to the next step.



NOTE

Note: If your ping does not return successfully within a reasonable amount of time (2 - 3 minutes), try these additional steps

- g. Power your system off with this command: poweroff.h.
- h. Unplug the power cords from the back of the system. Wait 30 seconds and then apply power to boot BMC.

8.4. POWERING ON YOUR SERVER WITH OPENBMC COMMANDS



NOTE

After your system powers on, the Petitboot interface loads. If you do not interrupt the boot process by pressing any key within 10 seconds, Petitboot automatically boots the first option.

To power on your server from a PC or notebook that is running Linux, follow these steps:

- Default user name: root
- Default password: OpenBmc (where, OpenBMC is using a zero and not a capital O)
 - 1. Open a terminal program on your PC or notebook.
 - 2. Log in to the BMC by running the following commands.

ssh root@<BMC server_ip_address> root@<BMC server password>

Where BMC server_ip_address is the IP address of the BMC and BMC server password is the password to authenticate.

- 3. To power on your server, run the following command:
 - \$ root@witherspoon:~# obmcutil poweron
- 4. Connect to OS console and use the default password **OpenBmc**.
 - ssh -p 2200 root@<BMC server_ip_address> root@

Where BMC server_ip_address is the IP address of the BMC and BMC server password is the password to authenticate.

8.5. CHOOSE YOUR INSTALLATION METHOD ON IBM ACCELERATED SERVERS

You can either install Red Hat Enterprise Linux from a USB device or through the network.

8.6. CONFIGURING PETITBOOT FOR NETWORK INSTALLATION

After the system powers on, the Petitboot bootloader scans local boot devices and network interfaces to find boot options that are available to the system. To install Red Hat Enterprise Linux from a network server, you need to set up a network interface (that is not the BMC network interface).

Set up a network connection and provide the network boot detail to Petitboot by following these steps:

- a. Connect an Ethernet cable to the second Ethernet port on the back of your system. Connect the other end to your network.
- b. On the Petitboot main screen, select c to configure your system options.
- c. In the Network field of the configuration screen, enter your network information:
 - i. Select your network type
 - ii. Select your network device (remember the interface name and mac address)
 - iii. Specify your IP/mask, Gateway, and DNS server (remember these setting as you will need them in the next step)
 - iv. Select OK to return to the main menu.
- d. Back on the Petitboot main screen, select \mathbf{n} to create new options.
- e. Choose your boot device or select to Specify paths/URLs manually and then enter your boot options:
 - i. In the Kernel field, enter the path to the kernel. This field is mandatory. Enter a URL similar to this one for a network:
 - http://<http_server_ip>/ppc/ppc64/vmlinuz
 - ii. In the Initrd field, enter the path to the init ramdisk. Enter a URL similar to this one for a network:
 - http://<http_server_ip>/ppc/ppc64/initrd.img

iii. In the Boot parameter field, set up the set up the repository path and the IP address of the server where the operating system is installed. For example:

append repo=http://<http_server_ip>/
root=live:http://<http_server_ip>/os/images/install.img ipv6.disable=1 ifname=
<ethernet_interface_name>:<mac_addr> ip=<os ip>::<gateway>:<2 digit mask>:
<hostname>:<ethernet_interface_name>:none nameserver=<anem_server> inst.text

You can accept the defaults for the rest of the fields.

- f. After you set your netboot options, select OK and press **Enter**.
- g. On the Petitboot main window, select User Item 1 as your boot option and press Enter.

8.7. CONFIGURING PETITBOOT FOR INSTALLATION WITH USB DEVICE ON ACCELERATED SERVERS

After the system powers on, the Petitboot bootloader scans local boot devices and network interfaces to find boot options that are available to the system. For information about creating a bootable USB device, see Section 7.1.1, "Creating a bootable USB device on Linux".

Use one of the following USB devices:

- USB attached DVD player with a single USB cable to stay under 1.0 Amps
- 8 GB 2.0 USB flash drive

Follow these steps to configure Petitboot:

1. Insert your bootable USB device into the front USB port. Petitboot displays the following:

[USB: sdb1 / 2015-10-30-11-05-03-00]

Rescue a Red Hat Enterprise Linux system (64-bit kernel)
Test this media & install Red Hat Enterprise Linux 8.x (64-bit kernel)

* Install Red Hat Enterprise Linux 8.x (64-bit kernel)



NOTE

Select Rescan devices if the USB device does not appear. If your device is not detected, you may have to try a different type.

- 2. Record the UUID of the USB device. For example, the UUID of the USB device in the above example is 2015-10-30-11-05-03-00.
- 3. Select Install Red Hat Enterprise Linux 8.x (64-bit kernel) and press **e** (Edit) to open the Petitboot Option Editor window.
- 4. Move the cursor to the Boot arguments section and add the following information:

inst.text inst.stage2=hd:UUID=your_UUID where your_UUID is the UUID that you recorded. Petitboot Option Editor

Device: () sda2 [f8437496-78b8-4b11-9847-bb2d8b9f7cbd]

(*) sdb1 [2015-10-30-11-05-03-00] () Specify paths/URLs manually

Kernel: /ppc/ppc64/vmlinuz Initrd: /ppc/ppc64/initrd.img

Device tree:

Boot arguments: ro inst.text inst.stage2=hd:UUID=2015-10-30-11-05-03-00

[OK] [Help] [Cancel]

- 5. Select **OK** to save your options and return to the Main menu.
- 6. Verify that Install Red Hat Enterprise Linux 8.x (64-bit kernel) is selected and then press **Enter** to begin your installation.

8.8. COMPLETING YOUR ACCELERATED SERVER INSTALLATION

After you select to boot the Red Hat Enterprise Linux 8.x installer, the installation program walks you through the steps.

a. Complete the installation program for Red Hat Enterprise Linux to set up disk options, your user name and password, time zones, and so on. The last step is to restart your system.



NOTE

While your system is restarting, remove the USB device.

b. After the system restarts, Petitboot displays the option to boot Red Hat Enterprise Linux 8.x. Select this option and press Enter.

PART IV. INSTALLING RED HAT ENTERPRISE LINUX ON IBM POWER SYSTEM L SERVERS

This section describes how to install Red Hat Enterprise Linux on the IBM L servers.

CHAPTER 9. INSTALLING RED HAT ENTERPRISE LINUX ON IBM POWER SYSTEM L SERVER

This guide helps you install Red Hat Enterprise Linux on an IBM Power System L server. Use these instructions for the following IBM Power System servers:

- 88247-22L (IBM Power System S822L)
- 8247-21L (IBM Power System S812L)
- 8247-42L (IBM Power System S824L)

For a list of supported distributions, see Supported Linux distributions for POWER8 and POWER9 Linux on Power systems.

9.1. OVERVIEW

Use this information to install Red Hat Enterprise Linux on a non-virtualized or bare metal IBM Power System L server. This procedure follows these general steps:

- Complete prerequisites
- Connecting to the ASMI
 - Connect using DHCP
 - Connect using Static IP
- Enabling IPMI
- Powering on your server with IPMI
 - Connecting from Linux notebook
 - Connecting from Windows notebook
- Configuring Petitboot and installing Red Hat Enterprise Linux

9.2. COMPLETING THE PREREQUISITES AND BOOTING YOUR FIRMWARE ON L SERVER

Before you install Red Hat Enterprise Linux, ensure that you have the following items:

- Ethernet cable
- VGA monitor. The VGA resolution must be set to 1024x768-60Hz.
- USB Keyboard
- Power cords and outlet for your system

Before you power on the system, follow these steps:

• If your system belongs in a rack, install your system into that rack. For instructions, see IBM Power Systems information at https://www.ibm.com/support/knowledgecenter/.

- Remove the shipping brackets from the power supplies. Ensure that the power supplies are fully seated in the system
- Access the server control panel.
- Connect the power cords to the system and plug them into the outlets.

At this point, your firmware is booting. Wait for the green power LED on the control panel to start flashing, indicating that it is ready to use, and for the prompt 01 N OPAL T to appear on the display.

9.3. CONNECTING TO ASMI WITH DHCP

To connect to the Advanced System Management Interface (ASMI) you need to set up your network connection. You can set up DHCP or static IP.

Use this type of connection if you are using DHCP. Use these steps to find the IP address of the service processor and then connect with the ASMI web interface. If you know what IP address that your server is using, complete step 1 and then skip to Step 5: Enabling

- 1. Connect an Ethernet cable to the HMC1 or HMC2 port on the back of your Power system to your DHCP network.
- 2. Access the control panel for your server.
- 3. Scroll to function 02 using Increment (↑) or Decrement (↓) buttons (up and down arrows) and then press **Enter**.
- 4. Move the cursor to the N by pressing Enter. The display looks like this example: **02 A N< T**
- 5. Change N to M to start manual mode using the Increment (↑) or Decrement (↓) buttons. The display looks like this example: **02 A M< T**
- 6. Press **Enter** twice to exit the mode menu.
- 7. Scroll to function 30 using Increment or Decrement buttons
- 8. Press **Enter** to enter subfunction. The display looks like this example: **30****
- 9. Use the Increment (1) or Decrement (1) buttons to select a network device. 3000 displays the IP address that is assigned to ETH0 (HMC1). 3001 displays the IP address that is assigned to ETH1 (HMC2)
- 10. Press Enter to display the selected IP address. Be sure to record this IP address.
- 11. Use the Increment (1) or Decrement (1) buttons to select subfunction exit (30**).
- 12. Press **Enter** to exit subfunction mode.
- 13. Scroll to 02 using Increment (1) or Decrement (1) buttons and press **Enter**.
- 14. Change the mode to N. The display looks like this example: **02 A N< T**

9.4. CONNECTING TO ASMI WITH STATIC IP ADDRESS

Use this type of connection if you are using a static IP address. This connection configures a console interface to the ASMI.

- 1. Connect an Ethernet cable from the PC or notebook to the Ethernet port labeled HMC1 on the back of the managed system.
- 2. Set your IP address on your PC or notebook to match the default values on your Power system. IP address on PC or notebook:

169.254.2.140 Subnet mask: 255.255.255.0 The default IP address of HMC1: 169.254.2.147



NOTE

The default values of HMC1 are already set and you do not need to change them. If you want to verify the IP address, follow the steps in Connecting to ASMI with DHCP to find the IP addresses with the control panel.

If you are running Linux on your PC or notebook, set your IP address by following these steps:

- 1. Log in as root.
- 2. Start a terminal session.
- 3. Run the follow command: ifconfig -a. Record these values so that you can reset your network connection later.
- 4. Type **ifconfig** *ethx* **169.254.2.140 netmask 255.255.25.0**. Replace *ethx* with either eth0 or eth1, depending on what your PC or notepad is using.

If you are running Windows 7 on your PC or notebook, set your IP address by following these steps:

- 1. Click Start > Control Panel.
- 2. **Select** Network and Sharing Center.
- 3. **Click** the network that is displayed in Connections.
- 4. Click Properties.
- 5. If the Security dialog box is displayed, click Continue.
- 6. Select Internet Protocol Version 4.
- 7. Click Properties.
- 8. **Select** Use the following IP address.
- 9. Use 169.254.2.140 for IP address and 255.255.255.0 for Subnet mask.
- 10. Click OK > Close > Close



NOTE

If HMC1 is occupied, use HMC2. Use IP address 169.254.3.140 and Subnet mask: 255.255.255.0 on your PC or notebook. The default IP address of the HMC2 is 169.254.3.147.

9.5. ENABLING IPMI

- 1. The first time that you connect to the firmware, enter the admin ID admin and password admin. After you log in, you will be forced to change the password. Be sure to record this password!
- 2. From the main menu, select System Configuration→Firmware Configuration. Verify that OPAL is selected as your Hypervisor Mode.
- 3. Follow these steps to set a password for your IPMI session:
 - a. From the main menu, select Login Profile → Change Passwords.
 - b. Select IPMI from the list of user IDs.
 - c. Enter the current password for the administrator (set in step 2) and then enter and confirm a password for IPMI.
 - d. Click Continue.
- 4. If your Power system is not using DHCP, you need to configure Network access. From the main menu, select Network Services > Network Configuration. To configure network access, follow these steps:
 - a. From the Network Configuration display, select IPv4 and Continue.
 - b. Select Configure this interface?
 - c. Verify that IPv4 is enabled.
 - d. Select Static for the type of IP address.
 - e. Enter a name for the host system.
 - f. Enter an IP address for the system.
 - g. Enter a subnet mask.
 - h. At the bottom of the page, enter a default gateway, Domain name, and IP address for the DNS server.
 - i. After you set the values for the Network configuration, click Continue.
 - j. Click Save Settings.
 - k. If you connected with a PC or notebook, you can remove the Ethernet cable from your PC or notebook and connect it to the network switch. To continue with a console connection, change the default IP address to the IP address that you assigned to the service processor.

9.6. POWERING ON YOUR L SERVER WITH IPMI

Intelligent Platform Management Interface (IPMI) is the default console to use when you configure your Power system. If you are using a Linux notebook or PC, use the **ipmitool** utility. If you are using a Windows notebook or PC, use the **ipmiutil** utility.

As the system powers up, you might notice the following actions:

• System reference codes appear on the control panel display while the system is being started.

- The system cooling fans are activated after approximately 30 seconds and accelerate to operating speed.
- The power LED on the control panel stops flashing and remains on, indicating that system power
 is on.



NOTE

After your system powers on, the Petitboot interface loads. If you do not interrupt the boot process by pressing any key within 10 seconds, Petitboot automatically boots the first option.

9.7. POWERING ON YOUR SYSTEM FROM A NOTEBOOK OR PC RUNNING LINUX

To power on your server from a notebook or PC running Linux, follow these steps:

- 1. Open a terminal program.
- 2. To power on your server, run the following command:
 - ipmitool -I lanplus -H fsp_ip_address -P _ipmi_password_ power on

Where *ipaddress* is the IP address of the Power system and *ipmi_password* is the password set up for IPMI.

3. Immediately activate your IPMI console by running this command:

ipmitool -I lanplus -H fsp_ip_address -P ipmi_password sol activate

TIP

Use your keyboard up arrow to display the previous **ipmitool** command. You can edit previous commands to avoid typing the entire command again.



NOTE

If you need to restart your system, follow these steps:

- 1. Deactivate the console by running this command:
 - ipmitool -I lanplus -H fsp_ip_address -P ipmi_password sol deactivate
- 2. Power your system off with this command:
 - ipmitool -I lanplus -H fsp_ip_address -P ipmi_password power off
- 3. Power your system on with this command:

ipmitool -I lanplus -H fsp_ip_address -P ipmi_password power on



NOTE

If you have not already done so, insert your DVD into the DVD drive or confirm the installer image in your network

9.8. POWERING ON YOUR SYSTEM FROM A NOTEBOOK OR PC RUNNING WINDOWS

To power on your server from a notebook or PC running Windows, follow these steps:

- 1. Open a command prompt and change the directory to C:\Program Files\sourceforge\ipmiutil
- 2. To power on your server, run the following command

ipmiutil power -u -N ipaddress -P ipmi_password

Where *ipaddress* is the IP address of the Power system and *ipmi_password* is the password set up for IPMI.

3. Immediately activate your IPMI console by running this command:

ipmiutil sol -a -r -N ipaddress -P ipmi_password

TIP

Use your keyboard up arrow to display the previous **ipmiutil** command. You can edit previous commands to avoid typing the entire command again.



NOTE

If you need to restart your system, follow these steps: . Deactivate the console by running this command:

ipmiutil sol -d -N ipaddress -P ipmi_password

1. Power your system off with this command:

ipmiutil power -d -N ipaddress -P ipmi_password

1. Power your system on with this command:

ipmiutil power -u -N ipaddress -P ipmi_password



NOTE

If you have not already done so, insert your DVD into the DVD drive or confirm the installer image in your network.

9.9. CONFIGURING PETITBOOT AND INSTALLING RED HAT ENTERPRISE LINUX

After the system powers on, the Petitboot boot loader scans local boot devices and network interfaces to find boot options that are available to the system. If you do not have network connectivity or the installation DVD in the disk drive, no boot options are listed.



IMPORTANT

As shown in the following example, for Power Systems servers, Petitboot lists any identified encrypted devices in the boot menu, even if there are no boot options associated with the devices. The encrypted devices can be ignored when selecting the boot device. In this example, *disk*, *system*, and *device* indicate your specific disk, system, and device information.

Petitboot (v1.11)
[Disk: sda2 / disk]
Red Hat Enterprise Linux (system) 8.x
*[Encrypted Device: rhel device / device

System information System configuration System status log Language Rescan devices Retrieve config from URL Plugins (0) Exit to shell

Procedure

- 1. From the Petitboot window, verify that you are booting Red Hat Enterprise Linux 8.x from the DVD drive.
- 2. Select the required Red Hat Enterprise Linux installer boot option and press **Enter**.
- 3. The installation process starts.



NOTE

If you do not interrupt the boot process by pressing any key within 10 seconds after the Petitboot window appears, Petitboot automatically boots the first option.

PART V. INSTALLING RED HAT ENTERPRISE LINUX ON IBM Z

This section describes how to install Red Hat Enterprise Linux on the IBM Z architecture.

CHAPTER 10. PREPARING FOR INSTALLATION ON IBM Z

10.1. OVERVIEW OF THE IBM Z INSTALLATION PROCESS

You can install Red Hat Enterprise Linux on IBM Z interactively or in unattended mode. Installation on IBM Z differs from other architectures as it is typically performed over a network, and not from local media. The installation consists of three phases:

1. Booting the installation

- Connect to the mainframe
- Perform an initial program load (IPL), or boot from the media containing the installation program

2. Starting the installation

 Connect to the system using SSH, and start the installation program using Virtual Network Computing (VNC)

3. Completing the installation

Use the Anaconda installation program to:

- Configure the network
- Specify language support
- Specify installation source
- Specify software packages to be installed
- Complete the installation

10.2. CUSTOMIZING BOOT PARAMETERS

Before the installation can begin, you must configure some mandatory boot parameters. When installing through z/VM, these parameters must be configured before you boot into the **generic.prm** file. When installing on an LPAR, the **rd.cmdline** parameter is set to **ask** by default, meaning that you will be given a prompt on which you can enter these boot parameters. In both cases, the required parameters are the same.



NOTE

All network configuration must now be specified by either by using a parameter file, or at the prompt.

Installation source

An installation source must always be configured. Use the **inst.repo=** option to specify the package source for the installation.

Network devices

Network configuration must be provided if network access will be required during the installation. If you plan to perform an unattended (Kickstart-based) installation using only local media such as a hard drive, network configuration can be omitted.

Use the **ip=** option for basic network configuration, and other options listed in as required.

Also use the **rd.znet=** kernel option, which takes a network protocol type, a comma delimited list of sub-channels, and, optionally, comma delimited **sysfs** parameter and value pairs. This parameter can be specified multiple times to activate multiple network devices.

For example:

rd.znet=qeth,0.0.0600,0.0.0601,0.0.0602,layer2=1,portname=foo

Storage devices

At least one storage device must always be configured for text mode installations.

The **rd.dasd=** option takes a Direct Access Storage Device (DASD) adapter device bus identifier. For multiple DASDs, specify the parameter multiple times, or use a comma separated list of bus IDs. To specify a range of DASDs, specify the first and the last bus ID. Example:

rd.dasd=0.0.0200 rd.dasd=0.0.0202(ro),0.0.0203(ro:failfast),0.0.0205-0.0.0207

The **rd.zfcp=** option takes a SCSI over FCP (zFCP) adapter device bus identifier, a world wide port name (WWPN), and a FCP LUN, then activates the device. This parameter can be specified multiple times to activate multiple zFCP devices. Example:

rd.zfcp=0.0.4000,0x5005076300C213e9,0x5022000000000000

Kickstart options

If you are using a Kickstart file to perform an automatic installation, you must always specify the location of the Kickstart file using the **inst.ks=** option. For an unattended, fully automatic Kickstart installation, the **inst.cmdline** option is also useful.

An example customized **generic.prm** file containing all mandatory parameters look similar to the following example:

Example 10.1. Customized generic.prm file

ro ramdisk_size=40000 cio_ignore=all,!condev inst.repo=http://example.com/path/to/repository rd.znet=qeth,0.0.0600,0.0.0601,0.0.0602,layer2=1,portno=0,portname=foo ip=192.168.17.115::192.168.17.254:24:foobar.systemz.example.com:enc600:none nameserver=192.168.17.1 rd.dasd=0.0.0200 rd.dasd=0.0.0202 rd.zfcp=0.0.4000,0x5005076300C213e9,0x5022000000000000 inst.ks=http://example.com/path/to/kickstart

Some installation methods also require a file with a mapping of the location of installation data in the file system of the DVD or FTP server and the memory locations where the data is to be copied. The file is typically named **generic.ins**, and contains file names for the initial RAM disk, kernel image, and parameter file (**generic.prm**) and a memory location for each file. An example **generic.ins** will look similar to the following example:

Example 10.2. Sample generic.ins file

images/kernel.img 0x00000000 images/initrd.img 0x02000000 images/genericdvd.prm 0x00010480 images/initrd.addrsize 0x00010408

A valid **generic.ins** file is provided by Red Hat along with all other files required to boot the installer. Modify this file only if you want to, for example, load a different kernel version than default.

10.3. BOOTING THE INSTALLATION

10.3.1. Booting the installation

After establishing a connection with the mainframe, you need to perform an initial program load (IPL), or boot, from the medium containing the installation program. This document describes the most common methods of installing Red Hat Enterprise Linux on IBM Z. In general, any method may be used to boot the Linux installation system, which consists of a kernel (**kernel.img**) and initial RAM disk (**initrd.img**) with parameters in the **generic.prm** file supplemented by user defined parameters. Additionally, a **generic.ins** file is loaded which determines file names and memory addresses for the initrd, kernel and **generic.prm**.

The Linux installation system is also called the *installation program* in this book.

The control point from where you can start the IPL process depends on the environment where your Linux is to run. If your Linux is to run as a z/VM guest operating system, the control point is the *control program* (CP) of the hosting z/VM. If your Linux is to run in LPAR mode, the control point is the mainframe's *Support Element* (SE) or an attached IBM Z *Hardware Management Console* (HMC).

You can use the following boot media only if Linux is to run as a guest operating system under z/VM:

z/VM reader

You can use the following boot media only if Linux is to run in LPAR mode:

- SE or HMC through a remote FTP server
- SE or HMC DVD

You can use the following boot media for both z/VM and LPAR:

- DASD
- SCSI device that is attached through an FCP channel
- FCP-attached SCSI DVD

If you use DASD and FCP-attached SCSI devices (except SCSI DVDs) as boot media, you must have a configured **zipl** boot loader.

10.4. PLANNING FOR INSTALLATION ON IBM Z

10.4.1. Pre-installation

Red Hat Enterprise Linux 8 runs on z13 or later IBM mainframe systems.

The installation process assumes that you are familiar with the IBM Z and can set up *logical partitions* (LPARs) and z/VM quest virtual machines.

For installation of Red Hat Enterprise Linux on IBM Z, Red Hat supports Direct Access Storage Device (DASD) and Fiber Channel Protocol (FCP) storage devices.



IMPORTANT

DASDs are hard disks that allow a maximum of three partitions per device. For example, **dasda** can have partitions **dasda1**, **dasda2**, and **dasda3**.

Pre-installation decisions

- Whether the operating system is to be run on an LPAR or as a z/VM guest operating system.
- If swap space is needed, and how much. Although it is recommended to assign enough memory to a z/VM guest virtual machine and let z/VM do the necessary swapping, there are cases where the amount of required RAM is hard to predict. Such instances should be examined on a caseby-case basis.
- Network configuration. Red Hat Enterprise Linux 8 for IBM Z supports the following network devices:
 - Real and virtual Open Systems Adapter (OSA)
 - Real and virtual HiperSockets
 - LAN channel station (LCS) for real OSA

Disk space

You will need to calculate and allocate sufficient disk space on DASDs or SCSI disks.

- A minimum of 10 GB is needed for a server installation, 20 GB if you want to install all packages.
- Disk space is also required for any application data. After the installation, you can add or delete more DASD or SCSI disk partitions.
- The disk space used by the newly installed Red Hat Enterprise Linux system (the Linux instance) must be separate from the disk space used by other operating systems you have installed on your system.

RAM

You will have to ensure enough RAM is available.

- 1 GB is recommended for the Linux instance. With some tuning, an instance might run with as little as 512 MB RAM.
- If installing from NFS, 1 GB is sufficient. However, if installing from an HTTP or FTP source, 1.5 GB is needed.
- Running at 512 MB in text mode can be done only when installing from NFS.



NOTE

When initializing swap space on a Fixed Block Architecture (FBA) DASD using the **SWAPGEN** utility, the **FBAPART** option must be used.

Additional resources

• For additional information on IBM Z, see http://www.ibm.com/systems/z.

10.5. INSTALLING UNDER Z/VM

Use the x3270 or c3270 terminal emulator, to log in to z/VM from other Linux systems, or use the IBM 3270 terminal emulator on the IBM Z Hardware Management Console (HMC). If you are running Microsoft Windows operating system, there are several options available, and can be found through an internet search. A free native Windows port of c3270 called wc3270 also exists.

When installing under z/VM, you can boot from:

- The z/VM virtual reader
- A DASD or an FCP-attached SCSI device prepared with the zipl boot loader
- An FCP-attached SCSI DVD drive
 - 1. Log on to the z/VM guest virtual machine chosen for the Linux installation.



NOTE

If your 3270 connection is interrupted and you cannot log in again because the previous session is still active, you can replace the old session with a new one by entering the following command on the z/VM logon screen:

logon user here

Replace *user* with the name of the z/VM guest virtual machine. Depending on whether an external security manager, for example RACF, is used, the logon command might vary.

If you are not already running **CMS** (single-user operating system shipped with z/VM) in your guest, boot it now by entering the command:

cp ipl cms

Be sure not to use CMS disks such as your A disk (often device number 0191) as installation targets. To find out which disks are in use by CMS, use the following query:

query disk

You can use the following CP (z/VM Control Program, which is the z/VM hypervisor) query commands to find out about the device configuration of your z/VM guest virtual machine:

• Query the available main memory, which is called *storage* in IBM Z terminology. Your guest should have at least 1 GB of main memory.

cp query virtual storage

• Query available network devices by type:

osa

OSA - CHPID type OSD, real or virtual (VSWITCH or GuestLAN), both in QDIO mode

hsi

HiperSockets - CHPID type IQD, real or virtual (GuestLAN type Hipers)

lcs

LCS - CHPID type OSE

For example, to query all of the network device types mentioned above, run:

cp query virtual osa

• Query available DASDs. Only those that are flagged **RW** for read-write mode can be used as installation targets:

cp query virtual dasd

Query available FCP channels:

cp query virtual fcp

10.6. USING THE Z/VM READER

Perform the following steps to boot from the z/VM reader:

Procedure

1. If necessary, add the device containing the z/VM TCP/IP tools to your CMS disk list. For example:

```
cp link tcpmaint 592 592 acc 592 fm
```

Replace fm with any **FILEMODE** letter.

2. Execute the command:

ftp host

Where **host** is the host name or IP address of the FTP server that hosts the boot images (**kernel.img** and **initrd.img**).

3. Log in and execute the following commands. Use the **(repl** option if you are overwriting existing **kernel.img**, **initrd.img**, **generic.prm**, or **redhat.exec** files:

cd /location/of/install-tree/images/ ascii get generic.prm (repl get redhat.exec (repl locsite fix 80 binary get kernel.img (repl get initrd.img (repl quit

4. Optionally, check whether the files were transferred correctly by using the CMS command **filelist** to show the received files and their format. It is important that **kernel.img** and **initrd.img** have a fixed record length format denoted by F in the Format column and a record length of 80 in the Lrecl column. For example:

VMUSER FILELIST A0 V 169 Trunc=169 Size=6 Line=1 Col=1 Alt=0 Cmd Filename Filetype Fm Format Lrecl Records Blocks Date Time REDHAT EXEC B1 V 22 1 1 4/15/10 9:30:40 GENERIC PRM B1 V 44 1 1 4/15/10 9:30:32 INITRD IMG B1 F 80 118545 2316 4/15/10 9:30:25 KERNEL IMG B1 F 80 74541 912 4/15/10 9:30:17

Press **PF3** to guit filelist and return to the CMS prompt.

5. Customize boot parameters in **generic.prm** as necessary. See Section 10.2, "Customizing boot parameters" for details.

Another way to configure storage and network devices is by using a CMS configuration file. In such a case, add the *CMSDASD=* and *CMSCONFFILE=* parameters to **generic.prm**. See Section 10.12, "IBM Z/VM configuration file" for more details.

6. Finally, execute the REXX script redhat.exec to boot the installation program:

redhat

10.7. USING A PREPARED DASD

Perform the following steps to use a Prepared DASD:

Procedure

1. Boot from the prepared DASD and select the **zipl** boot menu entry referring to the Red Hat Enterprise Linux installation program. Use a command of the following form:

cp ipl DASD_device_number loadparm boot_entry_number

Replace *DASD_device_number* with the device number of the boot device, and boot_entry_number with the **zipl** configuration menu for this device. For example:

cp ipl eb1c loadparm 0

10.8. USING A PREPARED FCP ATTACHED SCSI DISK

Perform the following steps to boot from a prepared FCP-attached SCSI disk:

Procedure

1. Configure the SCSI boot loader of z/VM to access the prepared SCSI disk in the FCP Storage Area Network. Select the prepared **zipI** boot menu entry referring to the Red Hat Enterprise Linux installation program. Use a command of the following form:

cp set loaddev portname WWPN lun LUN bootprog boot_entry_number

Replace WWPN with the World Wide Port Name of the storage system and LUN with the Logical Unit Number of the disk. The 16-digit hexadecimal numbers must be split into two pairs of eight digits each. For example:

cp set loaddev portname 50050763 050b073d lun 40204011 00000000 bootprog 0

2. Optionally, confirm your settings with the command:

query loaddev

3. Boot the FCP device connected with the storage system containing the disk with the following command:

cp ipl FCP_device

For example:

cp ipl fc00

10.9. USING AN FCP-ATTACHED SCSI DVD DRIVE

Perform the following steps to use a Prepared FCP attached SCSI DVD Drive:

Prerequistes

1. This requires a SCSI DVD drive attached to an FCP-to-SCSI bridge which is in turn connected to an FCP adapter in your IBM Z. The FCP adapter must be configured and available under z/VM.

Procedure

- 1. Insert your Red Hat Enterprise Linux for IBM Z DVD into the DVD drive.
- 2. Configure the SCSI boot loader of z/VM to access the DVD drive in the FCP Storage Area Network and specify **1** for the boot entry on the Red Hat Enterprise Linux for IBM Z DVD. Use a command of the following form:

cp set loaddev portname WWPN lun FCP_LUN bootprog 1

Replace WWPN with the WWPN of the FCP-to-SCSI bridge and FCP_LUN with the LUN of the DVD drive. The 16-digit hexadecimal numbers must be split into two pairs of eight characters each. For example:

cp set loaddev portname 20010060 eb1c0103 lun 00010000 00000000 bootprog 1

3. Optionally, confirm your settings with the command:

cp query loaddev

4. IPL on the FCP device connected with the FCP-to-SCSI bridge.

cp ipl FCP_device

For example:

cp ipl fc00

10.10. USING PARAMETER AND CONFIGURATION FILES ON IBM Z

The IBM Z architecture can use a customized parameter file to pass boot parameters to the kernel and the installation program.

You need to change the parameter file if you want to:

- Install unattended with Kickstart.
- Choose non-default installation settings that are not accessible through the installation program's interactive user interface, such as rescue mode.

The parameter file can be used to set up networking non-interactively before the installation program (**Anaconda**) starts.

The kernel parameter file is limited to 895 characters plus an end-of-line character. The parameter file can be variable or fixed record format. Fixed record format increases the file size by padding each line up to the record length. Should you encounter problems with the installation program not recognizing all specified parameters in LPAR environments, you can try to put all parameters in one single line or start and end each line with a space character.

The parameter file contains kernel parameters, such as **ro**, and parameters for the installation process, such as **vncpassword=test** or **vnc**.

10.11. REQUIRED CONFIGURATION FILE PARAMETERS ON IBM Z

Several parameters are required and must be included in the parameter file. These parameters are also provided in the file **generic.prm** in directory **images**/ of the installation DVD.

- ro
 Mounts the root file system, which is a RAM disk, read-only.
- ramdisk_size=size
 Modifies the memory size reserved for the RAM disk to ensure that the Red Hat
 Enterprise Linux installation program fits within it. For example: ramdisk_size=40000.

The **generic.prm** file also contains the additional parameter **cio_ignore=all,!condev**. This setting speeds up boot and device detection on systems with many devices. The installation program transparently handles the activation of ignored devices.



IMPORTANT

To avoid installation problems arising from **cio_ignore** support not being implemented throughout the entire stack, adapt the **cio_ignore**= parameter value to your system or remove the parameter entirely from your parameter file used for booting (IPL) the installation program.

10.12. IBM Z/VM CONFIGURATION FILE

Under z/VM, you can use a configuration file on a CMS-formatted disk. The purpose of the CMS configuration file is to save space in the parameter file by moving the parameters that configure the initial network setup, the DASD, and the FCP specification out of the parameter file.

Each line of the CMS configuration file contains a single variable and its associated value, in the following shell-style syntax: **variable=value**.

You must also add the **CMSDASD** and **CMSCONFFILE** parameters to the parameter file. These parameters point the installation program to the configuration file:

CMSDASD=cmsdasd_address

Where *cmsdasd_address* is the device number of a CMS-formatted disk that contains the configuration file. This is usually the CMS user's **A** disk.

For example: CMSDASD=191

CMSCONFFILE=configuration_file

Where *configuration_file* is the name of the configuration file. This value must be specified in lower case. It is specified in a Linux file name format: *CMS_file_name.CMS_file_type*.

The CMS file **REDHAT CONF** is specified as **redhat.conf**. The CMS file name and the file type can each be from one to eight characters that follow the CMS conventions.

For example: CMSCONFFILE=redhat.conf

10.13. INSTALLATION NETWORK PARAMETERS ON IBM Z

These parameters can be used to automatically set up the preliminary network, and can be defined in the CMS configuration file. These parameters are the only parameters that can also be used in a CMS configuration file. All other parameters in other sections must be specified in the parameter file.

NETTYPE="type"

Where *type* must be one of the following: **qeth**, **lcs**, or **ctc**. The default is **qeth**. Choose **lcs** for:

OSA-Express features

Choose **geth** for:

- OSA-Express features
- HiperSockets
- Virtual connections on z/VM, including VSWTICH and Guest LAN

SUBCHANNELS="device_bus_IDs"

Where device_bus_IDs is a comma-separated list of two or three device bus IDs. The IDs must be specified in lowercase.

Provides required device bus IDs for the various network interfaces:

qeth: SUBCHANNELS="read_device_bus_id,write_device_bus_id,data_device_bus_id" lcs or ctc: SUBCHANNELS="read_device_bus_id,write_device_bus_id"

For example (a sample qeth SUBCHANNEL statement):

SUBCHANNELS="0.0.f5f0,0.0.f5f1,0.0.f5f2"

PORTNAME="osa_portname" PORTNAME="lcs_portnumber"

This variable supports OSA devices operating in gdio mode or in non-gdio mode.

When using qdio mode (**NETTYPE="qeth"**), osa_portname is the portname specified on the OSA device when operating in qeth mode.

When using non-qdio mode (**NETTYPE="Ics"**), *Ics_portnumber* is used to pass the relative port number as a decimal integer in the range of 0 through 15.

PORTNO="portnumber"

You can add either **PORTNO="0"** (to use port 0) or **PORTNO="1"** (to use port 1 of OSA features with two ports per CHPID) to the CMS configuration file to avoid being prompted for the mode.

LAYER2="value"

Where *value* can be **0** or **1**.

Use **LAYER2="0"** to operate an OSA or HiperSockets device in layer 3 mode (**NETTYPE="qeth"**). Use **LAYER2="1"** for layer 2 mode. For virtual network devices under z/VM this setting must match the definition of the GuestLAN or VSWITCH to which the device is coupled.

To use network services that operate on layer 2 (the Data Link Layer or its MAC sublayer) such as DHCP, layer 2 mode is a good choice.

The qeth device driver default for OSA devices is now layer 2 mode. To continue using the previous default of layer 3 mode, set **LAYER2="0"** explicitly.

VSWITCH="value"

Where value can be **0** or **1**.

Specify **VSWITCH="1"** when connecting to a z/VM VSWITCH or GuestLAN, or **VSWITCH="0"** (or nothing at all) when using directly attached real OSA or directly attached real HiperSockets.

MACADDR="MAC_address"

If you specify **LAYER2="1"** and **VSWITCH="0"**, you can optionally use this parameter to specify a MAC address. Linux requires six colon-separated octets as pairs lower case hex digits - for example, **MACADDR=62:a3:18:e7:bc:5f**. Note that this is different from the notation used by z/VM. If you specify **LAYER2="1"** and **VSWITCH="1"**, you must not specify the **MACADDR**, because z/VM assigns a unique MAC address to virtual network devices in layer 2 mode.

CTCPROT="value"

Where value can be 0, 1, or 3.

Specifies the CTC protocol for **NETTYPE="ctc"**. The default is **0**.

HOSTNAME="string"

Where *string* is the host name of the newly-installed Linux instance.

IPADDR="IP"

Where IP is the IP address of the new Linux instance.

NETMASK="netmask"

Where *netmask* is the netmask.

The netmask supports the syntax of a prefix integer (from 1 to 32) as specified in IPv4 *classless interdomain routing* (CIDR). For example, you can specify **24** instead of **255.255.255.0**, or **20** instead of **255.255.240.0**.

GATEWAY="gw"

Where gw is the gateway IP address for this network device.

MTU="mtu"

Where mtu is the Maximum Transmission Unit (MTU) for this network device.

DNS="server1:server2:additional_server_terms:serverN"

Where "server1:server2:additional_server_terms:serverN" is a list of DNS servers, separated by colons. For example:

DNS="10.1.2.3:10.3.2.1"

SEARCHDNS="domain1:domain2:additional_dns_terms:domainN"

Where "domain1:domain2:additional_dns_terms:domainN" is a list of the search domains, separated by colons. For example:

SEARCHDNS="subdomain.domain:domain"

You only need to specify **SEARCHDNS**= if you specify the **DNS**= parameter.

DASD=

Defines the DASD or range of DASDs to configure for the installation.

The installation program supports a comma-separated list of device bus IDs, or ranges of device bus IDs with the optional attributes **ro**, **diag**, **erplog**, and **failfast**. Optionally, you can abbreviate device bus IDs to device numbers with leading zeros stripped. Any optional attributes should be separated by colons and enclosed in parentheses. Optional attributes follow a device bus ID or a range of device bus IDs.

The only supported global option is **autodetect**. This does not support the specification of non-existent DASDs to reserve kernel device names for later addition of DASDs. Use persistent DASD device names such as /dev/disk/by-path/name to enable transparent addition of disks later. Other global options such as probeonly, nopav, or nofcx are not supported by the installation program.

Only specify those DASDs that need to be installed on your system. All unformatted DASDs specified here must be formatted after a confirmation later on in the installation program.

Add any data DASDs that are not needed for the root file system or the /boot partition after installation.

For example:

DASD="eb1c,0.0.a000-0.0.a003,eb10-eb14(diag),0.0.ab1c(ro:diag)"

For FCP-only environments, remove the **DASD=** option from the CMS configuration file to indicate no DASD is present.

FCP_n="device_bus_ID WWPN FCP_LUN"

Where:

• *n* is typically an integer value (for example **FCP_1** or **FCP_2**) but could be any string with alphabetic or numeric characters or underscores.

- device_bus_ID specifies the device bus ID of the FCP device representing the host bus adapter (HBA) (for example 0.0.fc00 for device fc00).
- *WWPN* is the world wide port name used for routing (often in conjunction with multipathing) and is as a 16-digit hex value (for example **0x50050763050b073d**).
- FCP_LUN refers to the storage logical unit identifier and is specified as a 16-digit
 hexadecimal value padded with zeroes to the right (for example 0x4020400100000000).
 These variables can be used on systems with FCP devices to activate FCP LUNs such as
 SCSI disks. Additional FCP LUNs can be activated during the installation interactively or by
 means of a Kickstart file. An example value looks similar to the following:

FCP 1="0.0.fc00 0x50050763050b073d 0x4020400100000000"



IMPORTANT

Each of the values used in the FCP parameters (for example **FCP_1** or **FCP_2**) are site-specific and are normally supplied by the FCP storage administrator.

The installation program prompts you for any required parameters not specified in the parameter or configuration file except for FCP_n.

10.14. PARAMETERS FOR KICKSTART INSTALLATIONS ON IBM Z

The following parameters can be defined in a parameter file but do not work in a CMS configuration file.

inst.ks=URL

References a Kickstart file, which usually resides on the network for Linux installations on IBM Z. Replace *URL* with the full path including the file name of the Kickstart file. This parameter activates automatic installation with Kickstart.

inst.cmdline

When this option is specified, output on line-mode terminals (such as 3270 under z/VM or operating system messages for LPAR) becomes readable, as the installation program disables escape terminal sequences that are only applicable to UNIX-like consoles. This requires installation with a Kickstart file that answers all questions, because the installation program does not support interactive user input in cmdline mode.

Ensure that your Kickstart file contains all required parameters before you use the **inst.cmdline** option. If a required command is missing, the installation will fail.

10.15. MISCELLANEOUS PARAMETERS ON IBM Z

The following parameters can be defined in a parameter file but do not work in a CMS configuration file.

rd.live.check

Turns on testing of an ISO-based installation source; for example, when booted from an FCP-attached DVD or using **inst.repo=** with an ISO on local hard disk or mounted with NFS.

inst.nompath

Disables support for multipath devices.

inst.proxy=[protocol://][username[:password]@]host[:port]

Specify a proxy to use with installation over HTTP, HTTPS or FTP.

inst.rescue

Boot into a rescue system running from a RAM disk that can be used to fix and restore an installed system.

inst.stage2=URL

Specifies a path to a tree containing **install.img**, not to the **install.img** directly. Otherwise, follows the same syntax as **inst.repo=**. If **inst.stage2** is specified, it typically takes precedence over other methods of finding **install.img**. However, if **Anaconda** finds **install.img** on local media, the **inst.stage2** URL will be ignored.

If **inst.stage2** is not specified and **install.img** cannot be found locally, **Anaconda** looks to the location given by **inst.repo=** or **method=**.

If only **inst.stage2=** is given without **inst.repo=** or **method=**, **Anaconda** uses whatever repos the installed system would have enabled by default for installation.

Use the option multiple times to specify multiple HTTP, HTTPS or FTP sources. The HTTP, HTTPS or FTP paths are then tried sequentially until one succeeds:

```
inst.stage2=http://hostname/path_to_install_tree/inst.stage2=http://hostname/path_to_install_tree/inst.stage2=http://hostname/path_to_install_tree/
```

inst.syslog=IP/hostname[:port]

Sends log messages to a remote syslog server.

The boot parameters described here are the most useful for installations and trouble shooting on IBM Z, but only a subset of those that influence the installation program.

10.16. SAMPLE PARAMETER FILE AND CMS CONFIGURATION FILE ON IBM Z

To change the parameter file, begin by extending the shipped **generic.prm** file.

Example of **generic.prm** file:

```
ro ramdisk_size=40000 cio_ignore=all,!condev
CMSDASD="191" CMSCONFFILE="redhat.conf"
inst.vnc
inst.repo=http://example.com/path/to/dvd-contents
```

Example of **redhat.conf** file configuring a QETH network device (pointed to by **CMSCONFFILE** in **generic.prm**):

```
NETTYPE="qeth"
SUBCHANNELS="0.0.0600,0.0.0601,0.0.0602"
PORTNAME="FOOBAR"
PORTNO="0"
LAYER2="1"
MACADDR="02:00:be:3a:01:f3"
HOSTNAME="foobar.systemz.example.com"
IPADDR="192.168.17.115"
NETMASK="255.255.255.0"
```

GATEWAY="192.168.17.254" DNS="192.168.17.1" SEARCHDNS="systemz.example.com:example.com" DASD="200-203"

CHAPTER 11. INSTALLING IN AN LPAR

11.1. INSTALLING IN AN LPAR

Use this procedure when installing Red Hat Enterprise Linux into an LPAR.

Procedure

When installing in a logical partition (LPAR), you can boot from:

- The FTP server
- The DASD or an FCP-attached SCSI drive prepared with the zipl boot loader
- The FCP-attached SCSI DVD drive

Perform these steps first.

- Log in on the IBM System Z Hardware Management Console (HMC) or the Support Element (SE) as a user with sufficient privileges to install a new operating system to an LPAR. The SYSPROG user is recommended.
- 2. On the Systemstab, select the mainframe you want to work with, then on the Partitions tab select the LPAR to which you wish to install.
- 3. At the bottom of the screen, under Daily, find the Operating System Messages. Double-click **Operating System Messages** to show the text console on which Linux boot messages will appear.

Continue with the procedure for your installation source.

11.2. INSTALLING IN AN LPAR USING AN FTP SERVER

Use this procedure when installing Red Hat Enterprise Linux into an LPAR using an FTP server.

Procedure

- 1. Double-click Load from Removable Media or Server.
- 2. In the dialog box that follows, select **FTP Server**, and enter the following information:
 - **Host Computer** Host name or IP address of the FTP server you want to install from, for example **ftp.redhat.com**
 - User ID Your user name on the FTP server. Or, specify anonymous.
 - Password Your password. Use your email address if you are logging in as anonymous.
 - **File location (optional)** Directory on the FTP server holding the Red Hat Enterprise Linux for System z, for example **/rhel/s390x/**.
- 3. Click Continue.
- 4. In the dialog that follows, keep the default selection of generic.ins and click Continue.

11.3. INSTALLING IN AN LPAR USING A PREPARED DASD

Use this procedure when installing Red Hat Enterprise Linux into an LPAR using an already prepared DASD.

Procedure

- 1. Double-click Load.
- 2. In the dialog box that follows, select Normal as the Load type.
- 3. As Load address, fill in the device number of the DASD.
- 4. As **Load parameter**, fill in the number corresponding the **zipl** boot menu entry that you prepared for booting the Red Hat Enterprise Linux installation program.
- 5. Click the **OK** button.

11.4. INSTALLING IN AN LPAR USING A PREPARED FCP ATTACHED SCSI DISK

Use this procedure when installing Red Hat Enterprise Linux into and LPAR using an already prepared FCP attached SCSI disk.

Procedure

- 1. Double-click Load.
- 2. In the dialog box that follows, select SCSI as the Load type.
- 3. As Load address, fill in the device number of the FCP channel connected with the SCSI disk.
- 4. As **World wide port name**, fill in the WWPN of the storage system containing the disk as a 16-digit hexadecimal number.
- 5. As Logical unit number, fill in the LUN of the disk as a 16-digit hexadecimal number.
- 6. As **Boot program selector**, fill in the number corresponding the **zipl** boot menu entry that you prepared for booting the Red Hat Enterprise Linux installation program.
- 7. Leave the **Boot record logical block address** as **0** and the **Operating system specific load parameters** empty.
- 8. Click the **OK** button.

11.5. INSTALLING IN AN LPAR USING AN FCP ATTACHED SCSI DVD DRIVE

This requires a SCSI DVD drive attached to an FCP-to-SCSI bridge which is in turn connected to an FCP adapter in your System Z machine. The FCP adapter must be configured and available in your LPAR.

Procedure

1. Insert your Red Hat Enterprise Linux for IBM Z DVD into the DVD drive.

- 2. Double-click **Load**.
- 3. In the dialog box that follows, select SCSI as the Load type.
- 4. As **Load address**, fill in the device number of the FCP channel connected with the FCP-to-SCSI bridge.
- 5. As **World wide port name**, fill in the WWPN of the FCP-to-SCSI bridge as a 16-digit hexadecimal number.
- 6. As **Logical unit number**, fill in the LUN of the DVD drive as a 16-digit hexadecimal number.
- 7. As **Boot program selector**, fill in the number **1** to select the boot entry on the Red Hat Enterprise Linux for IBM Z DVD.
- 8. Leave the **Boot record logical block address** as **0** and the **Operating system specific load parameters** empty.
- 9. Click the **OK** button.

CHAPTER 12. CONFIGURING A LINUX INSTANCE ON IBM Z

This section describes most of the common tasks for installing Red Hat Enterprise Linux on IBM Z.

12.1. ADDING DASDS

Direct Access Storage Devices (DASDs) are a type of storage commonly used with IBM Z. For more information, see Working with DASDs in the IBM Knowledge Center. The following example is how to set a DASD online, format it, and make the change persistent.

Verify that the device is attached or linked to the Linux system if running under z/VM.

CP ATTACH EB1C TO *

To link a mini disk to which you have access, run the following commands:

CP LINK RHEL7X 4B2E 4B2E MR DASD 4B2E LINKED R/W

12.2. DYNAMICALLY SETTING DASDS ONLINE

To set a DASD online, follow these steps:

 Use the cio_ignore utility to remove the DASD from the list of ignored devices and make it visible to Linux:

cio_ignore -r device_number

Replace device_number with the device number of the DASD. For example:

cio_ignore -r 4b2e

2. Set the device online. Use a command of the following form:

chccwdev -e device_number

Replace device_number with the device number of the DASD. For example:

chccwdev -e 4b2e

As an alternative, you can set the device online using sysfs attributes:

a. Use the **cd** command to change to the /sys/ directory that represents that volume:

```
# cd /sys/bus/ccw/drivers/dasd-eckd/0.0.4b2e/
# ls -l
total 0
-r--r--- 1 root root 4096 Aug 25 17:04 availability
-rw-r--r-- 1 root root 4096 Aug 25 17:04 cmb_enable
-r--r--- 1 root root 4096 Aug 25 17:04 cutype
-rw-r--- 1 root root 4096 Aug 25 17:04 detach_state
-r--r--- 1 root root 4096 Aug 25 17:04 devtype
```

```
-r--r-- 1 root root 4096 Aug 25 17:04 discipline

-rw-r--r- 1 root root 4096 Aug 25 17:04 online

-rw-r--r-- 1 root root 4096 Aug 25 17:04 readonly

-rw-r--r-- 1 root root 4096 Aug 25 17:04 use_diag
```

b. Check to see if the device is already online:

```
# cat online
0
```

c. If it is not online, enter the following command to bring it online:

```
# echo 1 > online
# cat online
1
```

3. Verify which block devnode it is being accessed as:

```
# Is -I
total 0
-r--r---- 1 root root 4096 Aug 25 17:04 availability
Irwxrwxrwx 1 root root 0 Aug 25 17:07 block -> ../../../block/dasdb
-rw-r--r-- 1 root root 4096 Aug 25 17:04 cmb_enable
-r--r--- 1 root root 4096 Aug 25 17:04 cutype
-rw-r--r-- 1 root root 4096 Aug 25 17:04 detach_state
-r--r--- 1 root root 4096 Aug 25 17:04 devtype
-r--r--- 1 root root 4096 Aug 25 17:04 discipline
-rw-r---- 1 root root 4096 Aug 25 17:04 online
-rw-r---- 1 root root 4096 Aug 25 17:04 readonly
-rw-r---- 1 root root 4096 Aug 25 17:04 use_diag
```

As shown in this example, device 4B2E is being accessed as /dev/dasdb.

These instructions set a DASD online for the current session, but this is not persistent across reboots. For instructions on how to set a DASD online persistently, see Section 12.4, "Persistently setting DASDs online". When you work with DASDs, use the persistent device symbolic links under /dev/disk/by-path/.

12.3. PREPARING A NEW DASD WITH LOW-LEVEL FORMATTING

Once the disk is online, change back to the /**root** directory and low-level format the device. This is only required once for a DASD during its entire lifetime:

```
# cd /root
# dasdfmt -b 4096 -d cdl -p /dev/disk/by-path/ccw-0.0.4b2e
Drive Geometry: 10017 Cylinders * 15 Heads = 150255 Tracks
```

I am going to format the device /dev/disk/by-path/ccw-0.0.4b2e in the following way:

Device number of device: 0x4b2e

Labelling device : yes
Disk label : VOL1
Disk identifier : 0X4B2E
Extent start (trk no) : 0
Extent end (trk no) : 150254
Compatible Disk Layout : yes

```
Blocksize : 4096

--->> ATTENTION! <<---
All data of that device will be lost.

Type "yes" to continue, no will leave the disk untouched: yes cyl 97 of 3338 |#------| 2%
```

When the progress bar reaches the end and the format is complete, **dasdfmt** prints the following output:

```
Rereading the partition table... Exiting...
```

Now, use **fdasd** to partition the DASD. You can create up to three partitions on a DASD. In our example here, we create one partition spanning the whole disk:

```
# fdasd -a /dev/disk/by-path/ccw-0.0.4b2e reading volume label ..: VOL1 reading vtoc ..........: ok
auto-creating one partition for the whole disk... writing volume label...
writing VTOC...
rereading partition table...
```

After a (low-level formatted) DASD is online, it can be used like any other disk under Linux. For instance, you can create file systems, LVM physical volumes, or swap space on its partitions, for example /dev/disk/by-path/ccw-0.0.4b2e-part1. Never use the full DASD device (dev/dasdb) for anything but the commands dasdfmt and fdasd. If you want to use the entire DASD, create one partition spanning the entire drive as in the fdasd example above.

To add additional disks later without breaking existing disk entries in, for example, /etc/fstab, use the persistent device symbolic links under /dev/disk/by-path/.

12.4. PERSISTENTLY SETTING DASDS ONLINE

The above instructions described how to activate DASDs dynamically in a running system. However, such changes are not persistent and do not survive a reboot. Making changes to the DASD configuration persistent in your Linux system depends on whether the DASDs belong to the root file system. Those DASDs required for the root file system need to be activated very early during the boot process by the **initramfs** to be able to mount the root file system.

The **cio_ignore** commands are handled transparently for persistent device configurations and you do not need to free devices from the ignore list manually.

12.5. DASDS THAT ARE PART OF THE ROOT FILE SYSTEM

The file you have to modify to add DASDs that are part of the root file system has changed in Red Hat Enterprise Linux 8. Instead of editing the /etc/zipl.conf file, the new file to be edited, and its location, may be found by running the following commands:

```
# machine_id=$(cat /etc/machine-id)
# kernel_version=$(uname -r)
# ls /boot/loader/entries/$machine_id-$kernel_version.conf
```

There is one boot option to activate DASDs early in the boot process: **rd.dasd=**. This option takes a Direct Access Storage Device (DASD) adapter device bus identifier. For multiple DASDs, specify the parameter multiple times, or use a comma separated list of bus IDs. To specify a range of DASDs, specify the first and the last bus ID. Below is an example of the

/boot/loader/entries/4ab74e52867b4f998e73e06cf23fd761-4.18.0-80.el8.s390x.conf file for a system that uses physical volumes on partitions of two DASDs for an LVM volume group **vg_devel1** that contains a logical volume **lv_root** for the root file system.

title Red Hat Enterprise Linux (4.18.0-80.el8.s390x) 8.0 (Ootpa)
version 4.18.0-80.el8.s390x
linux /boot/vmlinuz-4.18.0-80.el8.s390x
initrd /boot/initramfs-4.18.0-80.el8.s390x.img
options root=/dev/mapper/vg_devel1-lv_root crashkernel=auto rd.dasd=0.0.0200 rd.dasd=0.0.0207
rd.lvm.lv=vg_devel1/lv_root rd.lvm.lv=vg_devel1/lv_swap cio_ignore=all,!condev
rd.znet=qeth,0.0.0a00,0.0a01,0.0.0a02,layer2=1,portno=0
id rhel-20181027190514-4.18.0-80.el8.s390x
grub_users \$grub_users
grub_arg --unrestricted
grub_class kernel

To add another physical volume on a partition of a third DASD with device bus ID **0.0.202b**. To do this, add **rd.dasd=0.0.202b** to the parameters line of your boot kernel in

/boot/loader/entries/4ab74e52867b4f998e73e06cf23fd761-4.18.0-32.el8.s390x.conf:

title Red Hat Enterprise Linux (4.18.0-80.el8.s390x) 8.0 (Ootpa) version 4.18.0-80.el8.s390x linux /boot/vmlinuz-4.18.0-80.el8.s390x initrd /boot/initramfs-4.18.0-80.el8.s390x.img options root=/dev/mapper/vg_devel1-lv_root crashkernel=auto rd.dasd=0.0.0200 rd.dasd=0.0.0207 rd.dasd=0.0.202b rd.lvm.lv=vg_devel1/lv_root rd.lvm.lv=vg_devel1/lv_swap cio_ignore=all,!condev rd.znet=qeth,0.0.0a00,0.0.0a01,0.0.0a02,layer2=1,portno=0 id rhel-20181027190514-4.18.0-80.el8.s390x grub_users \$grub_users grub_users grub_arg --unrestricted grub_class kernel



WARNING

Make sure the length of the kernel command line in the configuration file does not exceed 896 bytes. Otherwise, the boot loader cannot be saved, and the installation fails.

Run **zipl** to apply the changes of the configuration file for the next IPL:

zipl -V
Using config file '/etc/zipl.conf'
Using BLS config file '/boot/loader/entries/4ab74e52867b4f998e73e06cf23fd761-4.18.0-80.el8.s390x.conf'
Target device information

Device..... 5e:00 Partition..... 5e:01 Device name...... dasda Device driver name.....: dasd DASD device number..... 0201 Type..... disk partition Disk layout..... ECKD/compatible disk layout Geometry - heads...... 15 Geometry - sectors...... 12 Geometry - cylinders....: 13356 Geometry - start..... 24 File system block size....: 4096 Physical block size....: 4096 Device size in physical blocks..: 262152 Building bootmap in '/boot' Building menu 'zipl-automatic-menu' Adding #1: IPL section '4.18.0-80.el8.s390x' (default) initial ramdisk...: /boot/initramfs-4.18.0-80.el8.s390x.img kernel image.....: /boot/vmlinuz-4.18.0-80.el8.s390x kernel parmline...: 'root=/dev/mapper/vg devel1-lv root crashkernel=auto rd.dasd=0.0.0200 rd.dasd=0.0.0207 rd.dasd=0.0.202b rd.lvm.lv=vg devel1/lv root rd.lvm.lv=vg devel1/lv swap cio ignore=all,lcondev rd.znet=qeth,0.0.0a00,0.0.0a01,0.0.0a02,layer2=1,portno=0' component address: kernel image....: 0x00010000-0x0049afff parmline.....: 0x0049b000-0x0049bfff initial ramdisk.: 0x004a0000-0x01a26fff internal loader.: 0x0000a000-0x0000cfff Preparing boot menu Interactive prompt.....: enabled Menu timeout.....: 5 seconds Default configuration...: '4.18.0-80.el8.s390x' Preparing boot device: dasda (0201). Syncing disks... Done.

12.6. FCP LUNS THAT ARE PART OF THE ROOT FILE SYSTEM

The only file you have to modify for adding FCP LUNs that are part of the root file system has changed in Red Hat Enterprise Linux 8. Instead of editing the /etc/zipl.conf file, the new file to be edited, and its location, may be found by running the following commands:

```
# machine_id=$(cat /etc/machine-id)
# kernel_version=$(uname -r)
# ls /boot/loader/entries/$machine_id-$kernel_version.conf
```

Red Hat Enterprise Linux provides a parameter to activate FCP LUNs early in the boot process: **rd.zfcp=**. The value is a comma-separated list containing the device bus ID, the WWPN as 16 digit hexadecimal number prefixed with **0x**, and the FCP LUN prefixed with **0x** and padded with zeroes to the right to have 16 hexadecimal digits.

Below is an example of the /boot/loader/entries/4ab74e52867b4f998e73e06cf23fd761-4.18.0-80.el8.s390x.conf file for a system that uses physical volumes on partitions of two FCP LUNs for an LVM volume group vg_devel1 that contains a logical volume lv_root for the root file system. For simplicity, the example shows a configuration without multipathing.

```
title Red Hat Enterprise Linux (4.18.0-32.el8.s390x) 8.0 (Ootpa) version 4.18.0-32.el8.s390x linux /boot/vmlinuz-4.18.0-32.el8.s390x initrd /boot/initramfs-4.18.0-32.el8.s390x.img options root=/dev/mapper/vg_devel1-lv_root crashkernel=auto rd.zfcp=0.0.fc00,0x5105074308c212e9,0x401040a000000000 rd.zfcp=0.0.fc00,0x5105074308c212e9,0x401040a1000000000 rd.lvm.lv=vg_devel1/lv_root rd.lvm.lv=vg_devel1/lv_swap cio_ignore=all,lcondev rd.znet=qeth,0.0.0a00,0.0.0a01,0.0.0a02,layer2=1,portno=0 id rhel-20181027190514-4.18.0-32.el8.s390x grub_users $grub_users grub_arg --unrestricted grub_class kernel
```

To add another physical volume on a partition of a third FCP LUN with device bus ID 0.0.fc00, WWPN 0x5105074308c212e9 and FCP LUN 0x401040a30000000, add

rd.zfcp=0.0.fc00,0x5105074308c212e9,0x401040a300000000 to the parameters line of your boot kernel in /boot/loader/entries/4ab74e52867b4f998e73e06cf23fd761-4.18.0-32.el8.s390x.conf. For example:

```
title Red Hat Enterprise Linux (4.18.0-32.el8.s390x) 8.0 (Ootpa) version 4.18.0-32.el8.s390x linux /boot/vmlinuz-4.18.0-32.el8.s390x initrd /boot/initramfs-4.18.0-32.el8.s390x.img options root=/dev/mapper/vg_devel1-lv_root crashkernel=auto rd.zfcp=0.0.fc00,0x5105074308c212e9,0x401040a000000000 rd.zfcp=0.0.fc00,0x5105074308c212e9,0x401040a100000000 rd.zfcp=0.0.fc00,0x5105074308c212e9,0x401040a300000000 rd.lvm.lv=vg_devel1/lv_root rd.lvm.lv=vg_devel1/lv_swap cio_ignore=all,!condev rd.znet=qeth,0.0.0a00,0.0.0a01,0.0.0a02,layer2=1,portno=0 id rhel-20181027190514-4.18.0-32.el8.s390x grub_users $grub_users $grub_users grub_arg --unrestricted grub_class kernel
```



WARNING

Make sure the length of the kernel command line in the configuration file does not exceed 896 bytes. Otherwise, the boot loader cannot be saved, and the installation fails.

Run **zipl** to apply the changes of the configuration file for the next IPL:

Partition....: 08:01 Device name....: sda Device driver name.....sd Type..... disk partition Disk layout..... SCSI disk layout

Geometry - start..... 2048 File system block size....: 4096 Physical block size..... 512

Device size in physical blocks..: 10074112

Building bootmap in '/boot/'

Building menu 'rh-automatic-menu'

Adding #1: IPL section '4.18.0-32.el8.s390x' (default) kernel image.....: /boot/vmlinuz-4.18.0-32.el8.s390x

kernel parmline...: 'root=/dev/mapper/vg devel1-lv root crashkernel=auto

rd.zfcp=0.0.fc00,0x5105074308c212e9,0x401040a000000000 rd.zfcp=0.0.fc00,0x5105074308c212e9,0x401040a100000000

rd.zfcp=0.0.fc00,0x5105074308c212e9,0x401040a300000000 rd.lvm.lv=vg devel1/lv root

rd.lvm.lv=vg devel1/lv swap cio ignore=all,!condev

rd.znet=qeth,0.0.0a00,0.0.0a01,0.0.0a02,layer2=1,portno=0'

initial ramdisk...: /boot/initramfs-4.18.0-32.el8.s390x.img

component address:

kernel image....: 0x00010000-0x007a21ff parmline.....: 0x00001000-0x000011ff initial ramdisk.: 0x02000000-0x028f63ff internal loader.: 0x0000a000-0x0000a3ff

Preparing boot device: sda.

Detected SCSI PCBIOS disk layout. Writing SCSI master boot record.

Syncing disks...

Done.

12.7. FCP LUNS THAT ARE NOT PART OF THE ROOT FILE SYSTEM

FCP LUNs that are not part of the root file system, such as data disks, are persistently configured in the file /etc/zfcp.conf. It contains one FCP LUN per line. Each line contains the device bus ID of the FCP adapter, the WWPN as 16 digit hexadecimal number prefixed with $\mathbf{0x}$, and the FCP LUN prefixed with $\mathbf{0x}$ and padded with zeroes to the right to have 16 hexadecimal digits, separated by a space or tab. Entries in /etc/zfcp.conf are activated and configured by udev when an FCP adapter is added to the system. At boot time, all FCP adapters visible to the system are added and trigger udev.

Example content of /etc/zfcp.conf:

0.0.fc00 0x5105074308c212e9 0x401040a000000000

0.0.fc00 0x5105074308c212e9 0x401040a100000000 0.0.fc00 0x5105074308c212e9 0x401040a300000000

0.0.fcd0 0x5105074308c2aee9 0x401040a000000000

0.0.fcd0 0x5105074308c2aee9 0x401040a100000000

0.0.fcd0 0x5105074308c2aee9 0x401040a300000000

Modifications of /etc/zfcp.conf only become effective after a reboot of the system or after the dynamic addition of a new FCP channel by changing the system's I/O configuration (for example, a channel is attached under z/VM). Alternatively, you can trigger the activation of a new entry in /etc/zfcp.conf for an FCP adapter which was previously not active, by executing the following commands:

1. Use the **cio_ignore** utility to remove the FCP adapter from the list of ignored devices and make it visible to Linux:

```
# cio_ignore -r device_number
```

Replace device_number with the device number of the FCP adapter. For example:

```
# cio_ignore -r fcfc
```

2. To trigger the uevent that activates the change, issue:

```
# echo add > /sys/bus/ccw/devices/device-bus-ID/uevent
```

For example:

echo add > /sys/bus/ccw/devices/0.0.fcfc/uevent

12.8. ADDING A QETH DEVICE

The **qeth** network device driver supports IBM Z OSA-Express features in QDIO mode, HiperSockets, z/VM quest LAN, and z/VM VSWITCH.

The qeth device driver assigns the same interface name for Ethernet and Hipersockets devices: **enc**<**device number>**. The bus ID is composed of the channel subsystem ID, subchannel set ID, and device number, separated by dots; the device number is the last part of the bus ID, without leading zeroes. For example, the interface name will be **enca00** for a device with the bus ID **0.0.0a00**.

12.9. DYNAMICALLY ADDING A QETH DEVICE

To add a **qeth** device dynamically, follow these steps:

1. Determine whether the **qeth** device driver modules are loaded. The following example shows loaded **qeth** modules:

```
# Ismod | grep qeth
qeth_I3 69632 0
qeth_I2 49152 1
qeth 131072 2 qeth_I3,qeth_I2
qdio 65536 3 qeth,qeth_I3,qeth_I2
ccwgroup 20480 1 qeth
```

If the output of the **Ismod** command shows that the **qeth** modules are not loaded, run the **modprobe** command to load them:

```
# modprobe qeth
```

2. Use the **cio_ignore** utility to remove the network channels from the list of ignored devices and make them visible to Linux:

```
# cio_ignore -r read_device_bus_id,write_device_bus_id,data_device_bus_id
```

Replace read_device_bus_id,write_device_bus_id,data_device_bus_id with the three device bus IDs representing a network device. For example, if the read_device_bus_id is **0.0.f500**, the write_device_bus_id is **0.0.f501**, and the data_device_bus_id is **0.0.f502**:

cio_ignore -r 0.0.f500,0.0.f501,0.0.f502

3. Use the **znetconf** utility to sense and list candidate configurations for network devices:

```
# znetconf -u
Scanning for network devices...

Device IDs Type Card Type CHPID Drv.

-------
0.0.f500,0.0.f501,0.0.f502 1731/01 OSA (QDIO) 00 qeth
0.0.f503,0.0.f504,0.0.f505 1731/01 OSA (QDIO) 01 qeth
0.0.0400,0.0.0401,0.0.0402 1731/05 HiperSockets 02 qeth
```

4. Select the configuration you want to work with and use **znetconf** to apply the configuration and to bring the configured group device online as network device.

```
# znetconf -a f500
Scanning for network devices...
Successfully configured device 0.0.f500 (encf500)
```

5. Optionally, you can also pass arguments that are configured on the group device before it is set online:

```
# znetconf -a f500 -o portname=myname
Scanning for network devices...
Successfully configured device 0.0.f500 (encf500)
```

Now you can continue to configure the **encf500** network interface.

Alternatively, you can use **sysfs** attributes to set the device online as follows:

1. Create a **qeth** group device:

```
# echo read_device_bus_id,write_device_bus_id,data_device_bus_id > /sys/bus/ccwgroup/drivers/qeth/group
```

For example:

```
# echo 0.0.f500,0.0.f501,0.0.f502 > /sys/bus/ccwgroup/drivers/qeth/group
```

2. Next, verify that the **geth** group device was created properly by looking for the read channel:

ls /sys/bus/ccwgroup/drivers/geth/0.0.f500

You can optionally set additional parameters and features, depending on the way you are setting up your system and the features you require, such as:

- portno
- layer2

portname

3. Bring the device online by writing 1 to the online **sysfs** attribute:

echo 1 > /sys/bus/ccwgroup/drivers/qeth/0.0.f500/online

4. Then verify the state of the device:

```
# cat /sys/bus/ccwgroup/drivers/qeth/0.0.f500/online 1
```

A return value of **1** indicates that the device is online, while a return value **0** indicates that the device is offline.

5. Find the interface name that was assigned to the device:

```
# cat /sys/bus/ccwgroup/drivers/qeth/0.0.f500/if_name encf500
```

Now you can continue to configure the encf500 network interface.

The following command from the **s390utils** package shows the most important settings of your **qeth** device:

```
# Isqeth encf500
Device name
                       : encf500
            : 00.__
: 0.0.f500
^ f501
               : OSD_1000
card_type
cdev0
cdev1
              : 0.0.f501
cdev2
               : 0.0.f502
chpid
              : 76
online
              : 1
                : OSAPORT
portname
portno
              : 0
state
              : UP (LAN ONLINE)
priority_queueing : always queue 0
buffer_count
                 : 16
layer2
               : 1
isolation
               : none
```

12.10. PERSISTENTLY ADDING A QETH DEVICE

To make your new **qeth** device persistent, you need to create the configuration file for your new interface. The network interface configuration files are placed in the /etc/sysconfig/network-scripts/ directory.

The network configuration files use the naming convention **ifcfg-device**, where *device* is the value found in the **if_name** file in the **qeth** group device that was created earlier, for example **enc9a0**. The **cio_ignore** commands are handled transparently for persistent device configurations and you do not need to free devices from the ignore list manually.

If a configuration file for another device of the same type already exists, the simplest way to add the config file is to copy it to the new name and then edit it:

cd /etc/sysconfig/network-scripts # cp ifcfg-enc9a0 ifcfg-enc600

To learn IDs of your network devices, use the **Isqeth** utility:

If you do not have a similar device defined, you must create a new file. Use this example of /etc/sysconfig/network-scripts/ifcfg-0.0.09a0 as a template:

IBM QETH
DEVICE=enc9a0
BOOTPROTO=static
IPADDR=10.12.20.136
NETMASK=255.255.255.0
ONBOOT=yes
NETTYPE=qeth
SUBCHANNELS=0.0.09a0,0.0.09a1,0.0.09a2
PORTNAME=OSAPORT
OPTIONS='layer2=1 portno=0'
MACADDR=02:00:00:23:65:1a
TYPE=Ethernet

Edit the new **ifcfg-0.0.0600** file as follows:

- 1. Modify the **DEVICE** statement to reflect the contents of the **if_name** file from your **ccw** group.
- 2. Modify the **IPADDR** statement to reflect the IP address of your new interface.
- 3. Modify the **NETMASK** statement as needed.
- 4. If the new interface is to be activated at boot time, then make sure **ONBOOT** is set to **yes**.
- 5. Make sure the **SUBCHANNELS** statement matches the hardware addresses for your qeth device.
- 6. Modify the **PORTNAME** statement or leave it out if it is not necessary in your environment.
- 7. You can add any valid **sysfs** attribute and its value to the **OPTIONS** parameter. The Red Hat Enterprise Linux installation program currently uses this to configure the layer mode (**layer2**) and the relative port number (**portno**) of **qeth** devices.

 The **geth** device driver default for OSA devices is new layer 2 mode. To continue using old. **if**
 - The **qeth** device driver default for OSA devices is now layer 2 mode. To continue using old **ifcfg** definitions that rely on the previous default of layer 3 mode, add **layer2=0** to the **OPTIONS** parameter.

/etc/sysconfig/network-scripts/ifcfg-0.0.0600

IBM QETH
DEVICE=enc600
BOOTPROTO=static
IPADDR=192.168.70.87

NETMASK=255.255.255.0
ONBOOT=yes
NETTYPE=qeth
SUBCHANNELS=0.0.0600,0.0.0601,0.0.0602
PORTNAME=OSAPORT
OPTIONS='layer2=1 portno=0'
MACADDR=02:00:00:b3:84:ef
TYPE=Ethernet

Changes to an **ifcfg** file only become effective after rebooting the system or after the dynamic addition of new network device channels by changing the system's I/O configuration (for example, attaching under z/VM). Alternatively, you can trigger the activation of a **ifcfg** file for network channels which were previously not active yet, by executing the following commands:

1. Use the **cio_ignore** utility to remove the network channels from the list of ignored devices and make them visible to Linux:

```
# cio_ignore -r read_device_bus_id,write_device_bus_id,data_device_bus_id
```

Replace read_device_bus_id,write_device_bus_id,data_device_bus_id with the three device bus IDs representing a network device. For example, if the read_device_bus_id is **0.0.0600**, the write_device_bus_id is **0.0.0601**, and the data_device_bus_id is **0.0.0602**:

```
# cio_ignore -r 0.0.0600,0.0.0601,0.0.0602
```

2. To trigger the uevent that activates the change, issue:

```
# echo add > /sys/bus/ccw/devices/read-channel/uevent
```

For example:

- # echo add > /sys/bus/ccw/devices/0.0.0600/uevent
- 3. Check the status of the network device:
 - # Isgeth
- 4. Now start the new interface:
 - # ifup enc600
- 5. Check the status of the interface:

```
# ip addr show enc600
3: enc600: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP group default qlen 1000
link/ether 3c:97:0e:51:38:17 brd ff:ff:ff:ff
inet 10.85.1.245/24 brd 10.34.3.255 scope global dynamic enc600
valid_lft 81487sec preferred_lft 81487sec
inet6 1574:12:5:1185:3e97:eff:fe51:3817/64 scope global noprefixroute dynamic
valid_lft 2591994sec preferred_lft 604794sec
inet6 fe45::a455:eff:d078:3847/64 scope link
valid_lft forever preferred_lft forever
```

6. Check the routing for the new interface:

```
# ip route
default via 10.85.1.245 dev enc600 proto static metric 1024
12.34.4.95/24 dev enp0s25 proto kernel scope link src 12.34.4.201
12.38.4.128 via 12.38.19.254 dev enp0s25 proto dhcp metric 1
192.168.122.0/24 dev virbr0 proto kernel scope link src 192.168.122.1
```

7. Verify your changes by using the **ping** utility to ping the gateway or another host on the subnet of the new device:

```
# ping -c 1 192.168.70.8
PING 192.168.70.8 (192.168.70.8) 56(84) bytes of data.
64 bytes from 192.168.70.8: icmp_seq=0 ttl=63 time=8.07 ms
```

8. If the default route information has changed, you must also update /etc/sysconfig/network accordingly.

12.11. CONFIGURING AN IBM Z NETWORK DEVICE FOR NETWORK ROOT FILE SYSTEM

To add a network device that is required to access the root file system, you only have to change the boot options. The boot options can be in a parameter file, however, the /etc/zipl.conf file no longer contains specifications of the boot records. The file that needs to be modified can be located using the following commands:

```
# machine_id=$(cat /etc/machine-id)
# kernel_version=$(uname -r)
# ls /boot/loader/entries/$machine_id-$kernel_version.conf
```

There is no need to recreate the initramfs.

Dracut, the **mkinitrd** successor that provides the functionality in the initramfs that in turn replaces **initrd**, provides a boot parameter to activate network devices on IBM Z early in the boot process: **rd.znet=**.

As input, this parameter takes a comma-separated list of the **NETTYPE** (qeth, lcs, ctc), two (lcs, ctc) or three (qeth) device bus IDs, and optional additional parameters consisting of key-value pairs corresponding to network device sysfs attributes. This parameter configures and activates the IBM Z network hardware. The configuration of IP addresses and other network specifics works the same as for other platforms. See the **dracut** documentation for more details.

The cio_ignore commands for the network channels are handled transparently on boot.

Example boot options for a root file system accessed over the network through NFS:

```
root=10.16.105.196:/nfs/nfs_root cio_ignore=all,!condev rd.znet=qeth,0.0.0a00,0.0.0a01,0.0.0a02,layer2=1,portno=0,portname=OSAPORT ip=10.16.105.197:10.16.105.196:10.16.111.254:255.255.248.0:nfs-server.subdomain.domain:enc9a0:n one rd_NO_LUKS rd_NO_LVM rd_NO_MD rd_NO_DM LANG=en_US.UTF-8 SYSFONT=latarcyrheb-sun16 KEYTABLE=us
```